

# Calyx Tutorial

## PLDI 2023

<https://calyxir.org/tutorial>

Time	Topic
9-9:20am	Introduction to Calyx, and setting up
9:20-9:55am	Your first Calyx program
9:55-10am	fud, the hardware tool composer
10-10:10am	MrXL, a map-reduce frontend
10:10-10:50am	Implement a map operation for MrXL
10:50-10:55am	Cider, the Calyx interactive debugger
10:55-11am	Pollen, a pangenome analysis DSL
11-11:20am	Break
11:20am-12:20pm	Contest: extensions to MrXL!
12:20-12:30pm	Award ceremony and closing remarks

# opening remarks

- Hardware is important and Verilog is bad... we can fix that
  - Compilers are cool (as opposed to hand-writing fixed-function code)
- Setting up: check out Docker container, run fud check, run runt

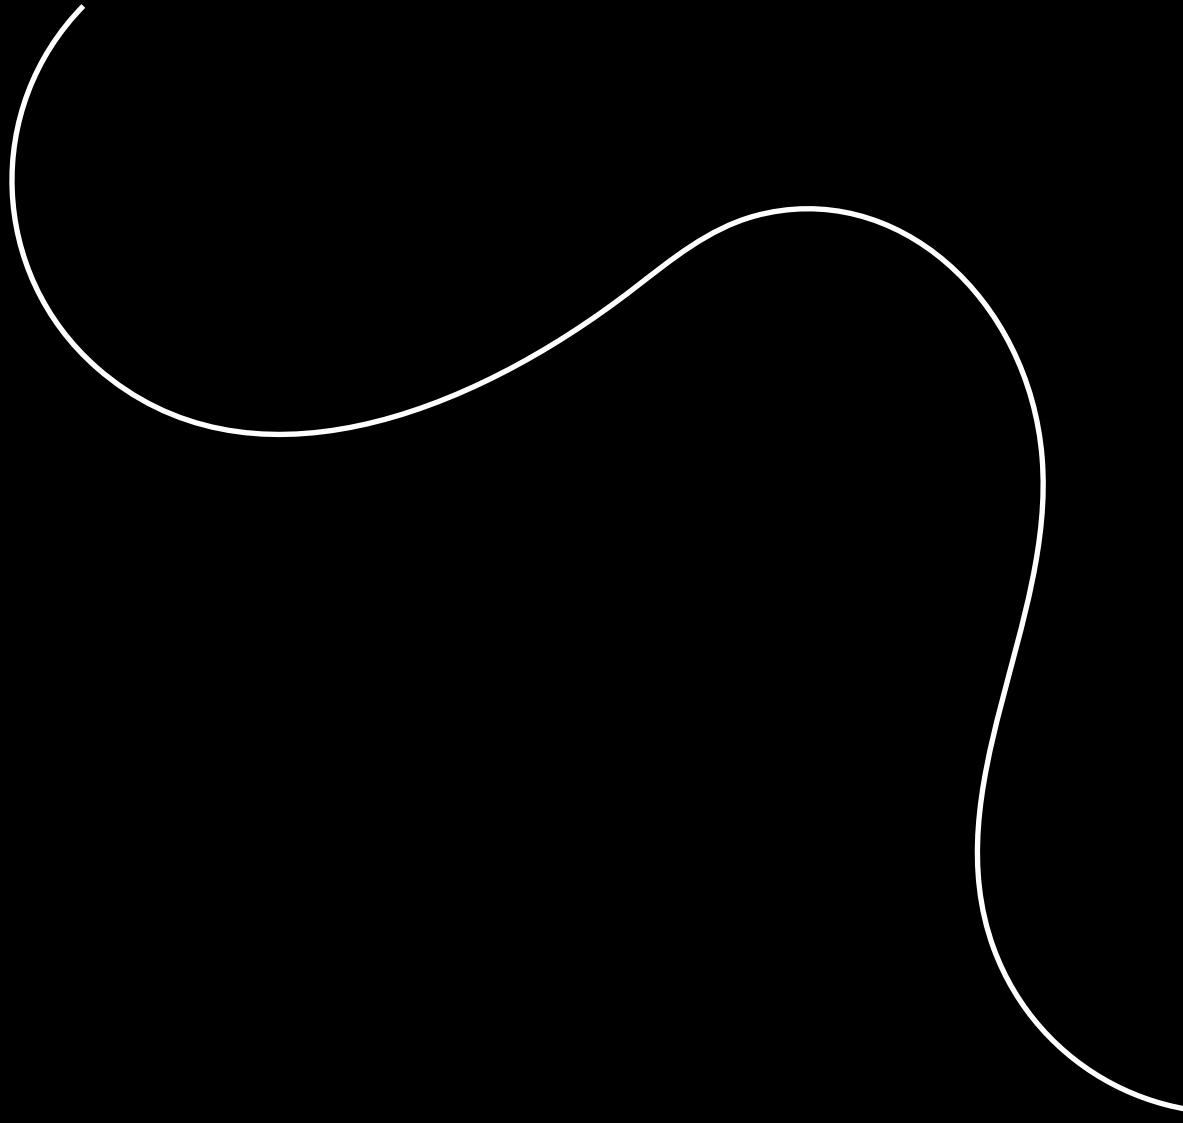
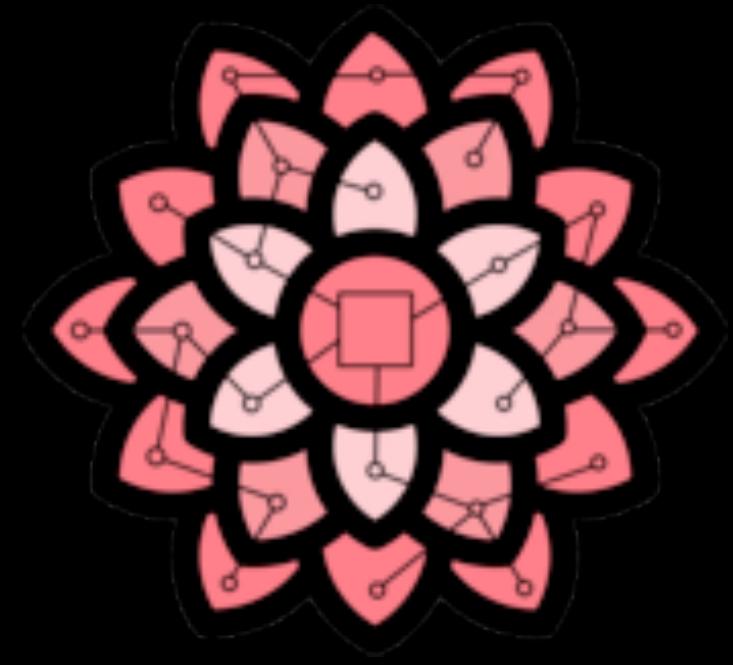
- Companies are investing in custom hardware
- Countries are investing in hardware design
- And mostly importantly, you can write papers about hardware

- But, using Verilog is a pain
  - Bad language
  - Bad tools
  - Complex constraints (name drop things like timing closure)

1. Install Docker
2. Get the Docker container
3. Run the following commands in the container
  - `fud check ...`
  - `runt -i ...`

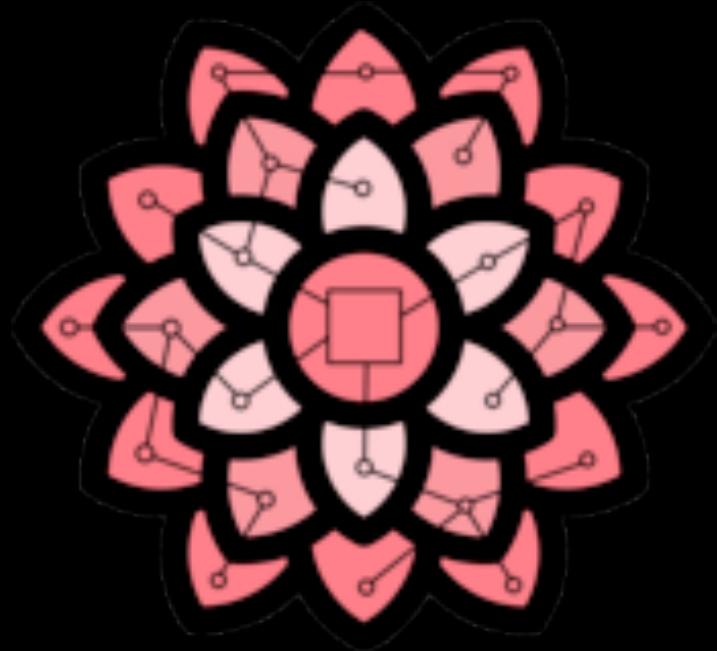
- Calyx programs have components
  - Which have cells, wires, and control
  - Wires are structured using groups
  - Control operators schedule them
- Link to language tutorial

Calix

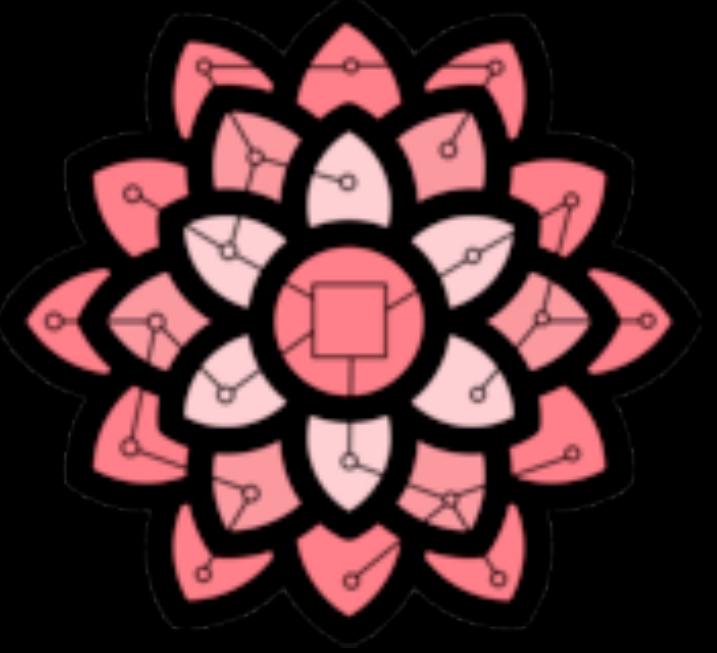


Calix

# MrXL



Calyx



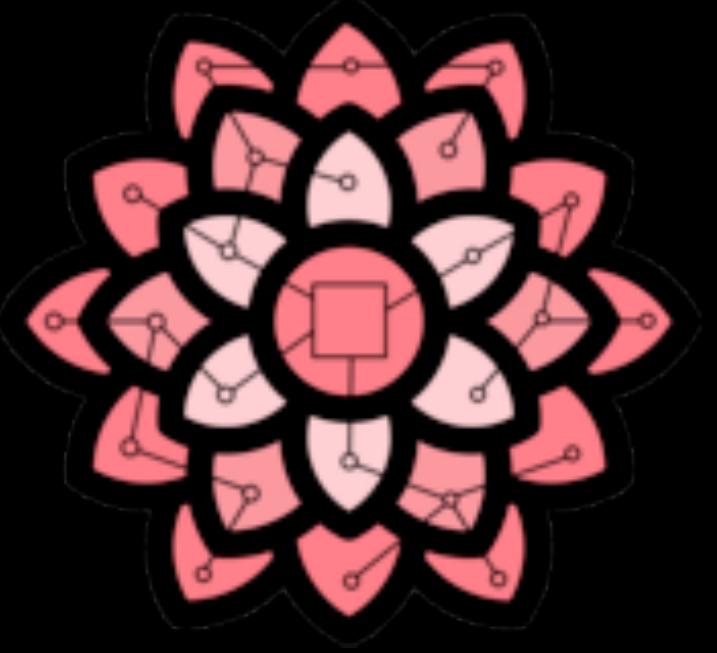
MrXL



tvm



Calyx



**MrXL**



**MyDSL**



# MrXL



fud, the Calyx driver

# fud, the Calyx driver

```
fud e squares.mrxl --from mrxl --to calyx
```

# fud, the Calyx driver

```
fud e squares.mrxl --from mrxl --to calyx
```

# fud, the Calyx driver

```
import "primitives/core.futil";
import "primitives/binary_operators.futil";
component main() -> () {
    cells {
        @external avec_b0 = std_mem_d1(32, 2, 32);
        @external avec_b1 = std_mem_d1(32, 2, 32);
        @external squares_b0 = std_mem_d1(32, 2, 32);
        @external squares_b1 = std_mem_d1(32, 2, 32);
        idx_b0_0 = std_reg(32);
        incr_b0_0 = std_add(32);
        lt_b0_0 = std_lt(32);
        mul_b0_0 = std_mult_pipe(32);
        idx_b1_0 = std_reg(32);
        incr_b1_0 = std_add(32);
        lt_b1_0 = std_lt(32);
        mul_b1_0 = std_mult_pipe(32);
    }
    wires {
        group incr_idx_b0_0 {
            incr_b0_0.left = idx_b0_0.out;
            incr_b0_0.right = 32'd1;
        }
    }
}
```

fud, the Calyx driver

# fud, the Calyx driver

```
fud e squares.mrxml --from mrxml --to verilog
```

# fud, the Calyx driver

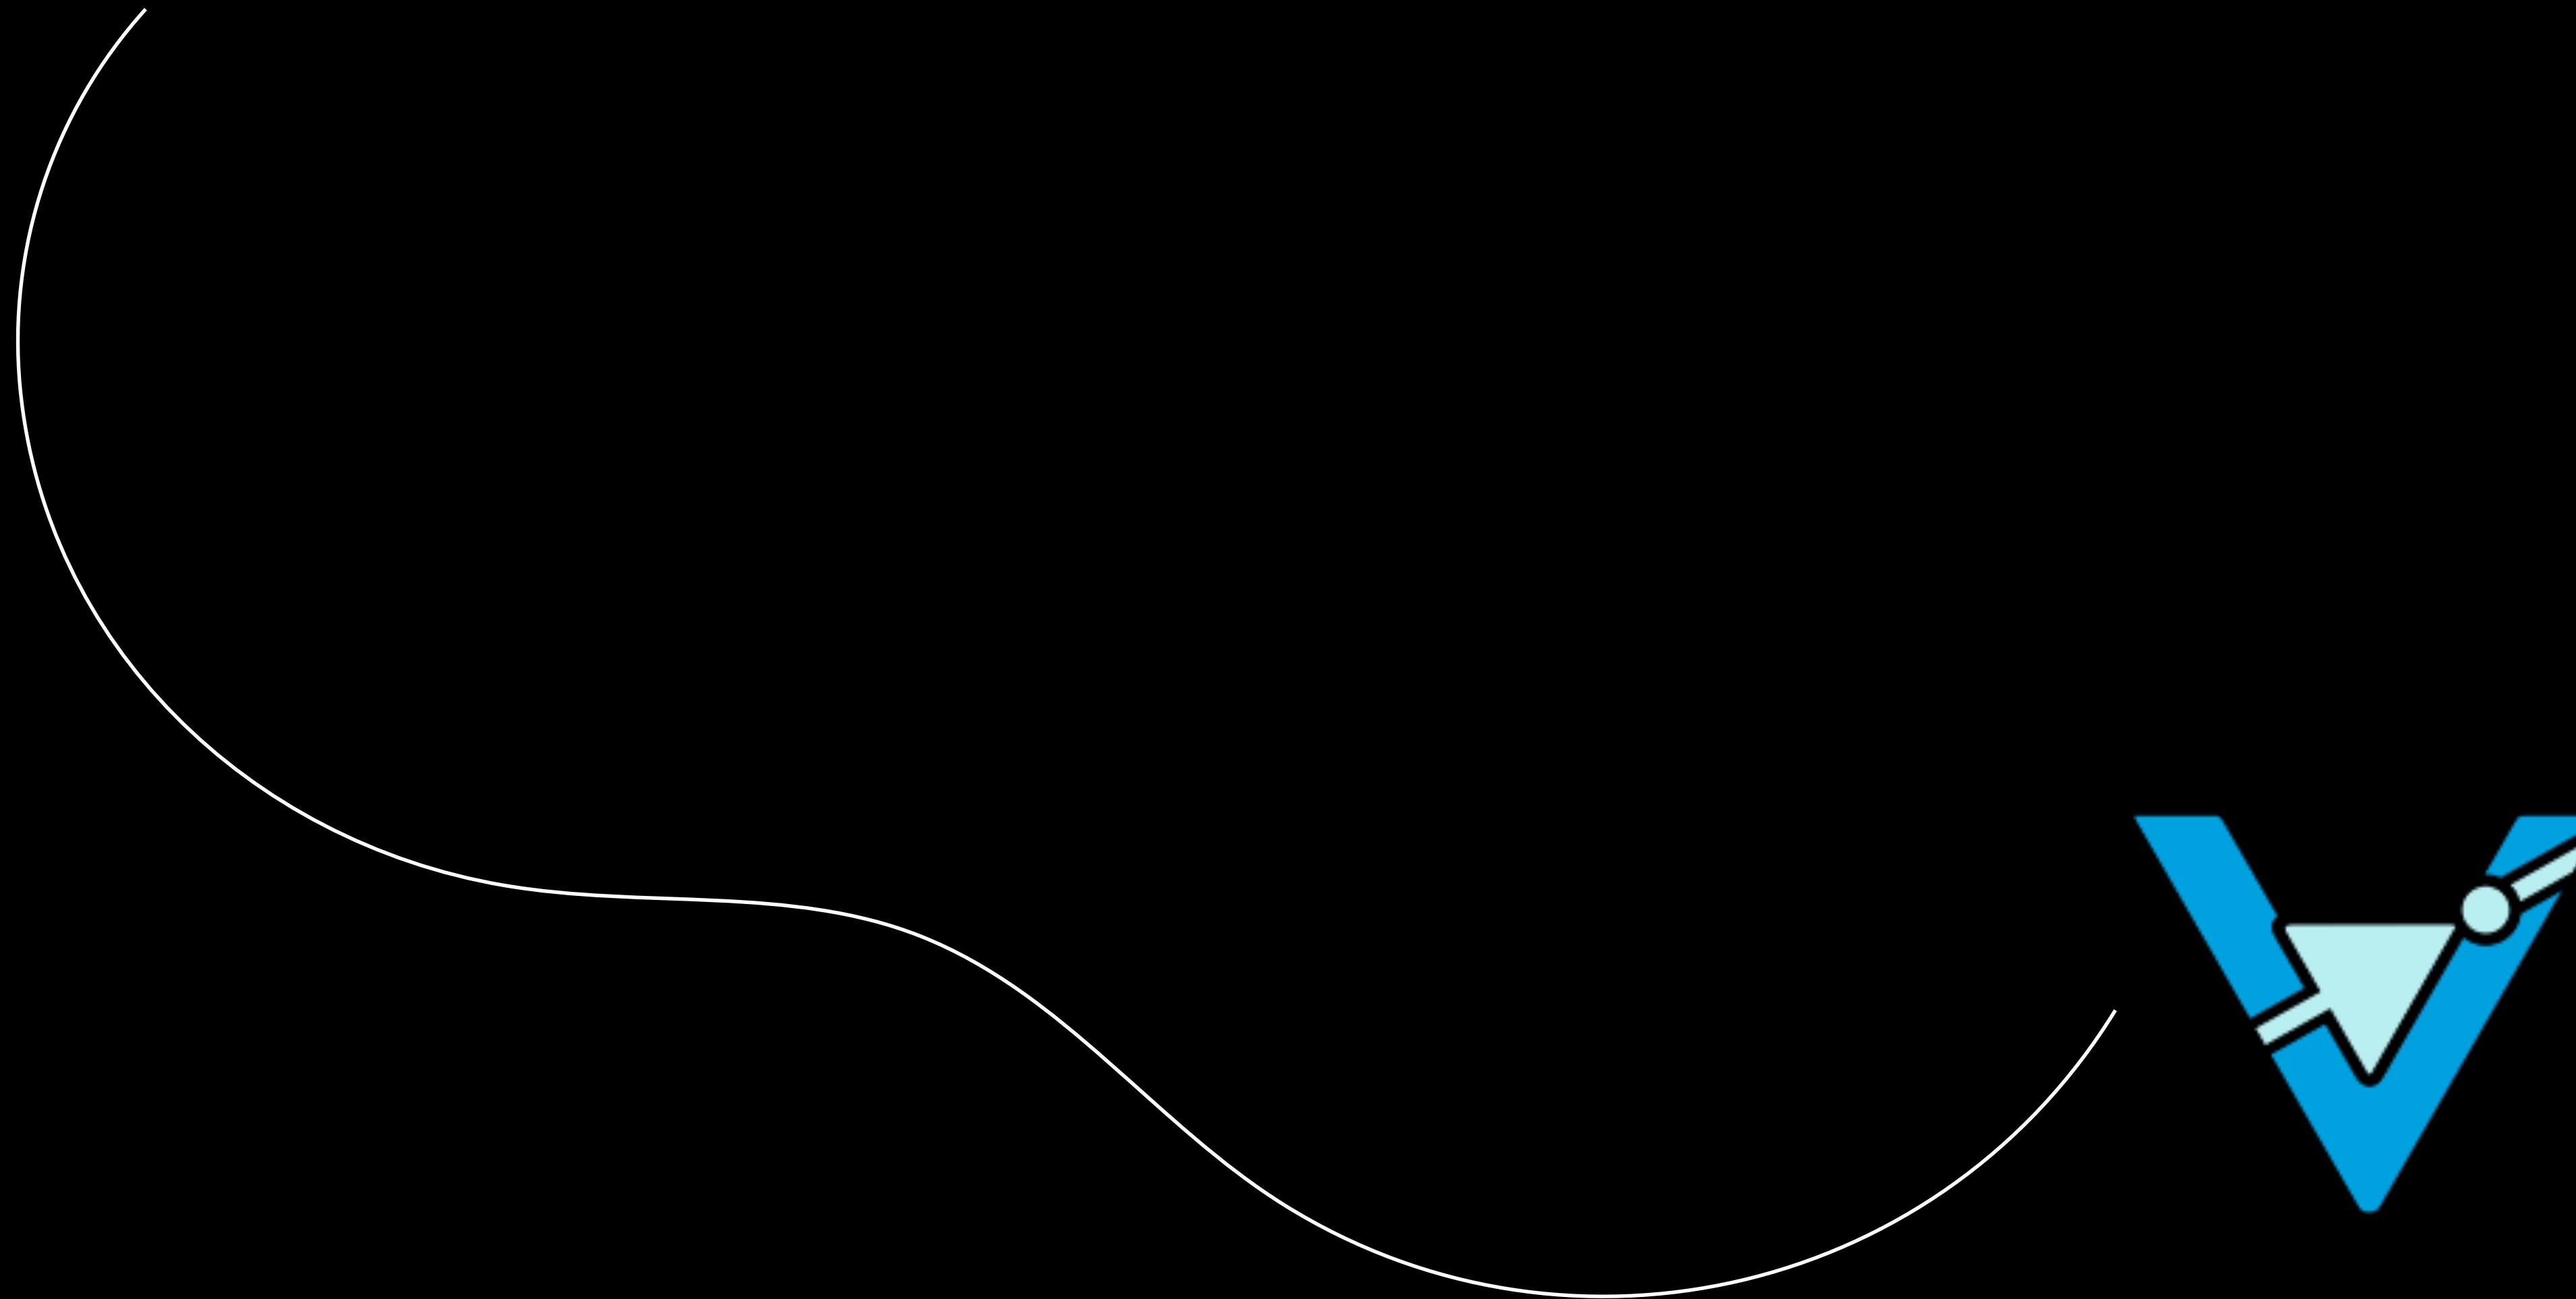
```
fud e squares.mrxml --from mrxml --to verilog
```

# fud, the Calyx driver

```
module main(
    input logic go,
    input logic clk,
    input logic reset,
    output logic done
);
// COMPONENT START: main
string DATA;
int CODE;
initial begin
    CODE = $value$plusargs("DATA=%s", DATA);
    $display("DATA (path to meminit files): %s", DATA);
    $readmemh({DATA, "/avec_b0.dat"}, avec_b0.mem);
    $readmemh({DATA, "/avec_b1.dat"}, avec_b1.mem);
    $readmemh({DATA, "/squares_b0.dat"}, squares_b0.mem);
    $readmemh({DATA, "/squares_b1.dat"}, squares_b1.mem);
end
final begin
    $writememh({DATA, "/avec_b0.out"}, avec_b0.mem);
    $writememh({DATA, "/avec_b1.out"}, avec_b1.mem);
    $writememh({DATA, "/squares_b0.out"}, squares_b0.mem);
    $writememh({DATA, "/squares_b1.out"}, squares_b1.mem);
end
logic [31:0] avec_b0_addr0;
logic [31:0] avec_b0_write_data;
```



# MrXL



# MrXL

mem



MrXL

Calyx

mem



MrXL



mem



MrXL



mem



MrXL



mem



MrXL

CalX

mem

.SV



MrXL

CalX

.SV

mem



MrXL

CalX

.SV

mem





```
fud e squares.mrxl --from mrxl --to dat \
--through verilog -s mrxl.data squares.mrxl.data
```

```
fud e squares.mrxml --from mrxml --to dat \
--through verilog -s mrxml.data squares.mrxml.data
```

```
{  
  "cycles": 17,  
  "memories": {  
    "avec_b0": [  
      0,  
      1  
    ],  
    "avec_b1": [  
      4,  
      5  
    ],  
    "squares_b0": [  
      0,  
      1  
    ],  
    "squares_b1": [  
      16,  
      25  
    ]  
  }  
}
```

```
fud e squares.mrxl --from mrxl --to dat \
--through verilog -s mrxl.data squares.mrxl.data -n
```

```
fud e squares.mrxl --from mrxl --to dat \
--through verilog -s mrxl.data squares.mrxl.data -n
```

```
fud will perform the following steps:
- mrxl.mktmp: Make temporary directory to store Verilator build files.
- mrxl.set_mrxl_prog: Set stages.mrxl.prog as `input`
- mrxl.mrxl-data.get_mrxl_prog: Dynamically retrieve the value of stages.mrxl.prog
- mrxl.mrxl-data.convert_mrxl_data_to_calyx_data: Converts MrXL input into calyx input
- transform: transform input to String
- mrxl.save_data: Save verilog.data in `tmpdir` and update stages.verilog.data
- mrxl.run_mrxl: mrxl
- calyx.run_futil: /Users/ps/Research/calyx-ref/calyx/target/debug/calyx -l /Users/ps/Re
search/calyx-ref/calyx -b verilog
- transform: transform input to Path
- verilog.mktmp: Make temporary directory to store Verilator build files.
- verilog.get_verilog_data: Dynamically retrieve the value of stages.verilog.data
- verilog.check_verilog_for_mem_read: Read input verilog to see if `verilog.data` needs
to be set.
- verilog.json_to_dat: Converts a `json` data format into a series of `dat` files insid
e the given
    temporary directory.
- verilog.compile_with_verilator: verilator --trace {input_path} /Users/ps/Research/caly
x-ref/calyx/fud/icarus/tb.sv --binary --top-module TOP --Mdir {tmpdir_name} -fno-inline
- verilog.simulate: Simulates compiled Verilator code.
- verilog.output_json: Convert .dat files back into a json and extract simulated cycles
from log.
- verilog.cleanup: Cleanup Verilator build files that we no longer need.
```



```
fud e squares.mrxml --from mrxml --to calyx
```

```
fud e squares.mrxml --from mrxml --to verilog
```

```
fud e squares.mrxl --from mrxl --to dat \
--through icarus-verilog -s mrxl.data squares.mrxl.data
```

```
fud e squares.mrxl --from mrxl --to dat \
--through verilog -s mrxl.data squares.mrxl.data
```

```
fud e squares.mrxml --from mrxml --to interpreter-out \
-s mrxml.data squares.mrxml.data
```

```
fud e squares.mrxl --from mrxl --to interpreter-out \
-s mrxl.data squares.mrxl.data
```

```
fud e squares.xclbin --from xclbin --to fpga \
-s fpga.data squares.mrxl.data
```

# MrXL: map-reduce accelerator

Real frontends are not built in a day,  
so let's work on a toy frontend for Calyx.

# MrXL: map-reduce accelerator

Real frontends are not built in a day,  
so let's work on a toy frontend for Calyx.

We will introduce MrXL, and then you will implement its map operation.

# MrXL: map-reduce accelerator

Real frontends are not built in a day,  
so let's work on a toy frontend for Calyx.

We will introduce MrXL, and then you will implement its map operation.

Watch your directories!  
Work in `frontends/mrxl`

# anatomy of a frontend

program.mrxml

program.futil

# anatomy of a frontend

program.mrxml

program.ast

program.futil

# anatomy of a frontend

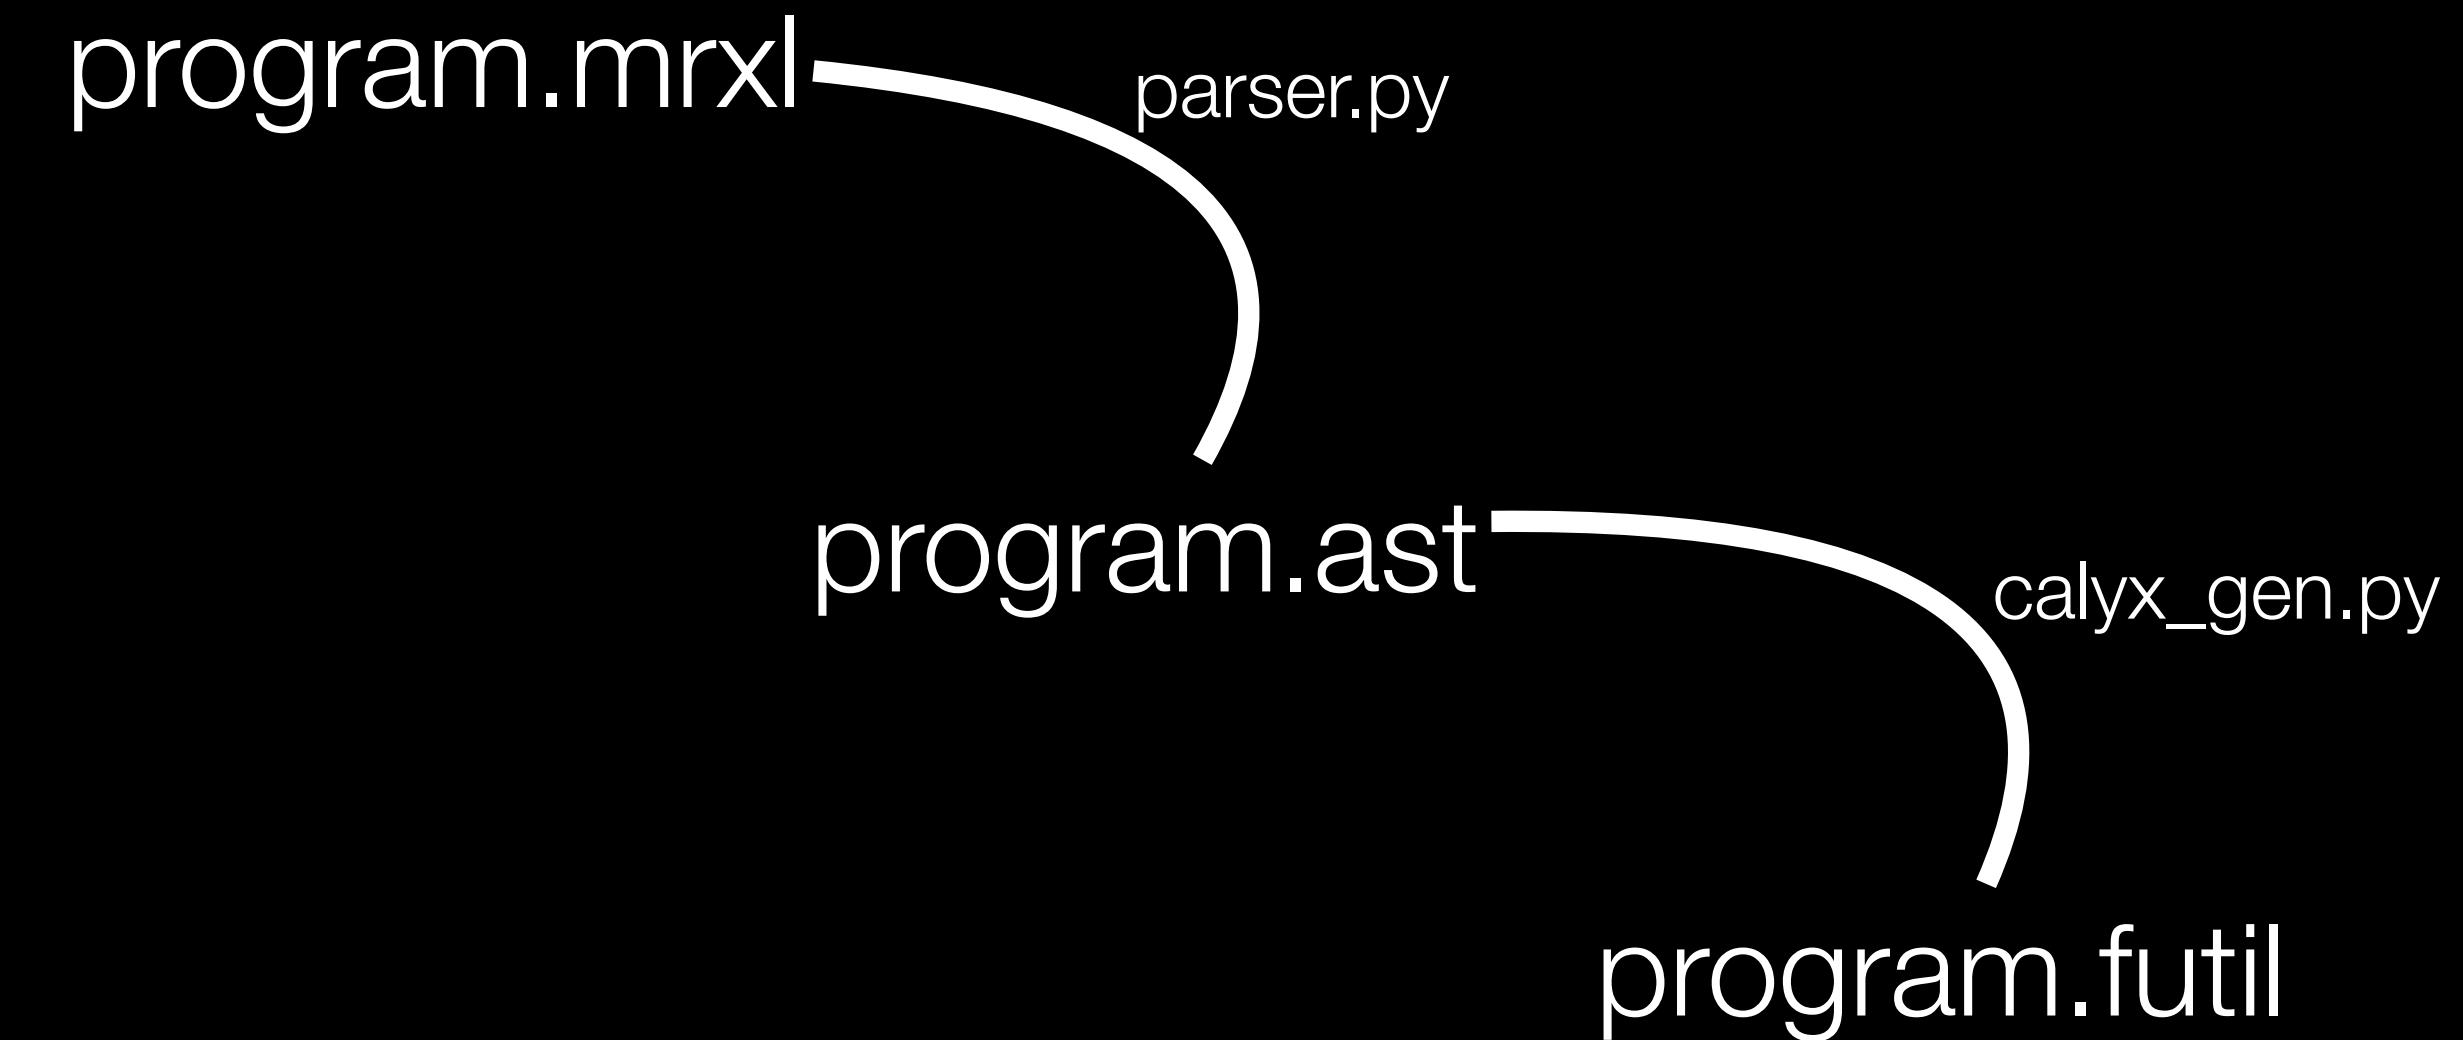
program.mrxml

parser.py

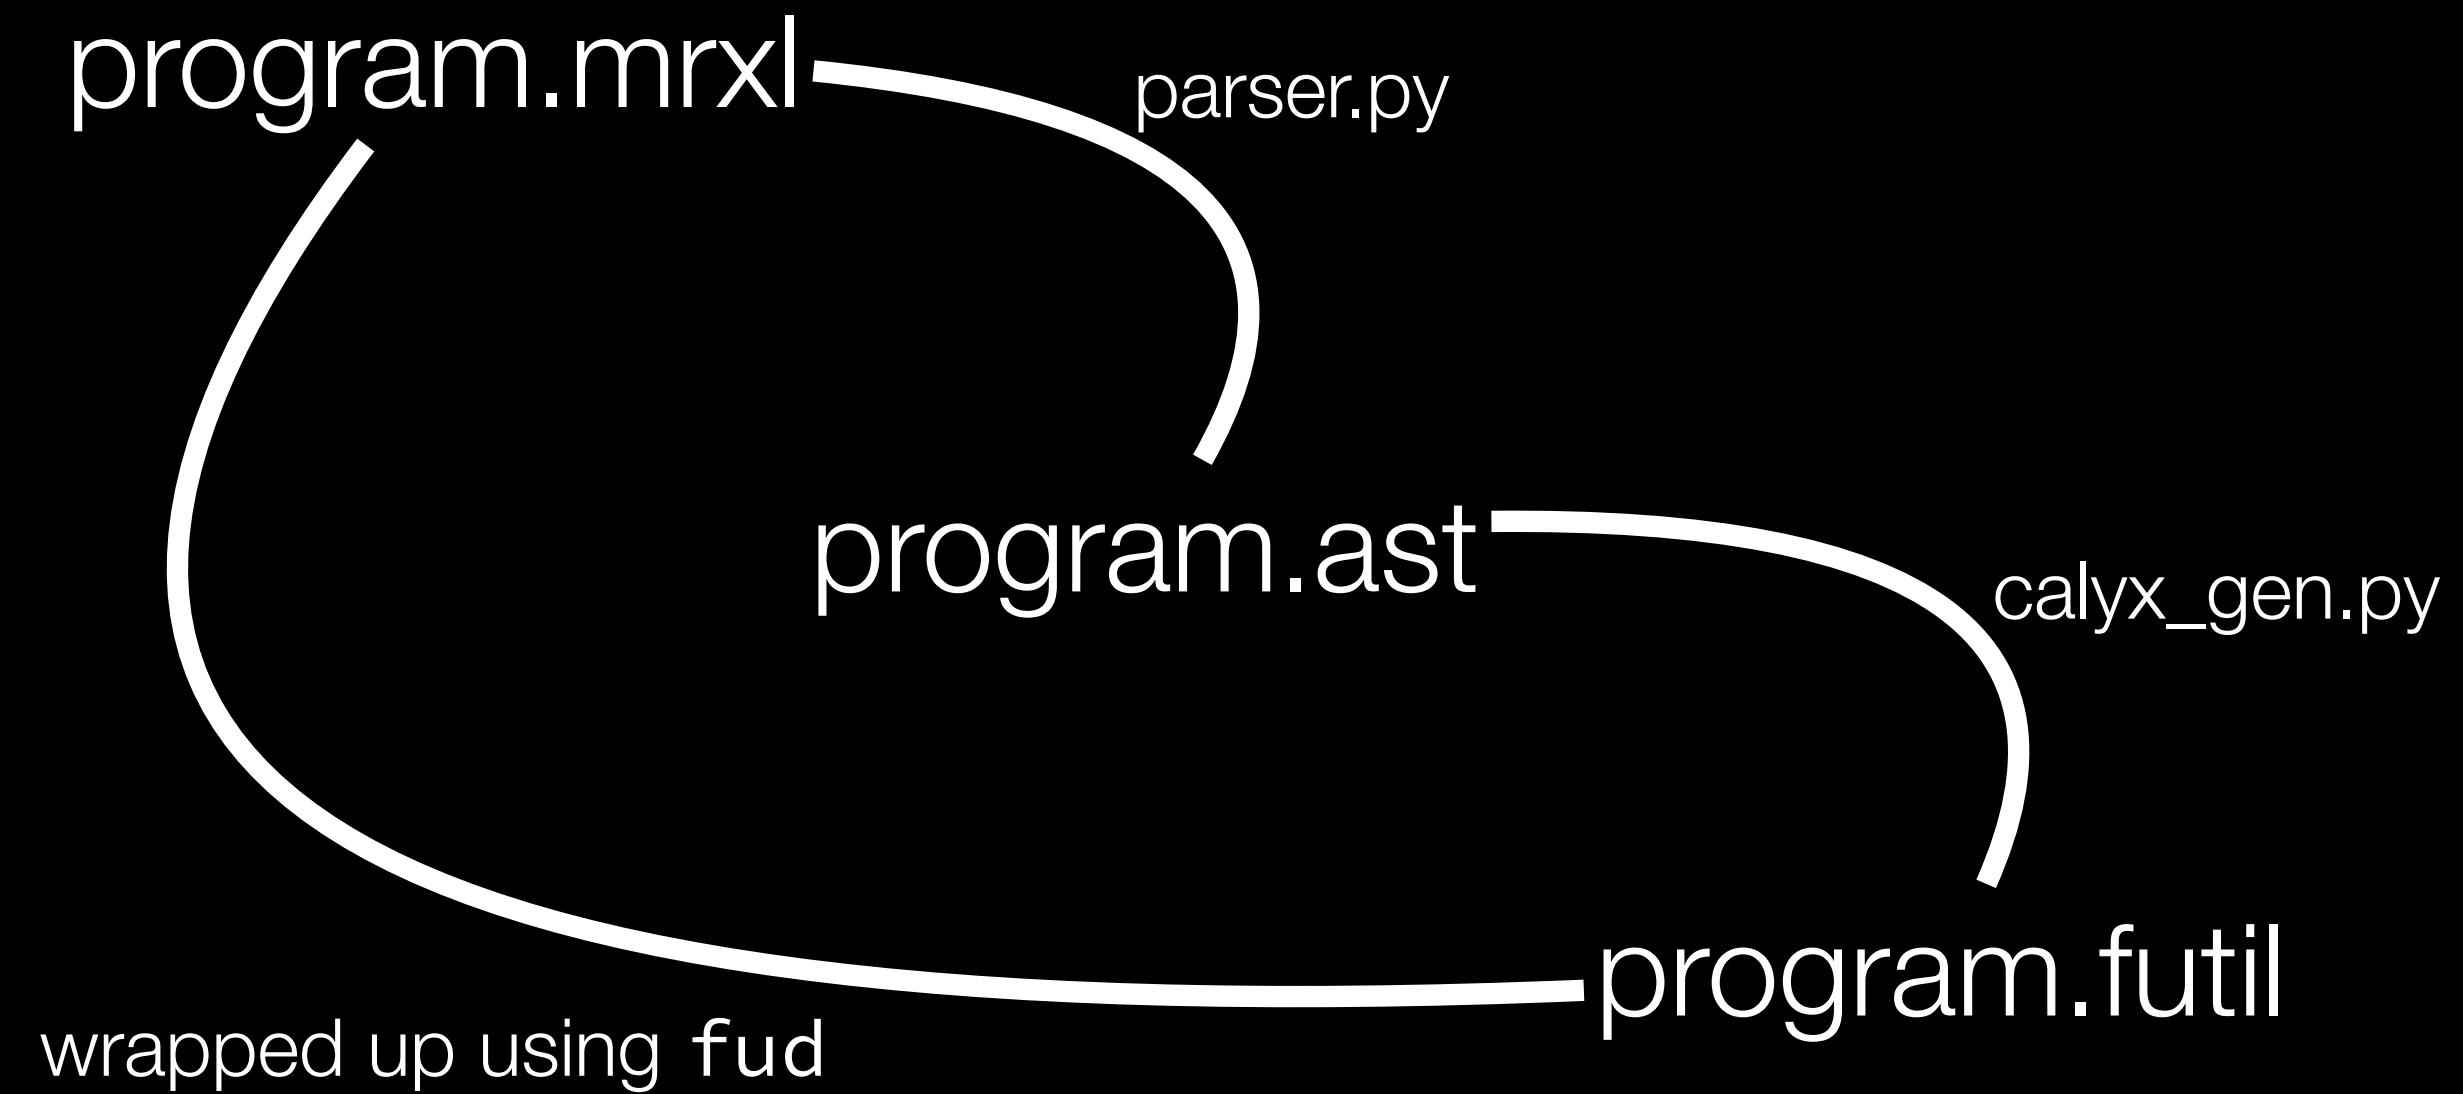
program.ast

program.futil

# anatomy of a frontend



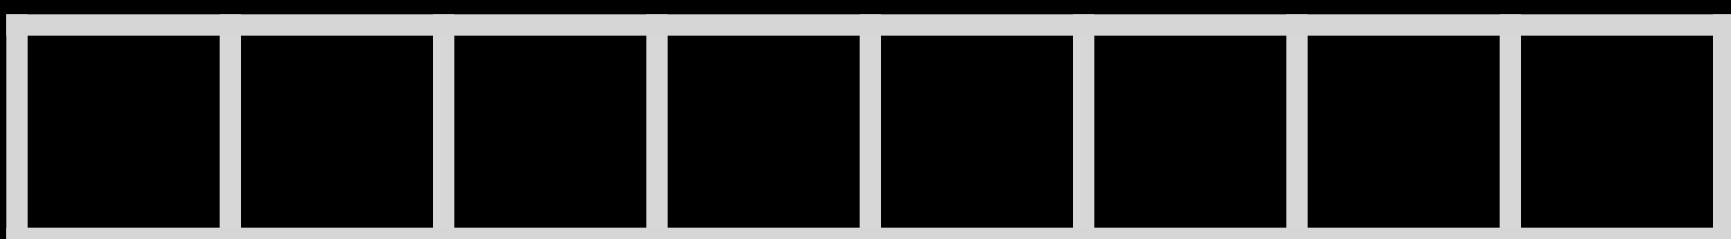
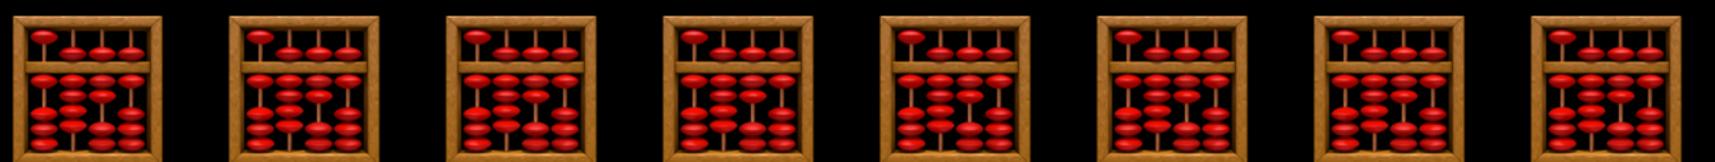
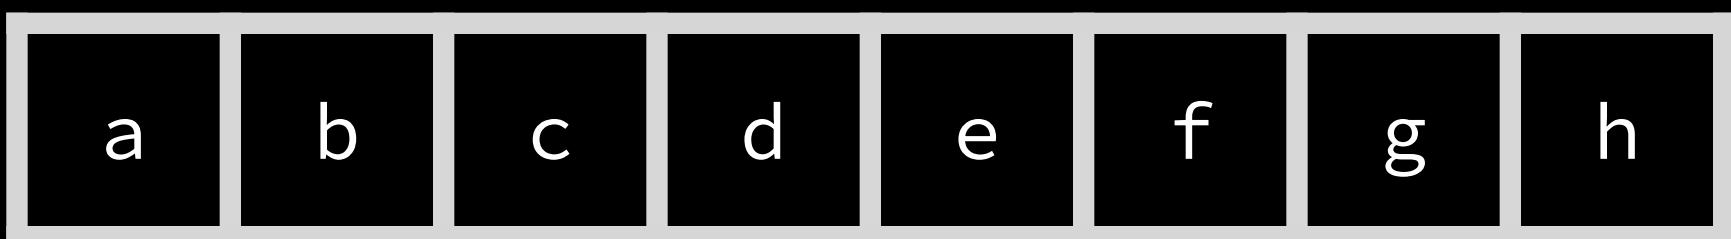
# anatomy of a frontend



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

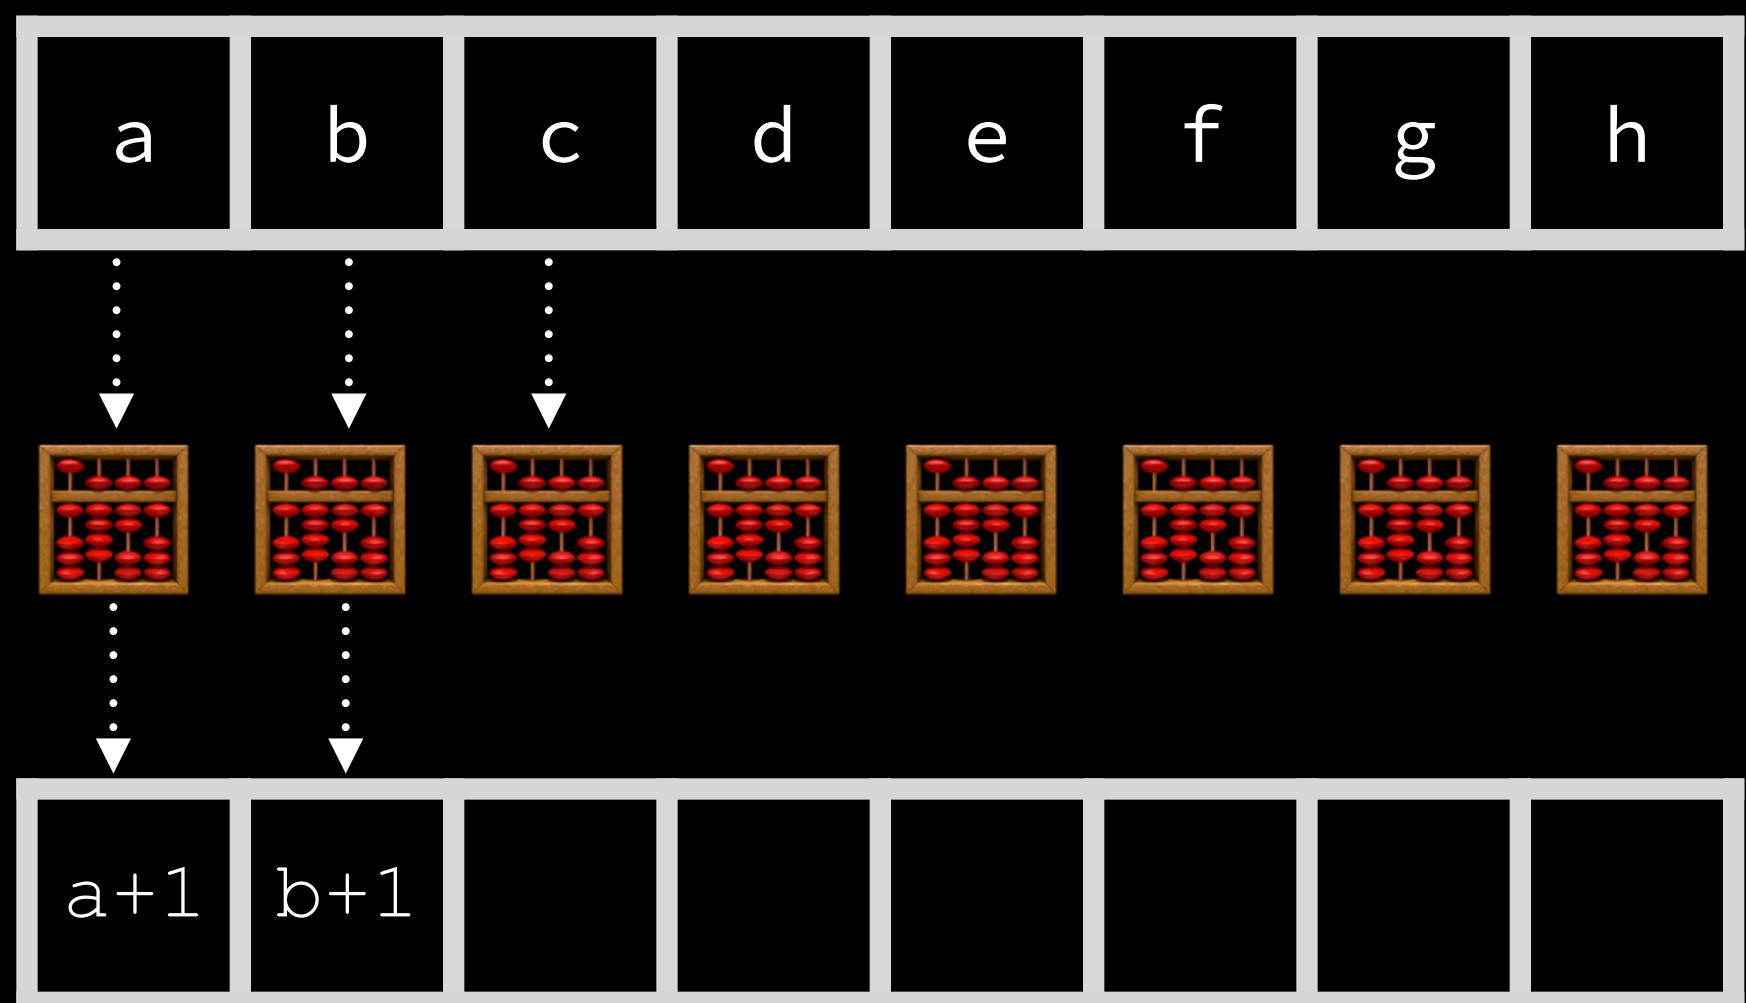
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

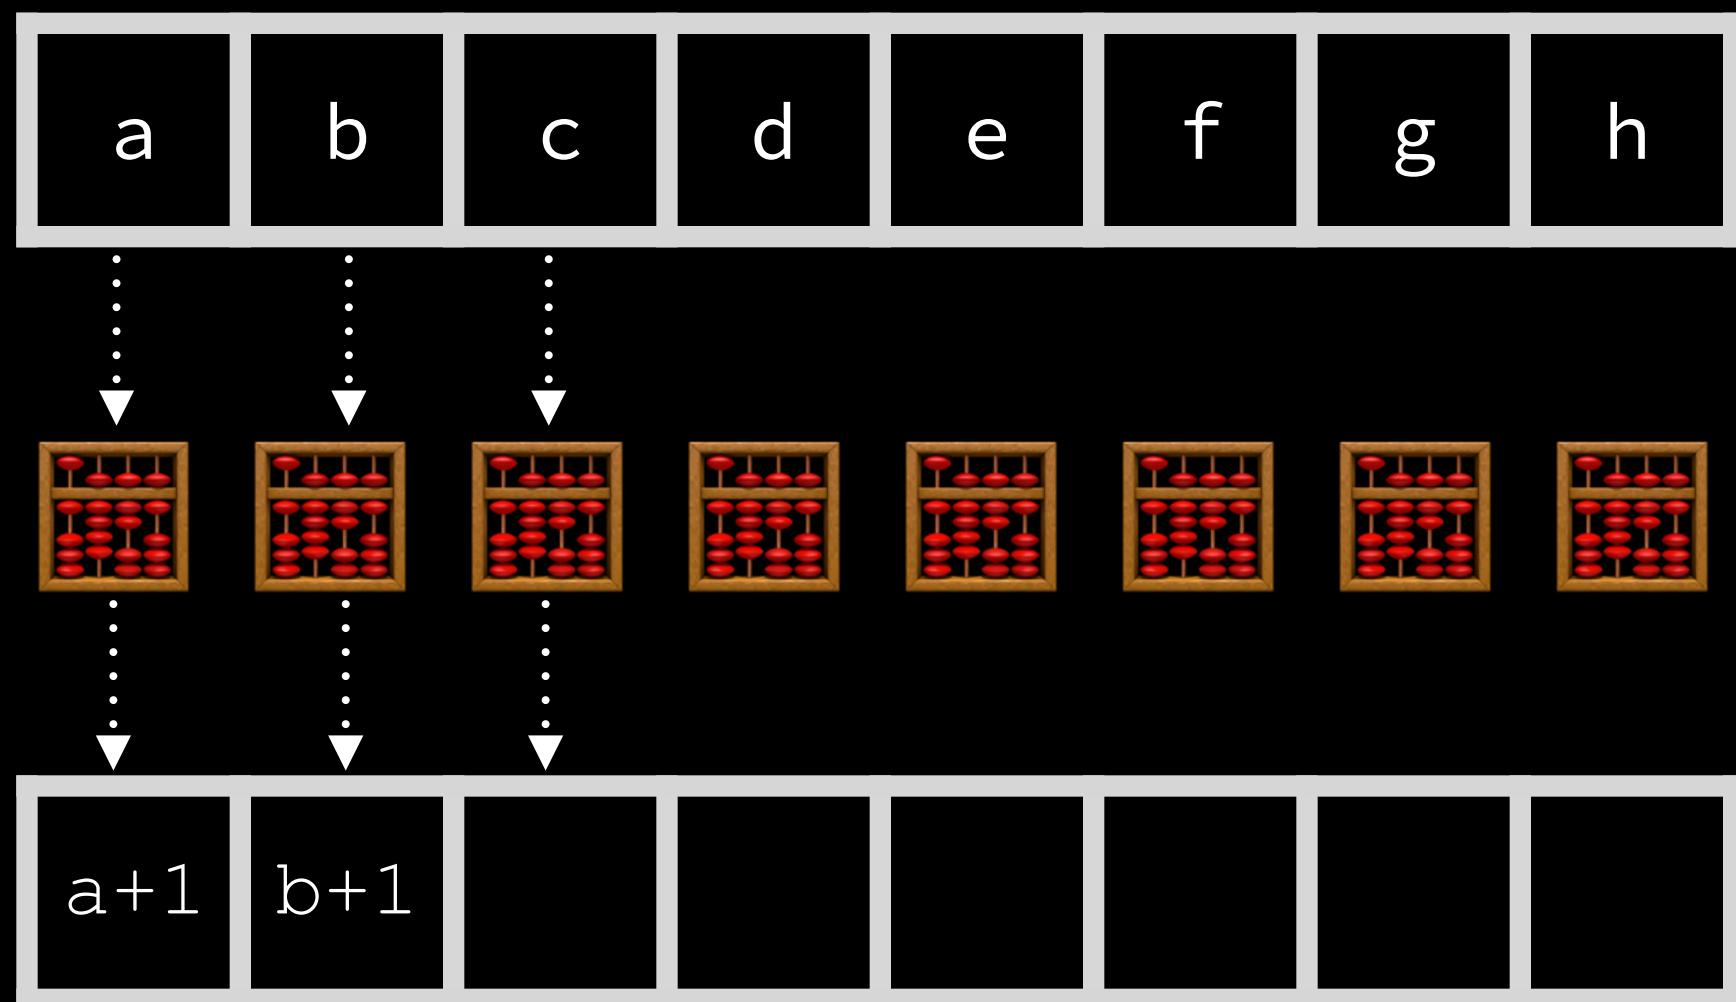
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

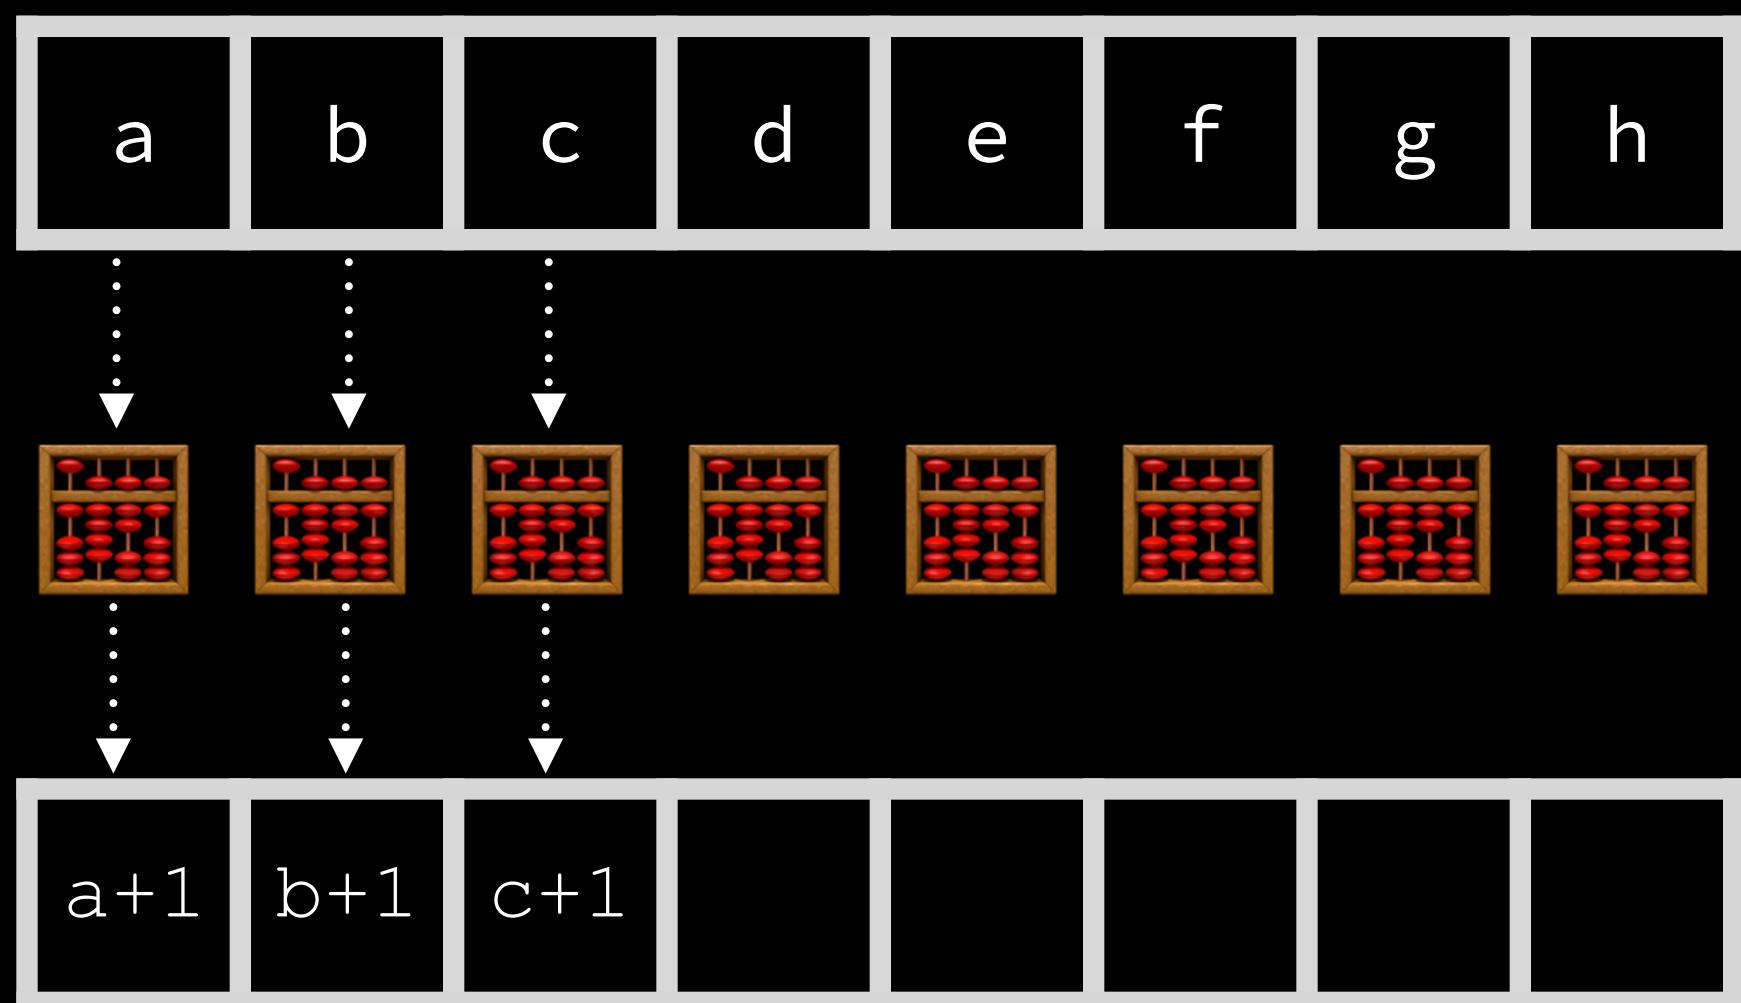
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

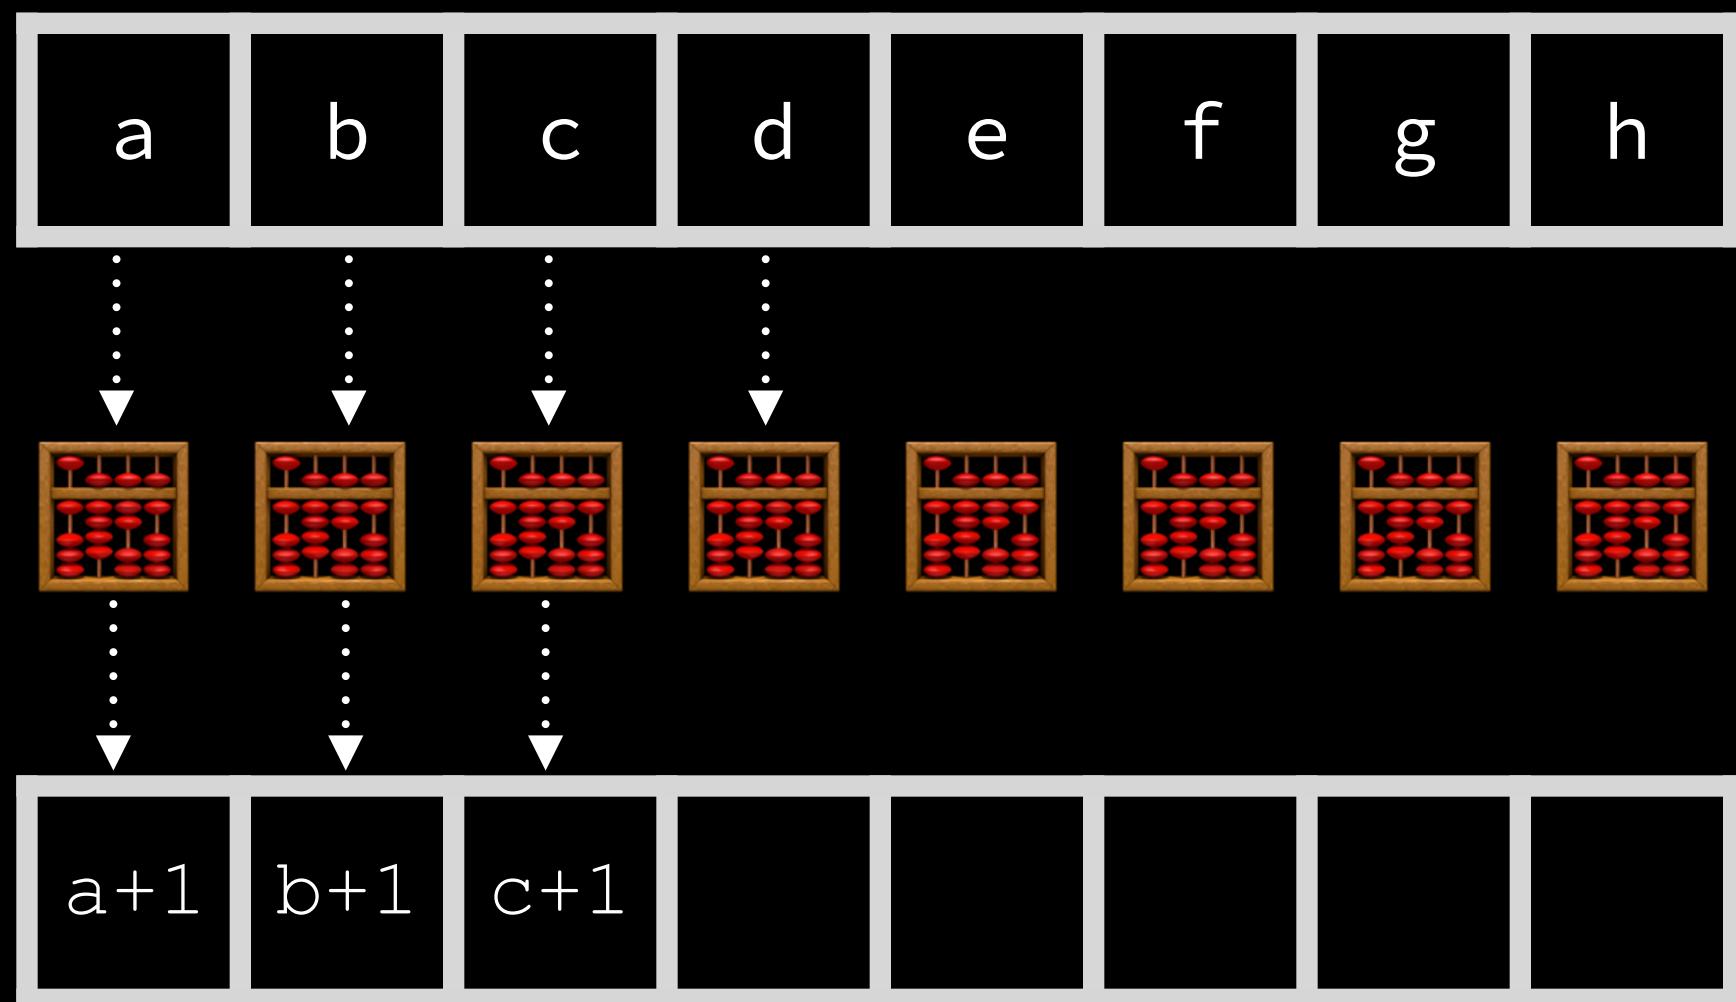
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

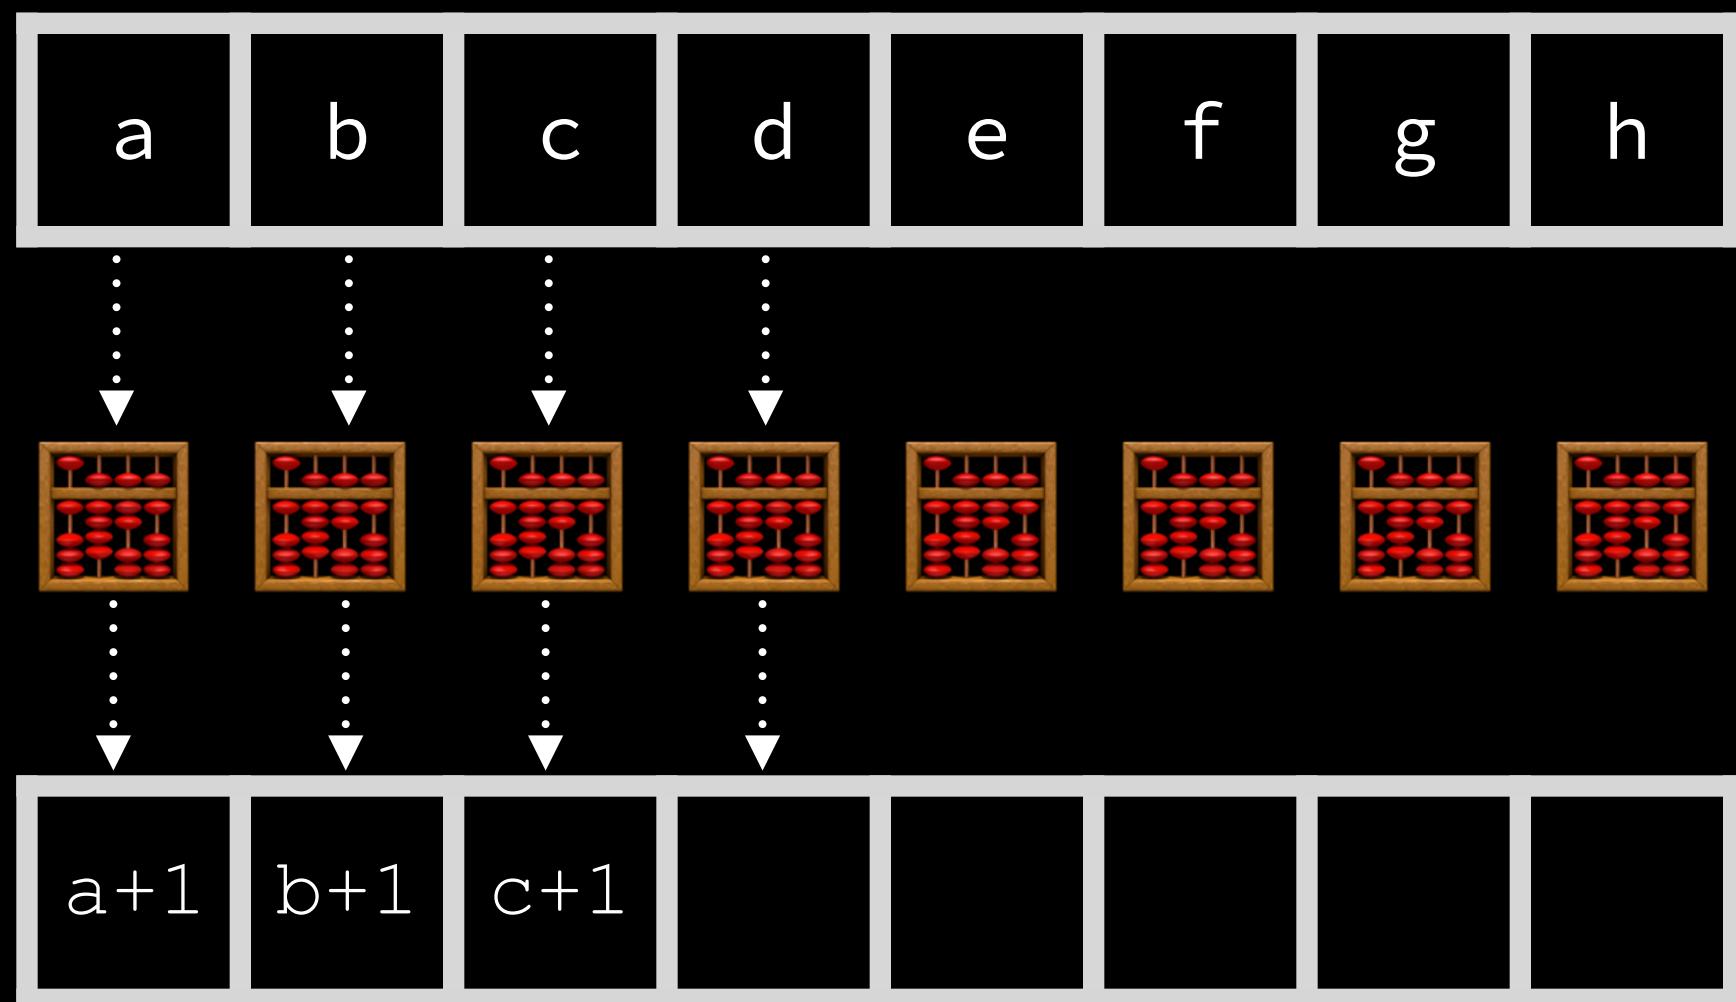
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

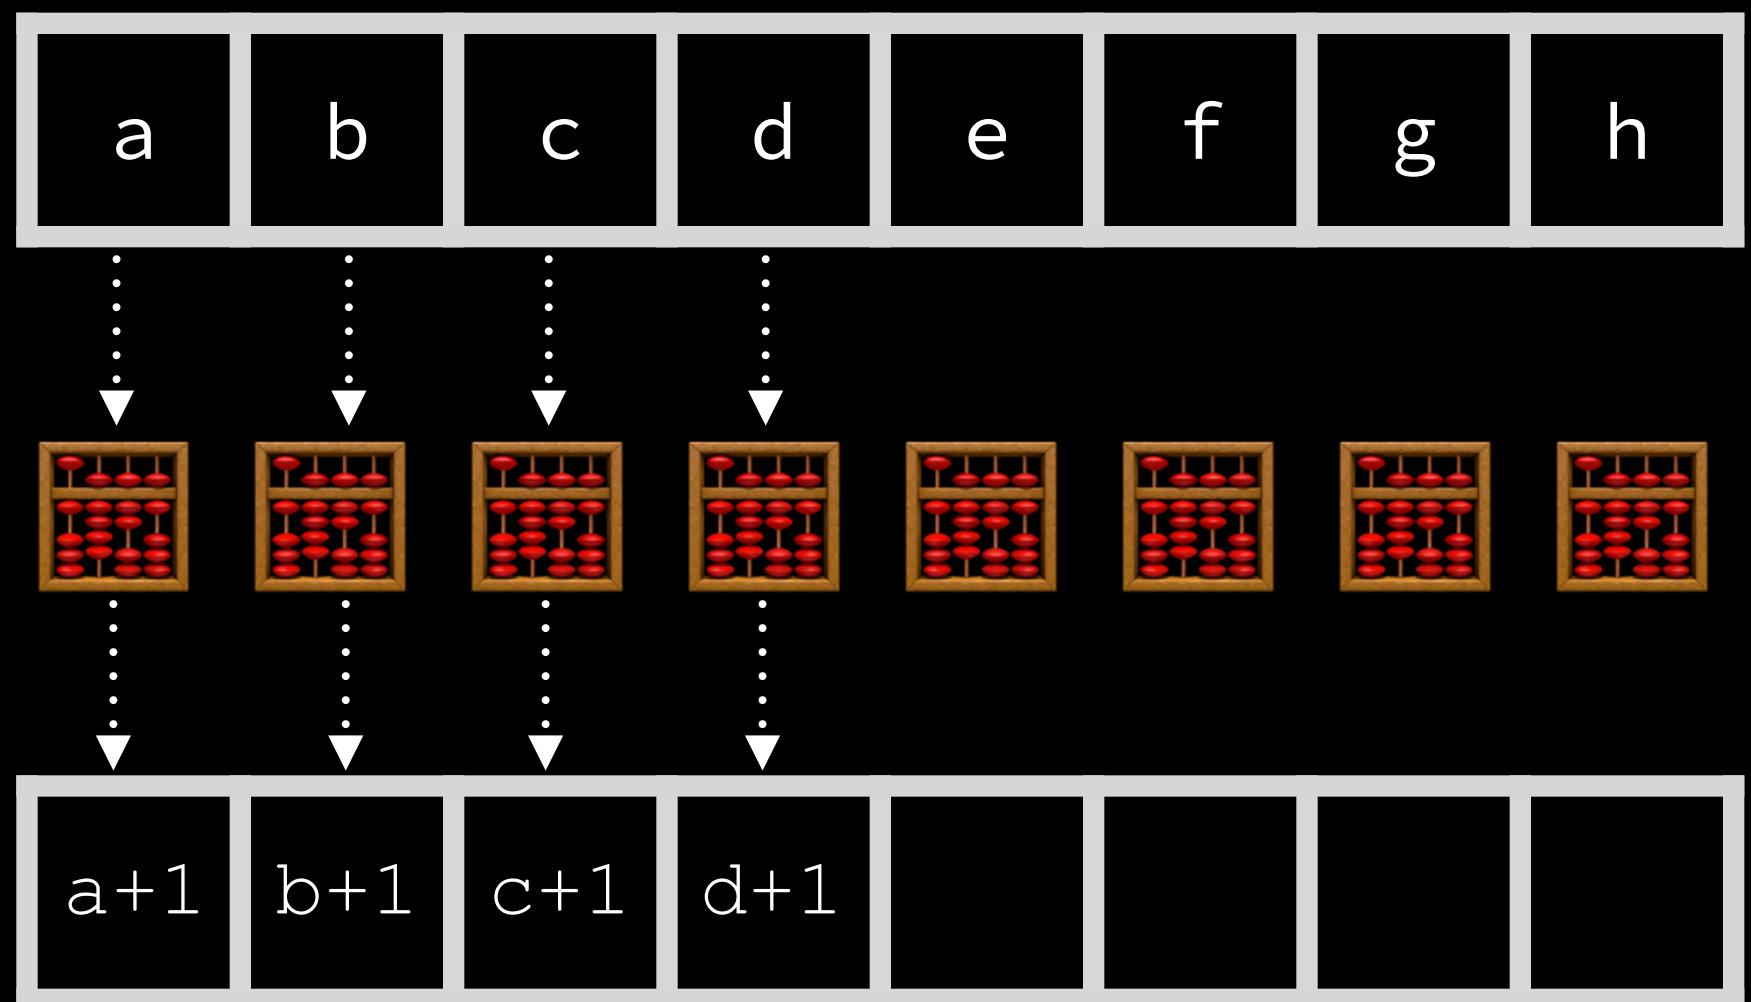
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

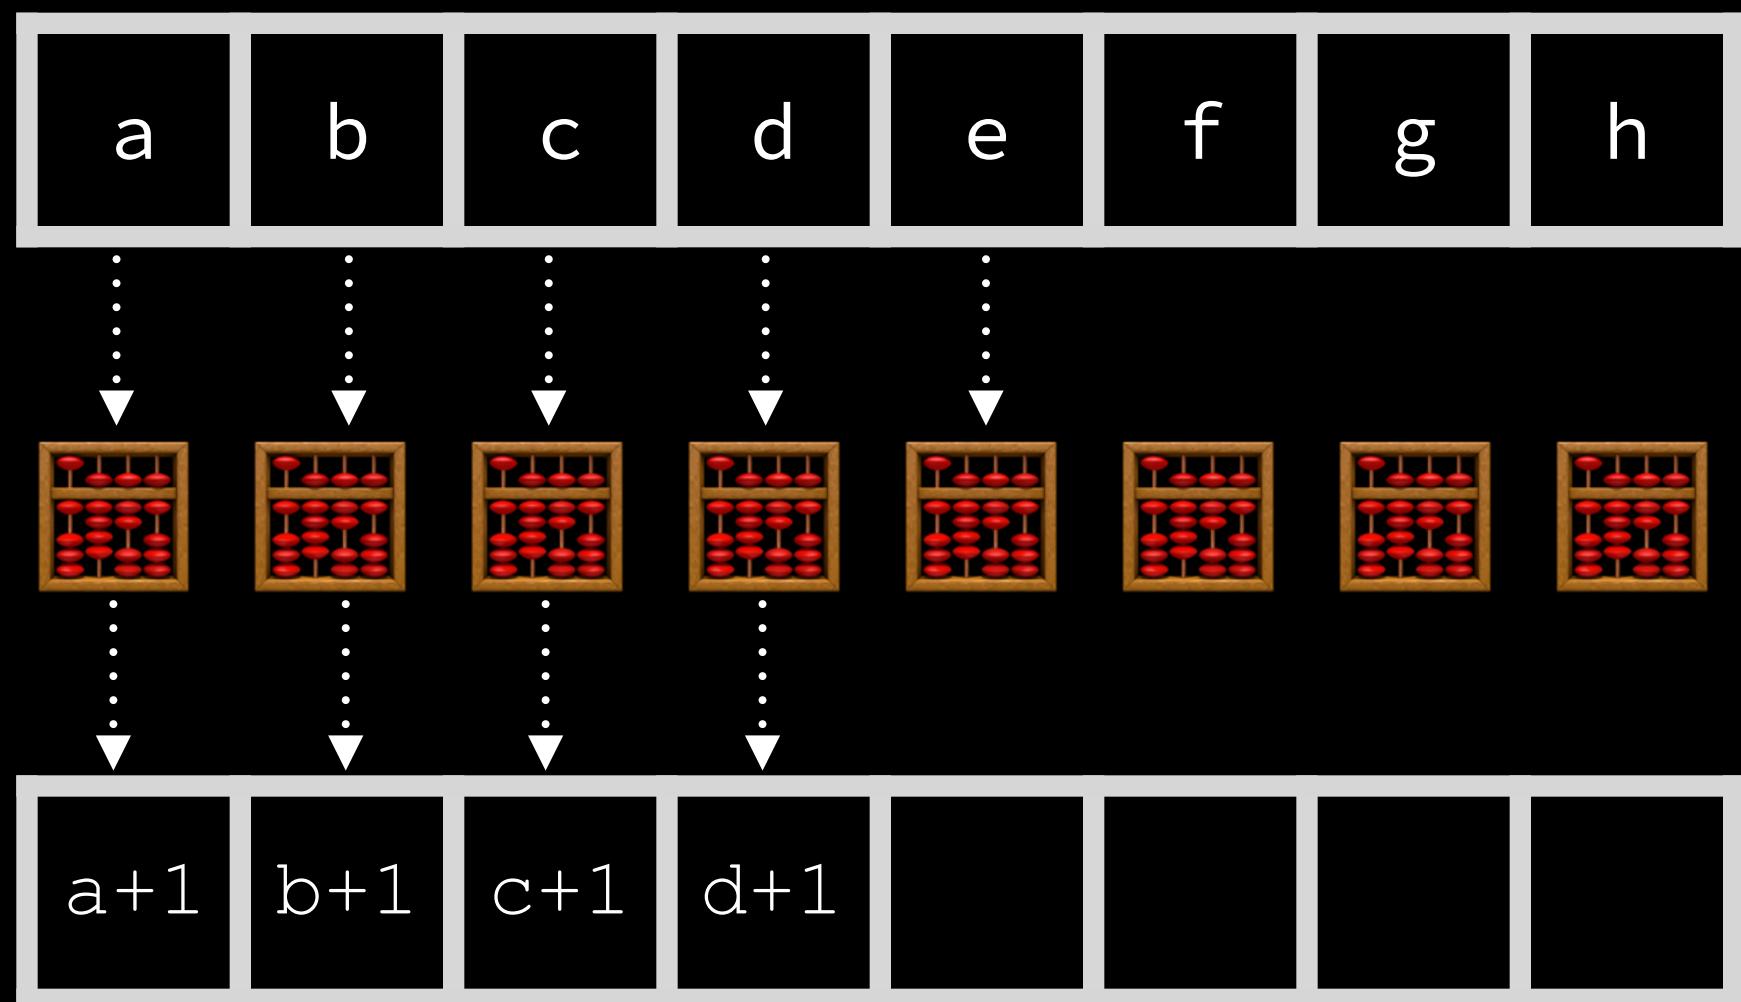
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

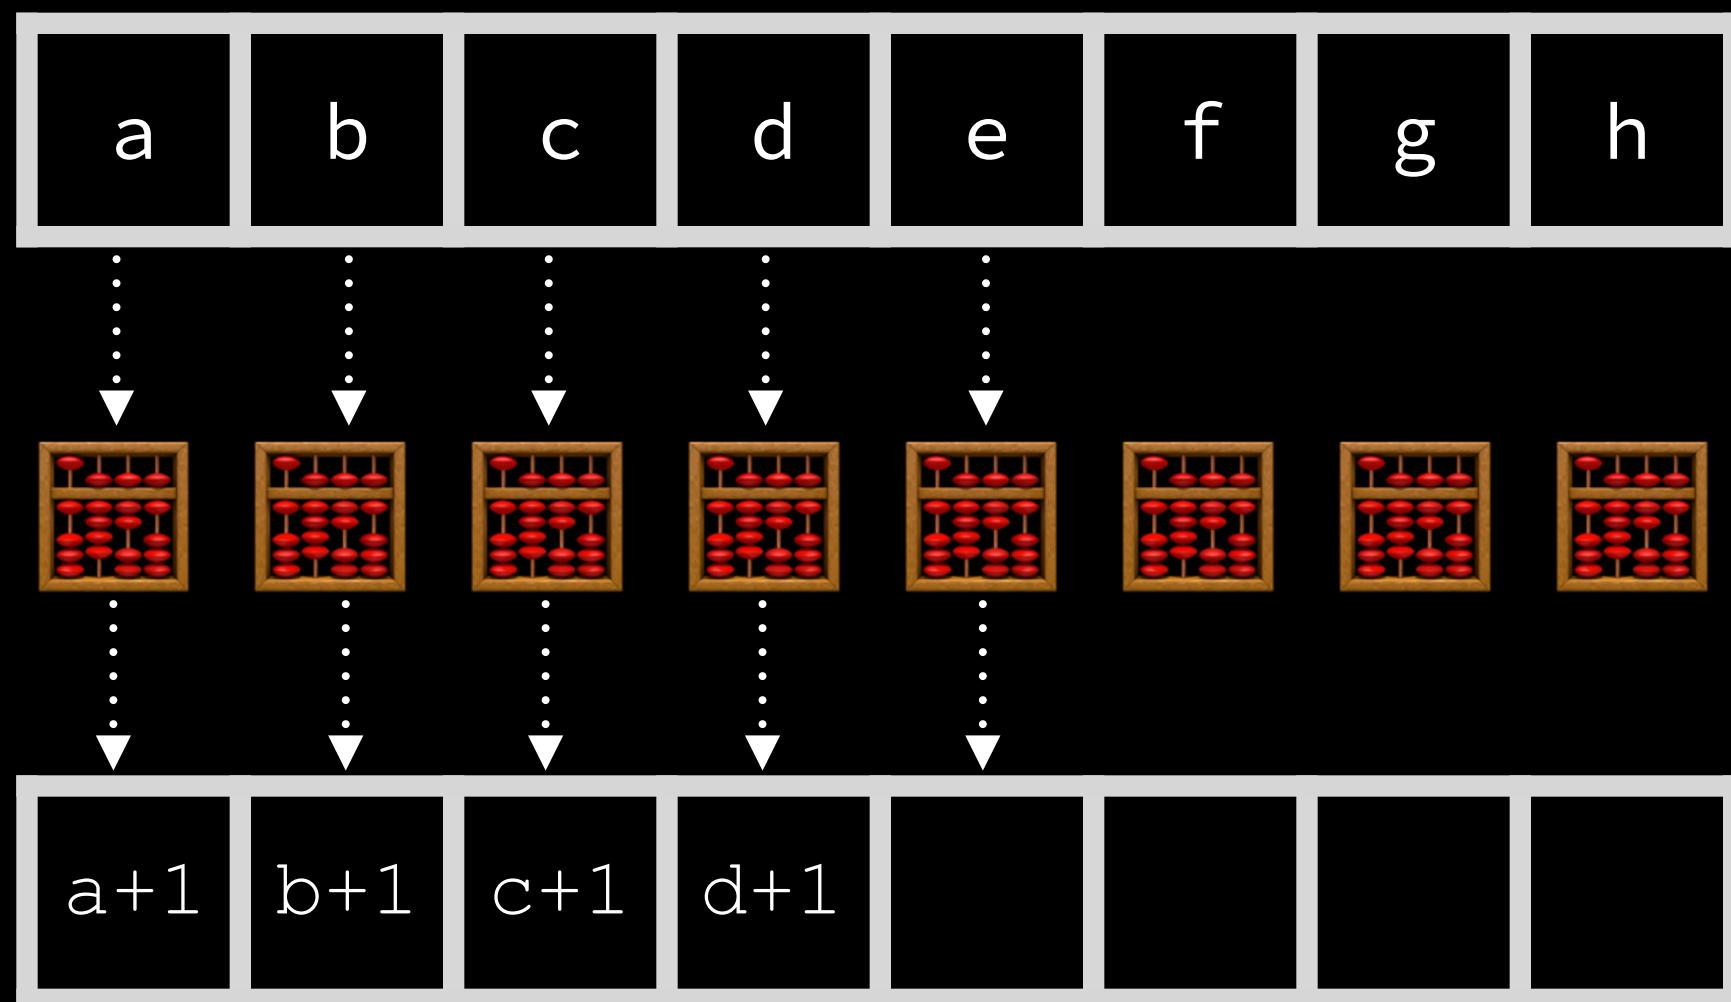
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

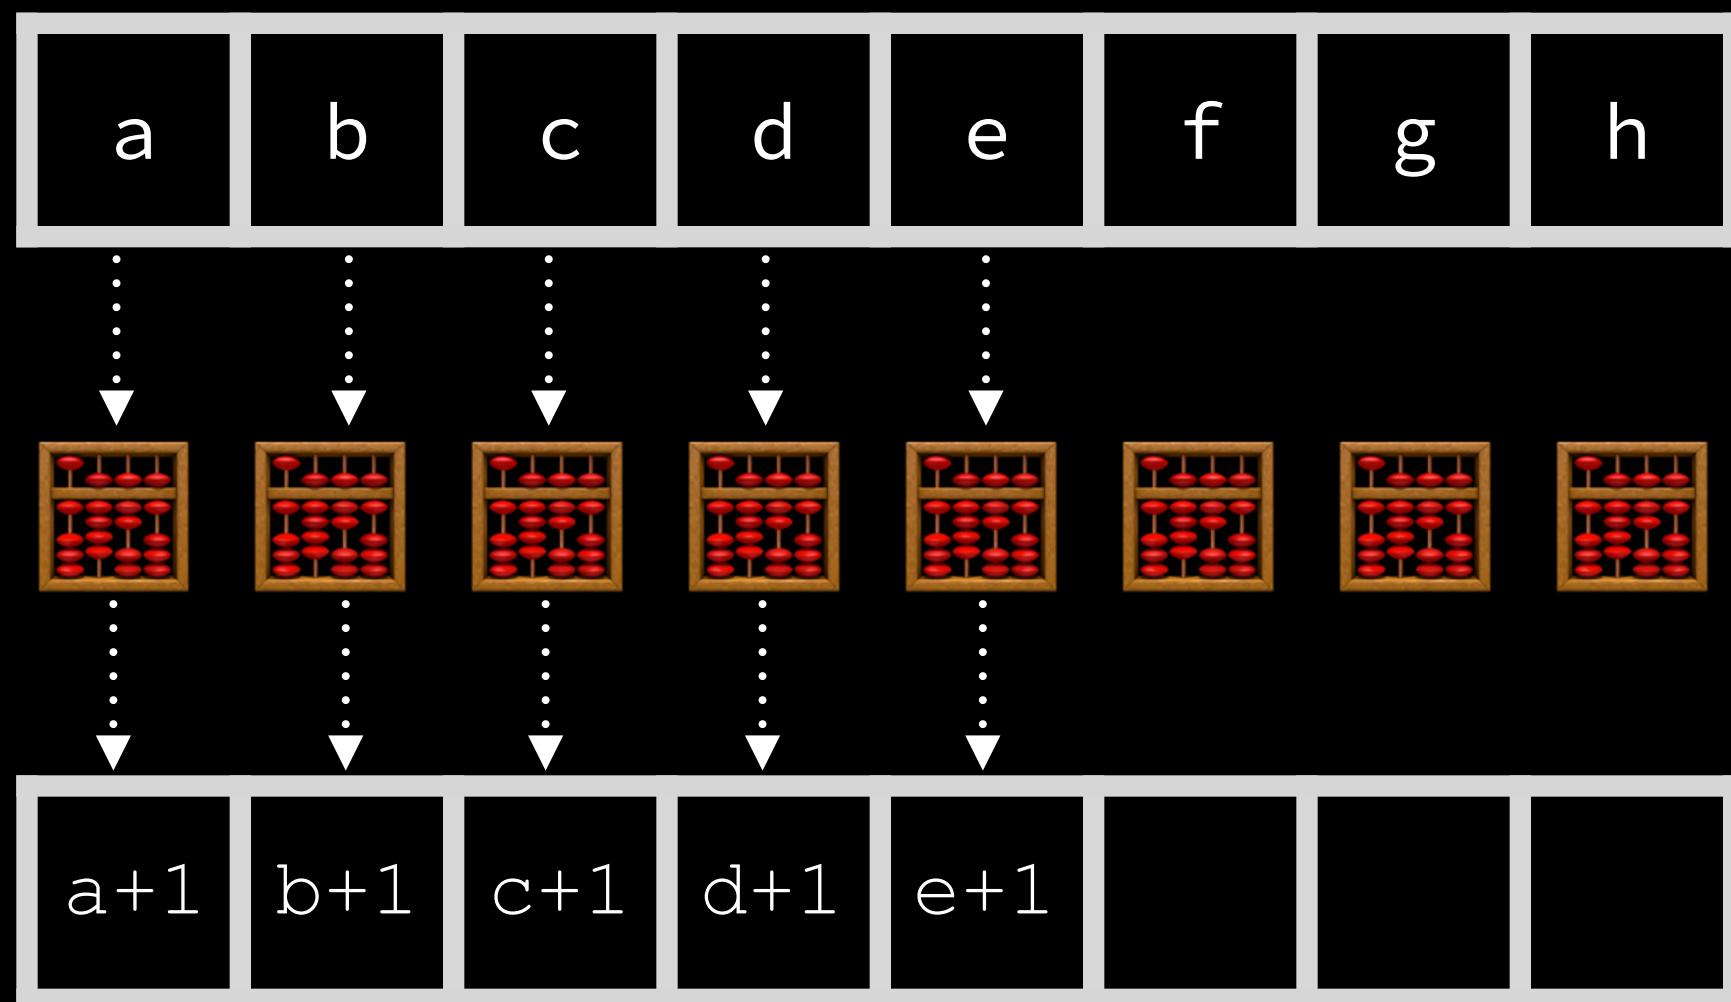
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

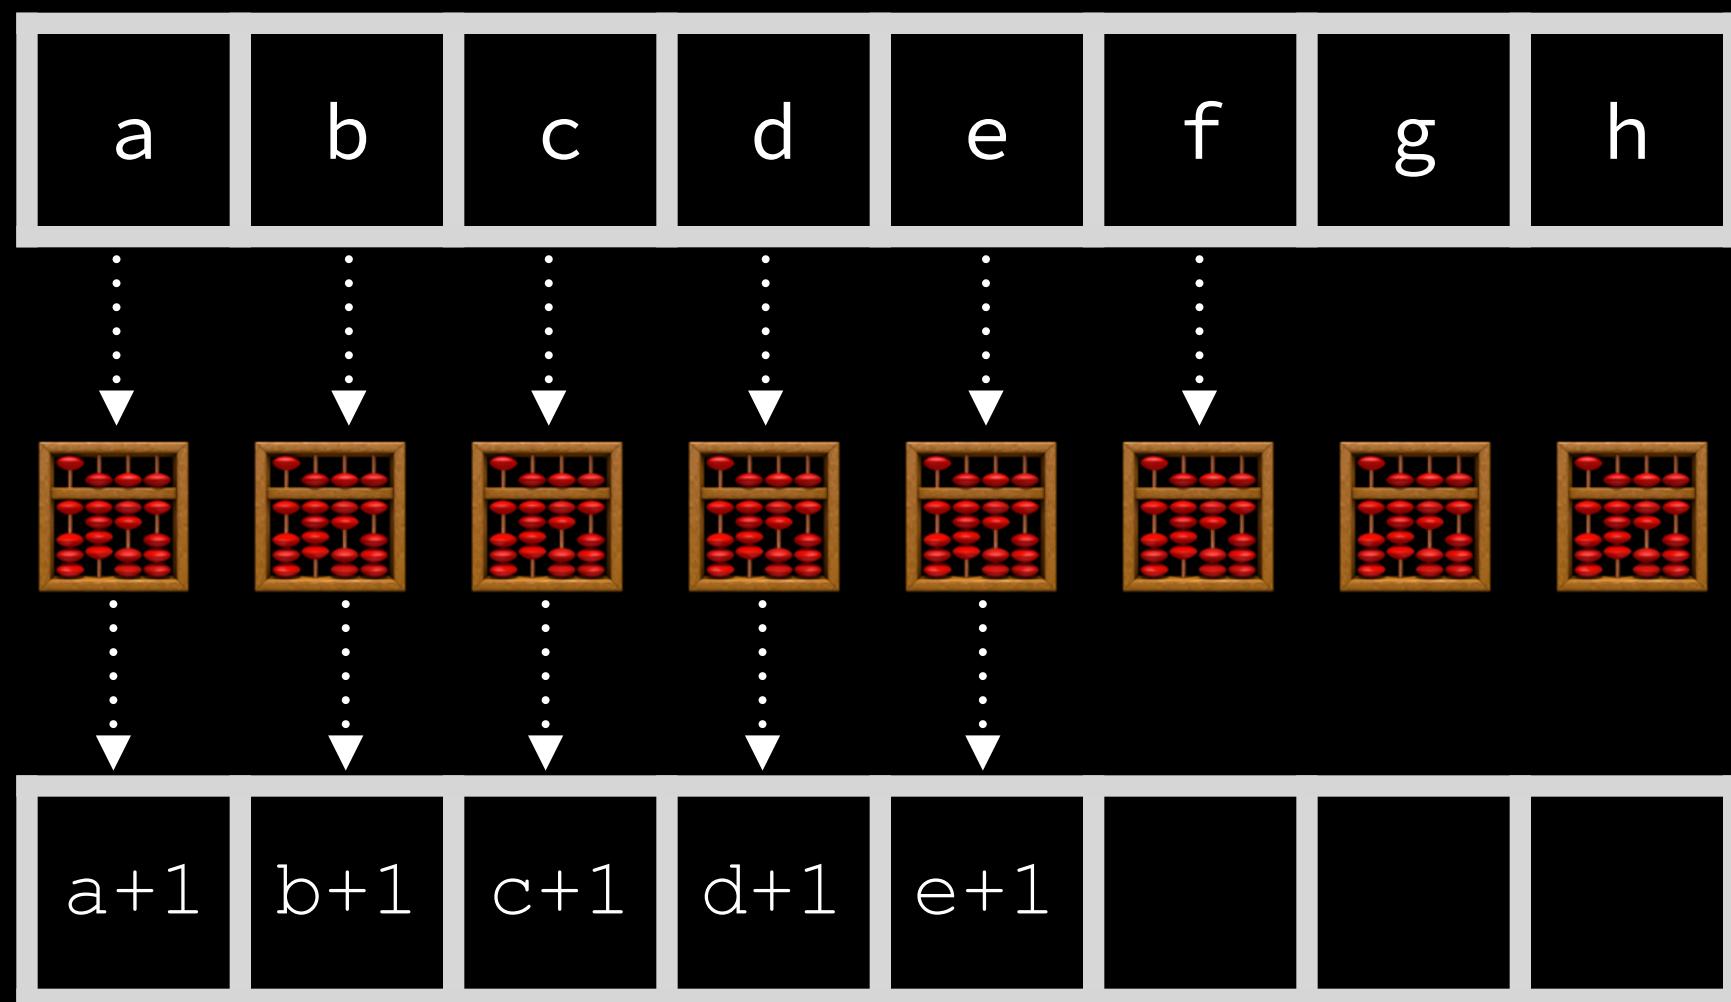
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

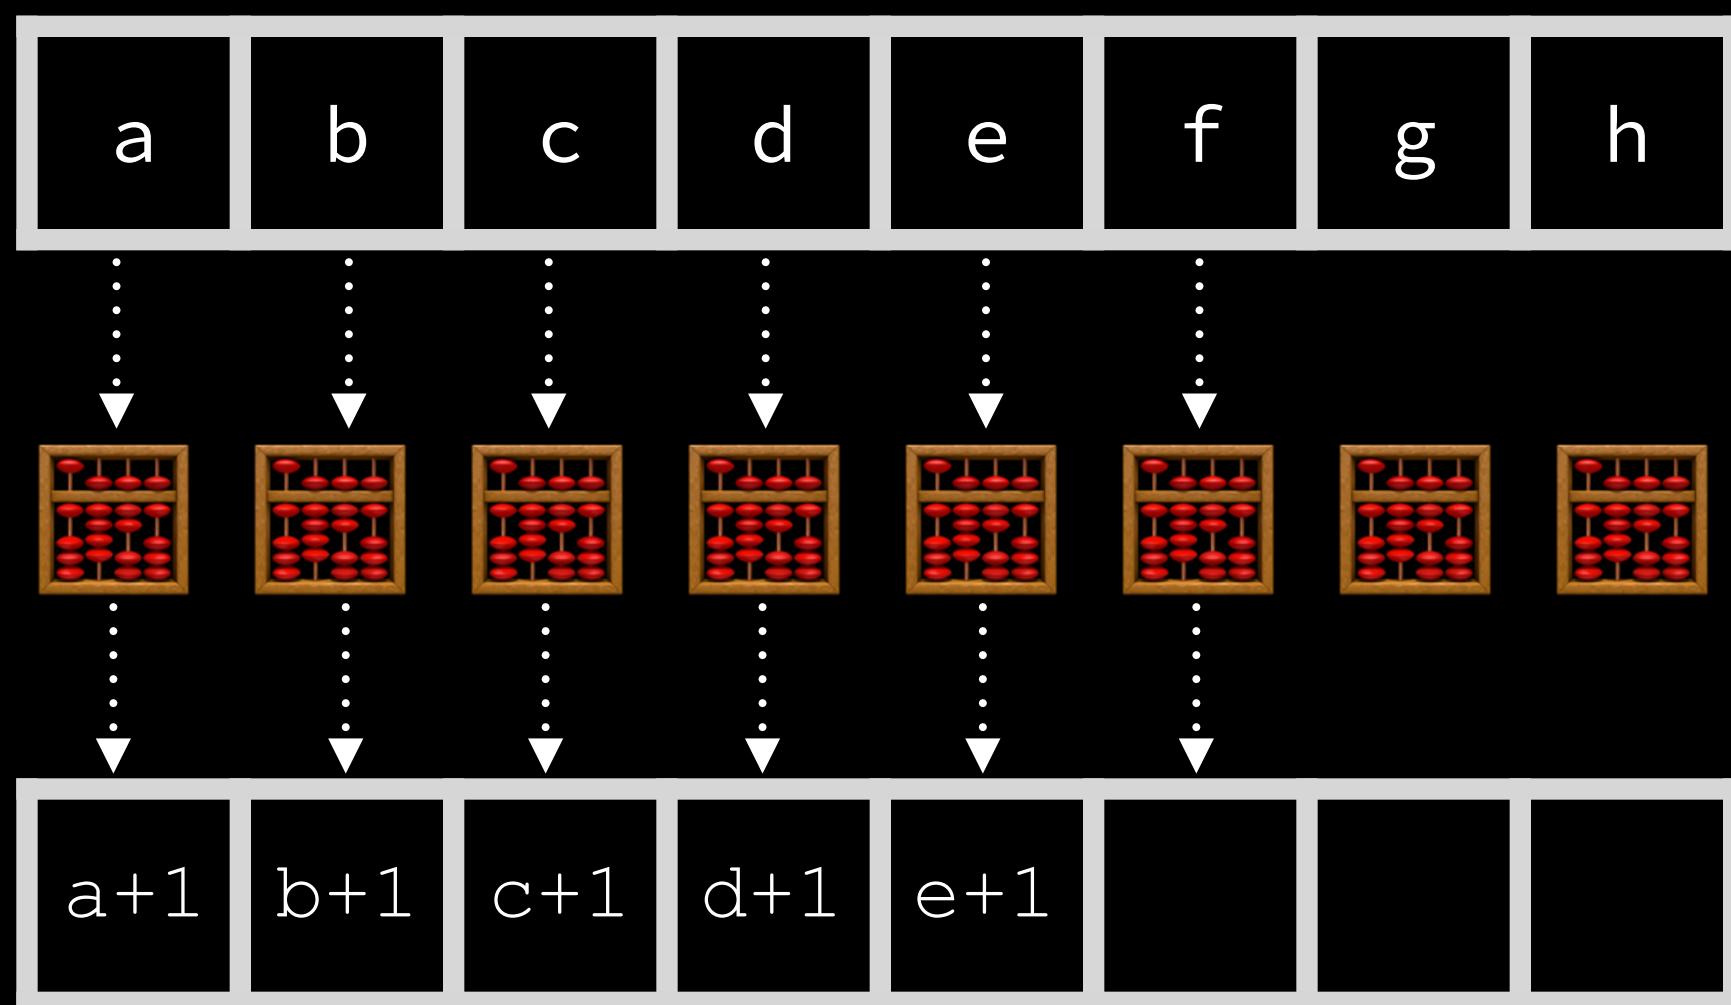
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

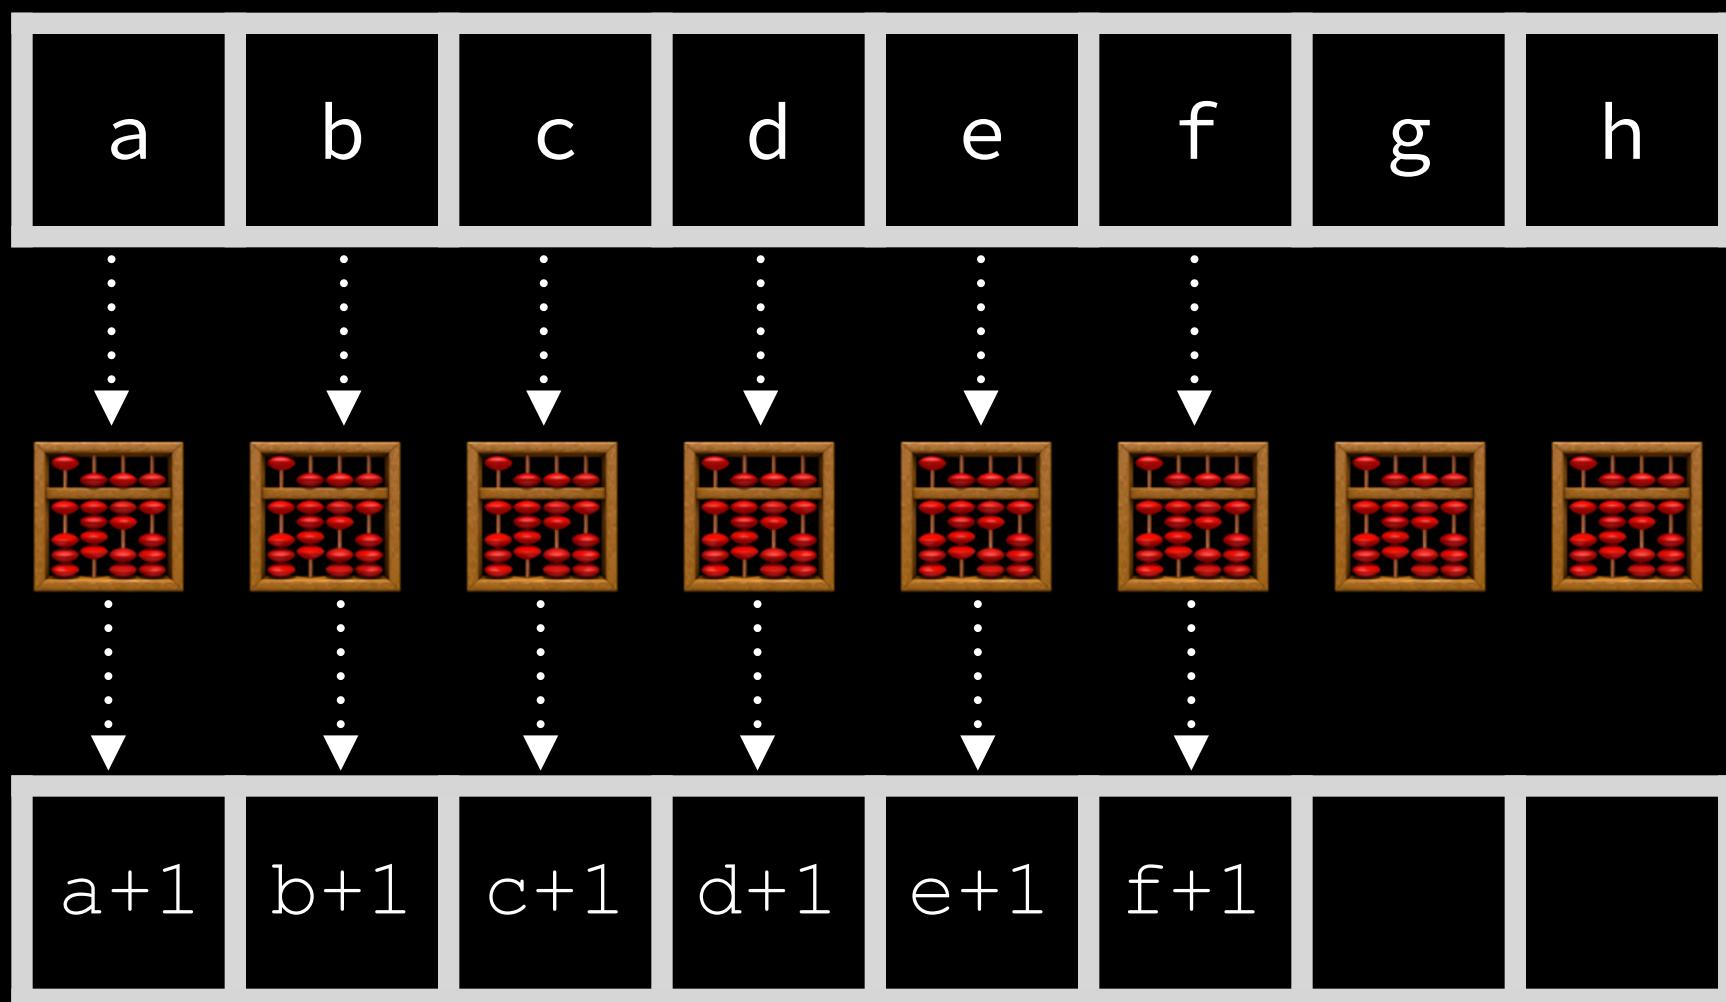
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

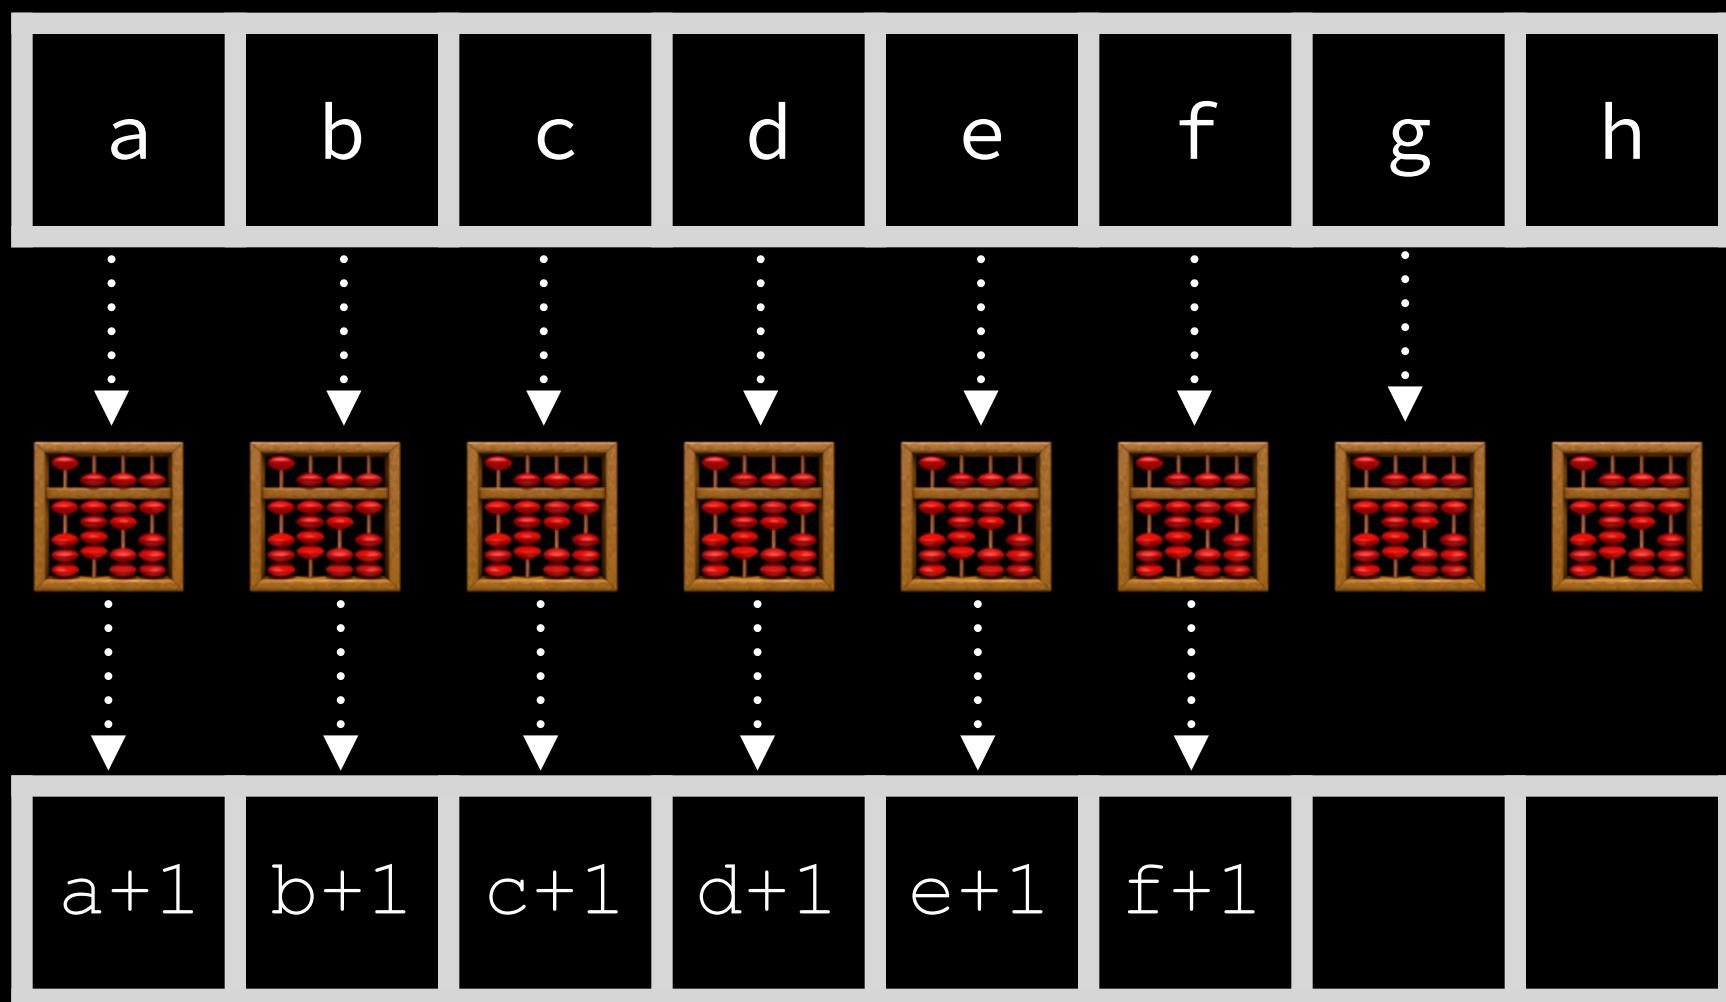
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

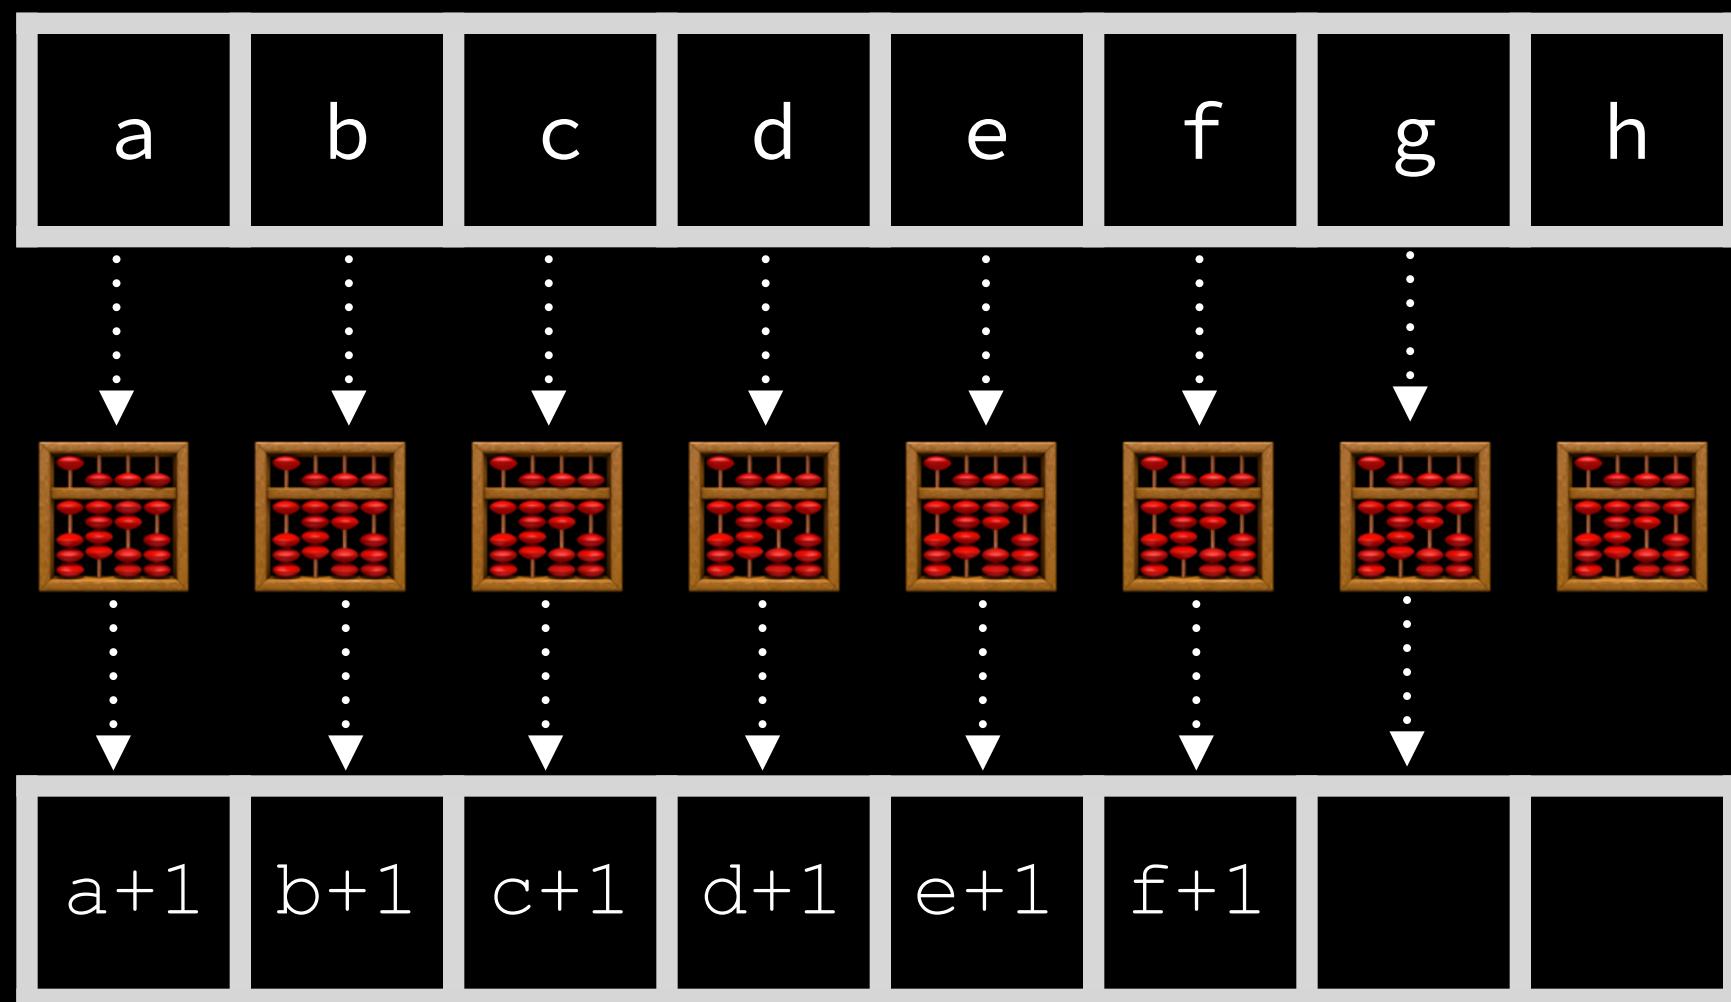
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

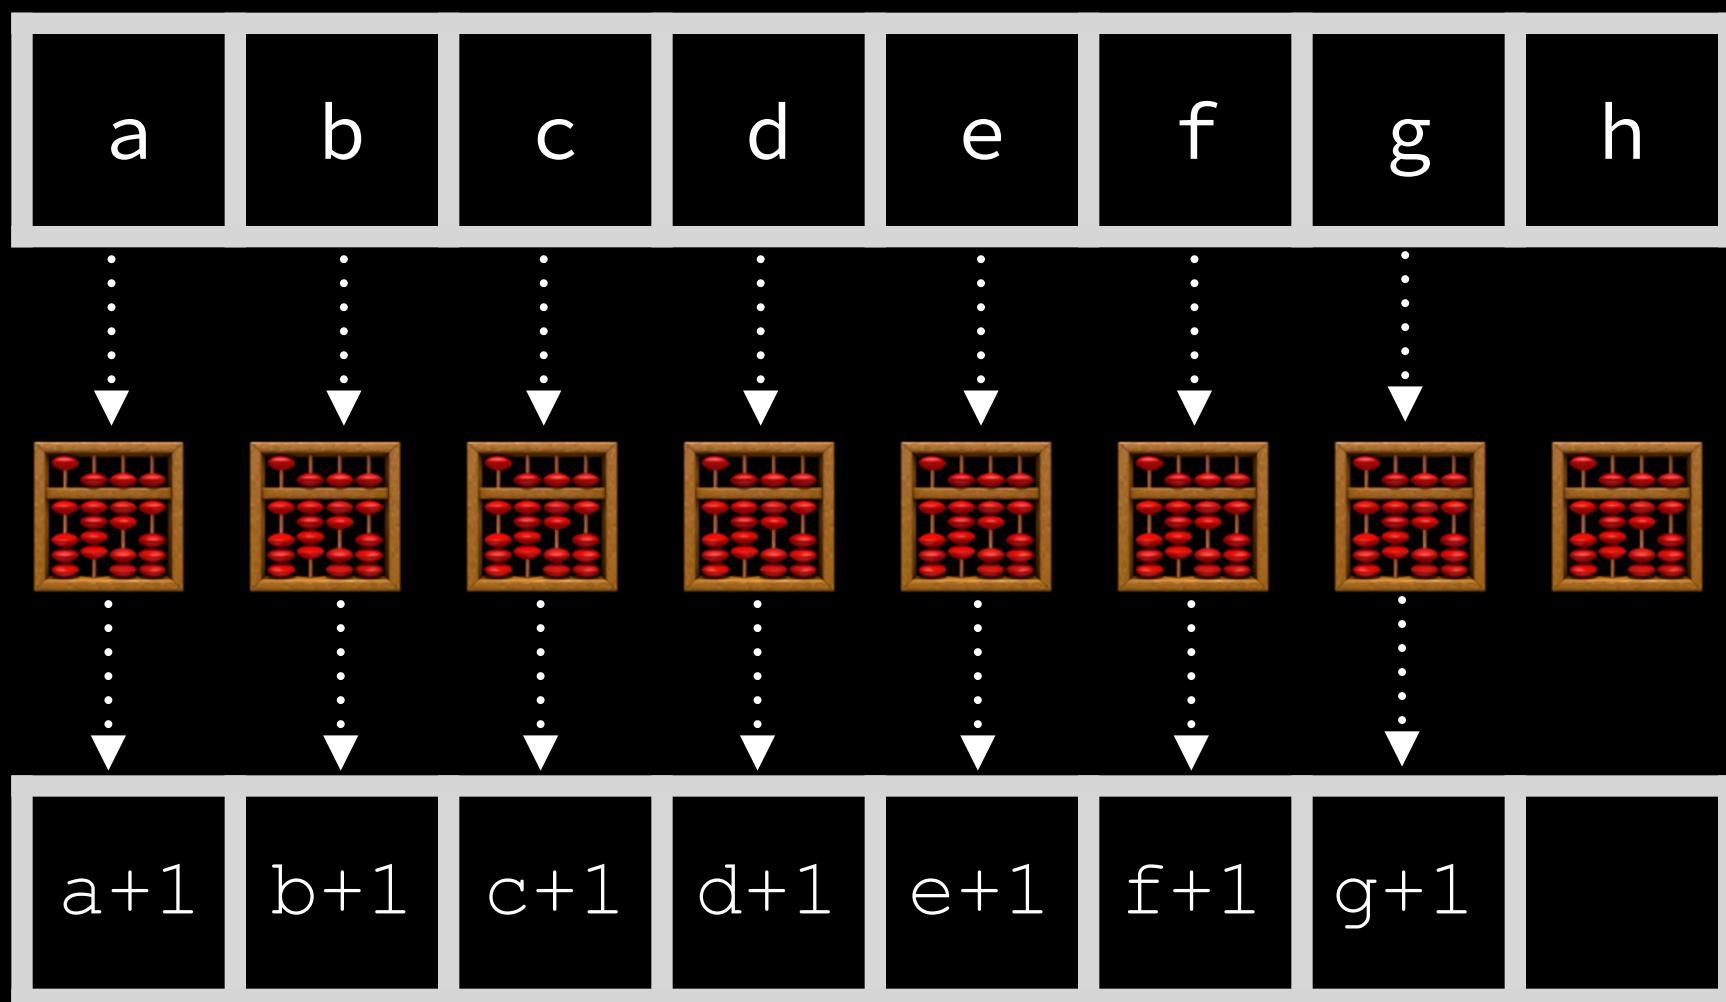
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

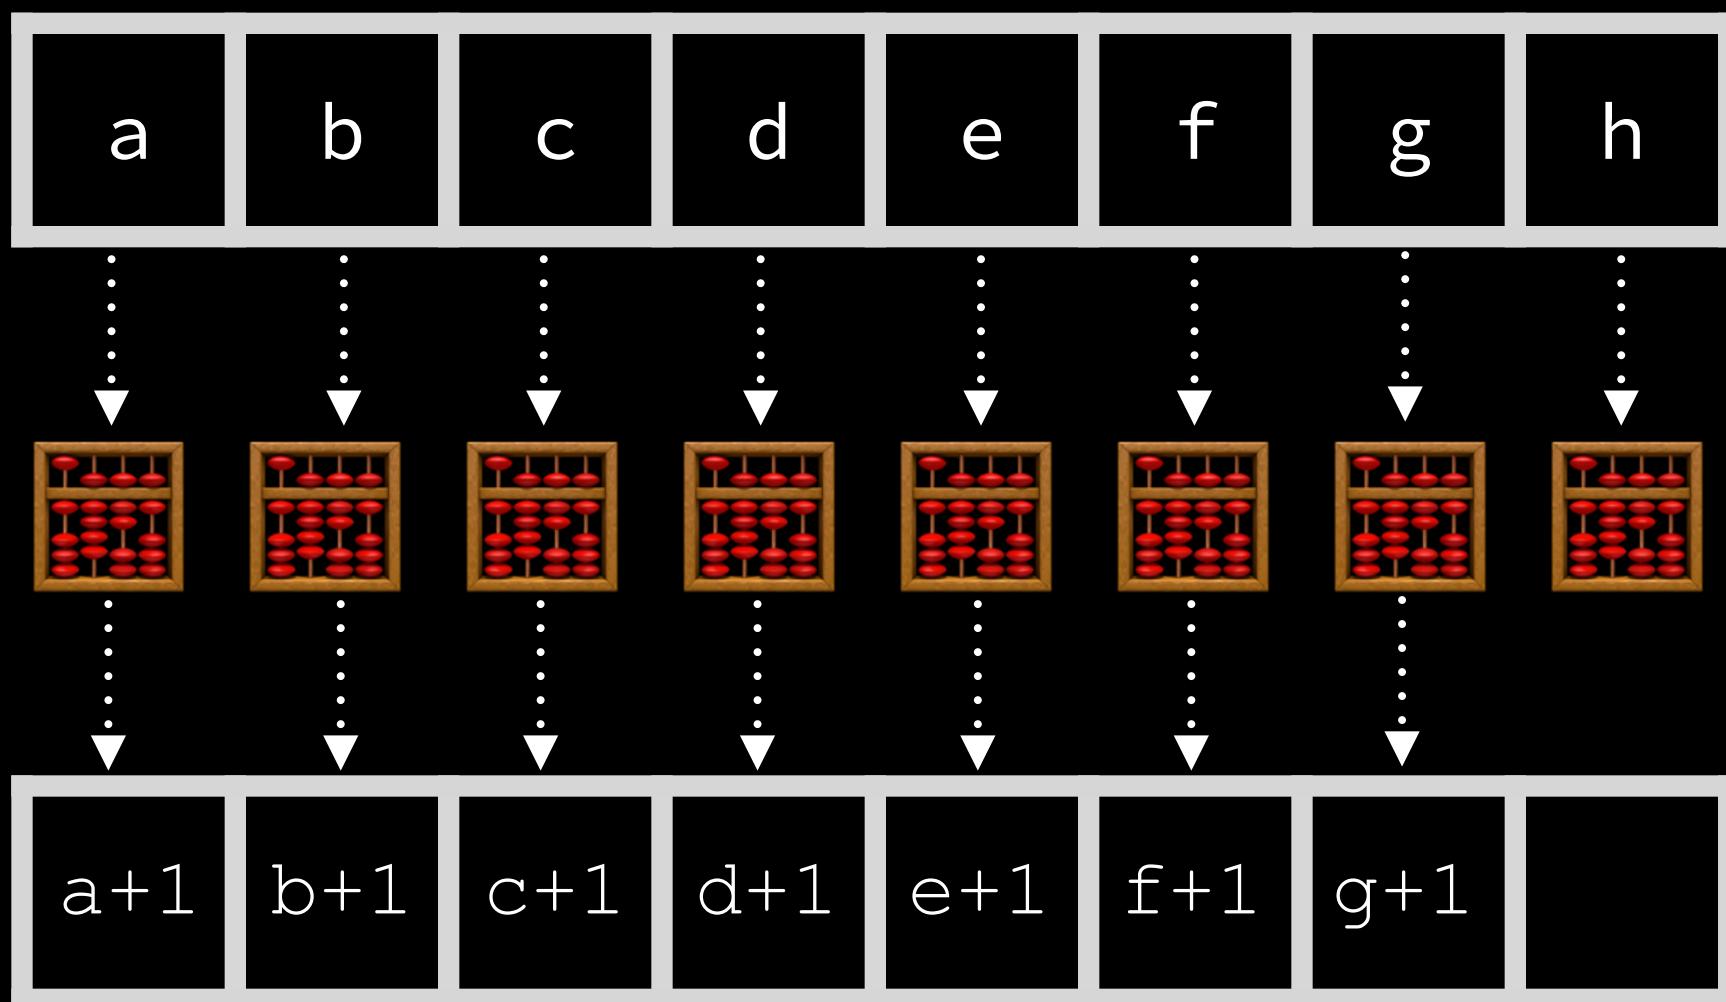
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

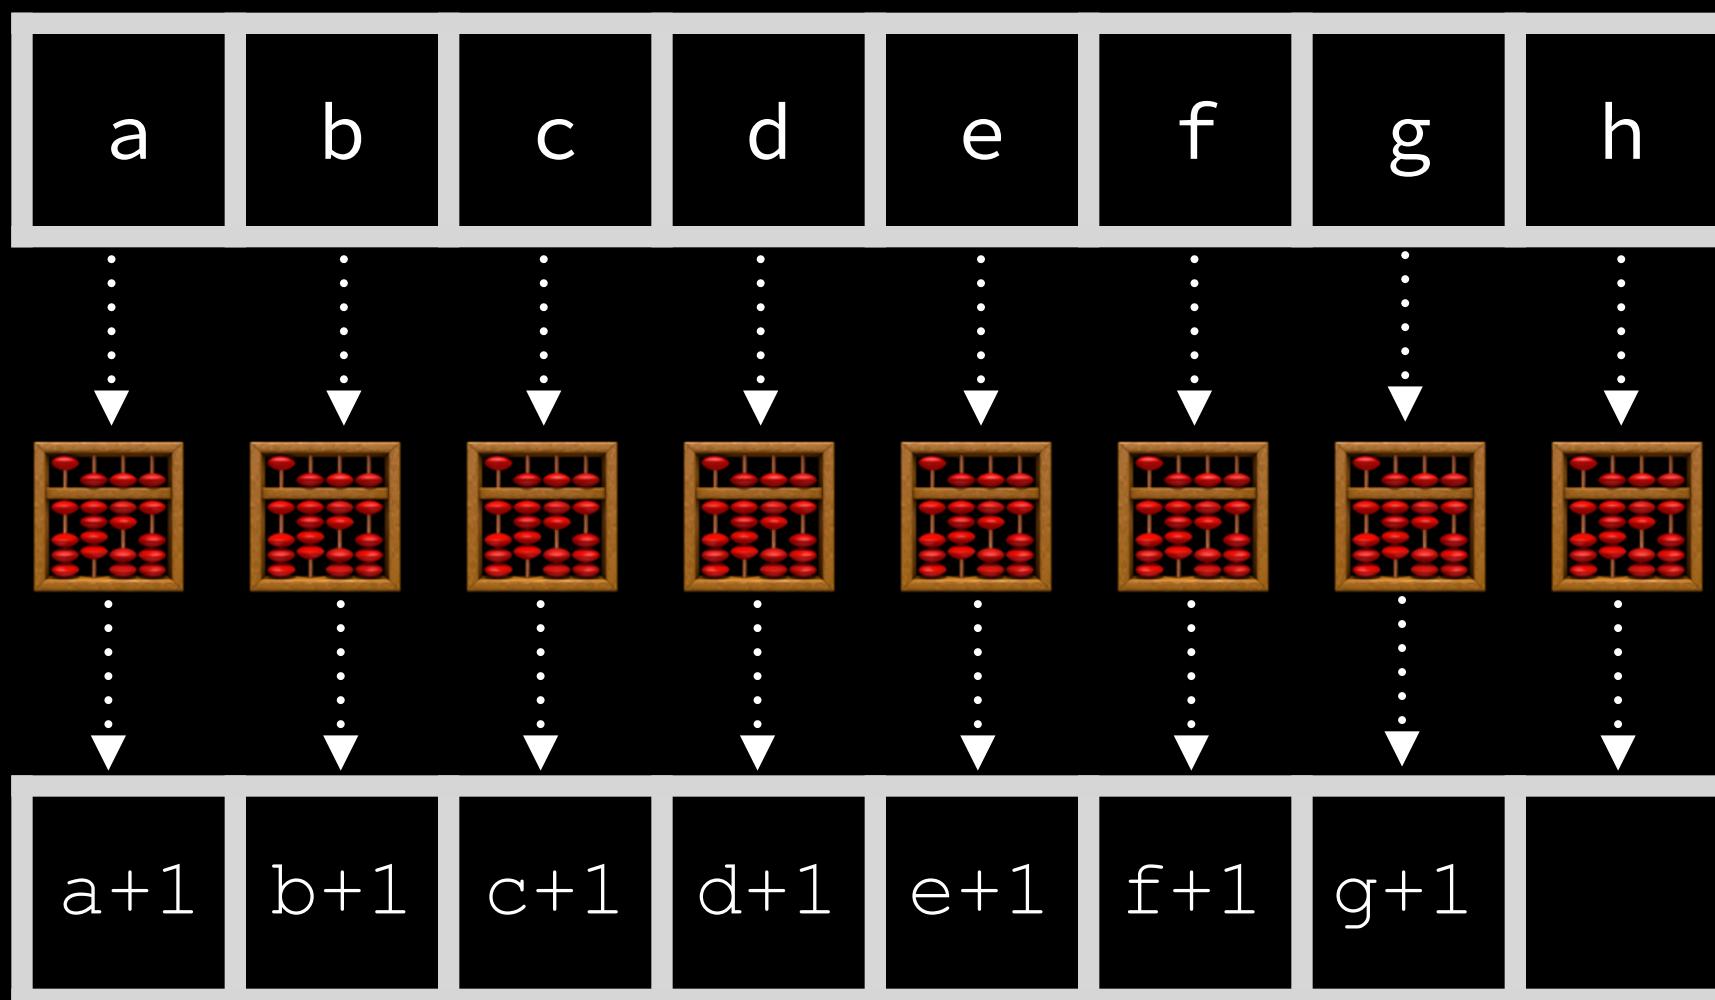
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

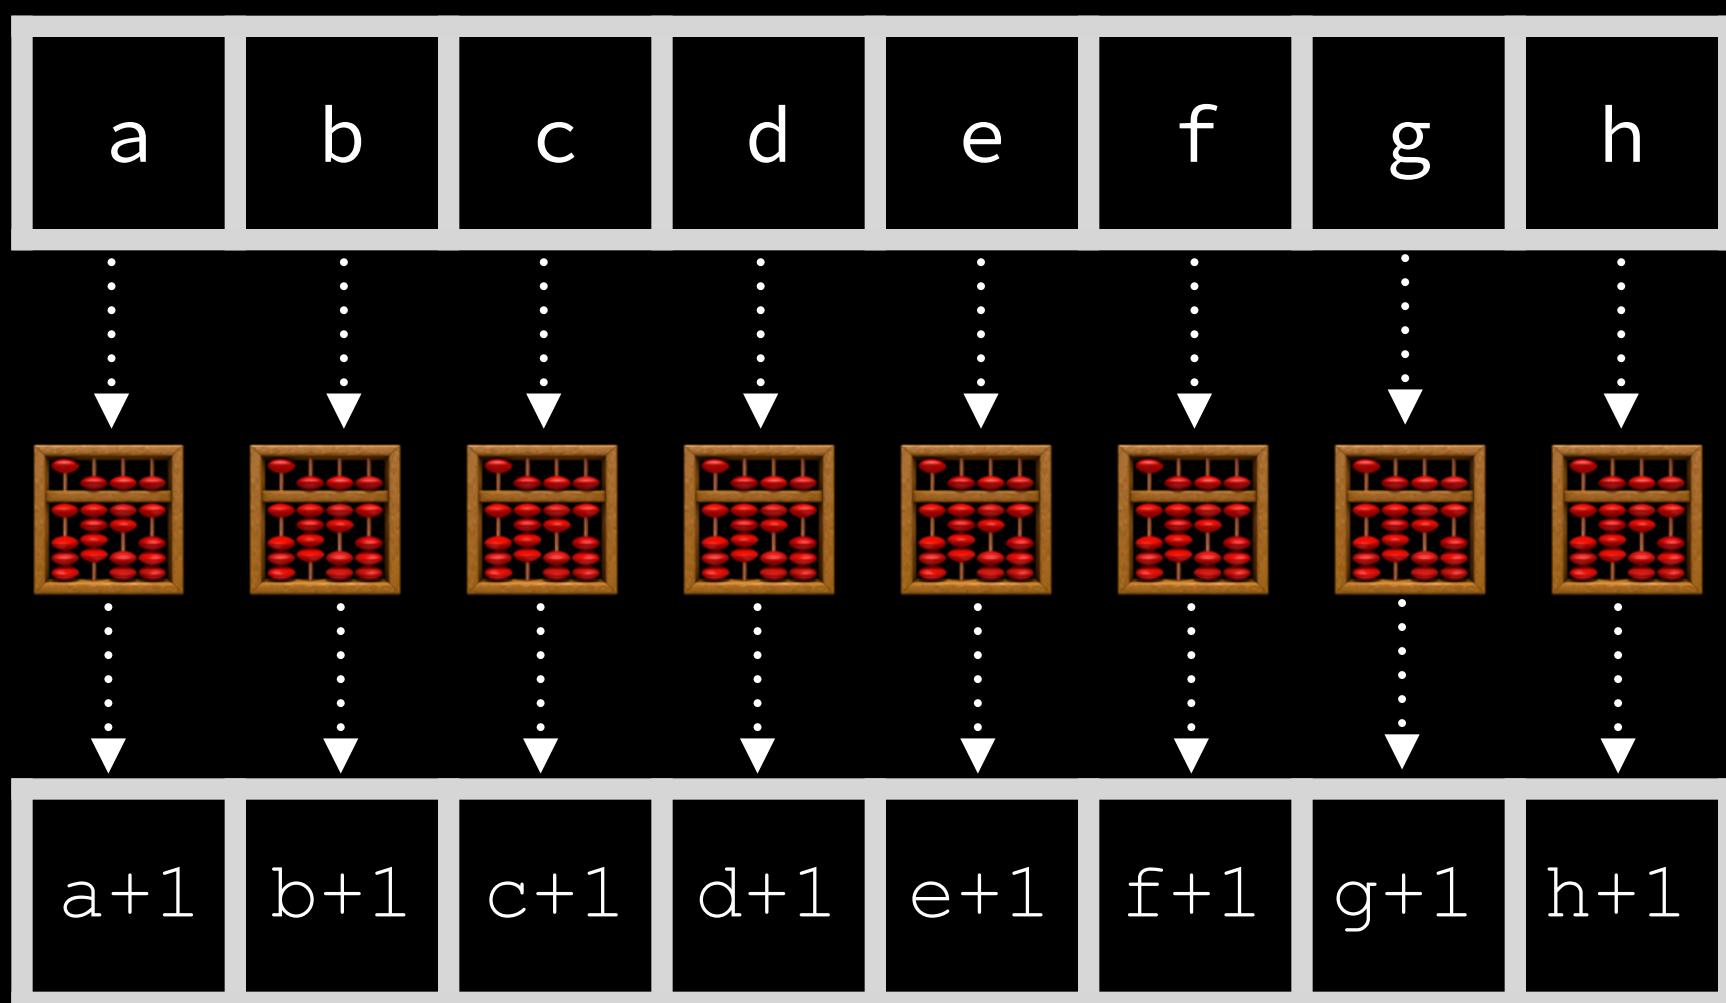
Say I give you all the compute  
you could ask for:



# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:

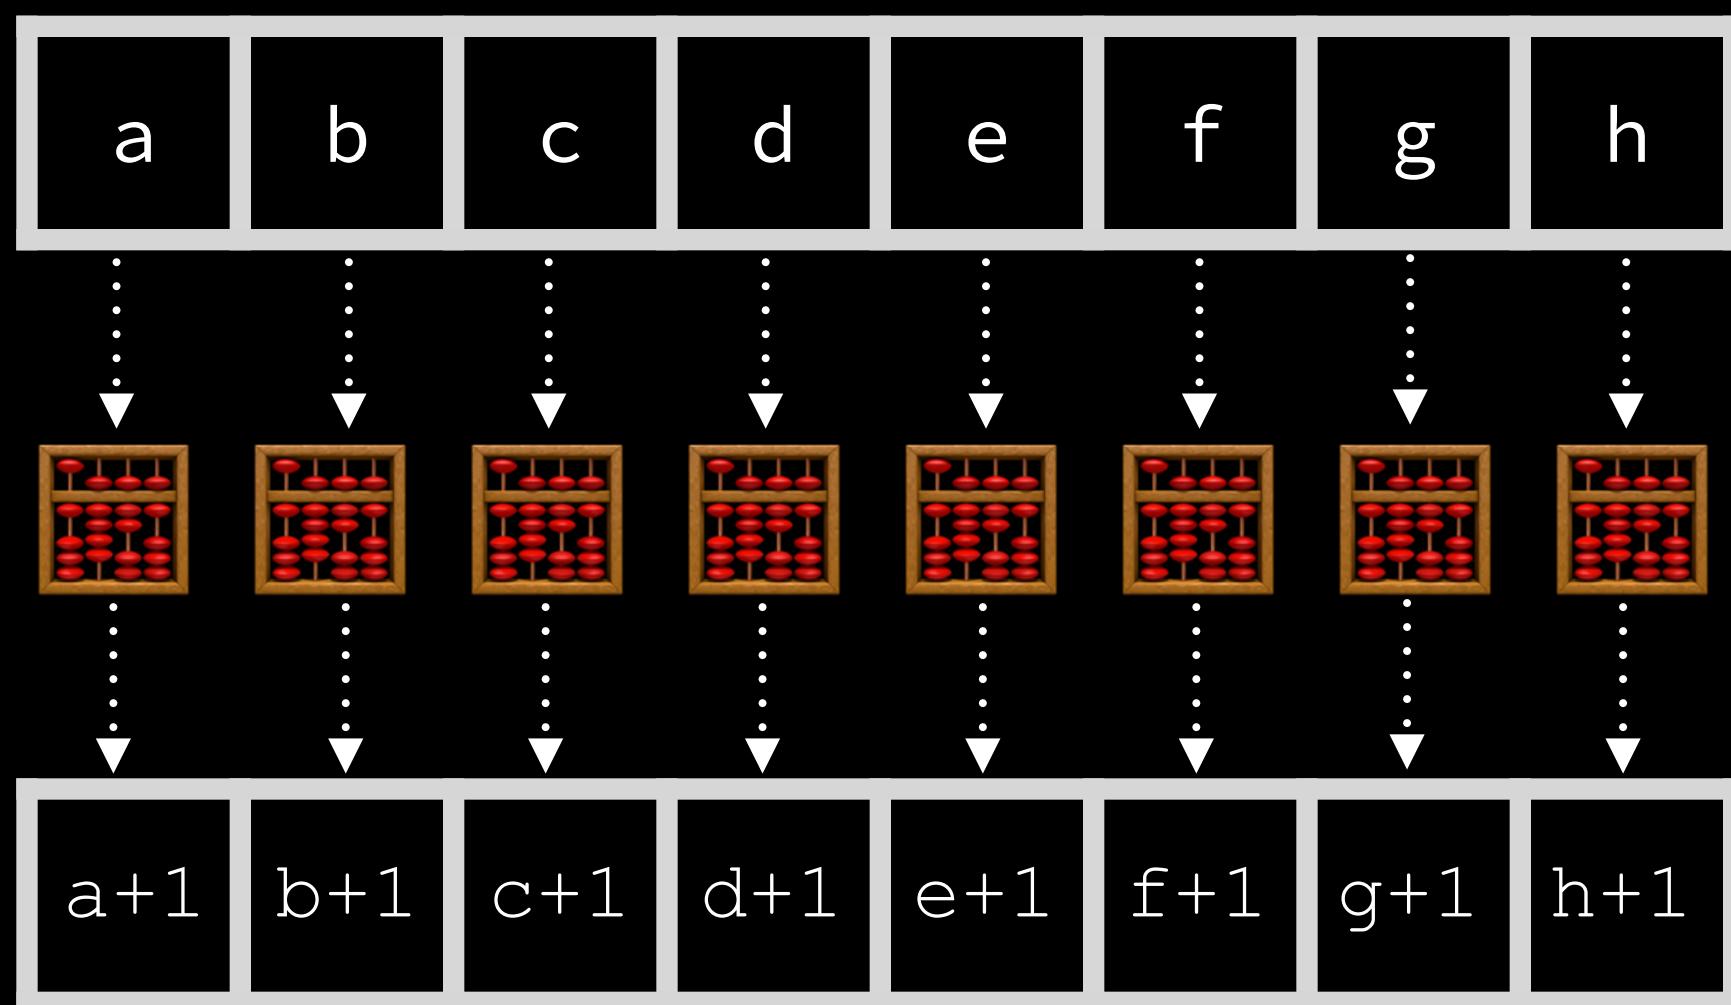


# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:

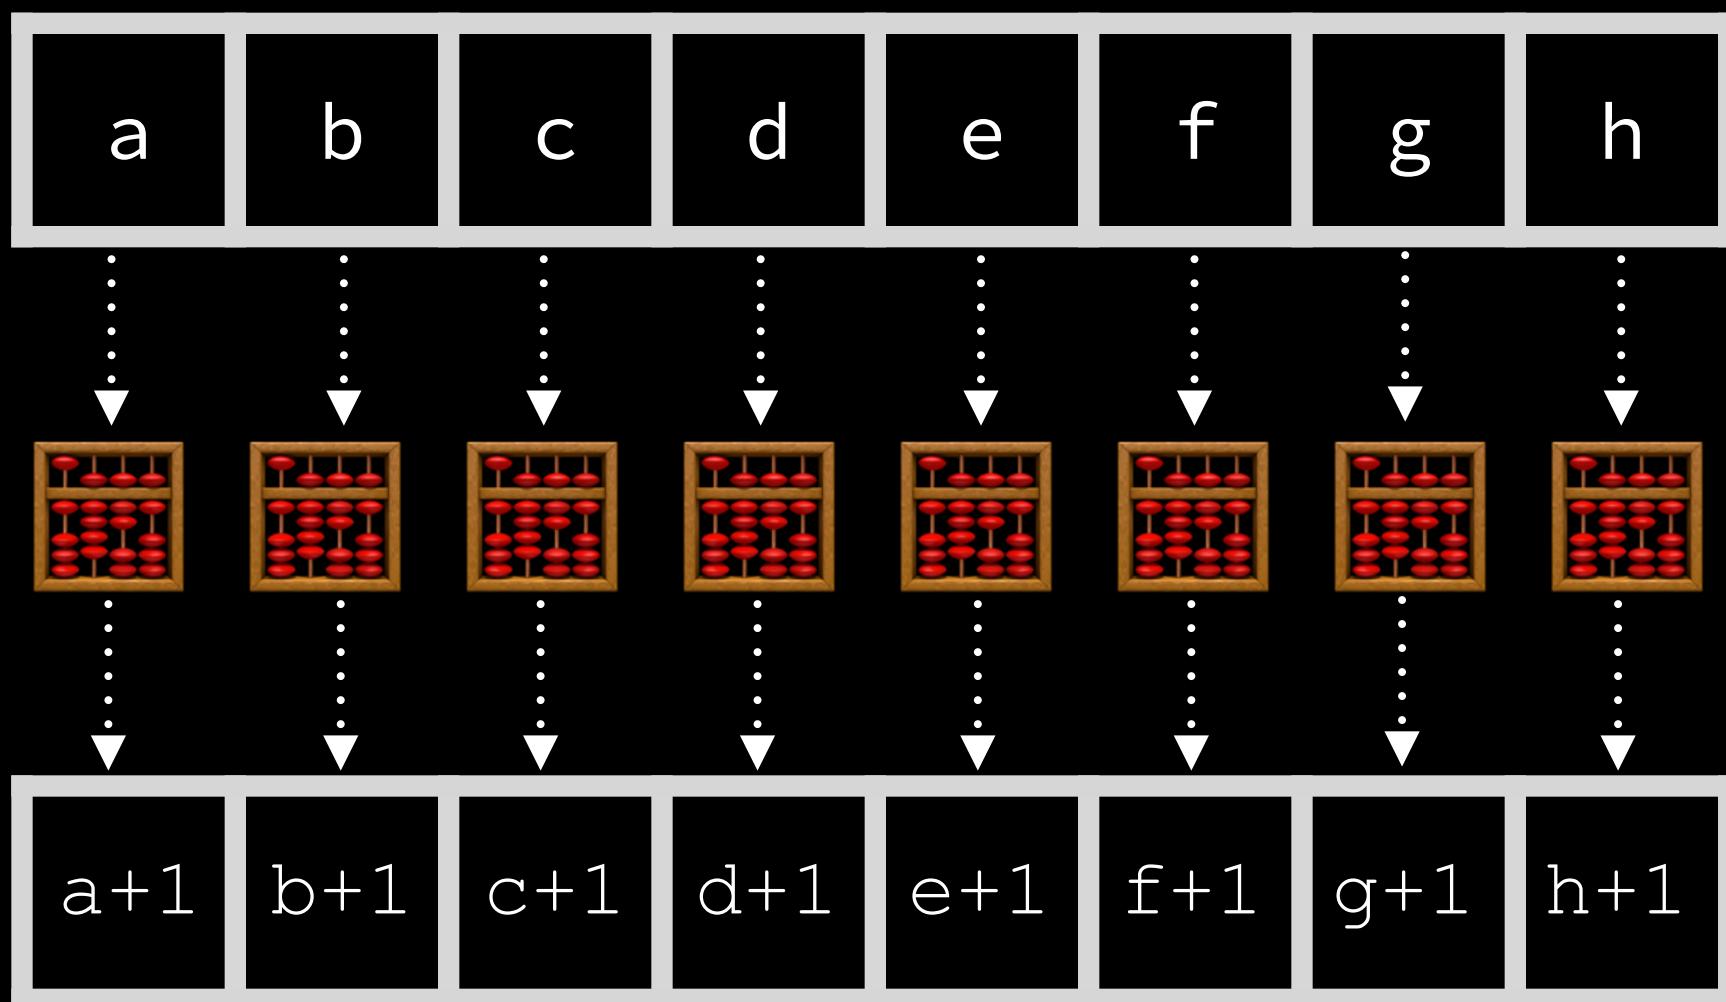
But if we *bank* the arrays,  
we really can parallelize:



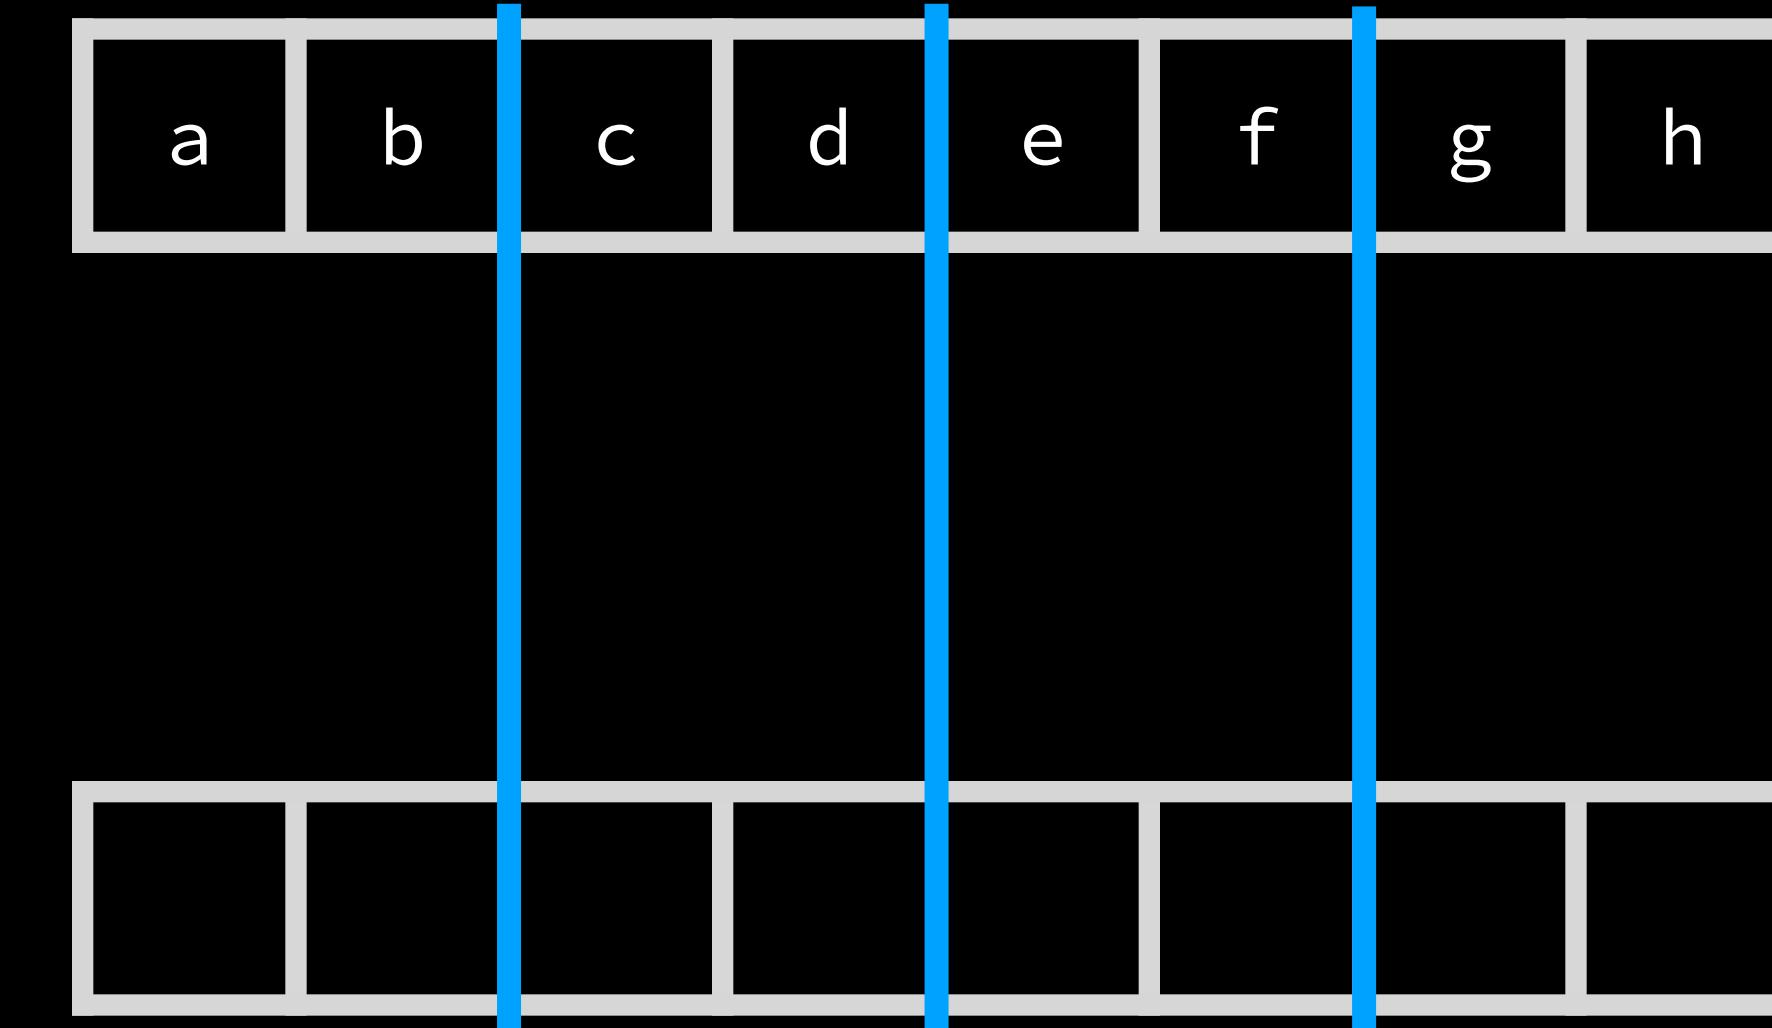
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



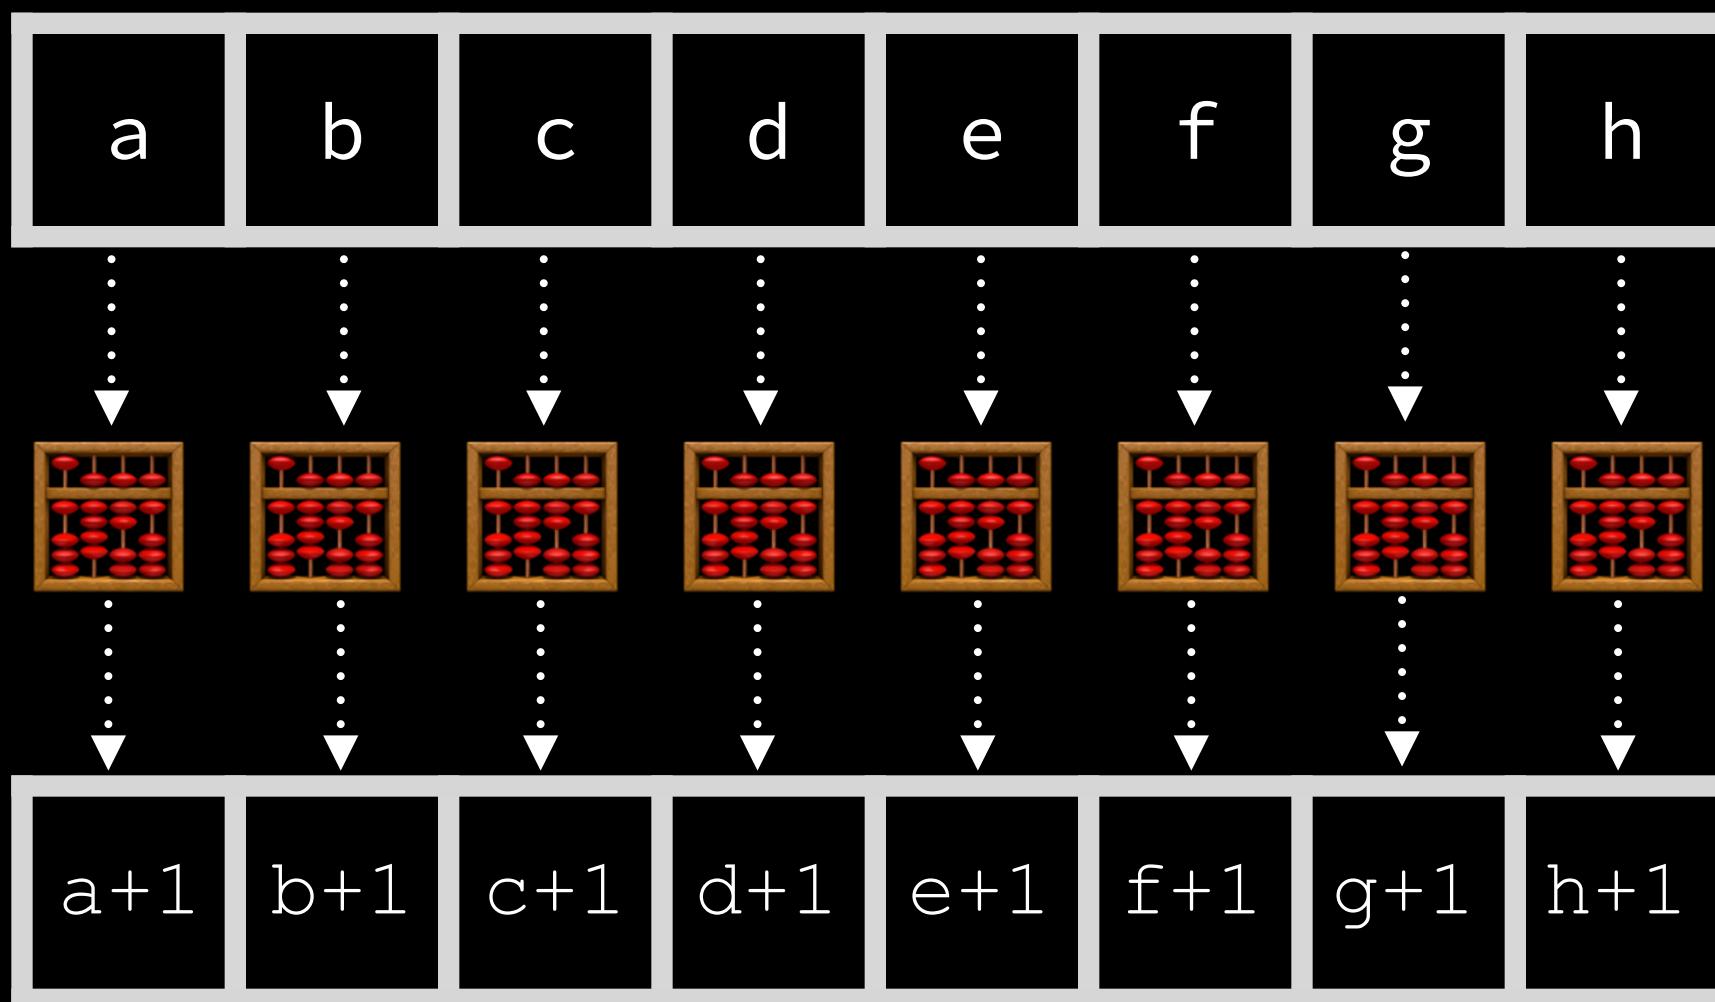
But if we *bank* the arrays,  
we really can parallelize:



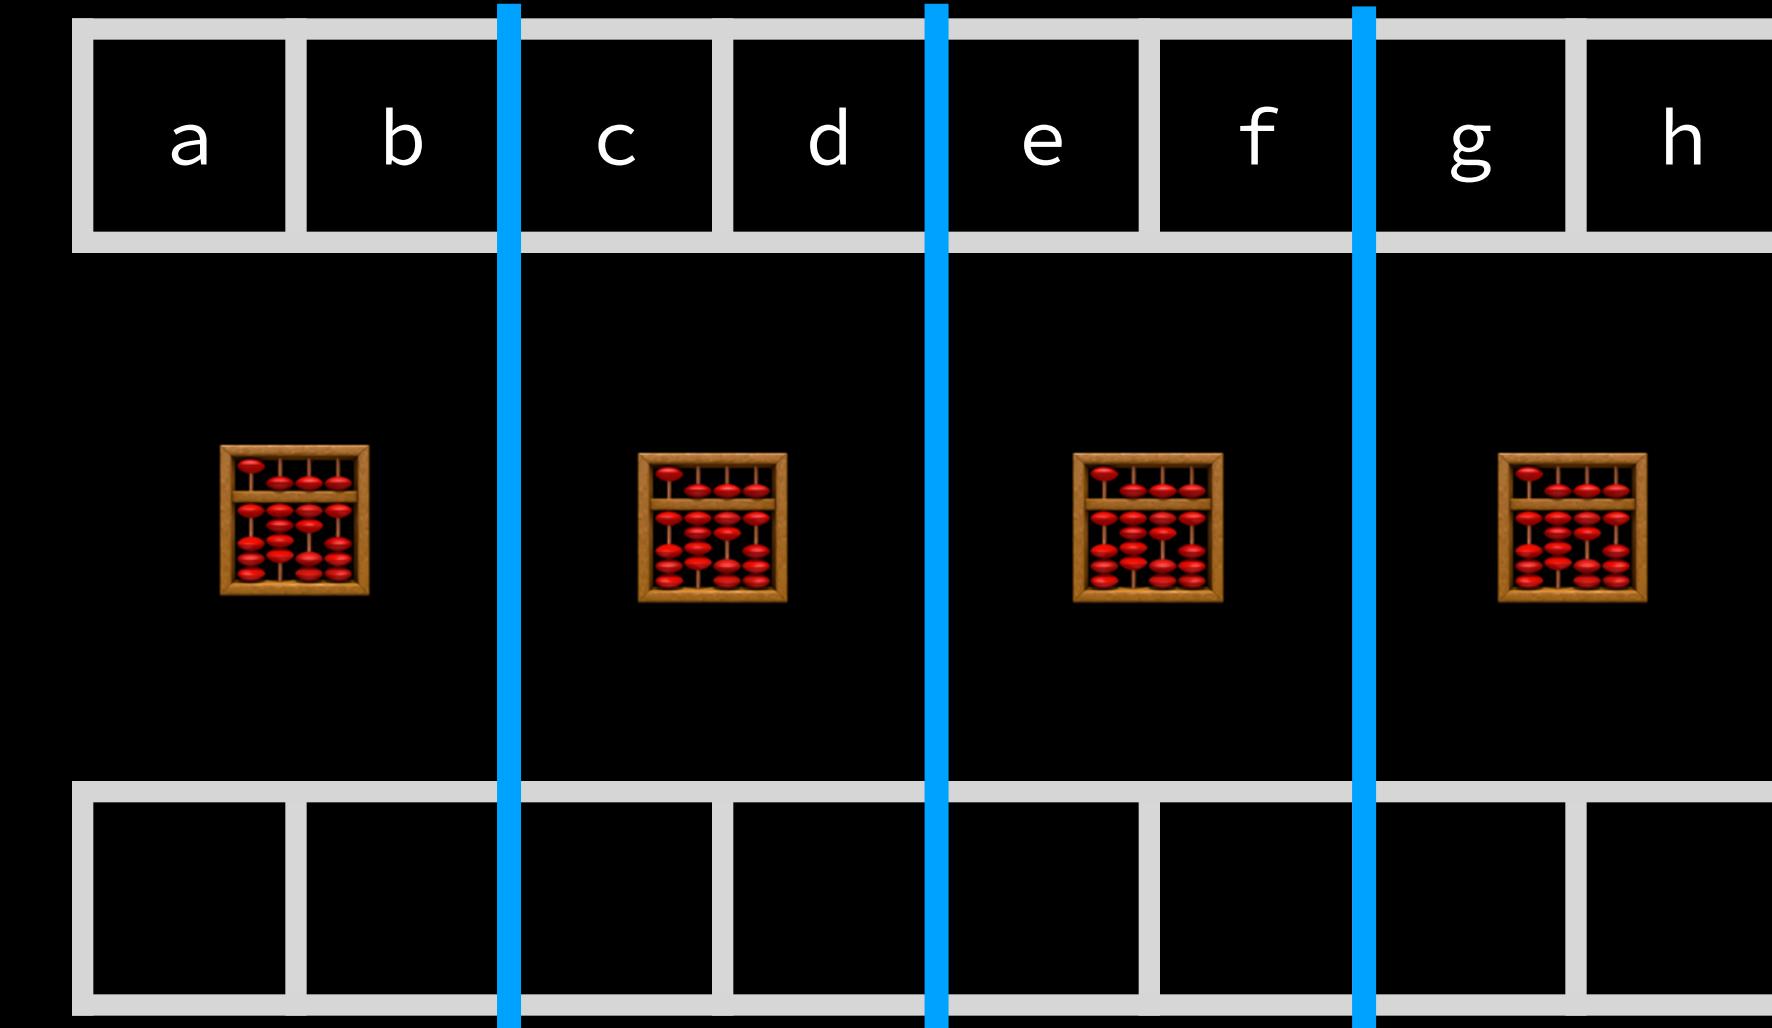
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



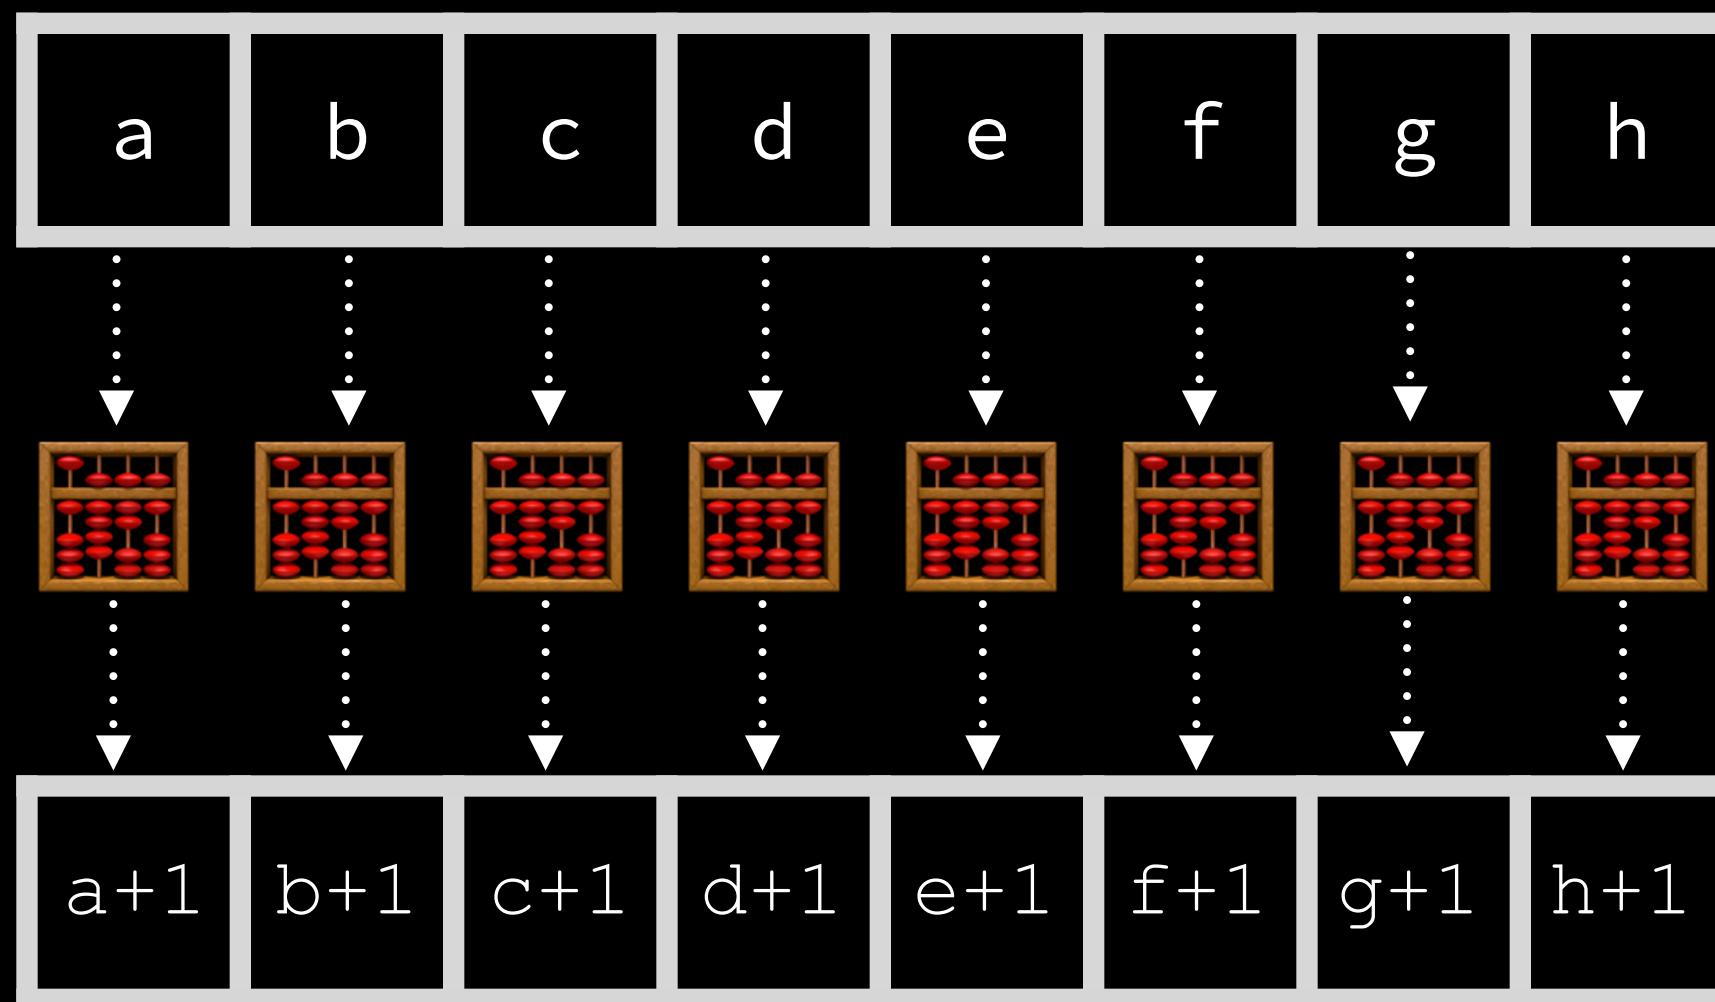
But if we *bank* the arrays,  
we really can parallelize:



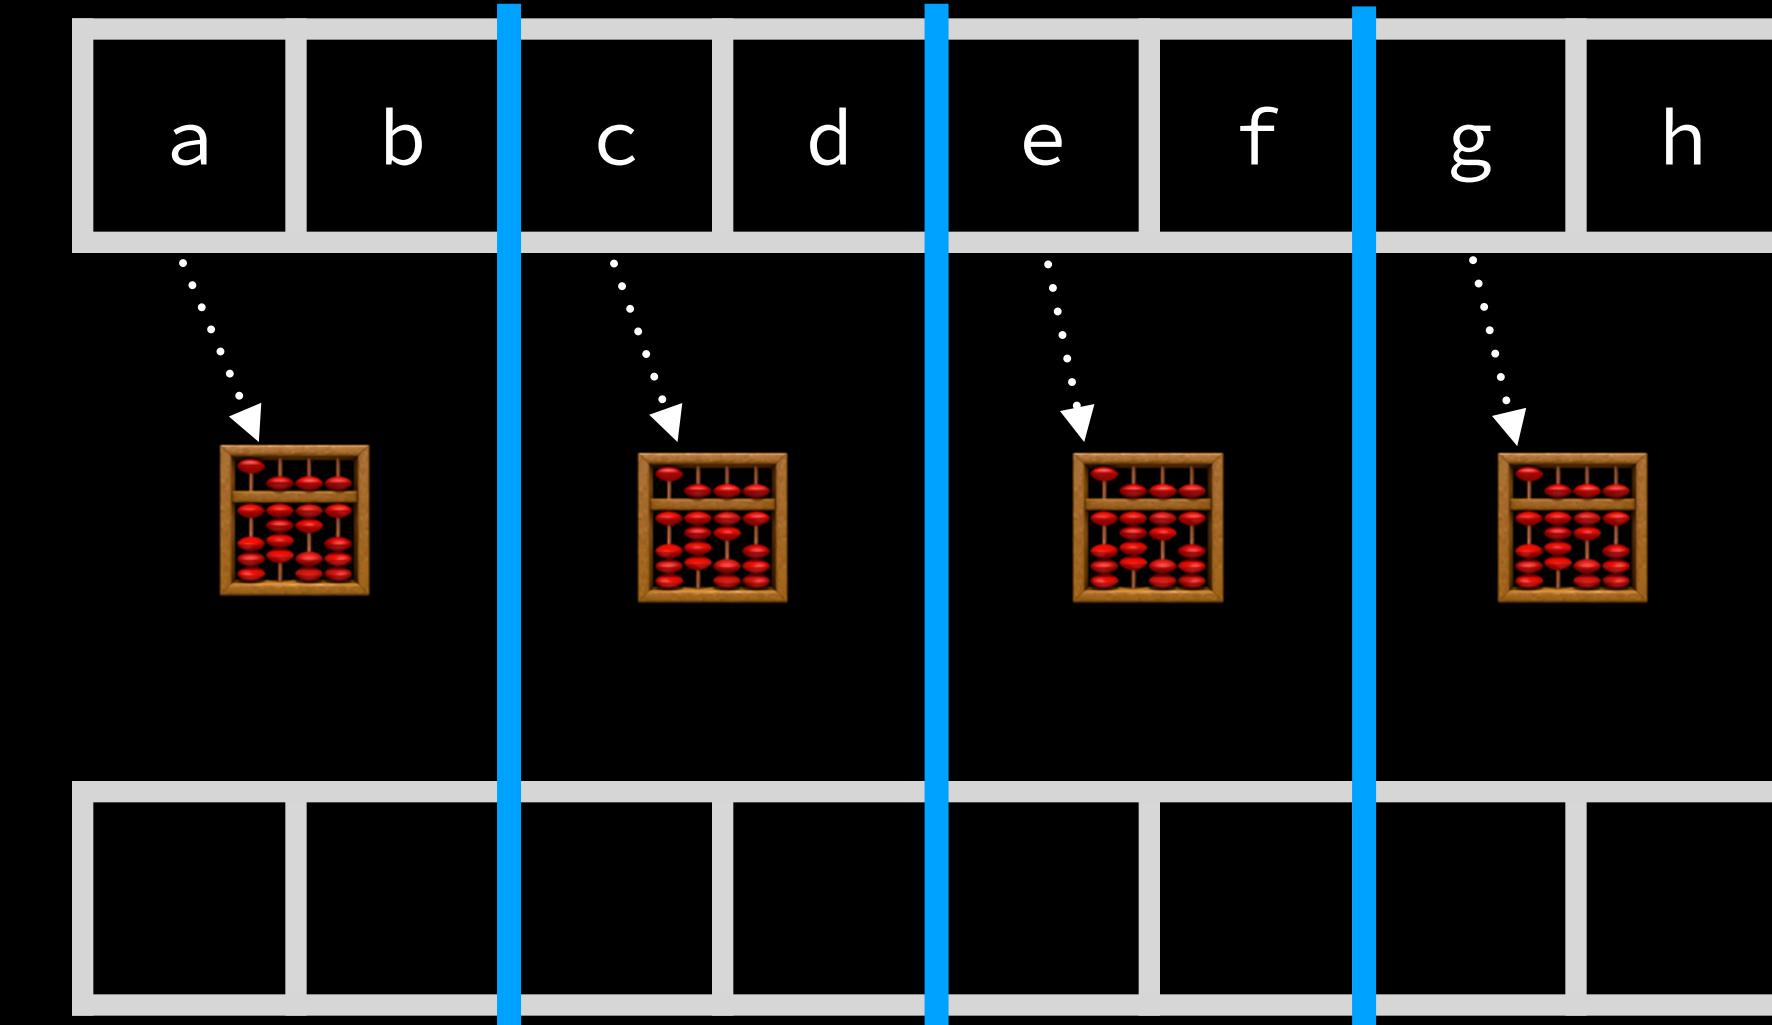
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



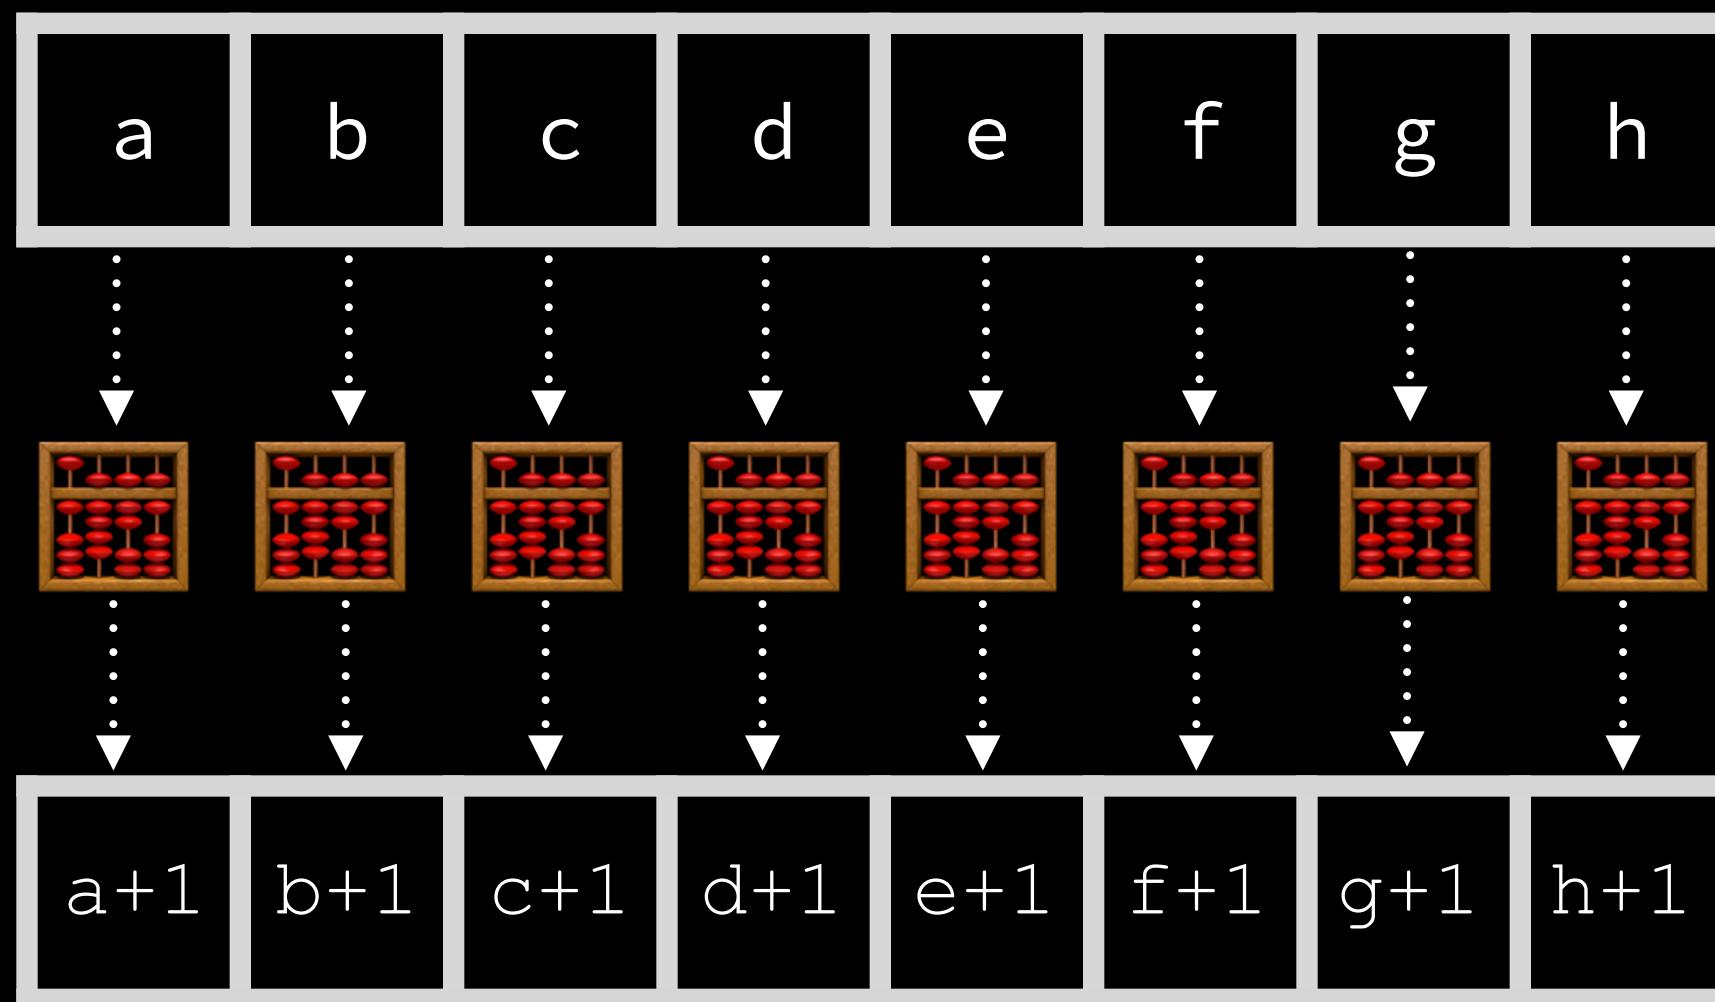
But if we *bank* the arrays,  
we really can parallelize:



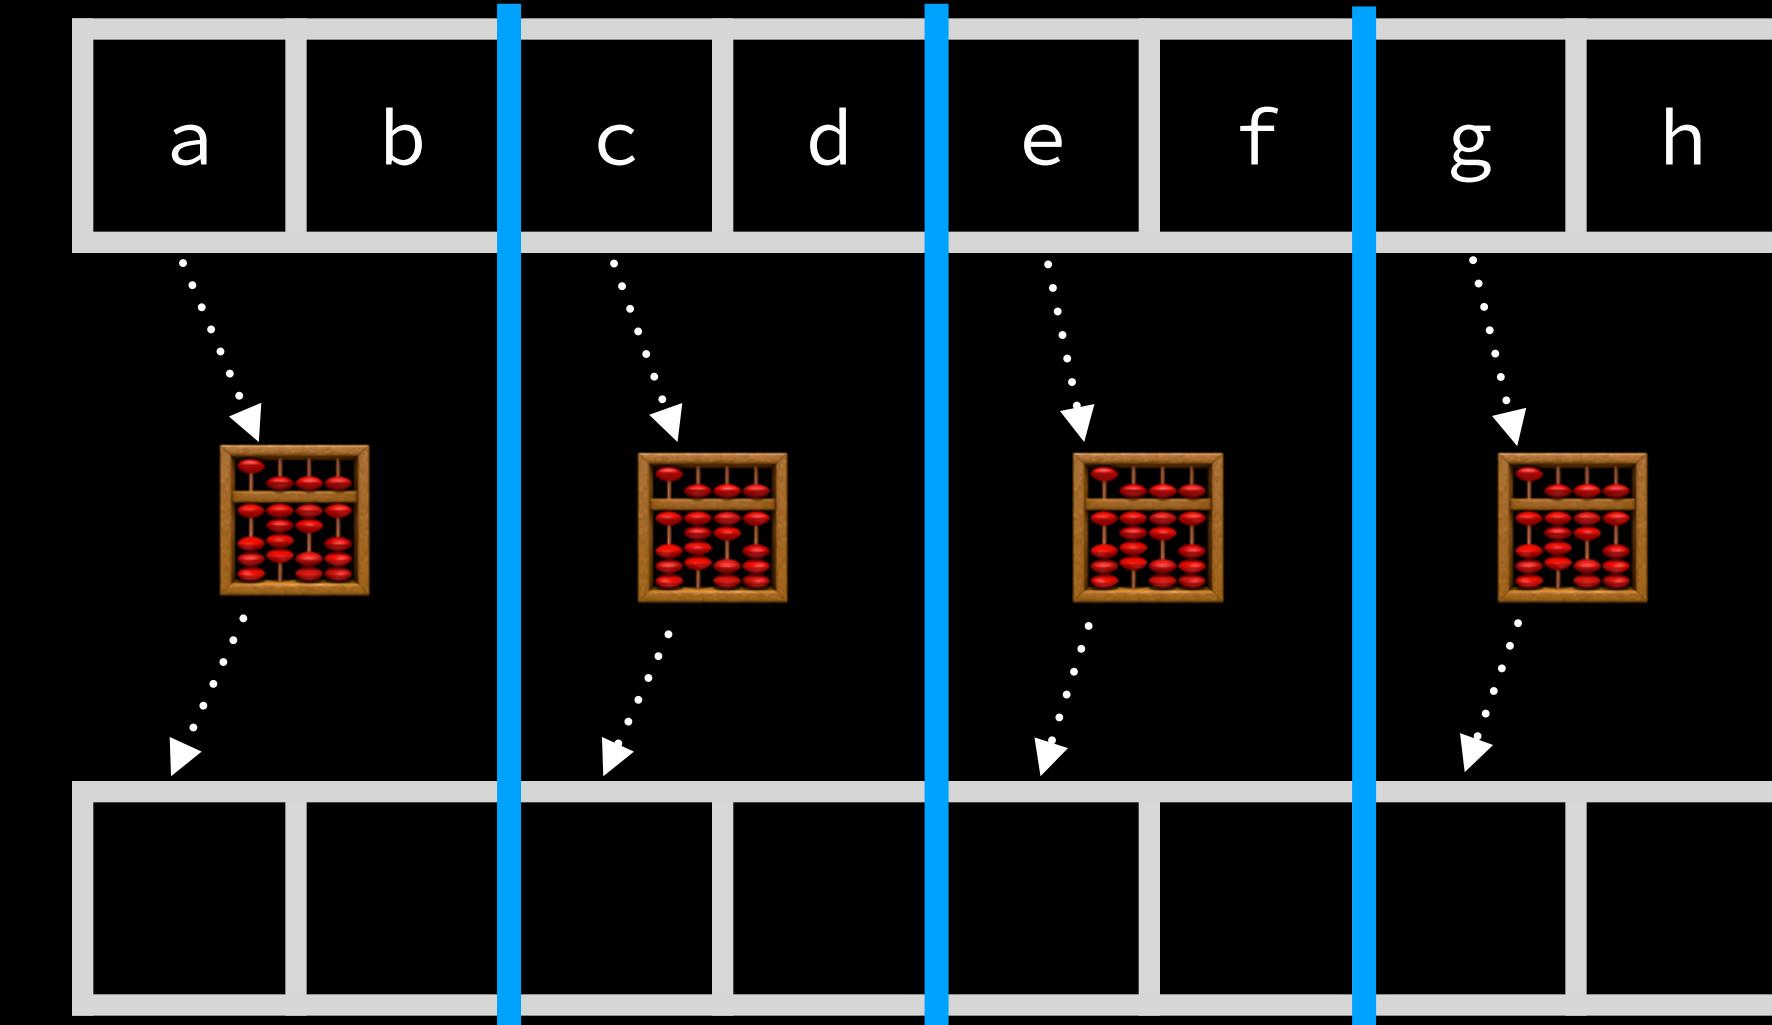
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



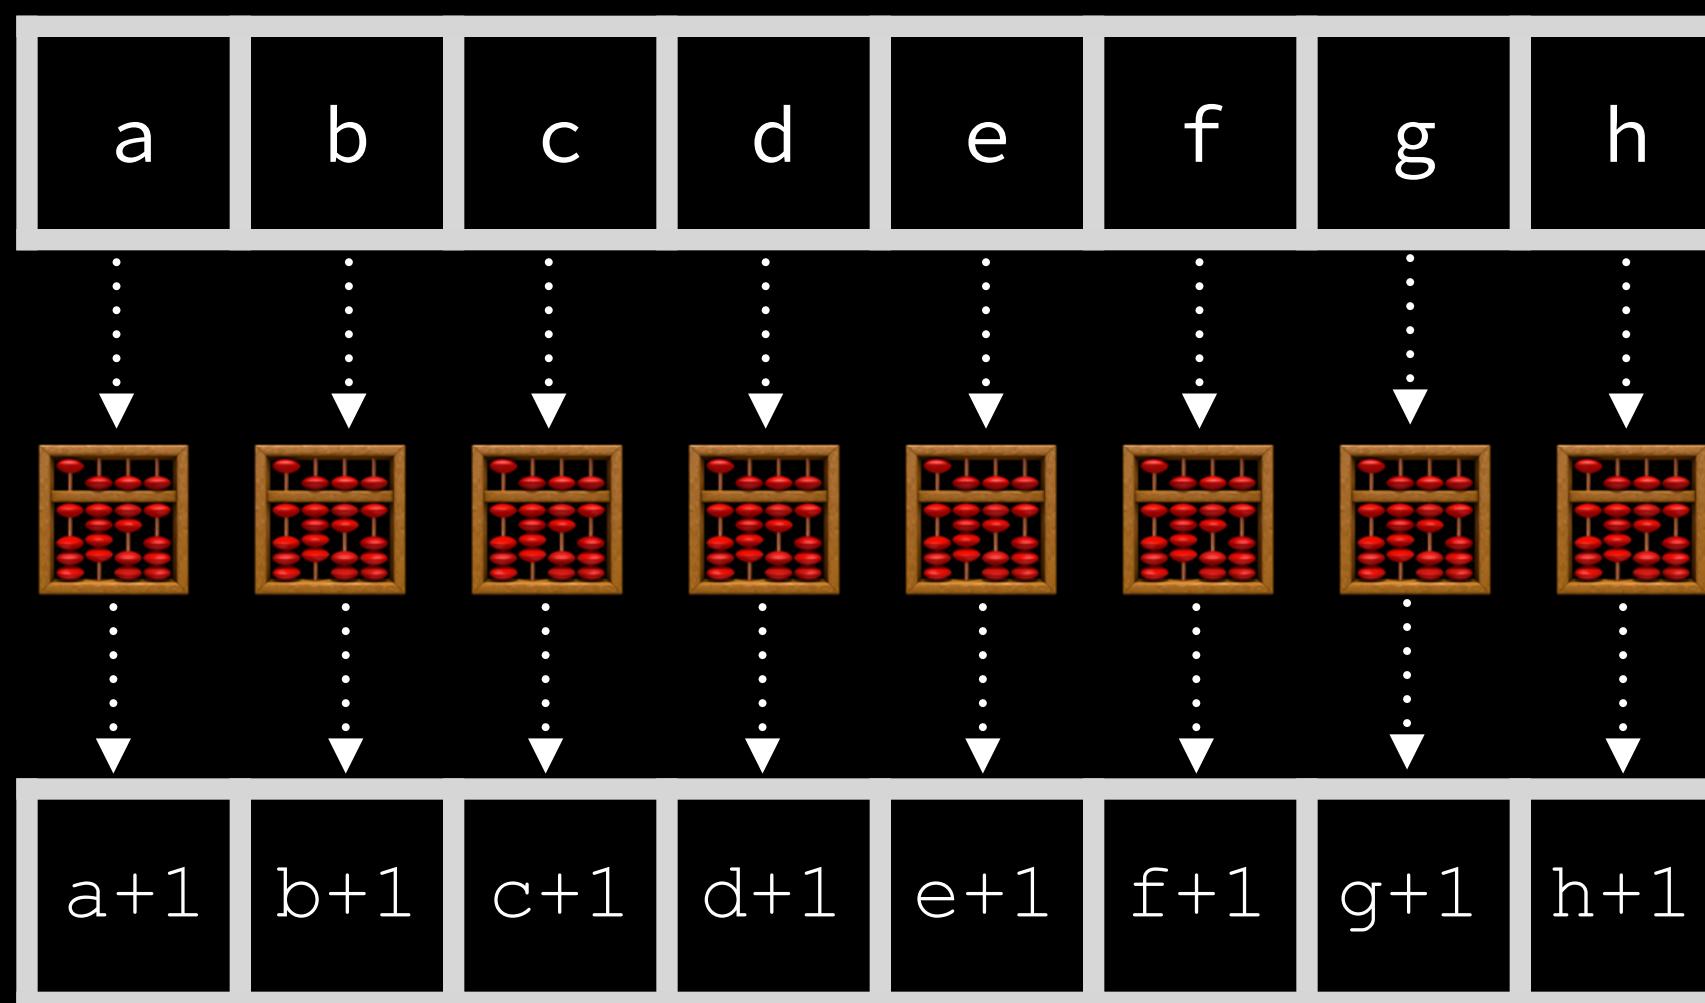
But if we *bank* the arrays,  
we really can parallelize:



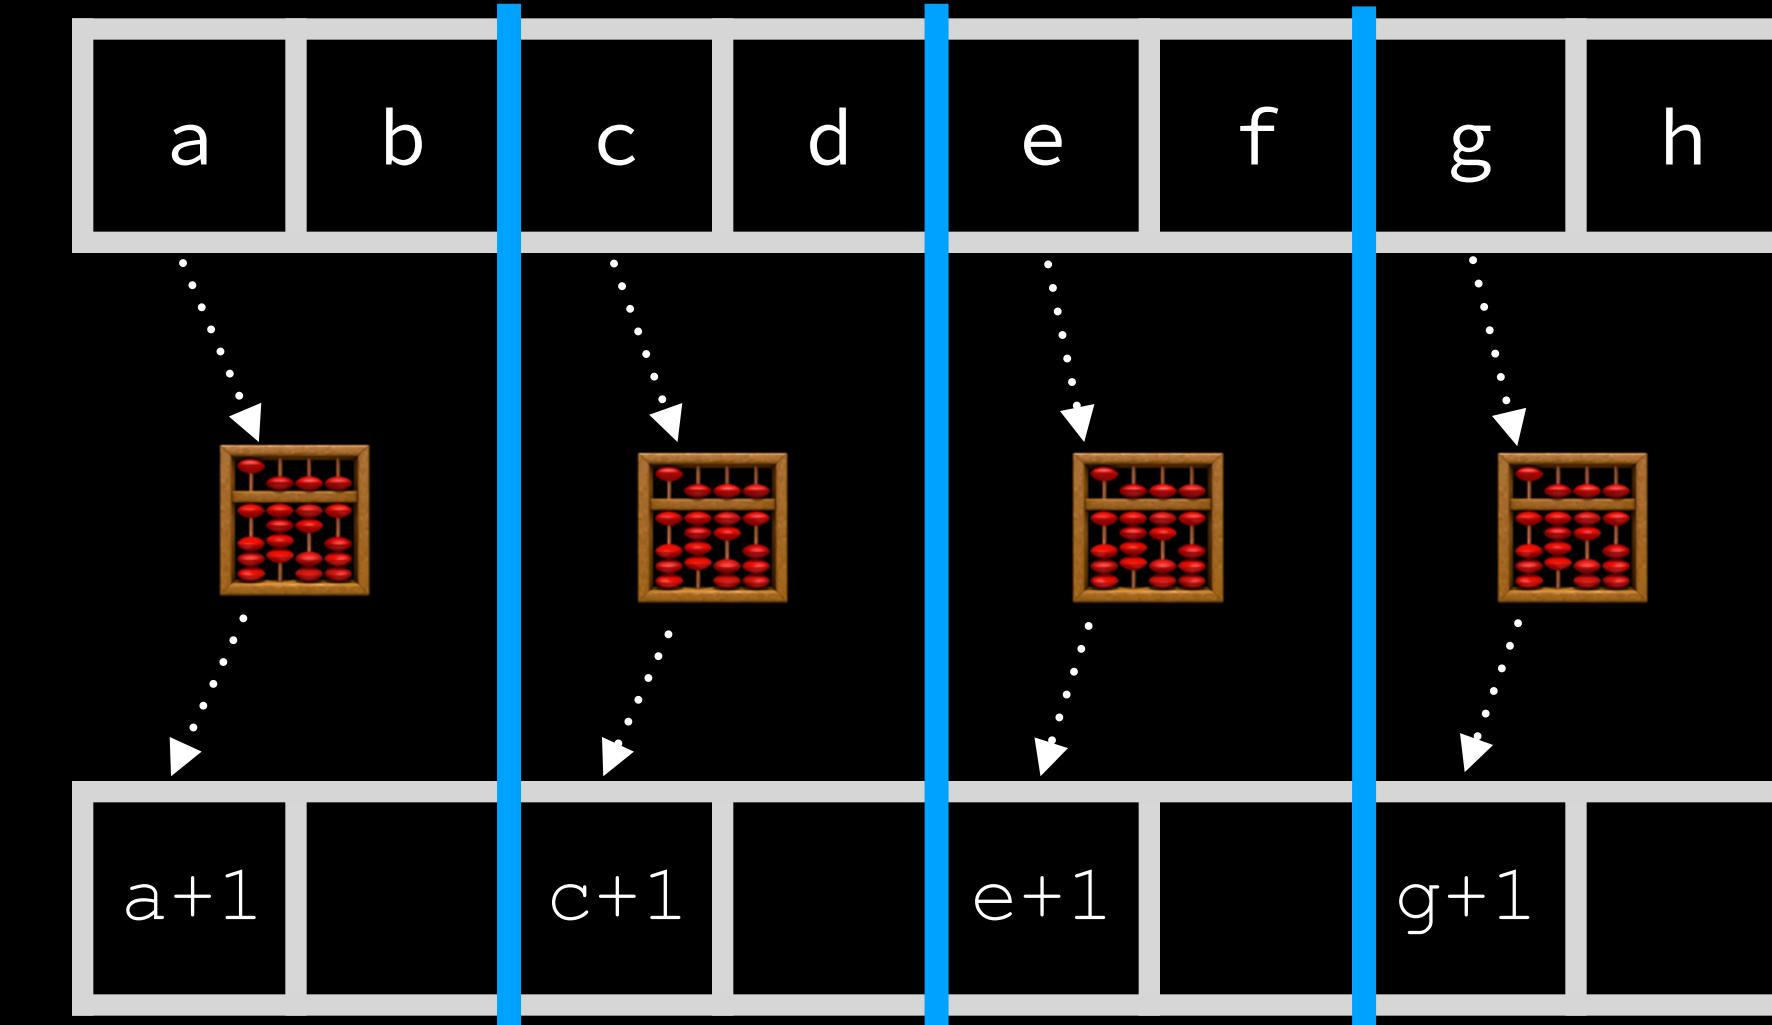
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



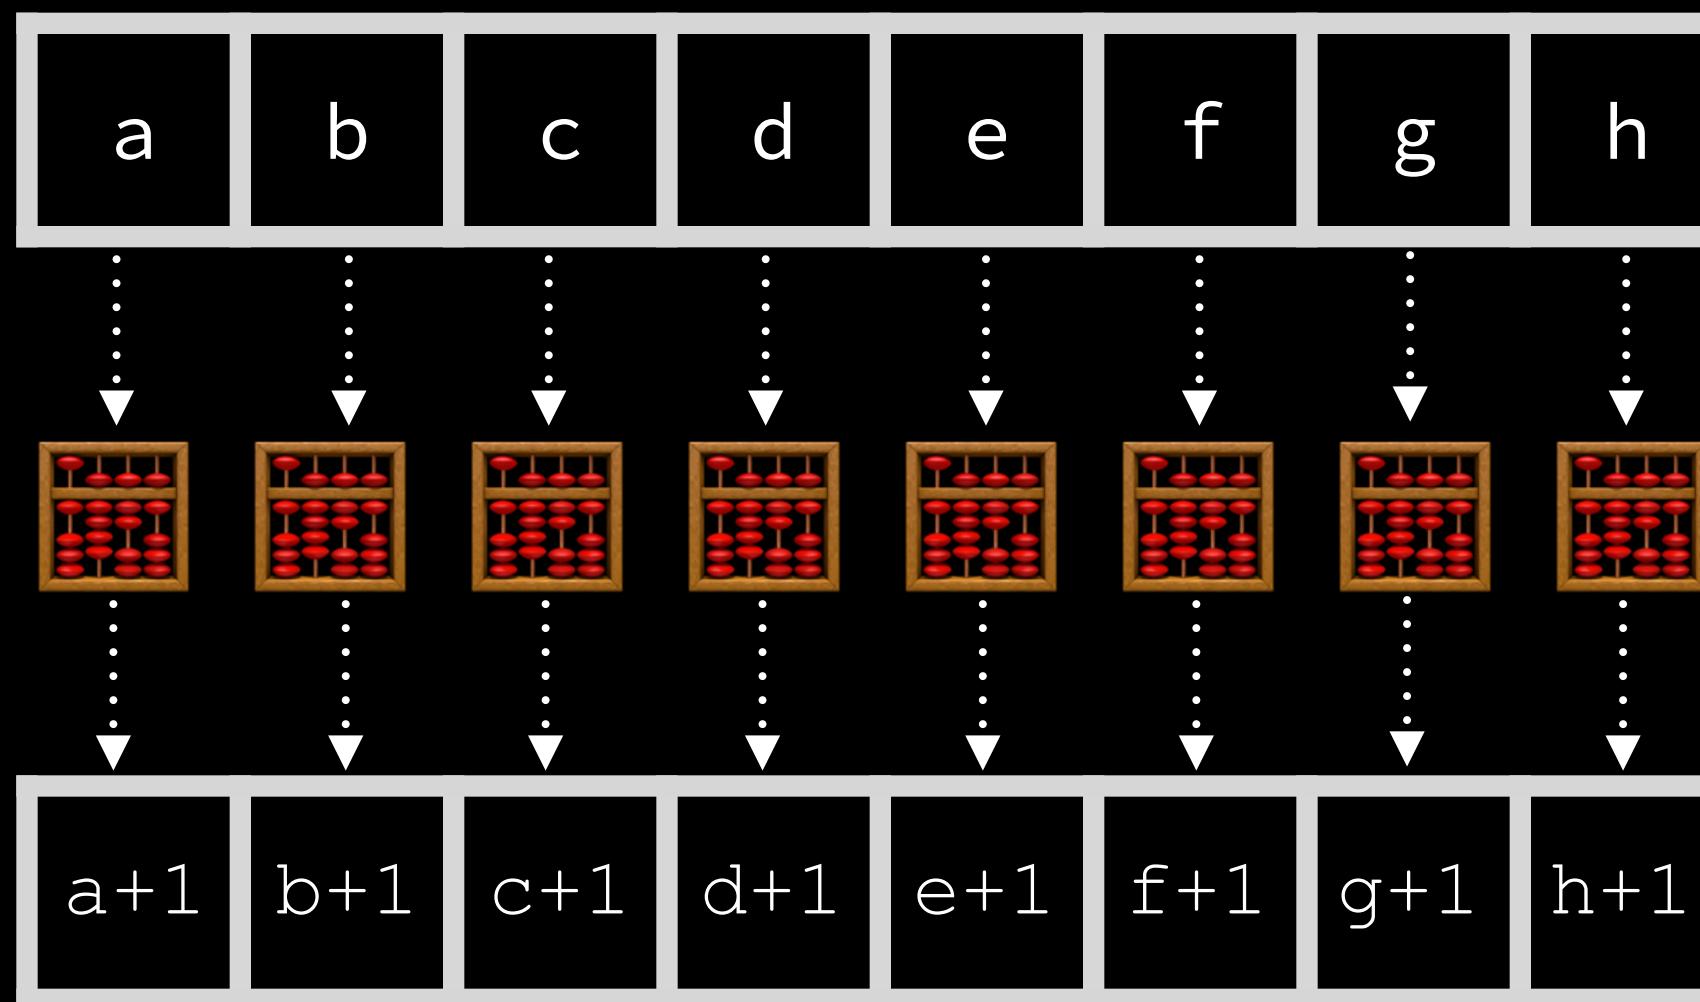
But if we *bank* the arrays,  
we really can parallelize:



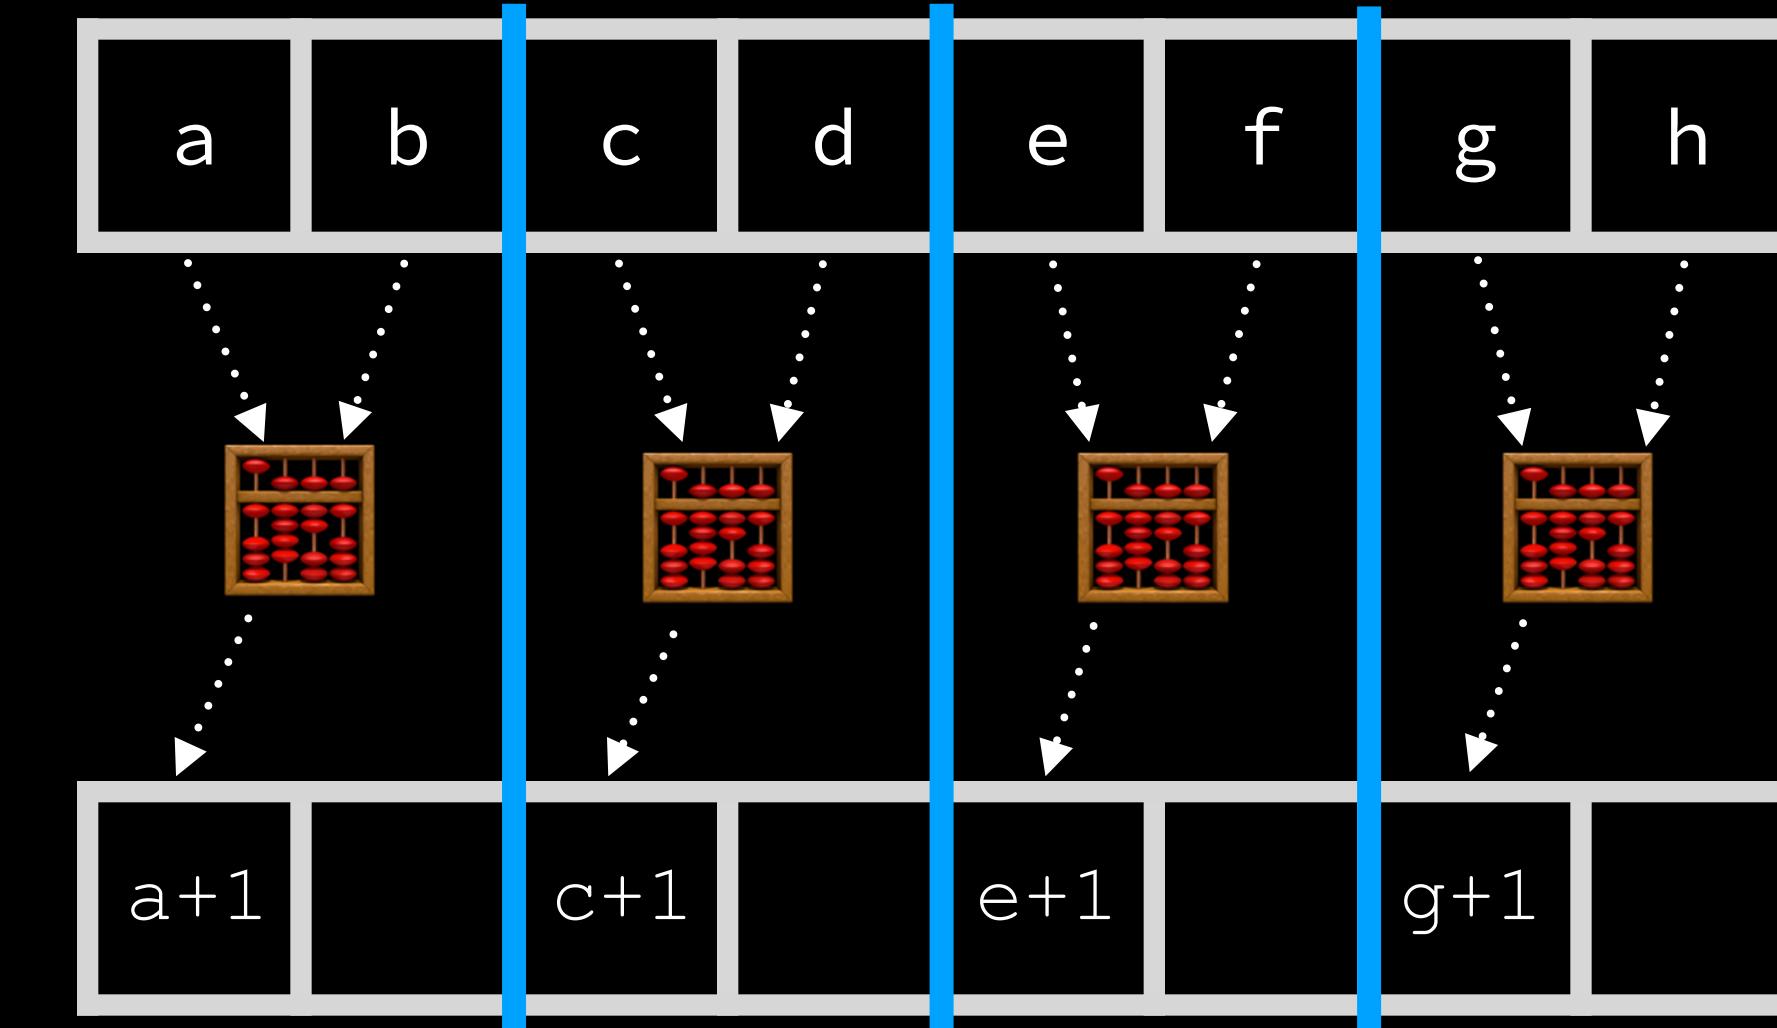
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



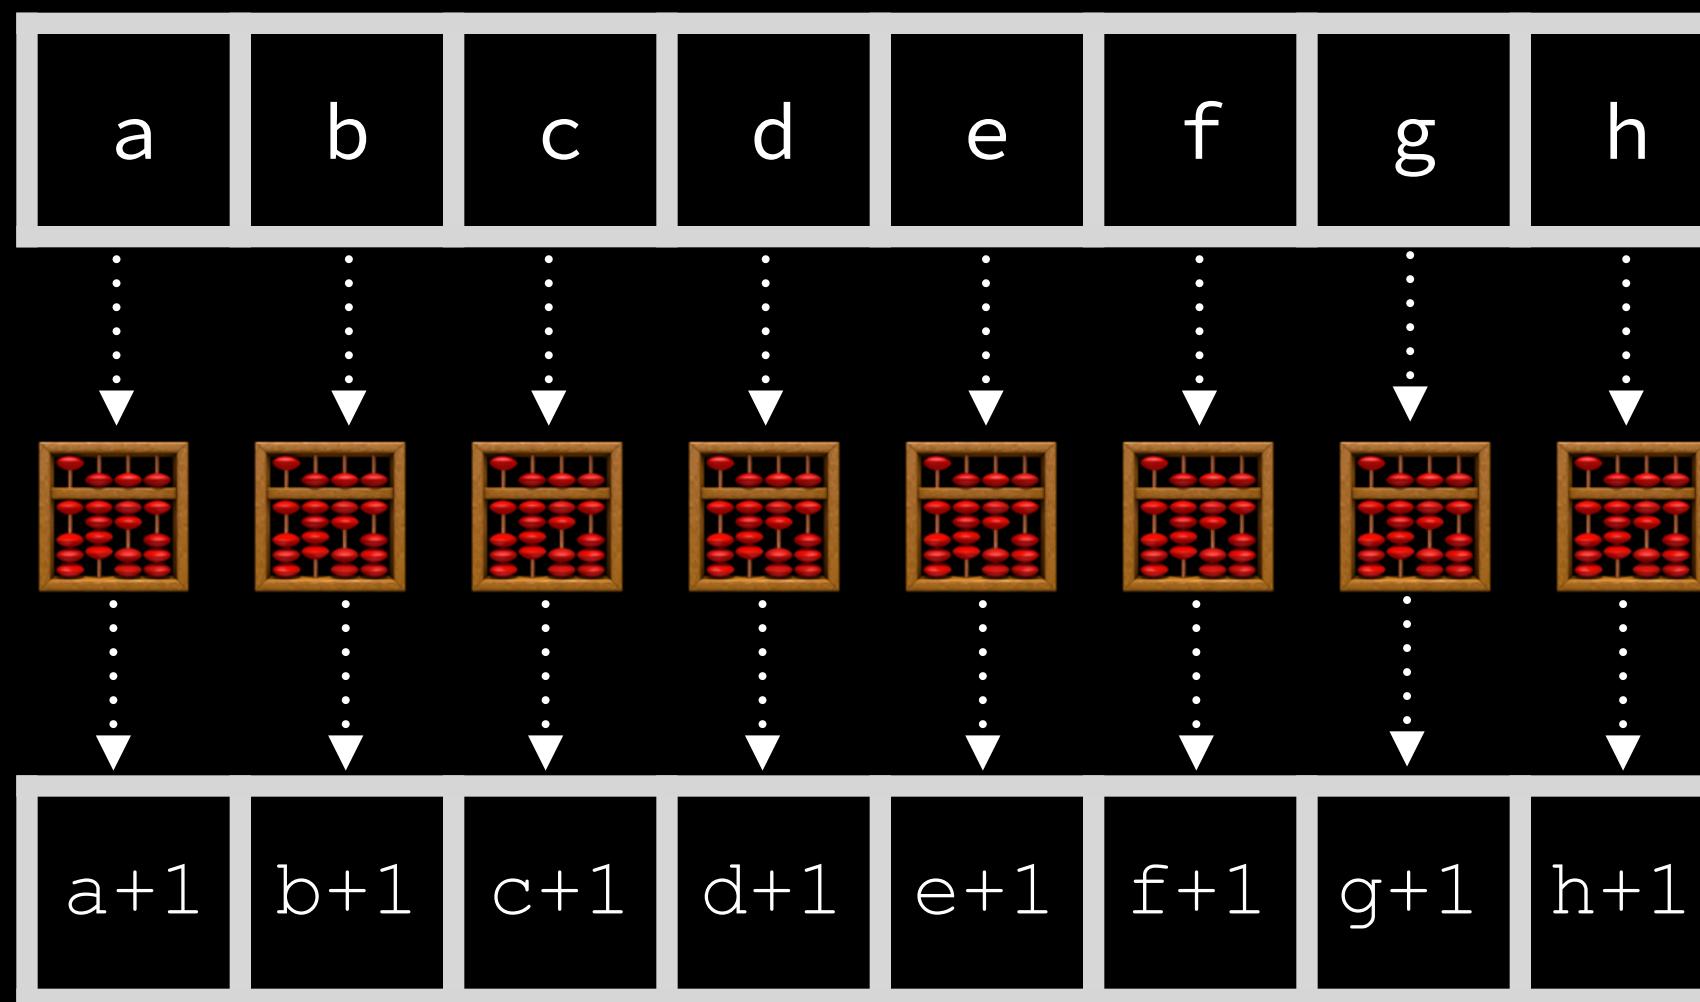
But if we *bank* the arrays,  
we really can parallelize:



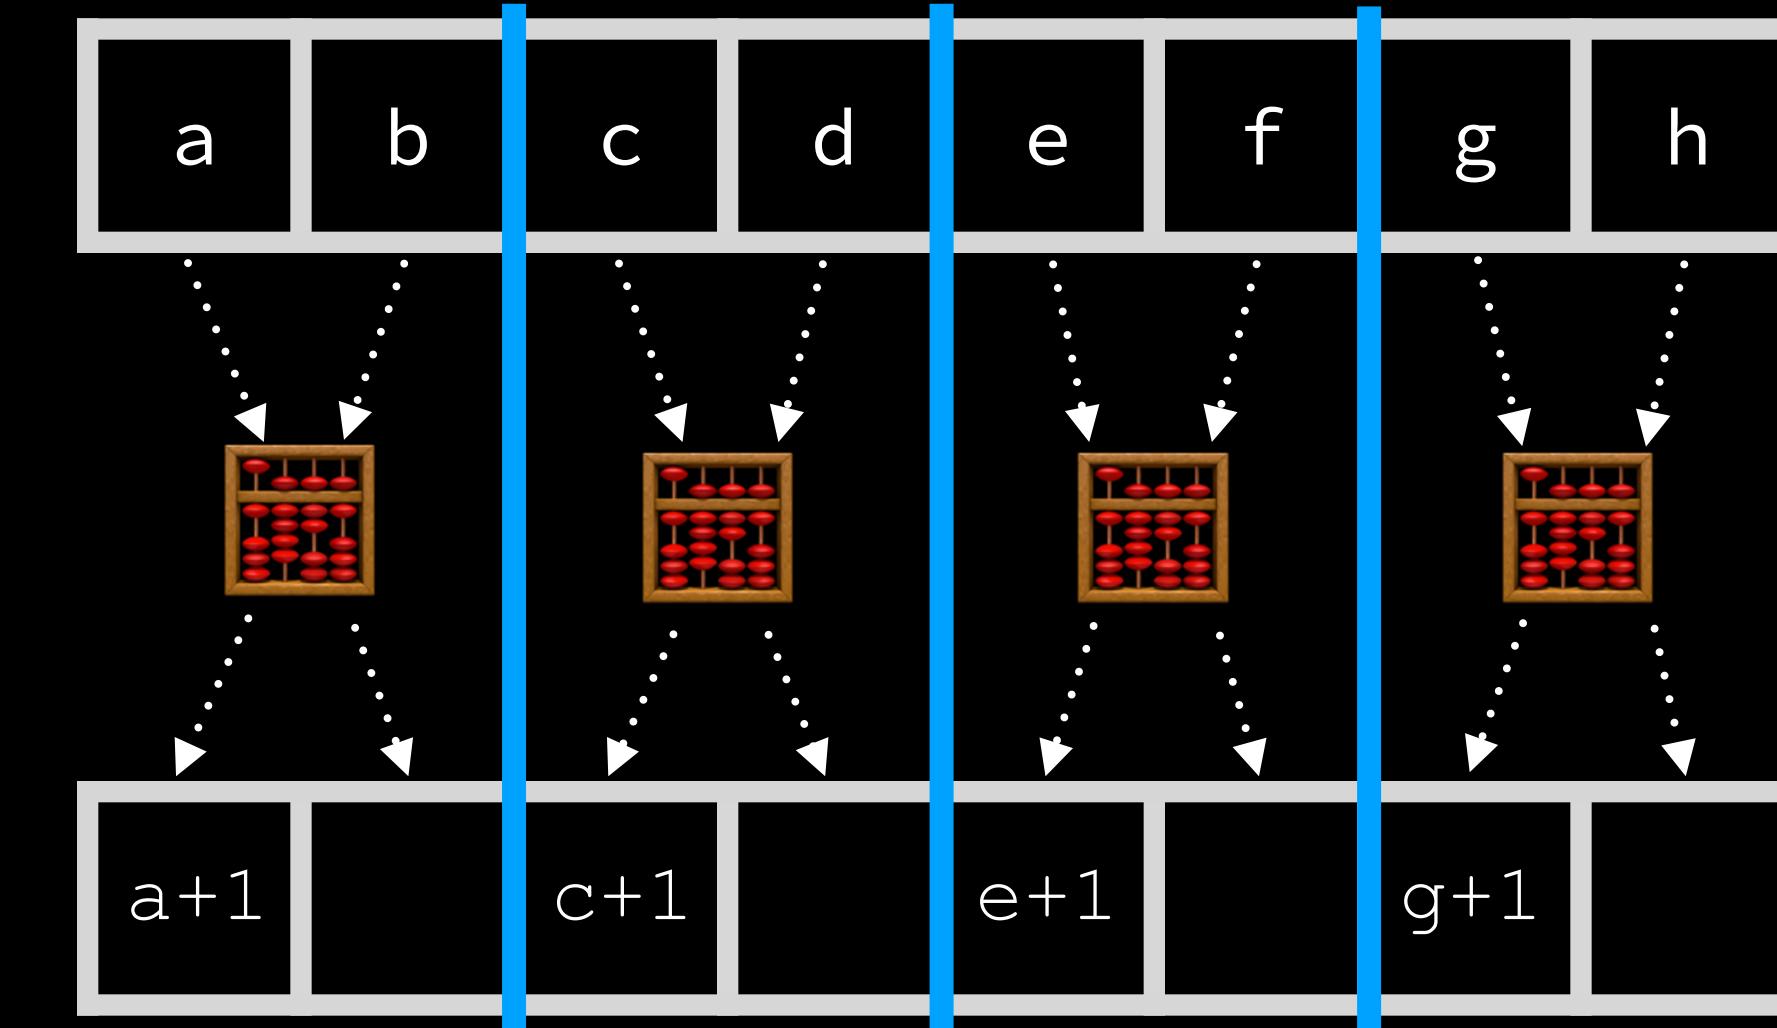
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



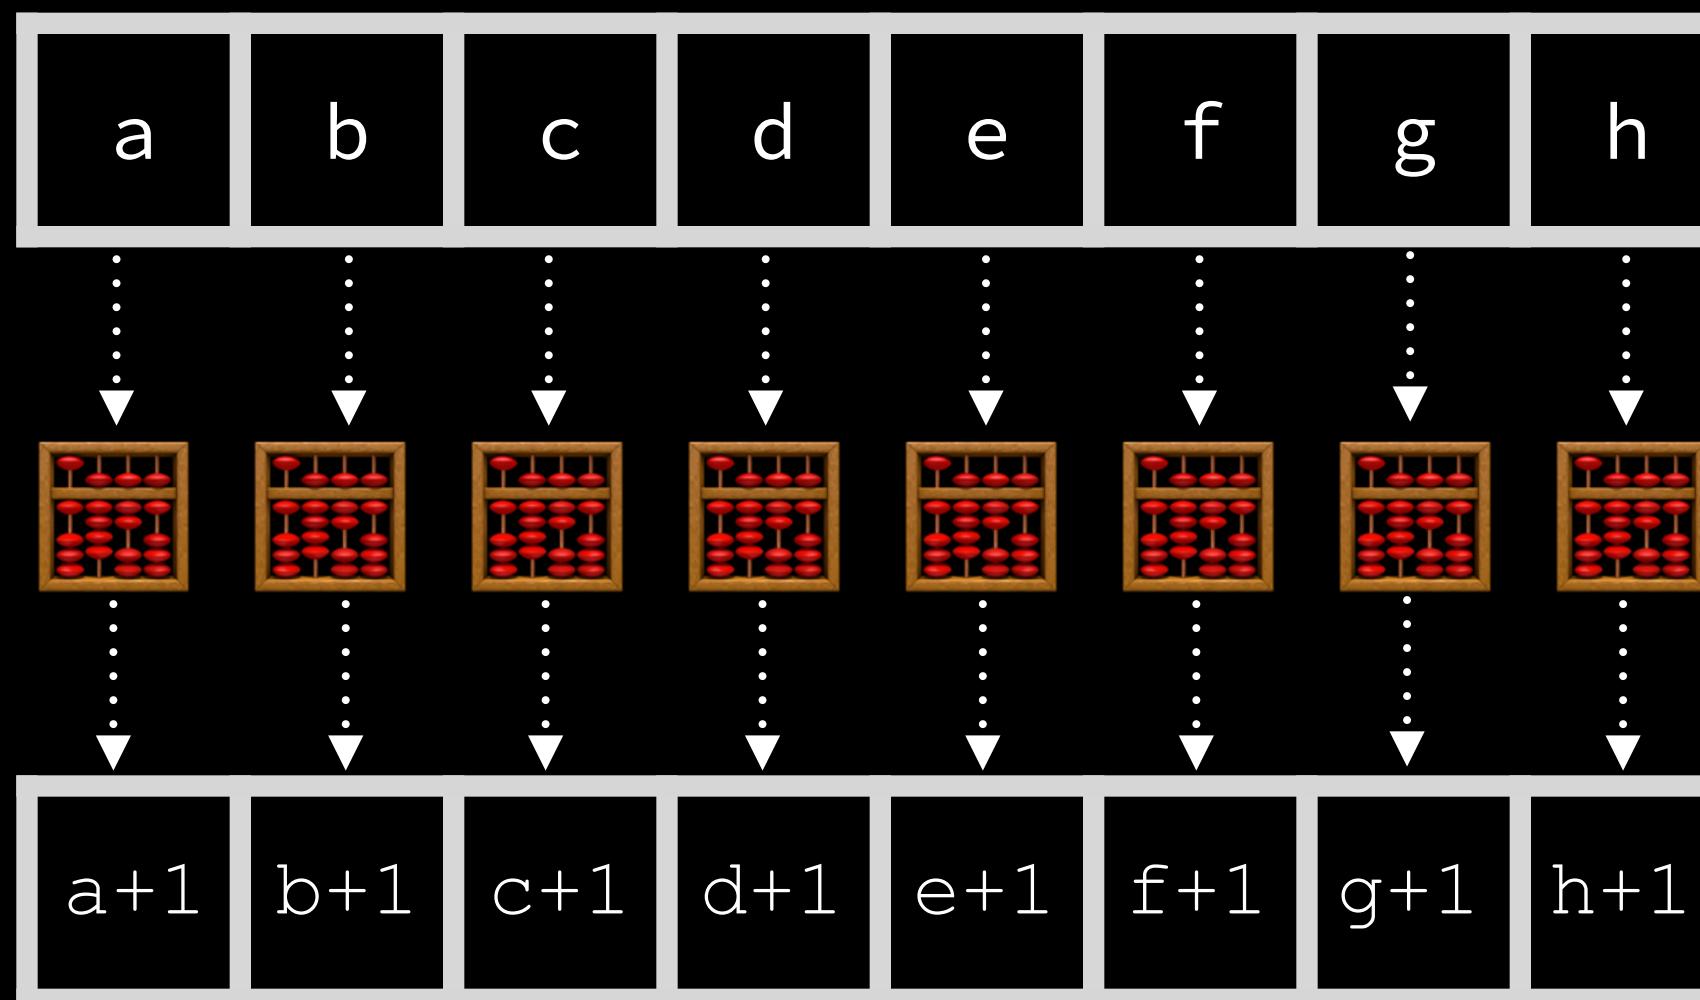
But if we *bank* the arrays,  
we really can parallelize:



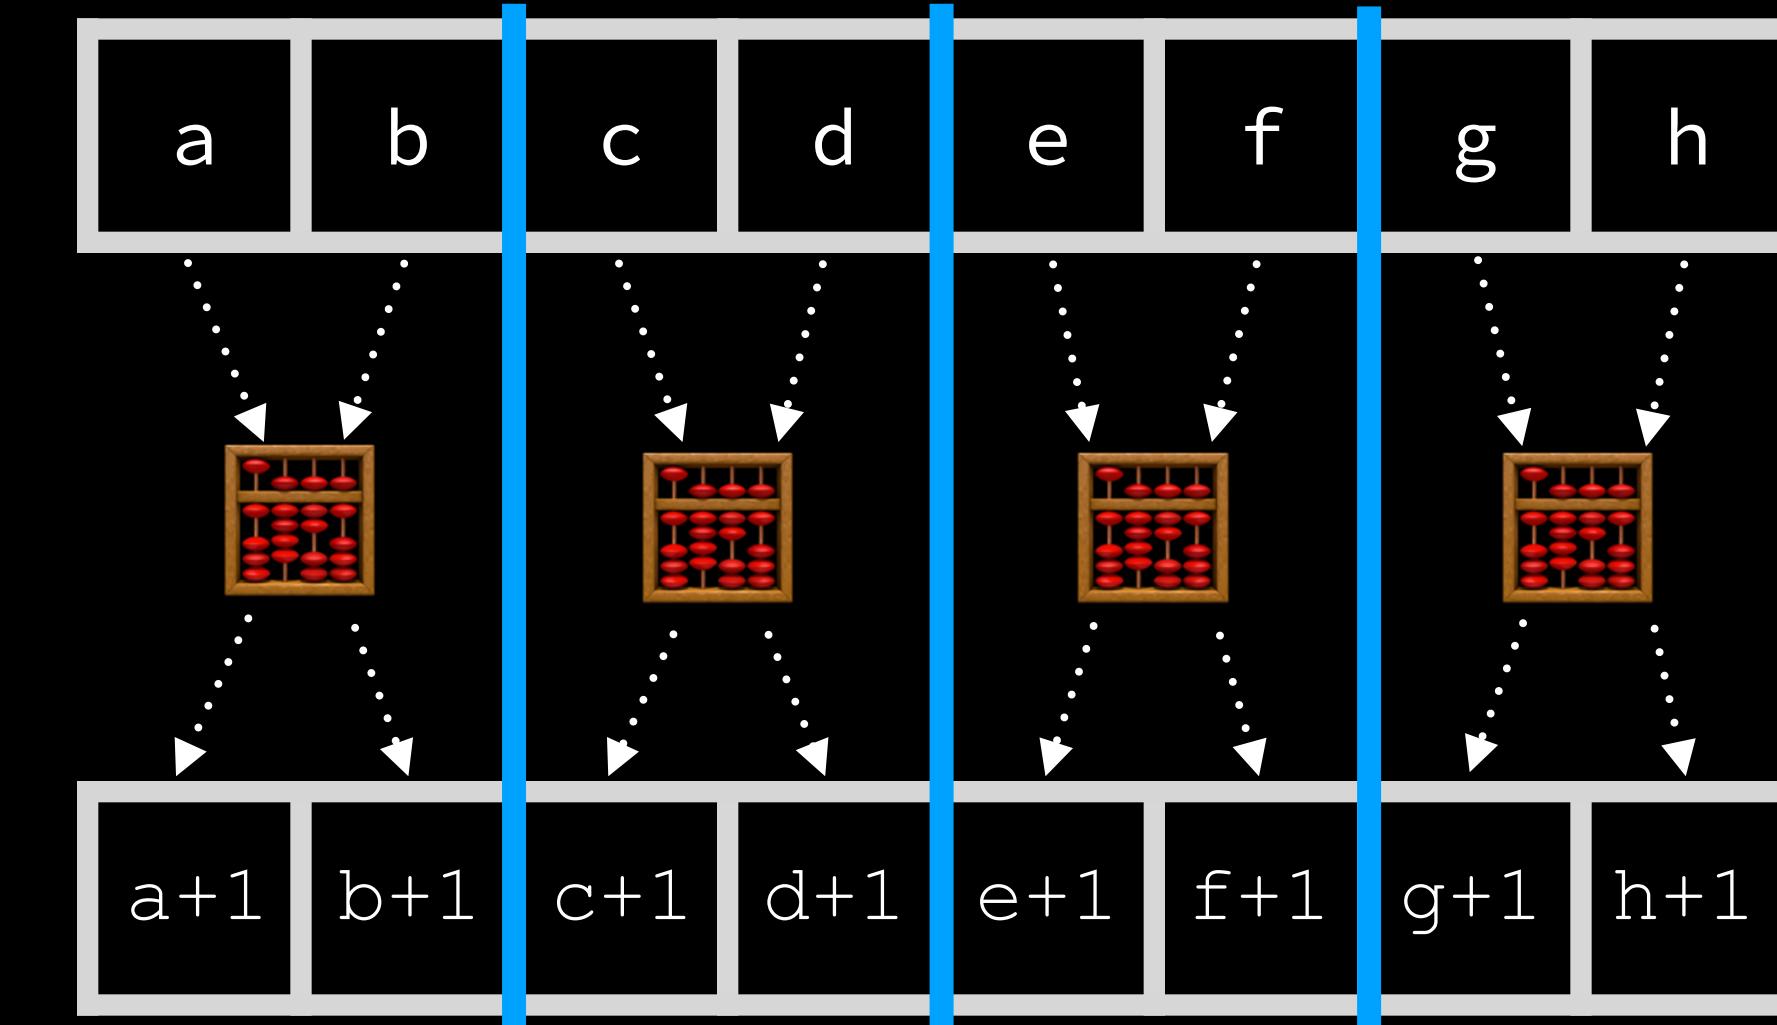
# memory banking

Reads from a memory have a limitation:  
one read/write per tick of the clock.

Say I give you all the compute  
you could ask for:



But if we *bank* the arrays,  
we really can parallelize:



give MrXL a map operation!

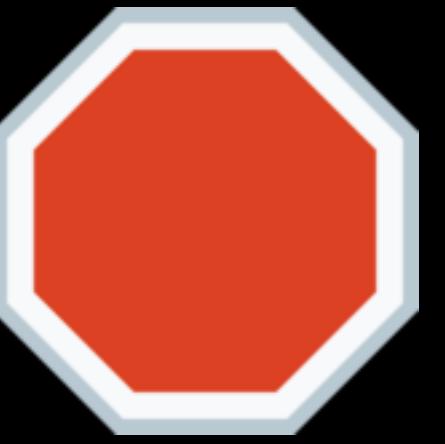
# give MrXL a map operation!

Psst: consider implementing just add first, end-to-end

# give MrXL a map operation!

Psst: consider implementing just add first, end-to-end

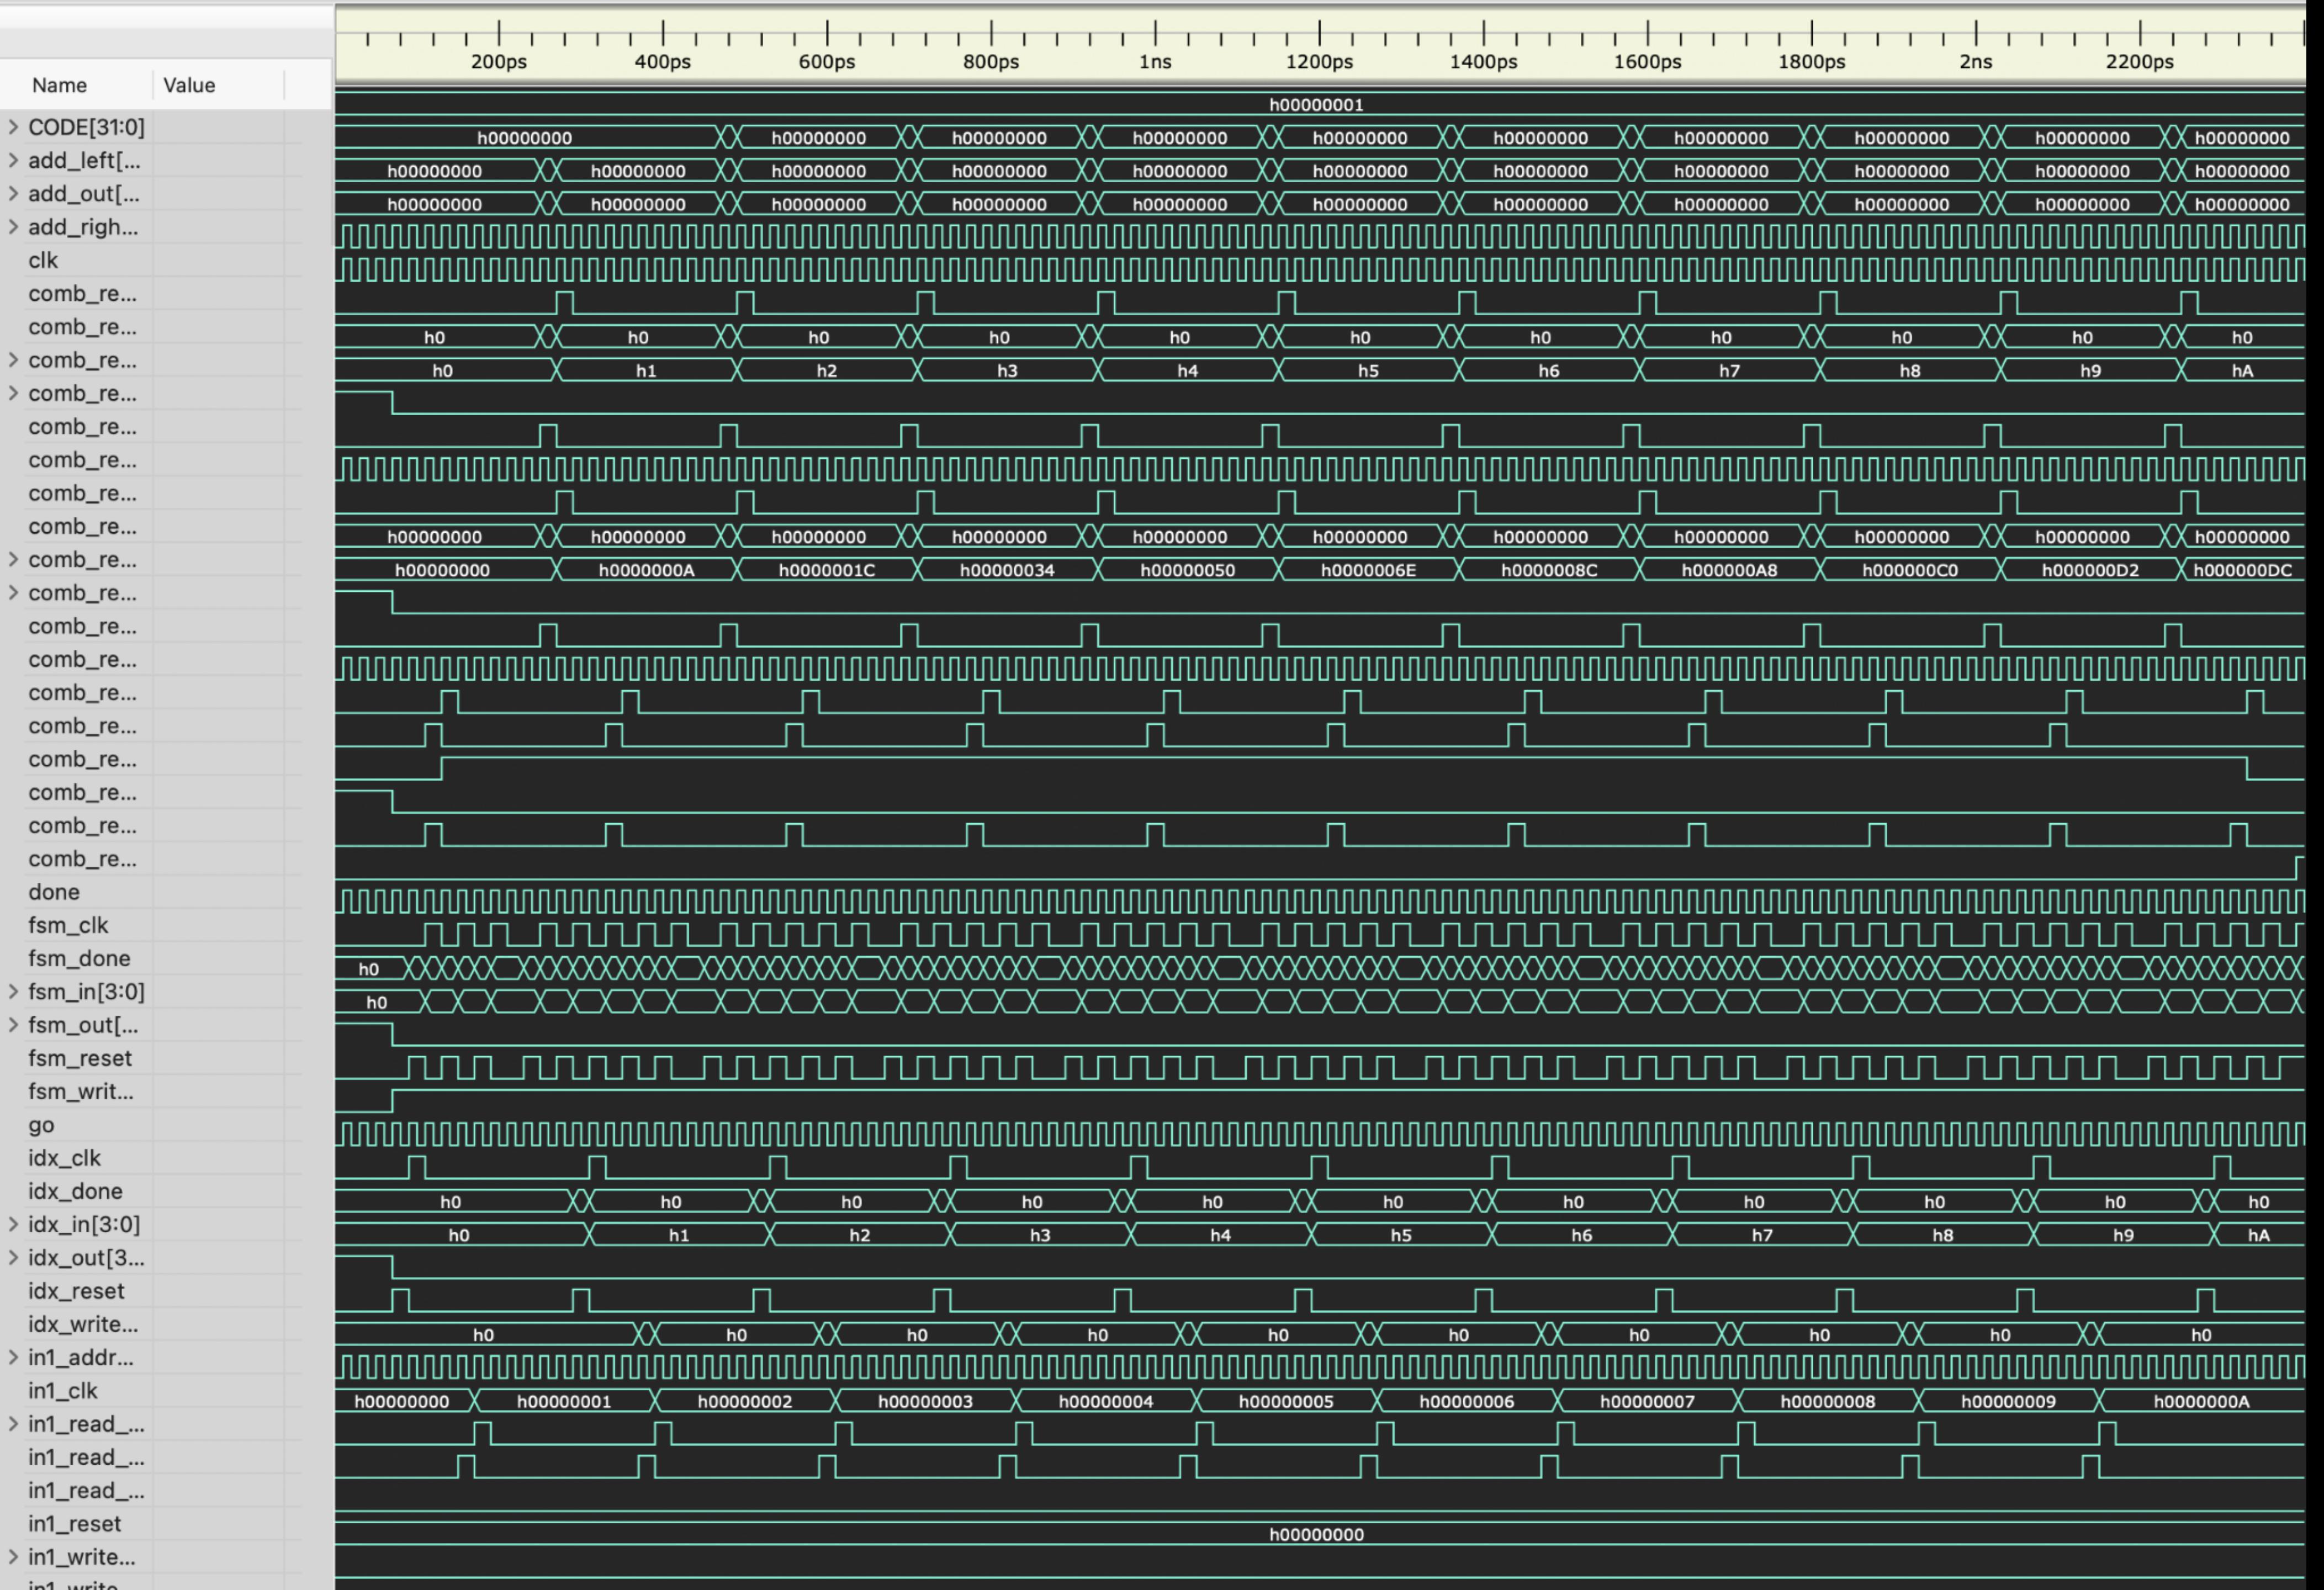
```
fud e --from mrxl test/sos.mrxl \
--to dat --through verilog \
pass this flag → -s mrxl.flags "--my-map" \
to run your version           \
-s mrxl.data test/sos.mrxl.data
```



study the implementation  
that's in place

# Cider





# Cider: Calyx Interpreter and Debugger

Provides a GDB-like debugging experience for Calyx programs

Insight: Use Calyx ***Groups*** as coarse time units

Insight: Use Calyx ***Groups*** as coarse time units

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Insight: Use Calyx ***Groups*** as coarse time units

```
read_mem  
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Insight: Use Calyx ***Groups*** as coarse time units

```
for i in 0..4:  
    read_mem  
    z[i] = x[i] * y[i]  
    do_mul
```

Insight: Use Calyx ***Groups*** as coarse time units

```
for i in 0..4:  
    read_mem  
    z[i] = x[i] * y[i]  
    do_mul  
    write_mem
```

Insight: Use Calyx ***Groups*** as coarse time units

```
for i in 0..4:  
    read_mem  
    z[i] = x[i] * y[i]  
    do_mul  
    write_mem  
    upd_idx
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg  
idx_reg = 1
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg  
idx_reg = 1  
idx_reg = 2
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg  
idx_reg = 1  
idx_reg = 2  
idx_reg = 3
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg  
idx_reg = 1  
idx_reg = 2  
idx_reg = 3  
WARN - Integer overflow, source: idx_adder
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg  
idx_reg = 1  
idx_reg = 2  
idx_reg = 3  
WARN - Integer overflow, source: idx_adder  
idx_reg = 0
```

```
for i in 0..4:  
    z[i] = x[i] * y[i]
```

Observed behavior: does not terminate

```
> watch after upd_idx with print-state \u idx_reg  
idx_reg = 1  
idx_reg = 2  
idx_reg = 3  
WARN - Integer overflow, source: idx_adder  
idx_reg = 0  
idx_reg = 1
```

# the bug

```
cells {  
    idx_reg = reg(2);  
    idx_adder = reg(2);  
}
```

```
while idx_reg <= 3 {  
    read_mem;  
    do_mul;  
    write_mem;  
    upd_idx;  
}
```

# the bug

```
cells {  
    idx_reg = reg(2);  
    idx_adder = reg(2);  
}
```



```
while idx_reg <= 3 {  
    read_mem;  
    do_mul;  
    write_mem;  
    upd_idx;  
}
```

# the bug

```
cells {  
    idx_reg = reg(2);  
    idx_adder = reg(2);  
}
```



```
cells {  
    idx_reg = reg(3);  
    idx_adder = reg(3);  
}
```

```
while idx_reg <= 3 {  
    read_mem;  
    do_mul;  
    write_mem;  
    upd_idx;  
}
```

# using Cider

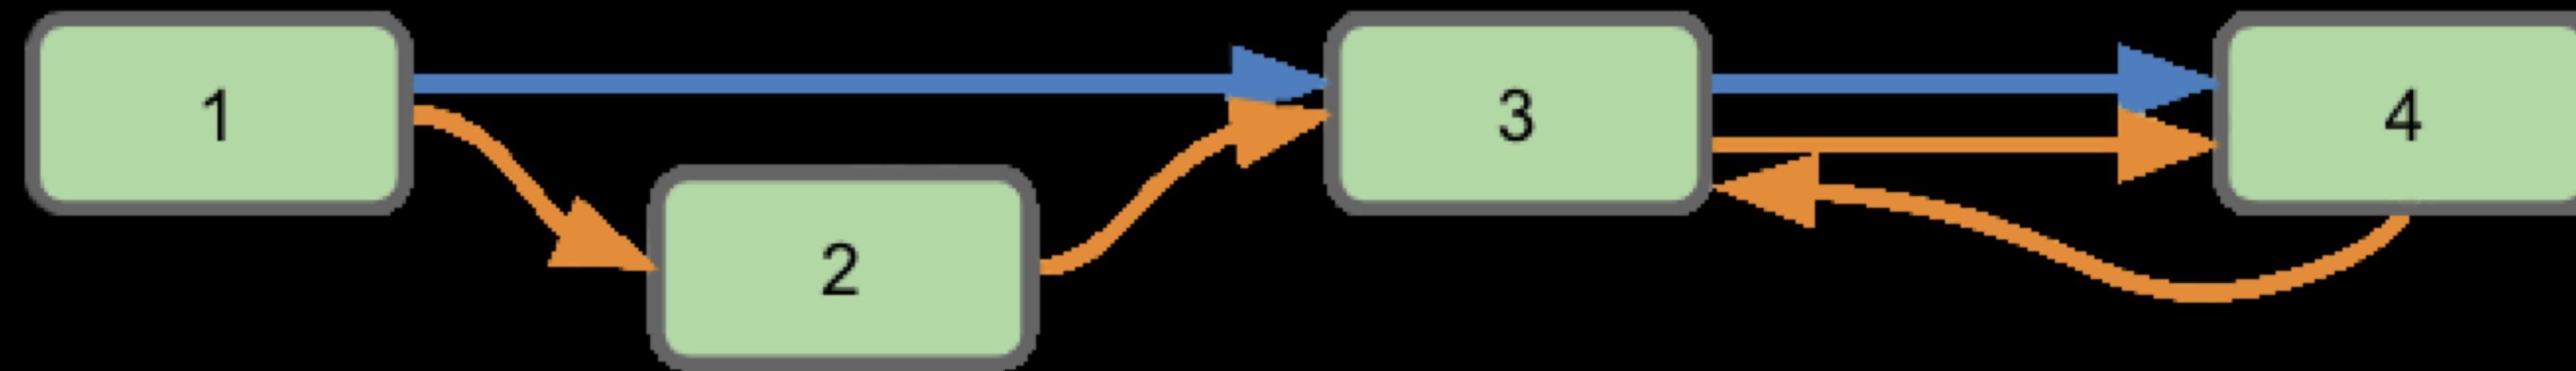
```
fud e --from mrxl test/sos.mrxl --to interpreter-out
```

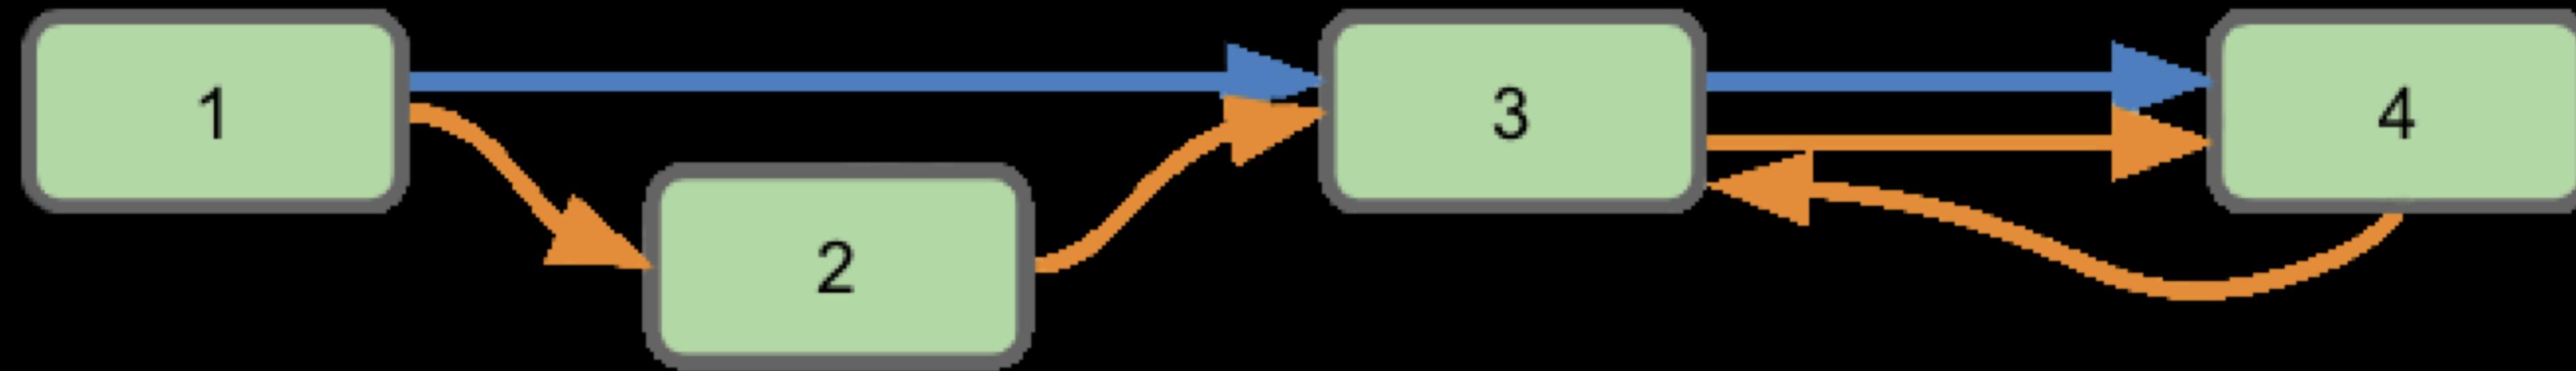
```
fud e --from mrxl test/sos.mrxl --to debugger
```



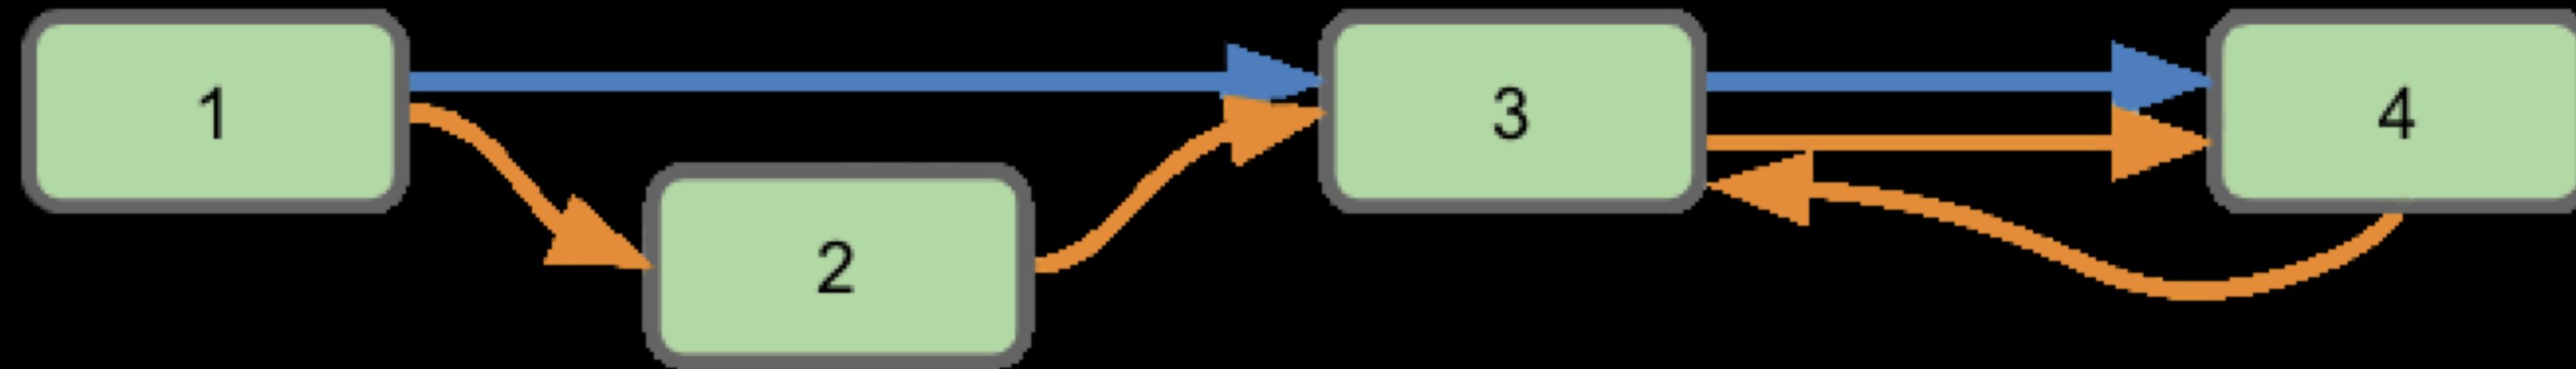
poen

The logo features the word "poen" in a lowercase, sans-serif font. The letters are primarily blue, except for the letter "e" which is orange. A stylized DNA double helix is positioned behind the letters, with its blue and orange segments interwoven through the text. The DNA structure has dark blue and orange horizontal bands.





```
out_graph g;
parset depth[int, g];
for segment in graph.segments {
    emit segment.steps.size() to depths;
}
```



```
out_graph g;
parset depth[int, g];
for segment in graph.segments {
    emit segment.steps.size() to depths;
}
```

```
out_graph g;
parset depth[int, g];
for segment in graph.segments {
    emit segment.steps.size() to depths;
}
```

```
exine depth -n 2
```

=

```
output depths : int[4]
depths := map 2 (s <- graph.segments) {s.steps.size()}
```

# takeaways

Domain experts with minimal hardware knowledge can make use of Calyx

Calyx can be a backend for complex DSLs

The skills you've gained generating map are broadly applicable to all sorts of language features

have a break,  
have a Kit Kat

# MrXL: extensions

Three options:

1. Implement the reduce operation.
2. Allow the same memory to be banked repeatedly.
3. Office hours!

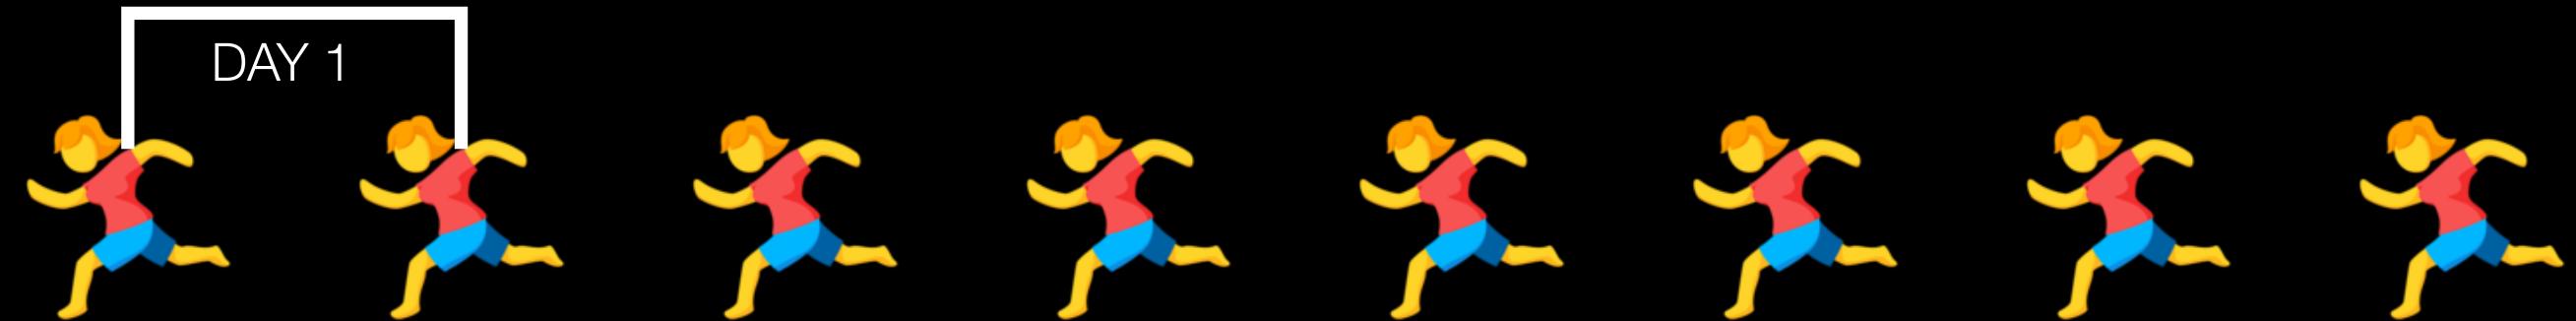
# reduction trees

How best to run a tennis tournament?



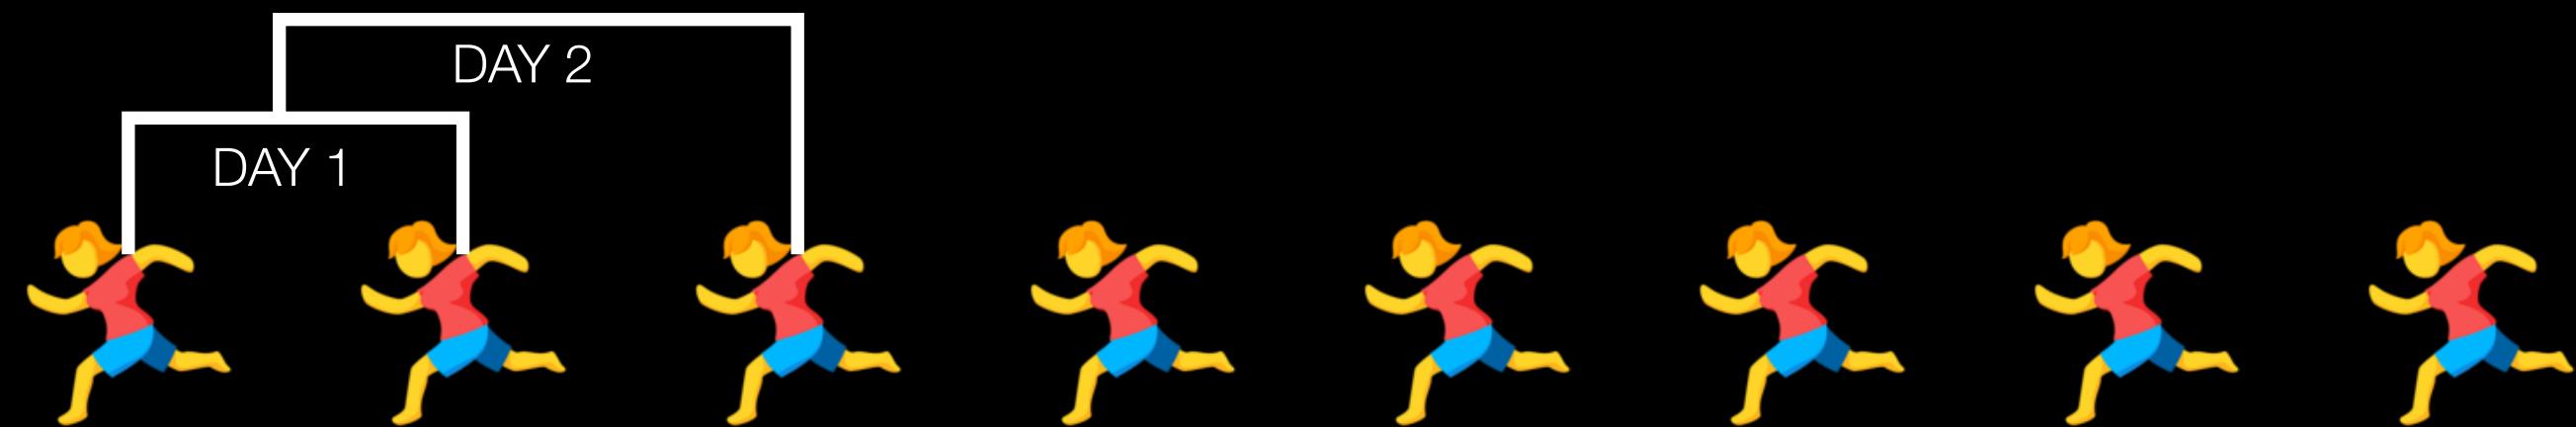
# reduction trees

How best to run a tennis tournament?



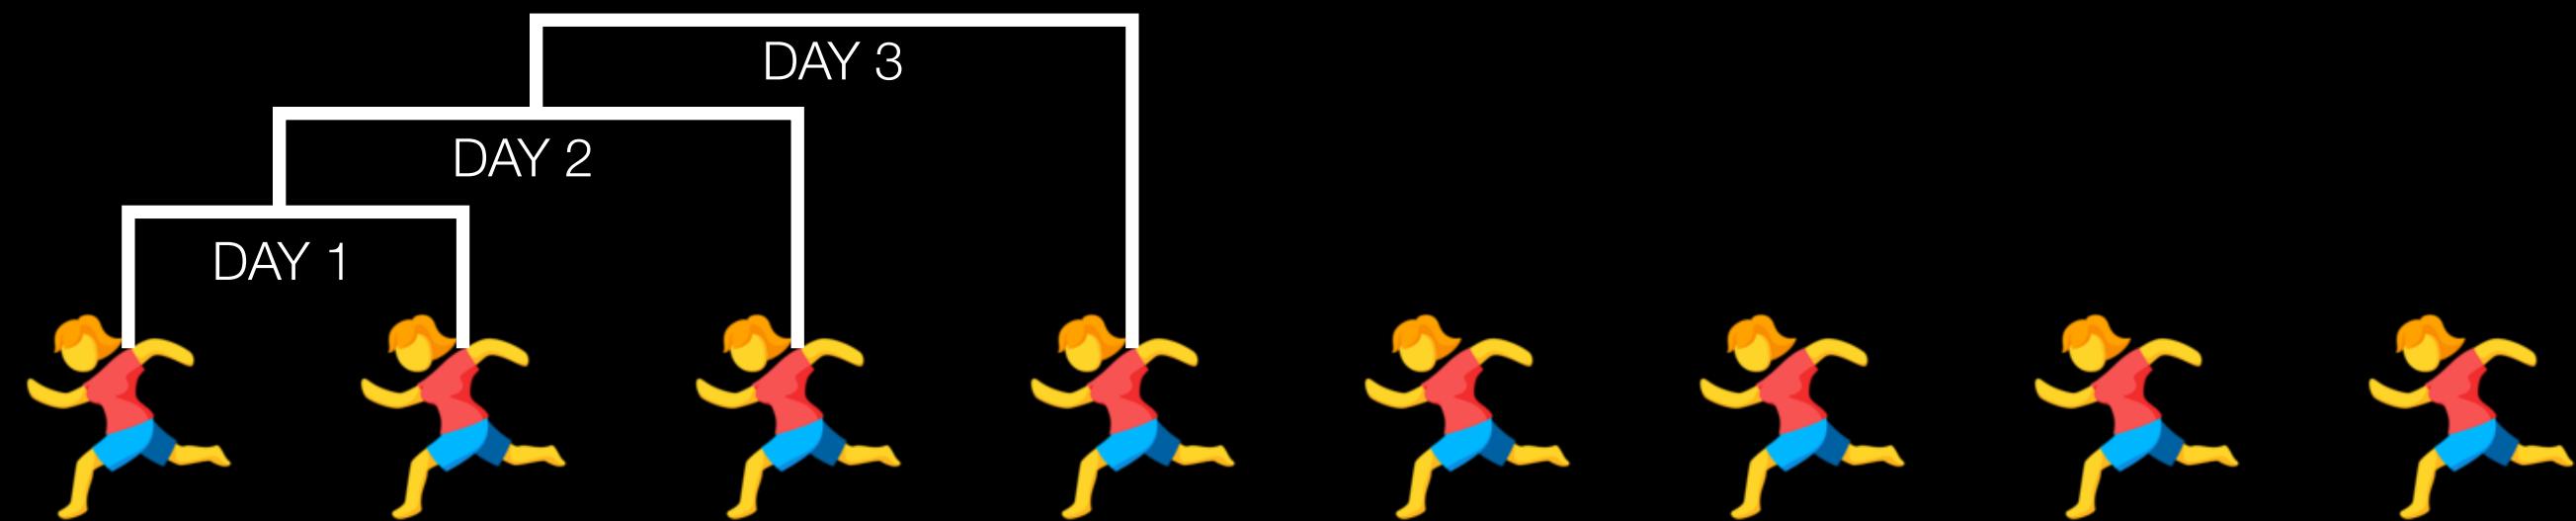
# reduction trees

How best to run a tennis tournament?



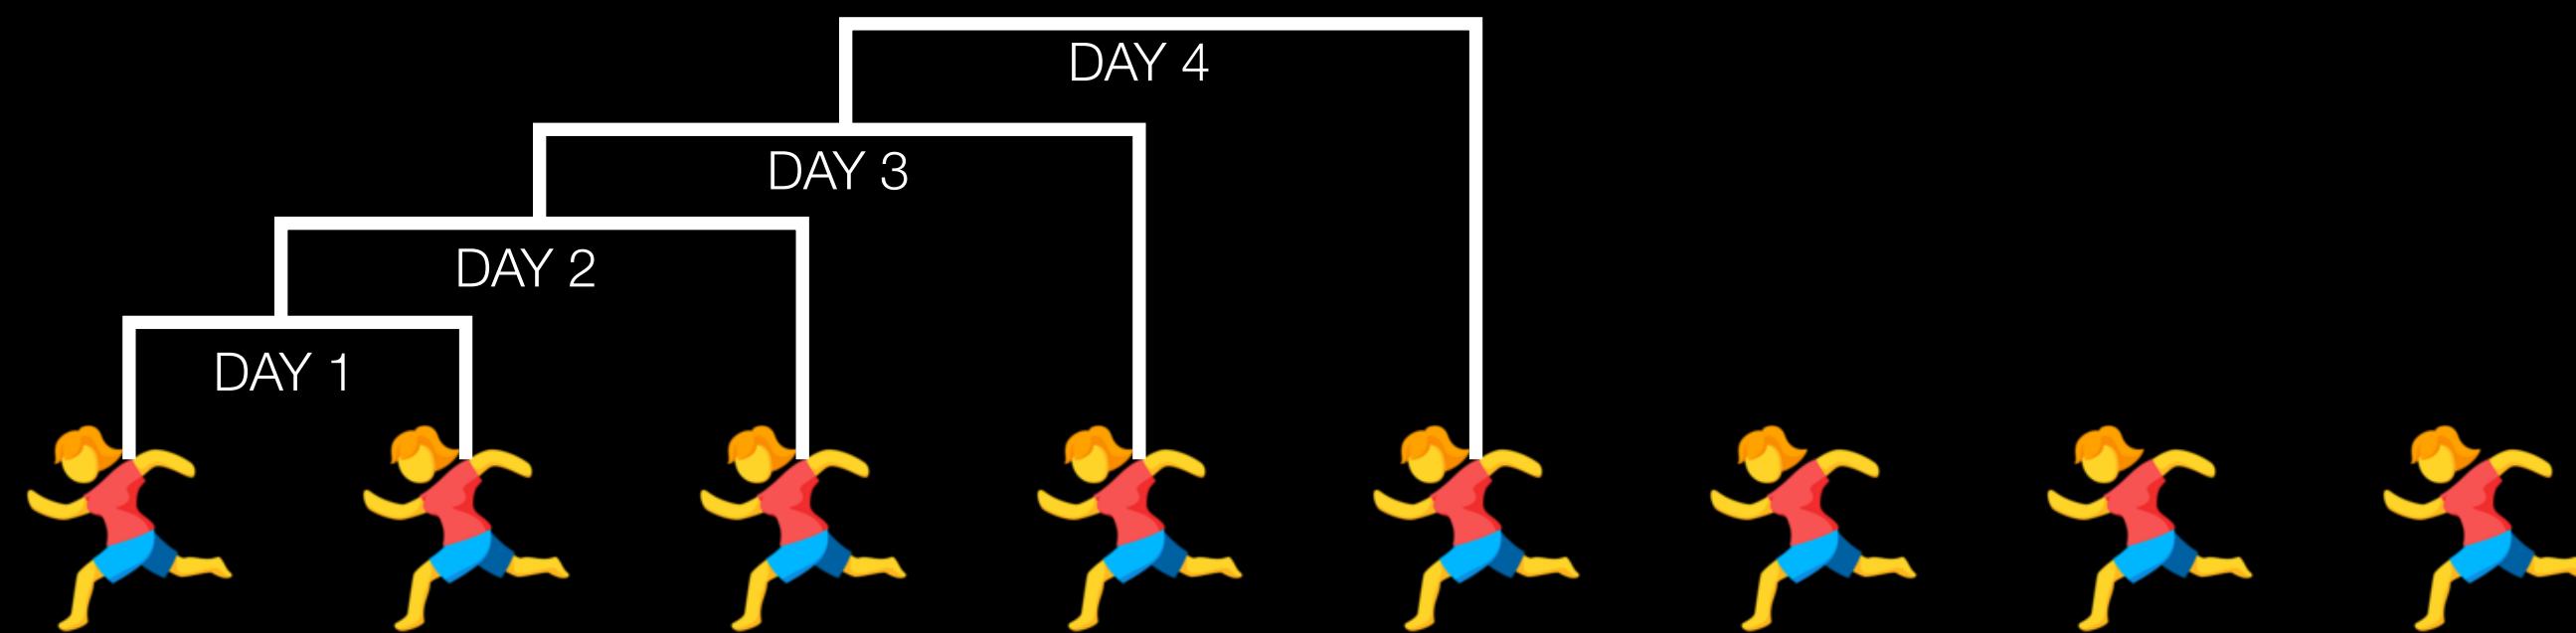
# reduction trees

How best to run a tennis tournament?



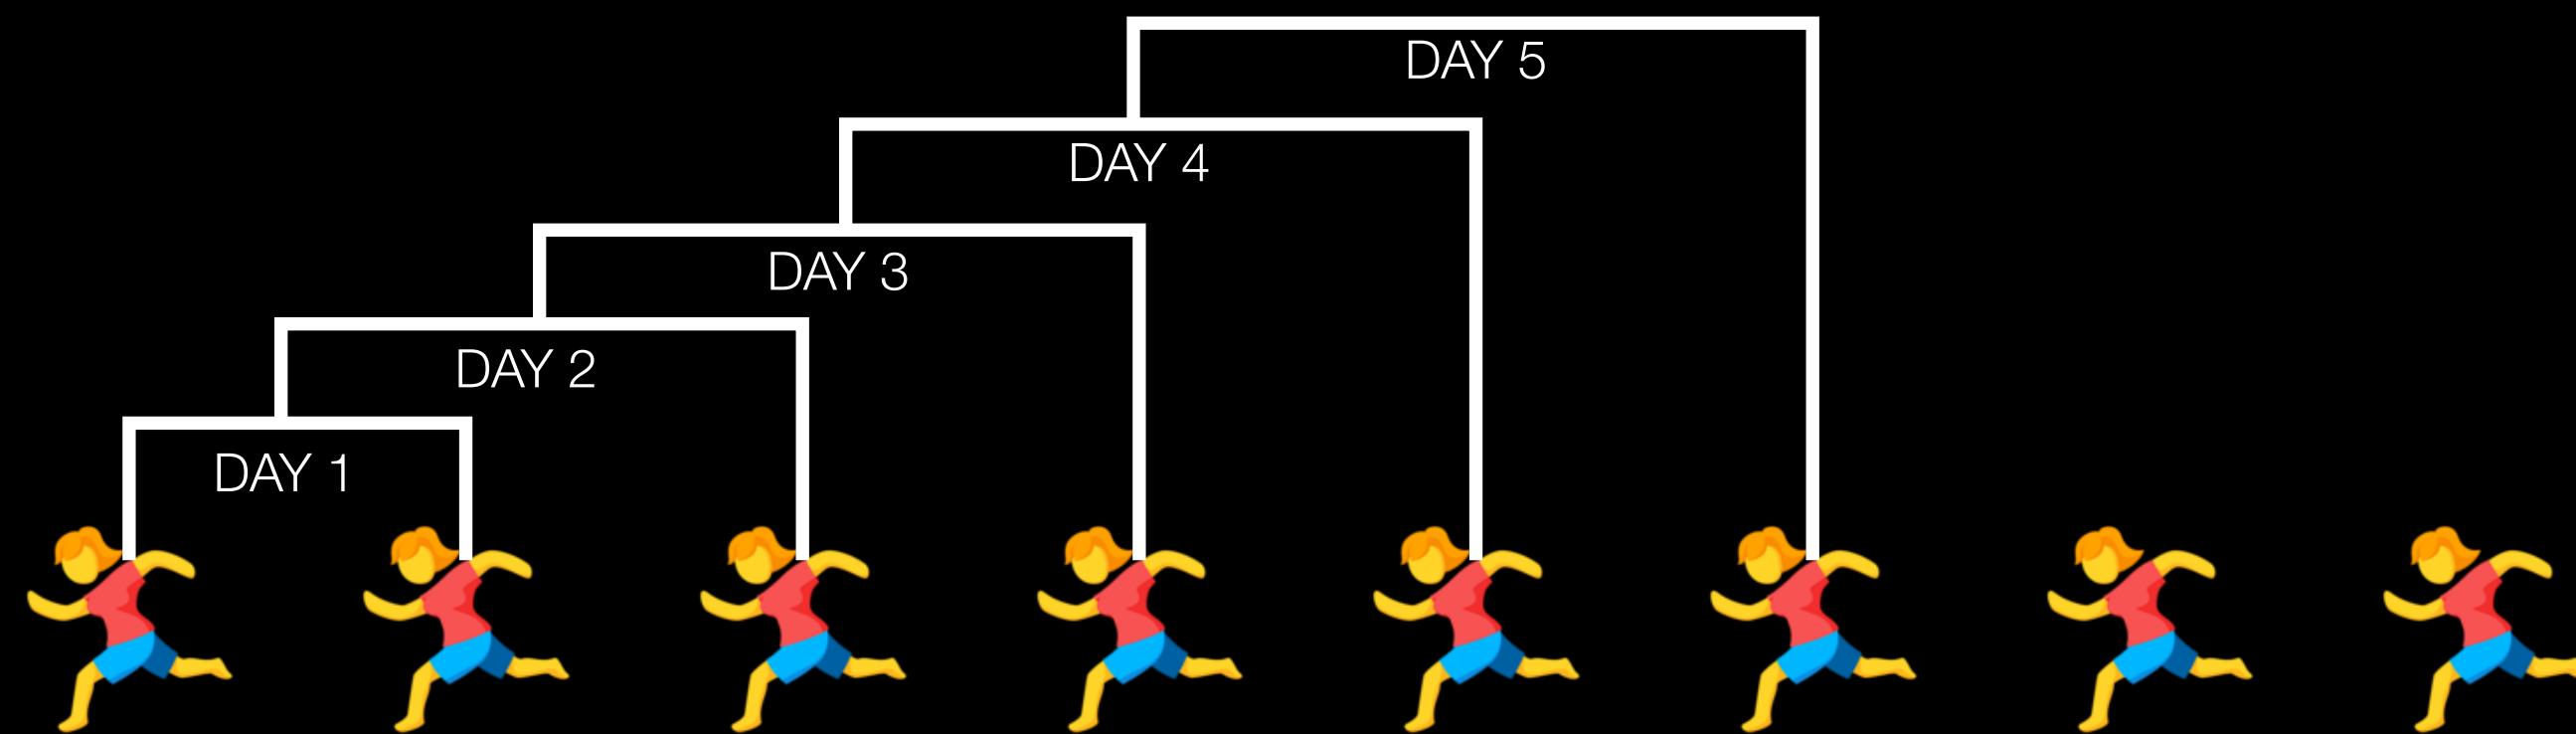
# reduction trees

How best to run a tennis tournament?



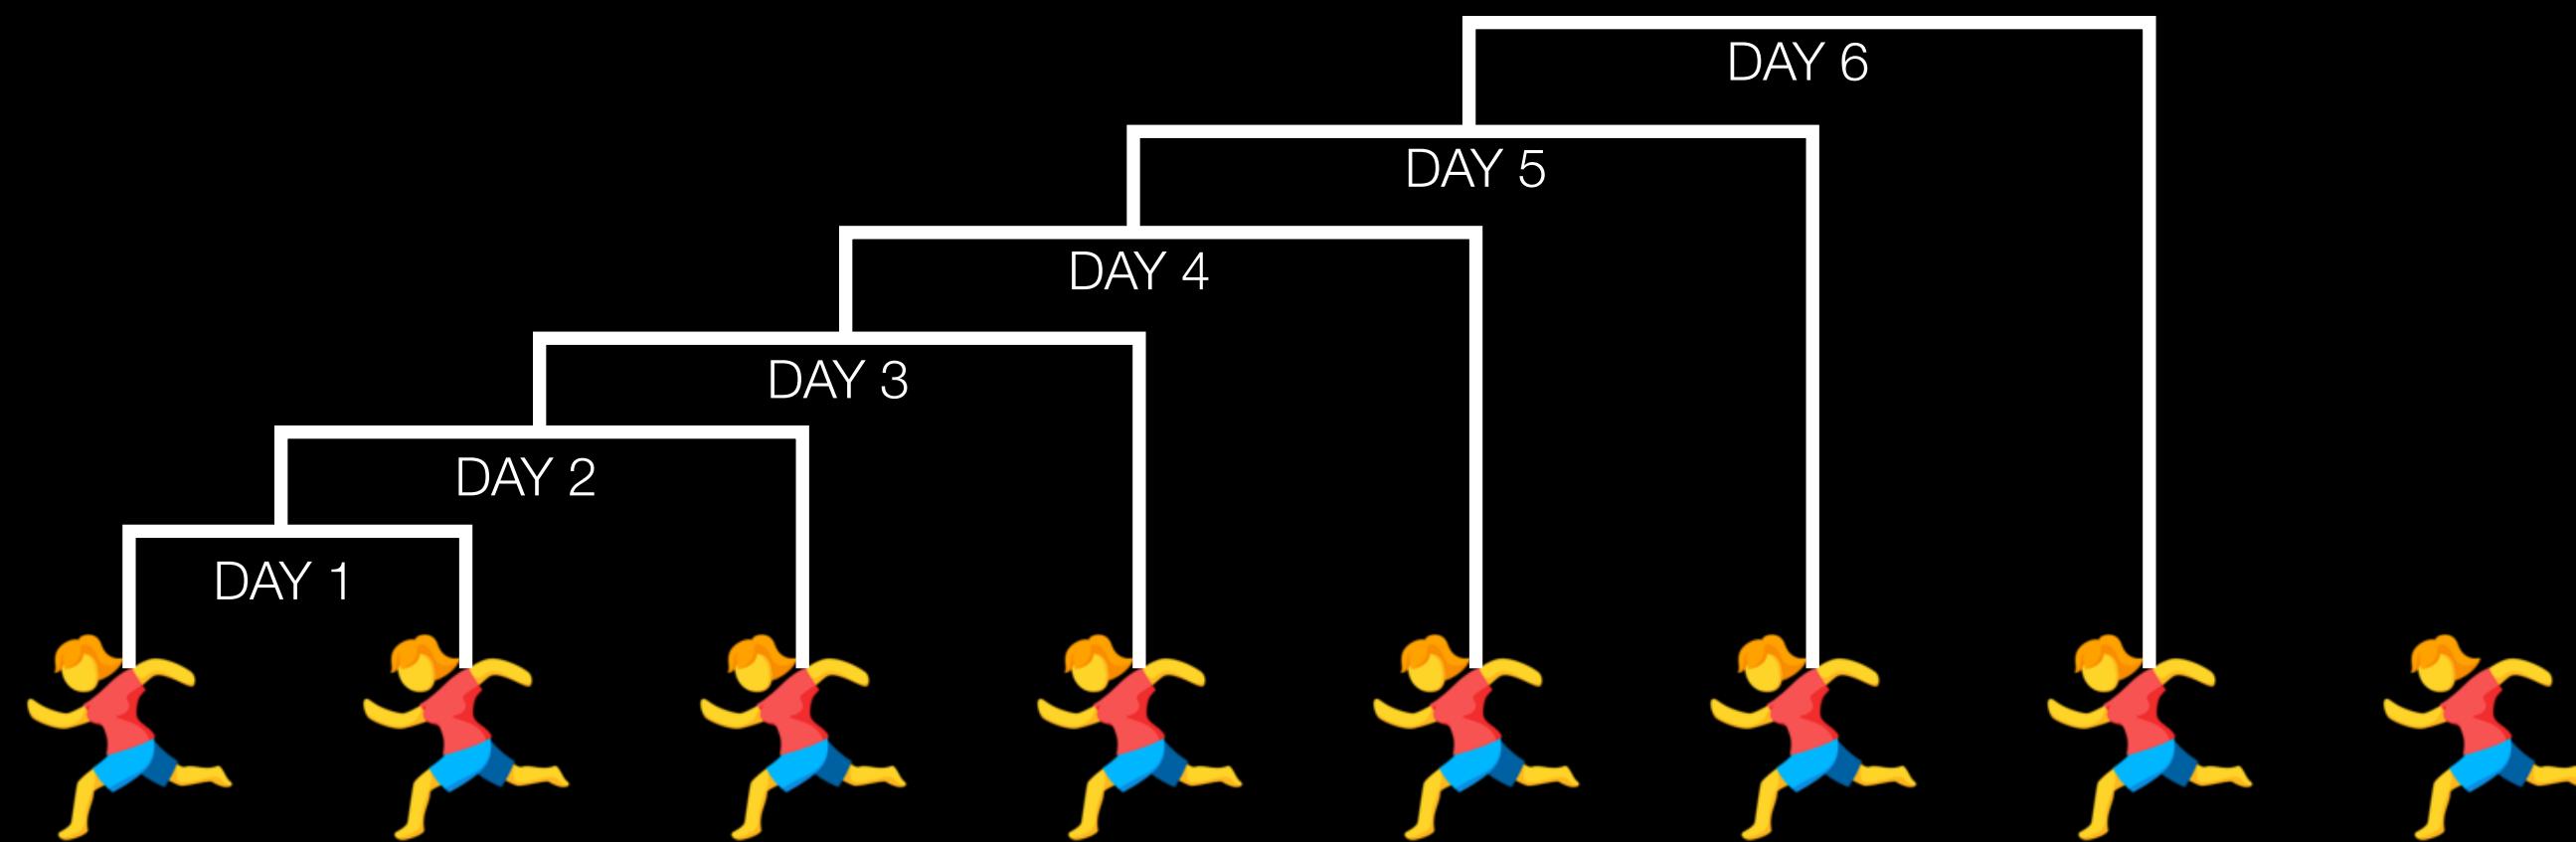
# reduction trees

How best to run a tennis tournament?



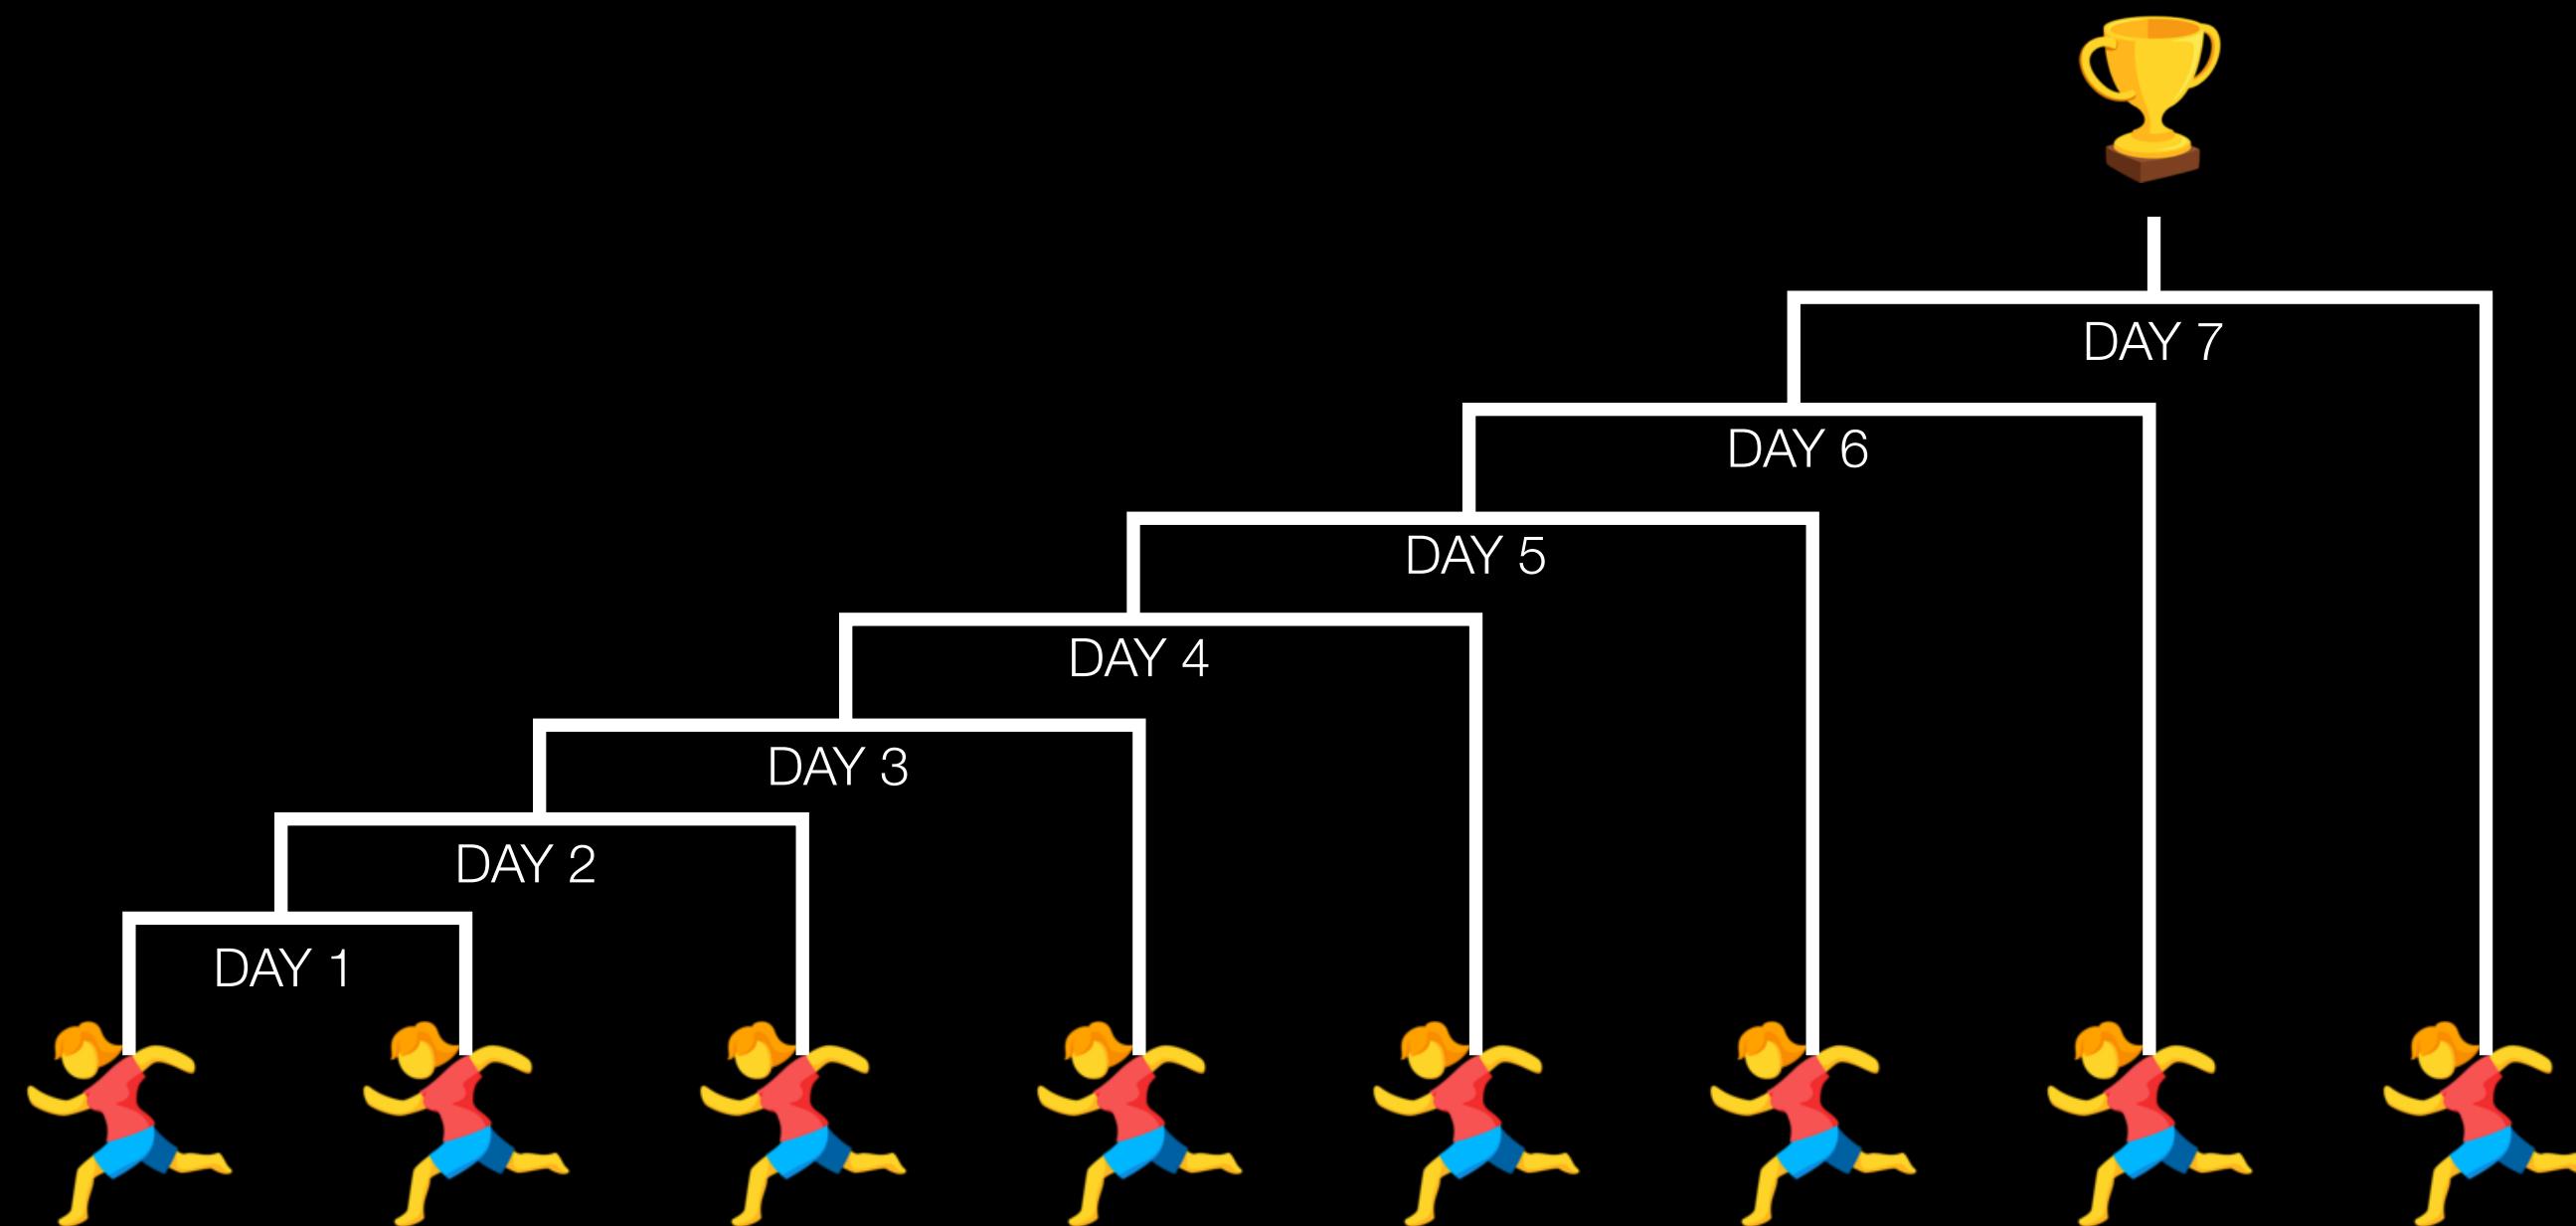
# reduction trees

How best to run a tennis tournament?



# reduction trees

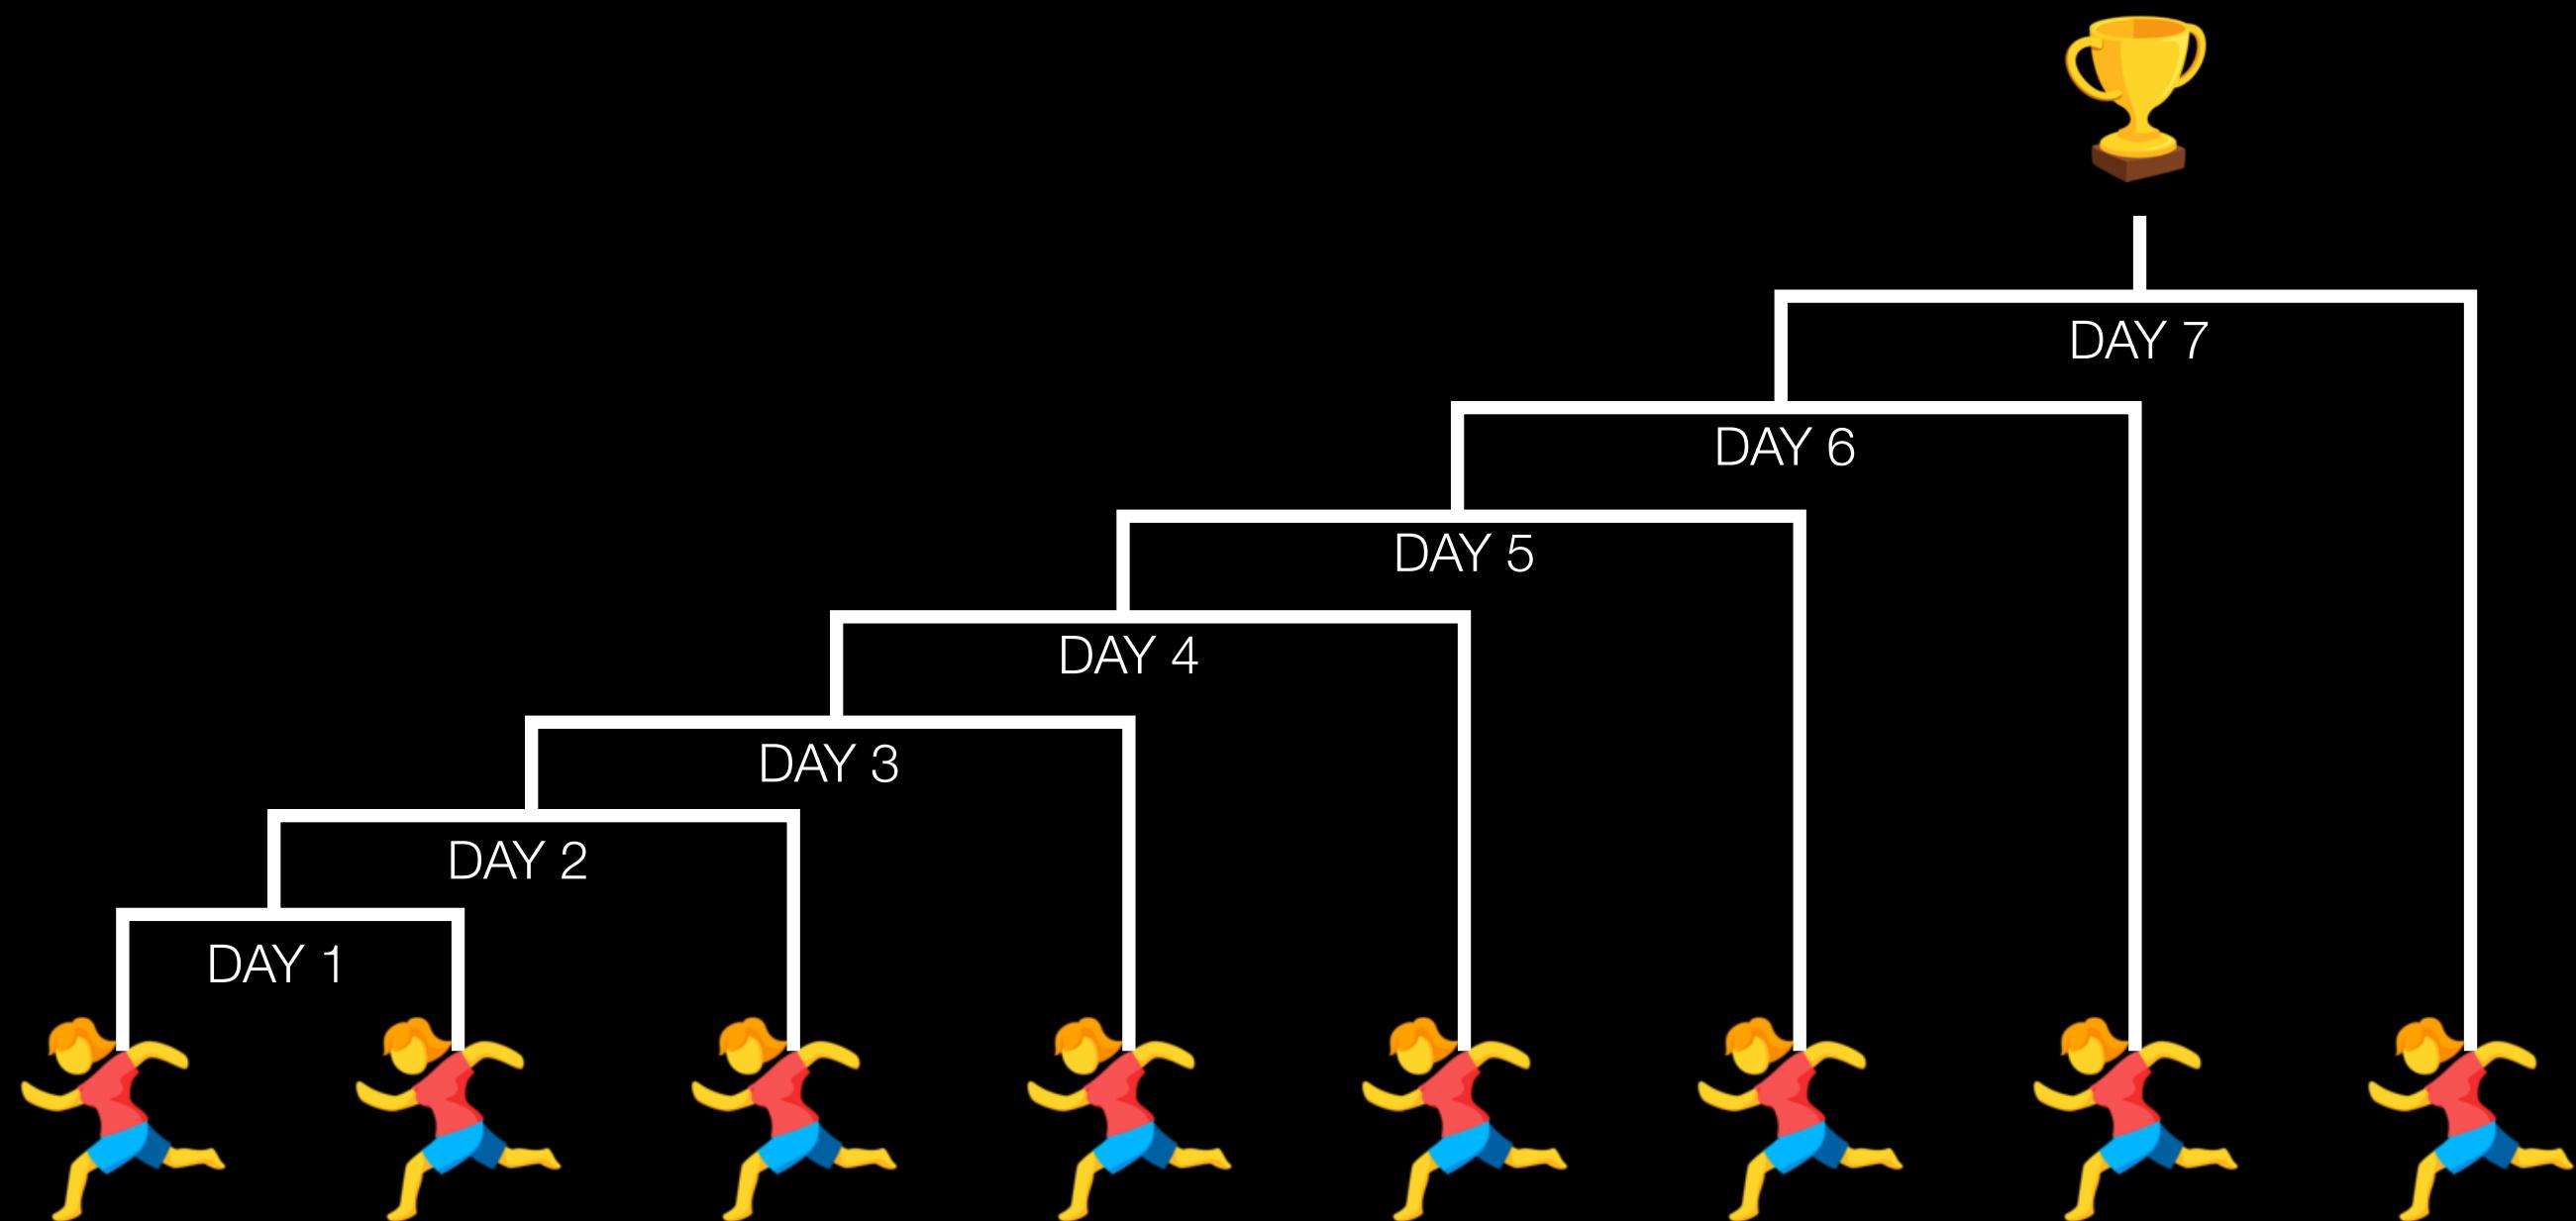
How best to run a tennis tournament?



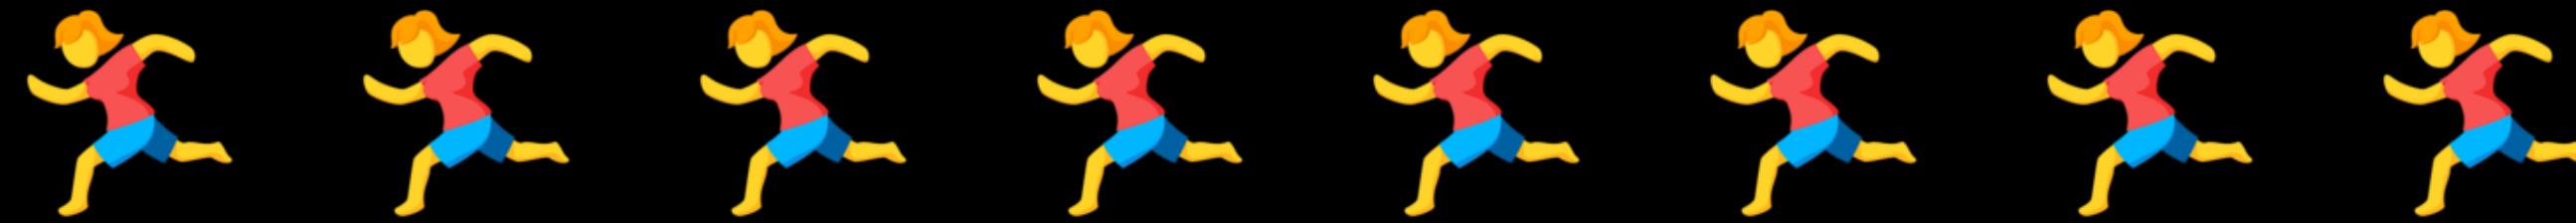
# reduction trees

How best to run a tennis tournament?

This clearly has problems...



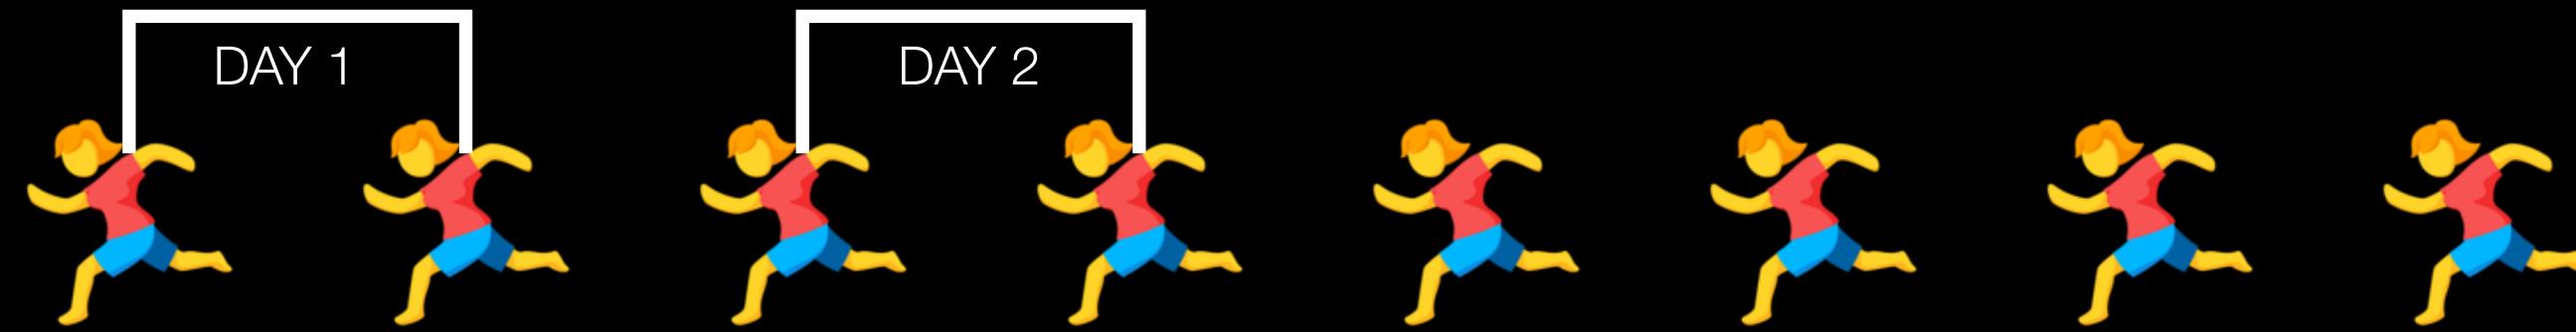
# reduction trees



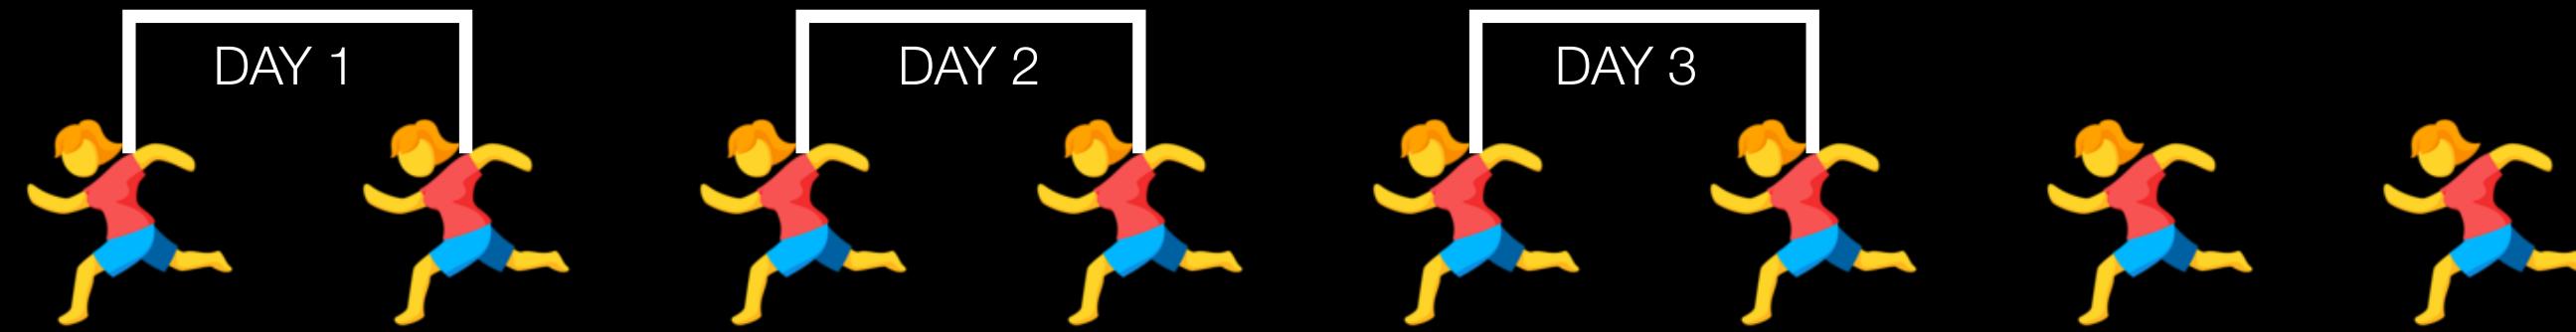
# reduction trees



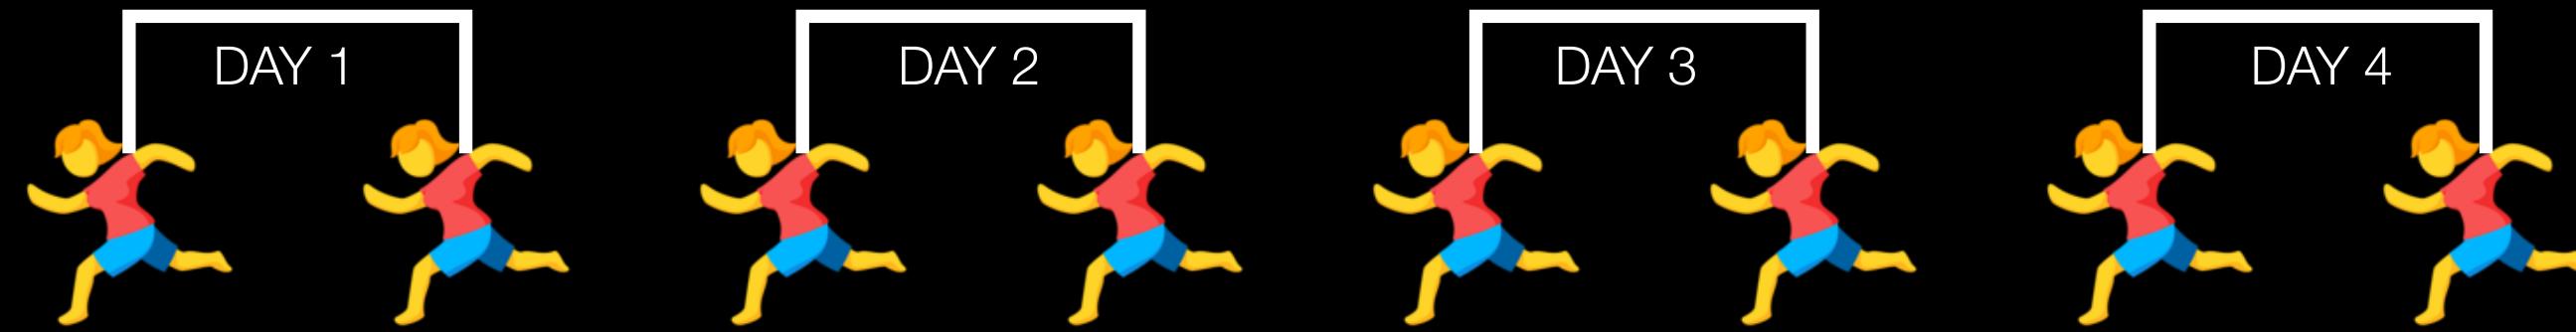
# reduction trees



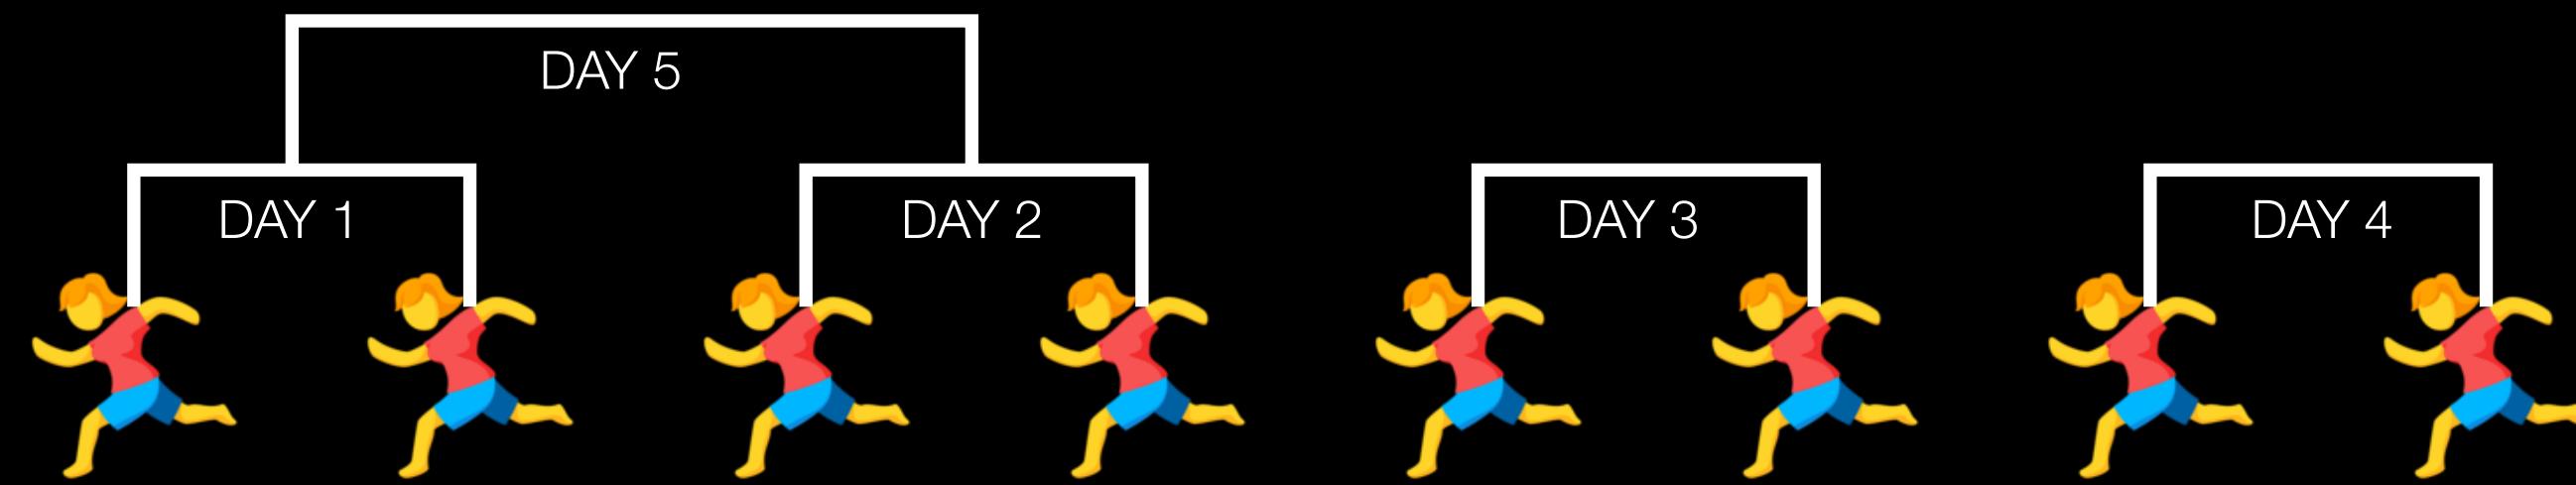
# reduction trees



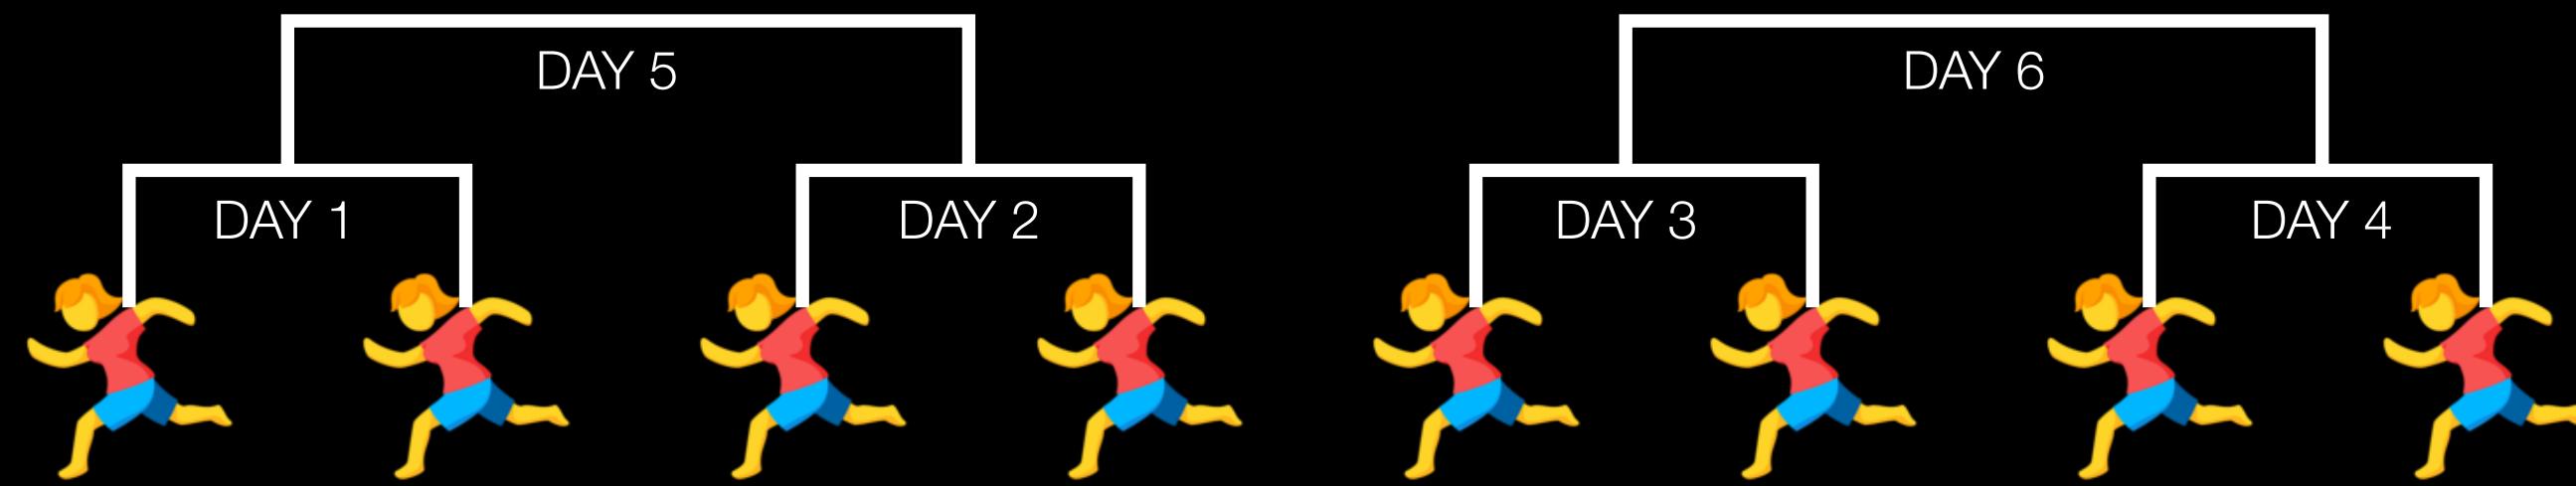
# reduction trees



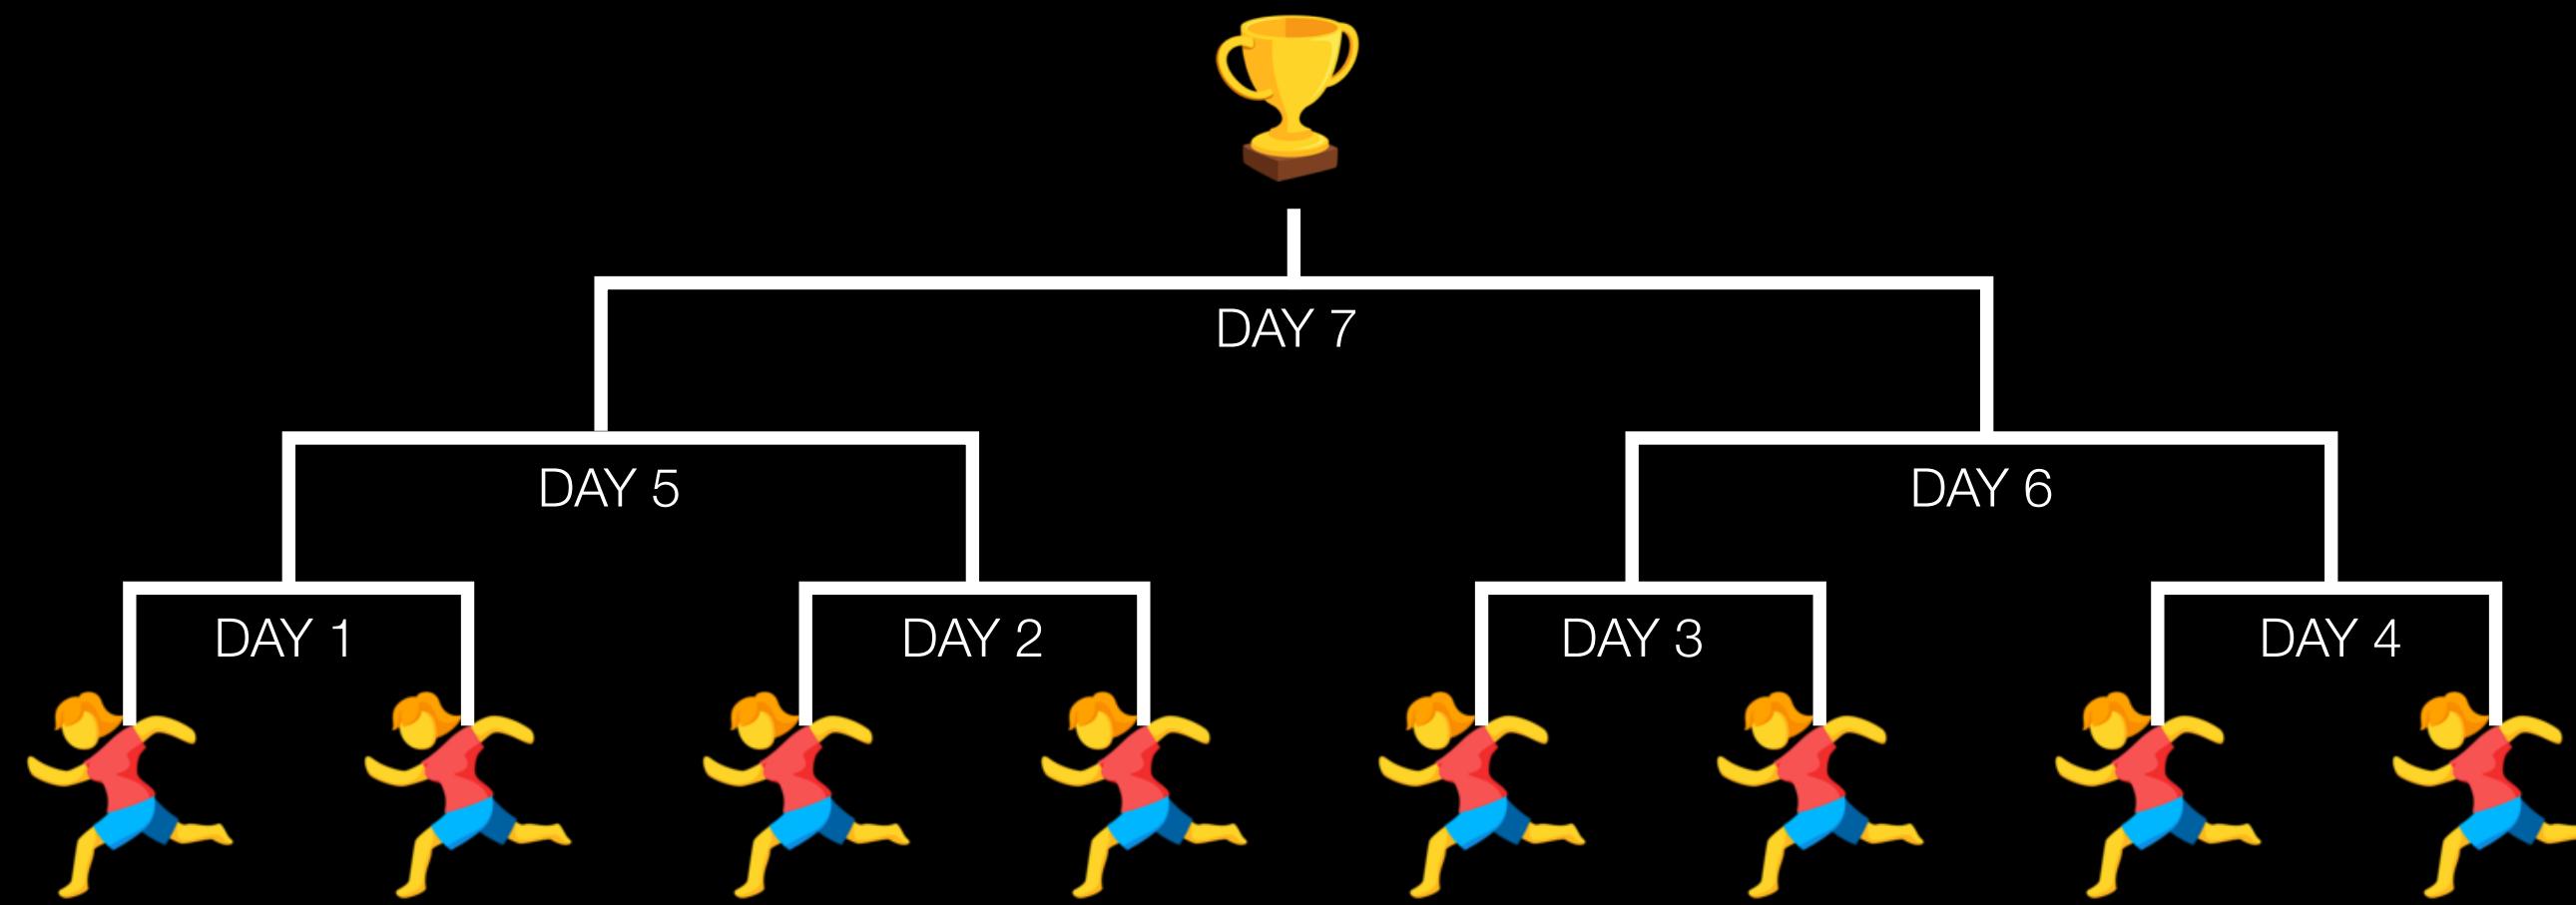
# reduction trees



# reduction trees

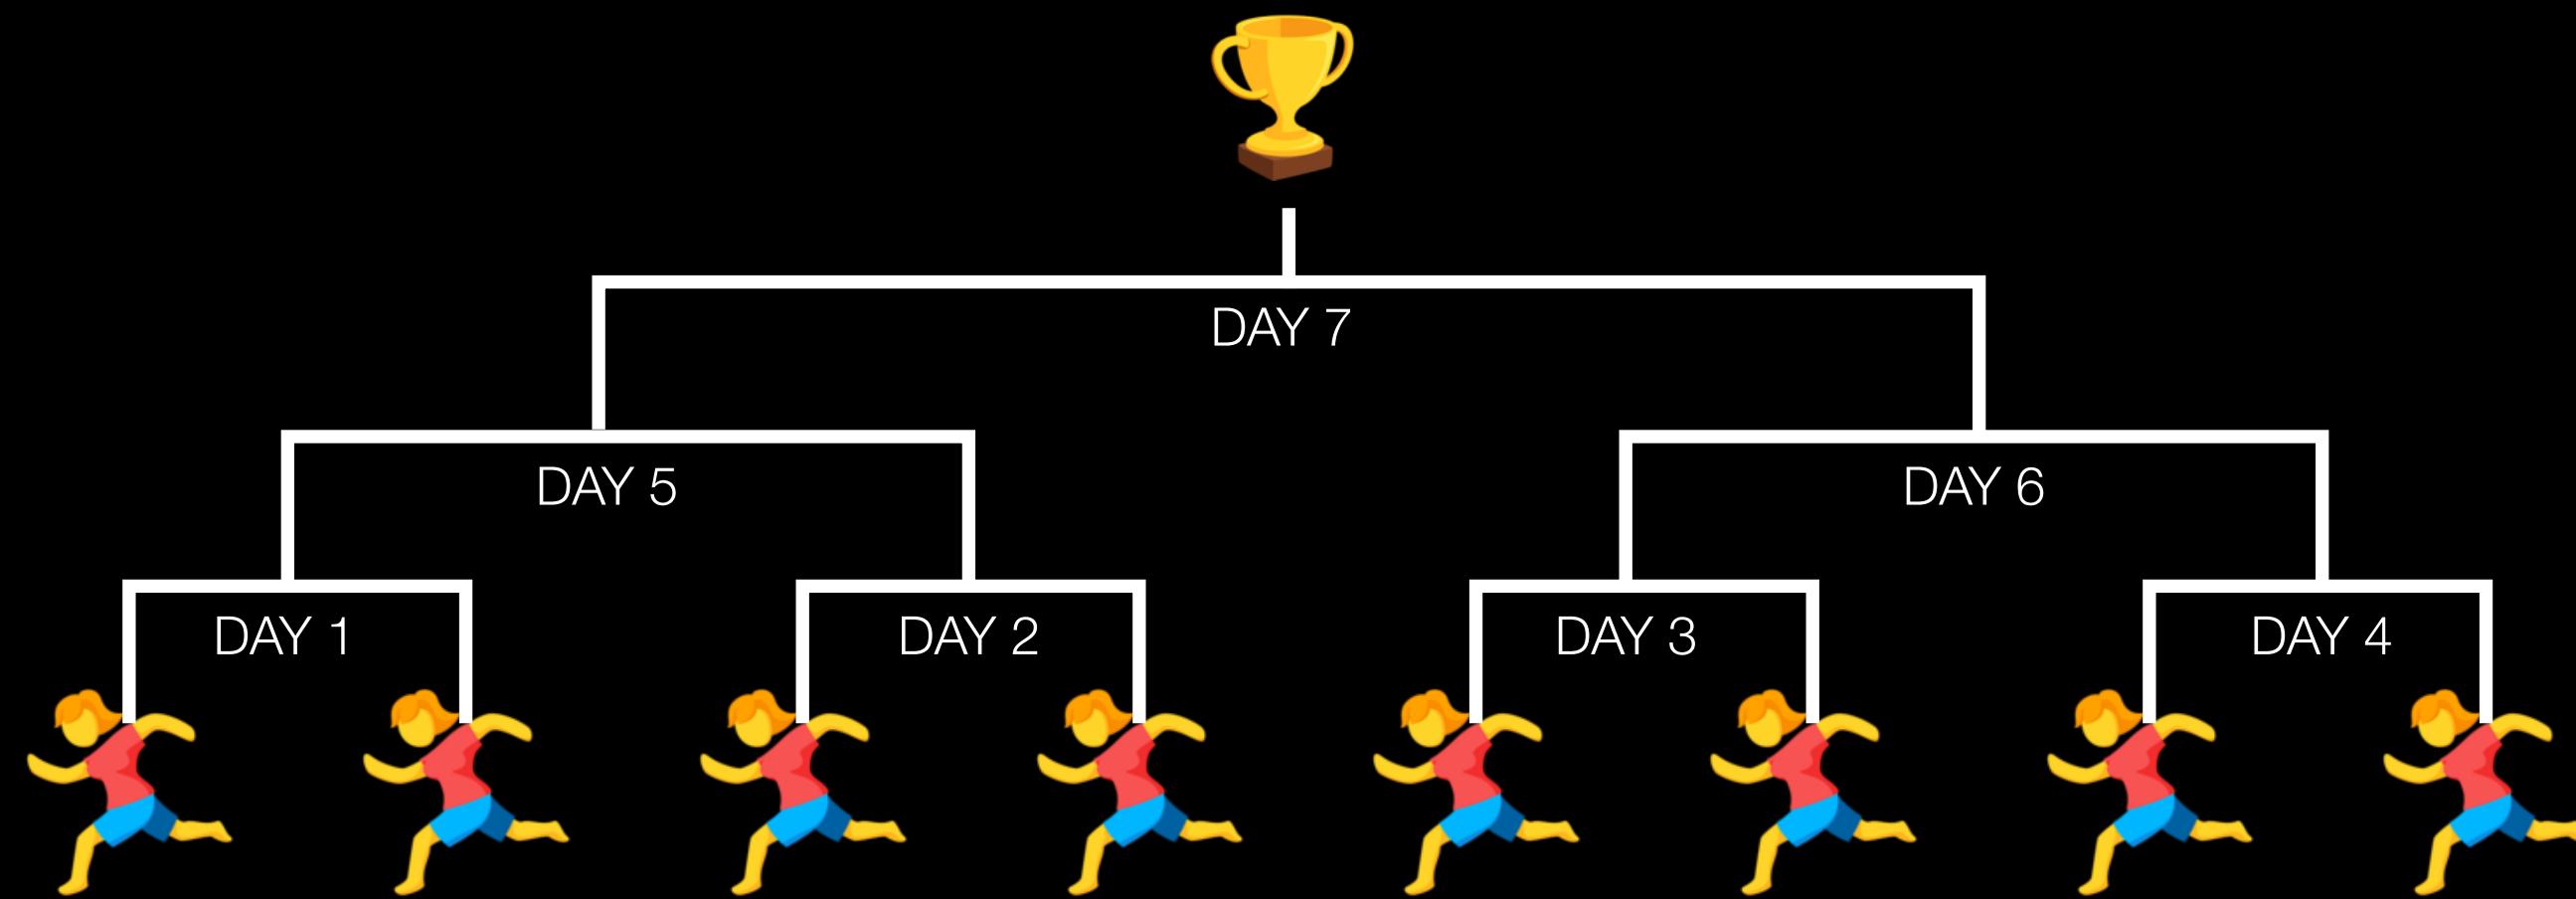


# reduction trees



# reduction trees

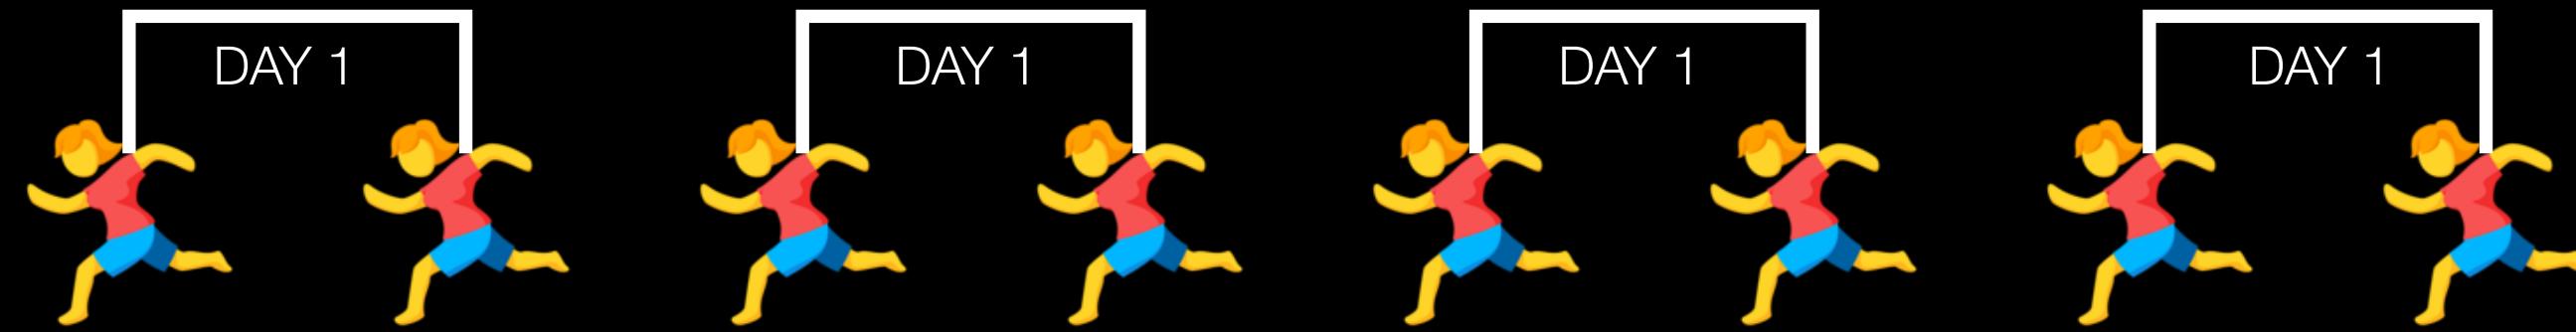
A little better!



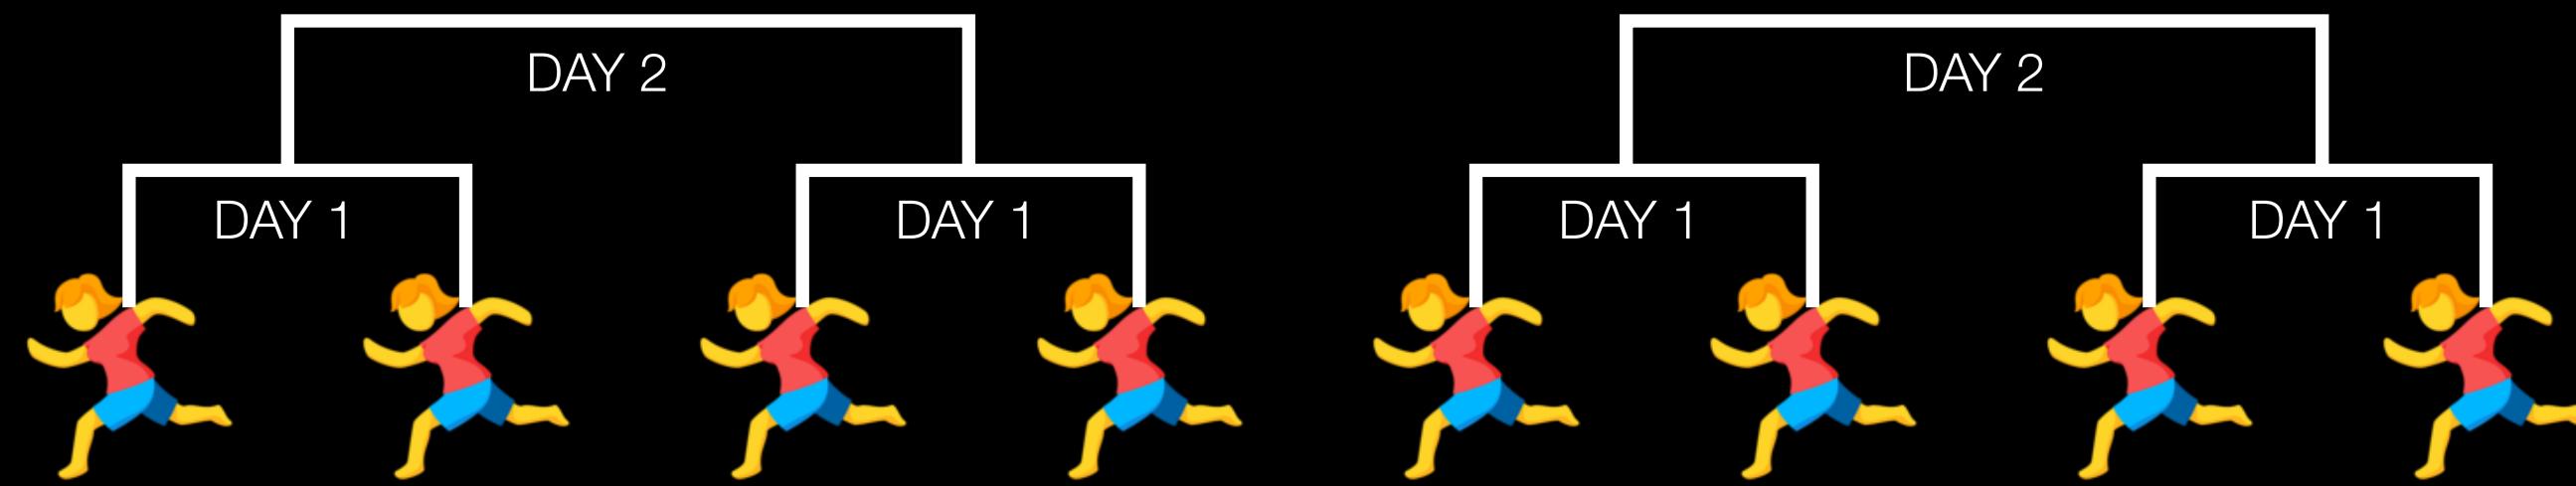
# reduction trees



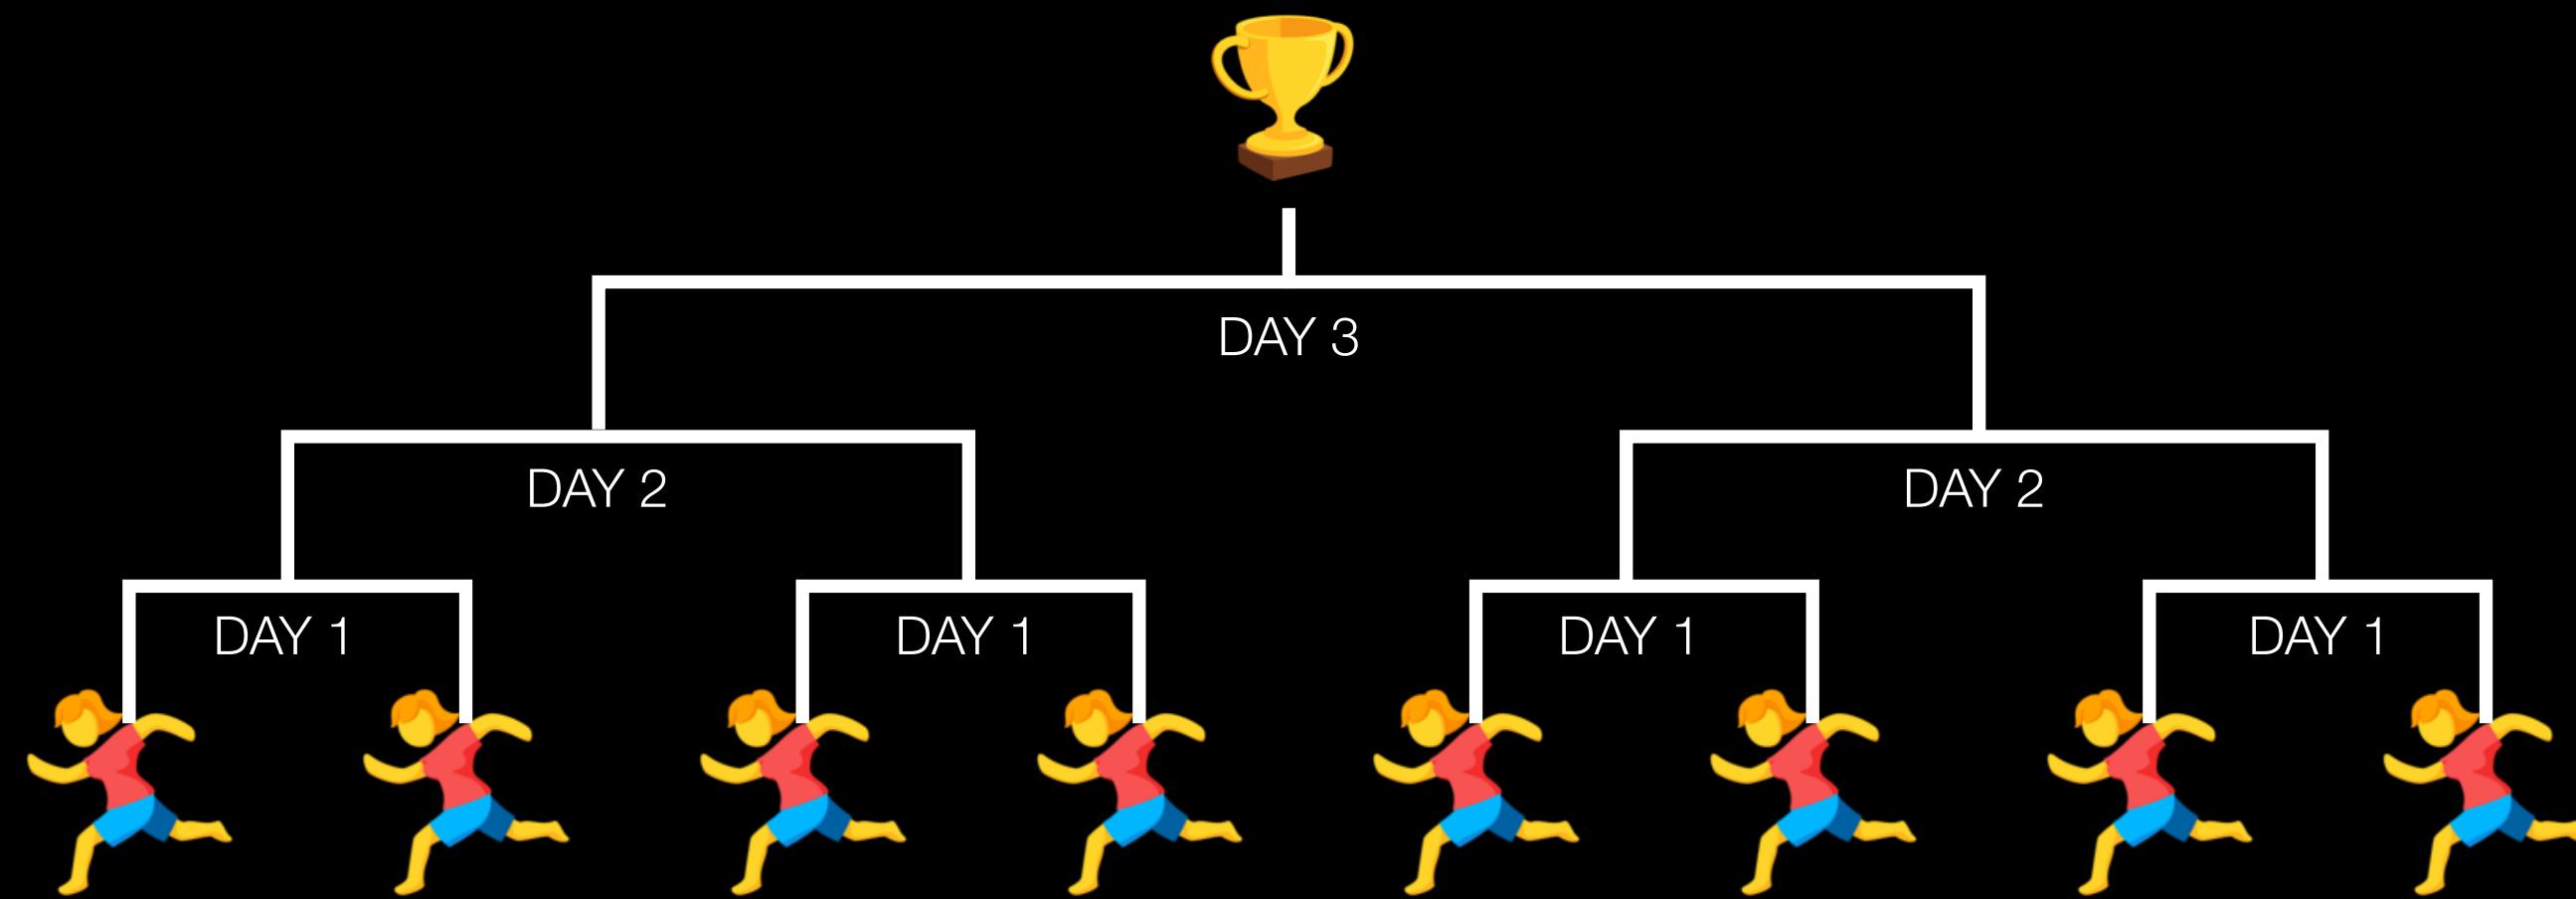
# reduction trees



# reduction trees

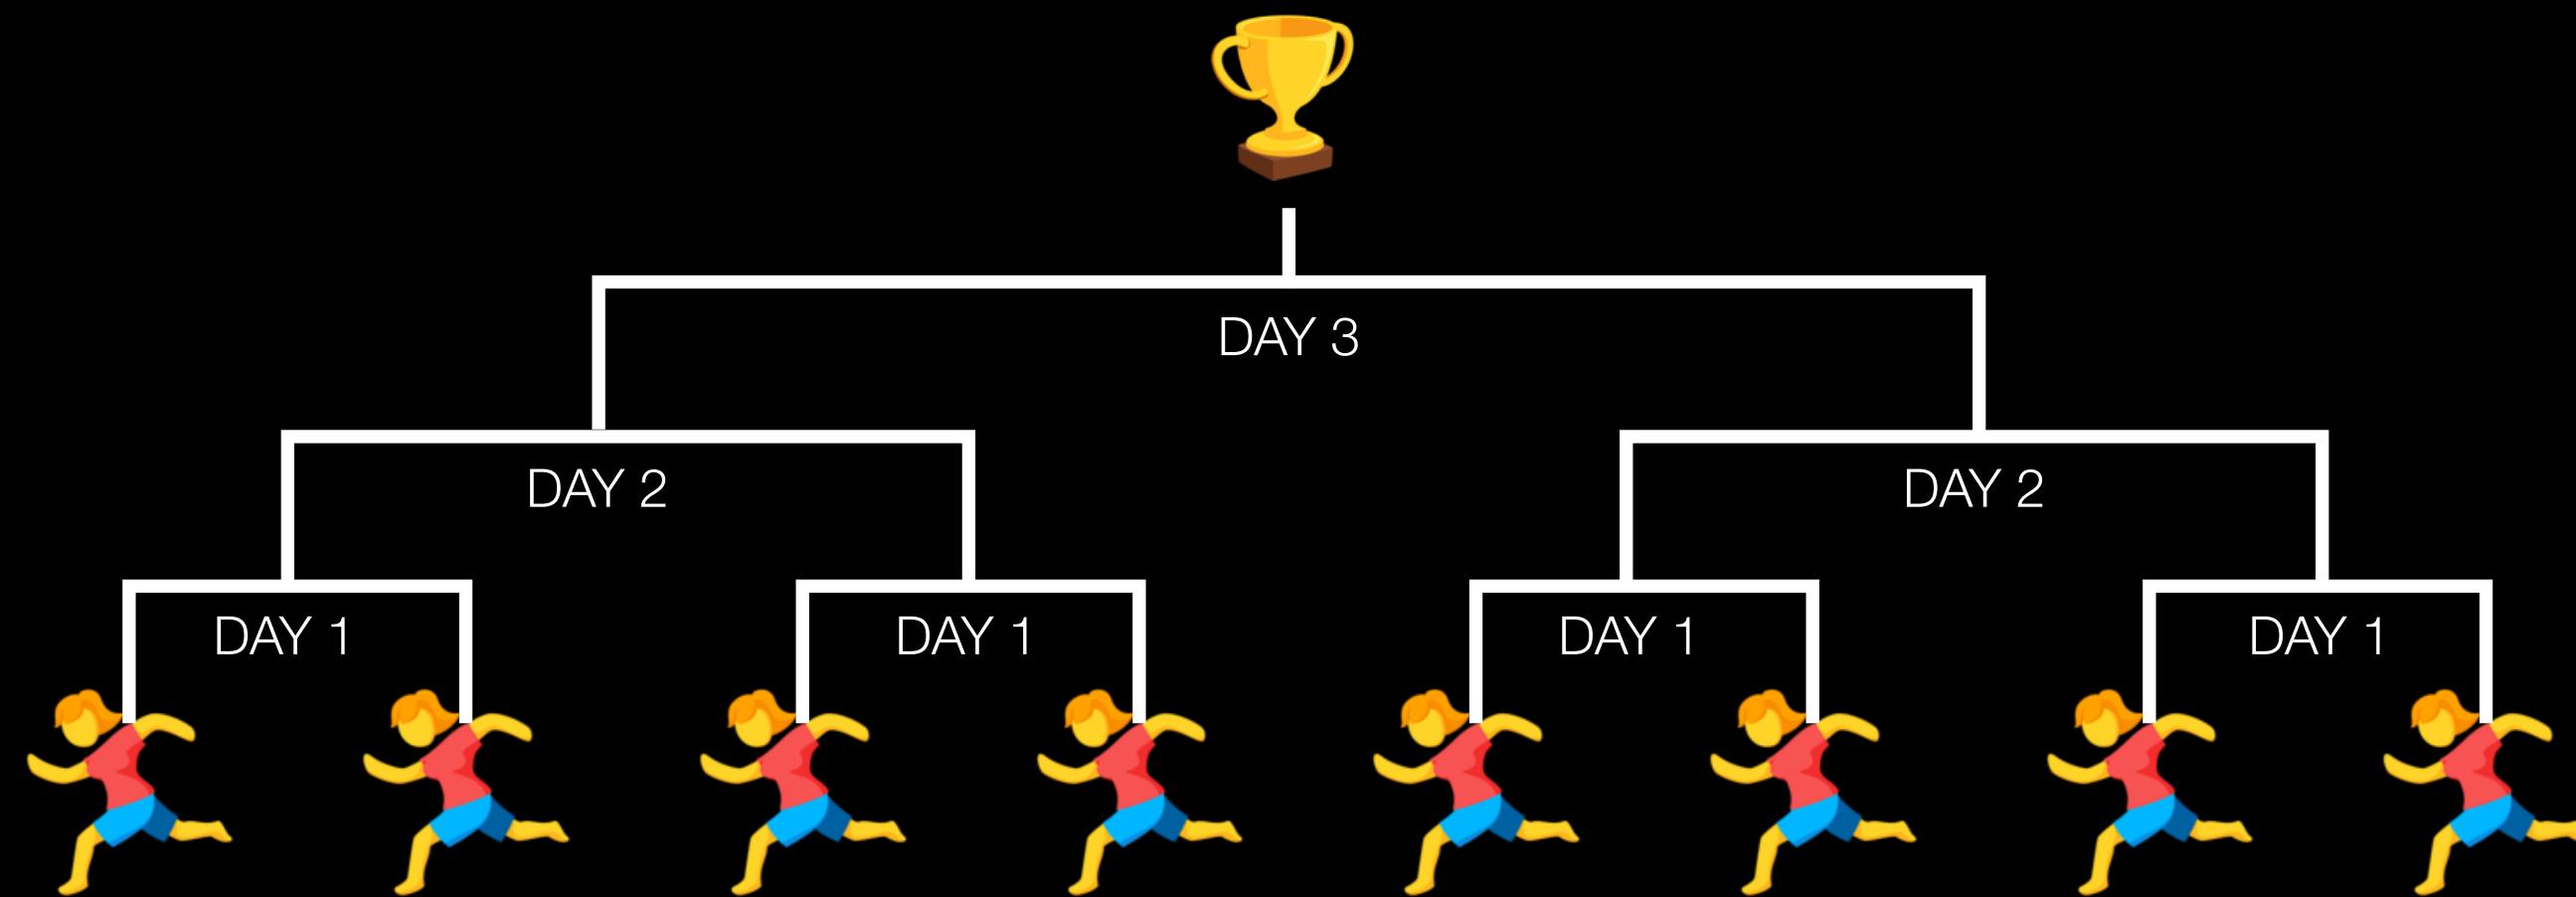


# reduction trees



# reduction trees

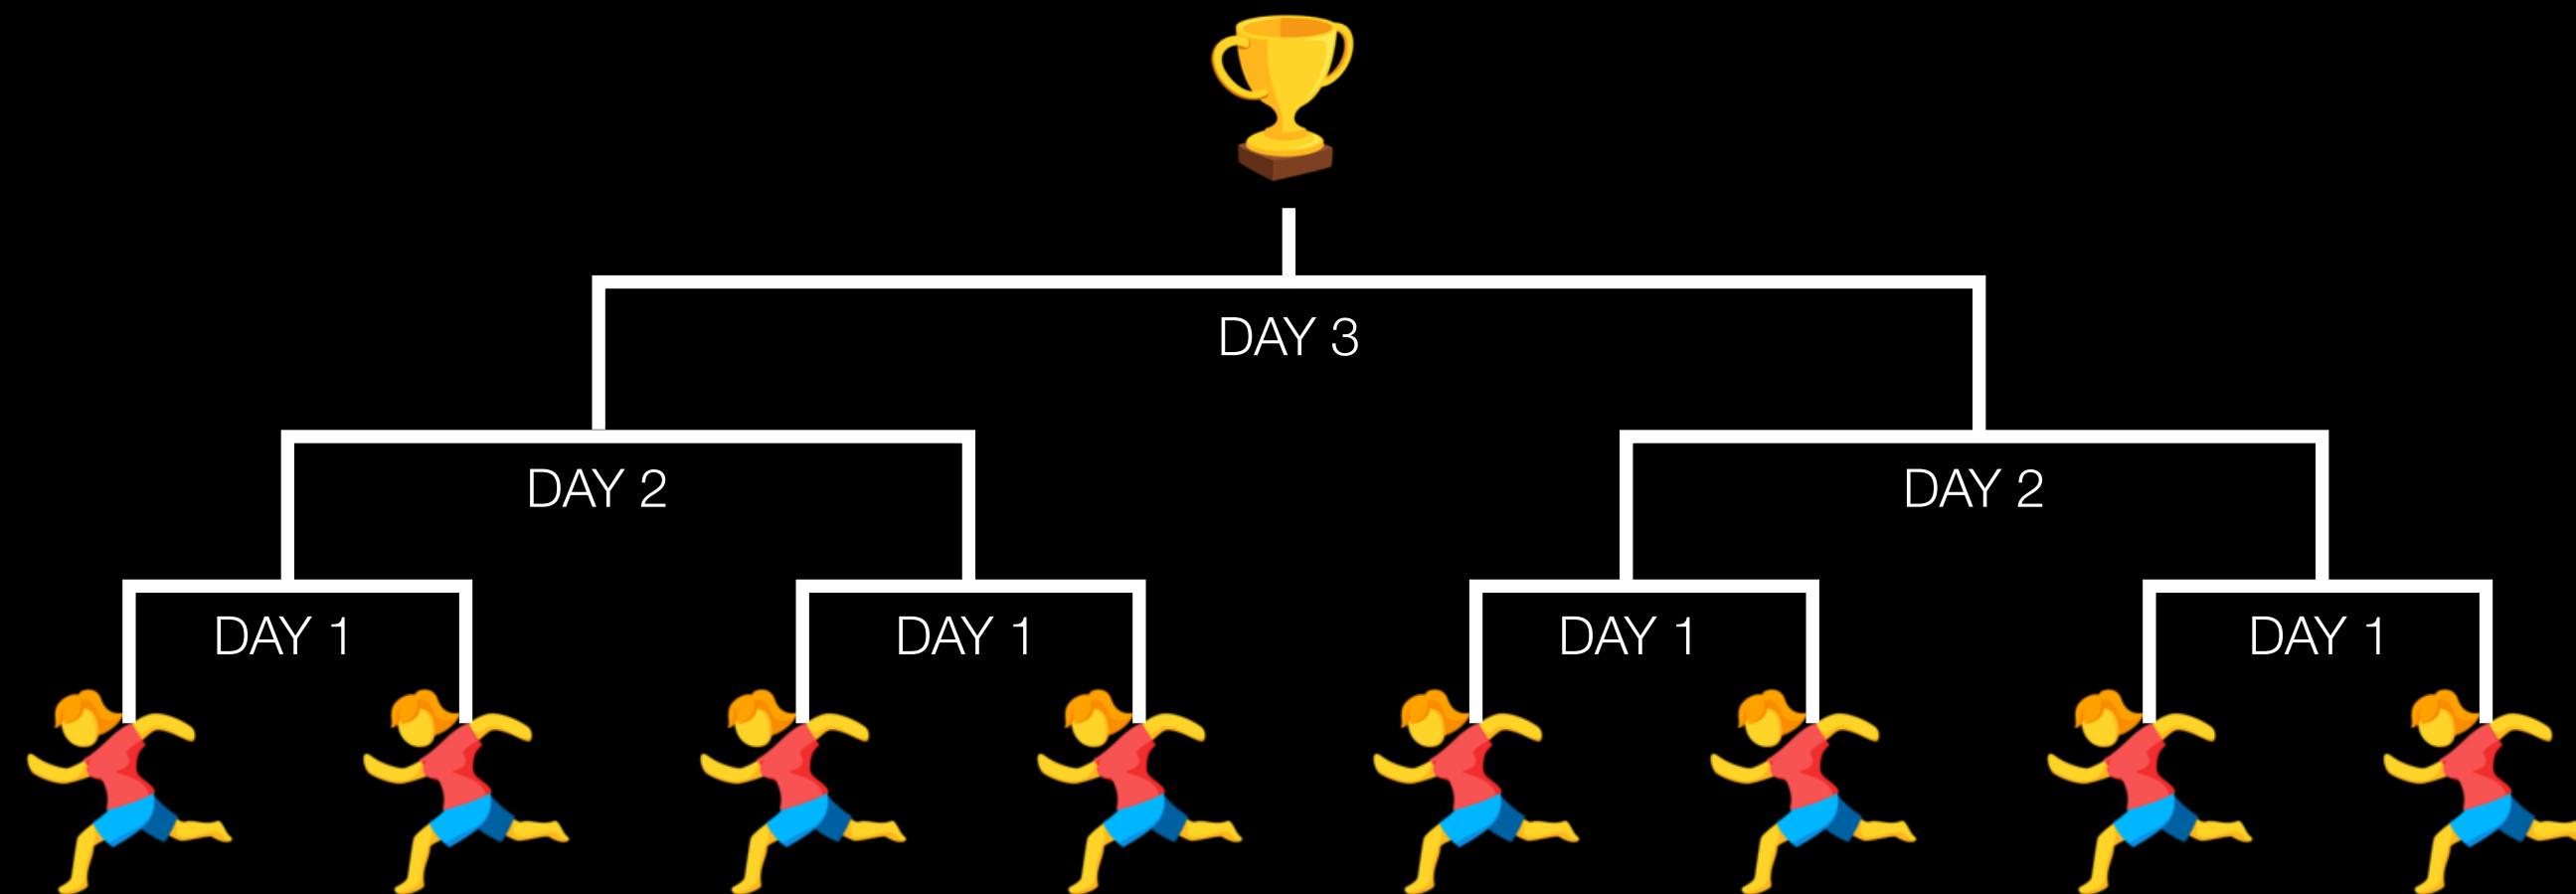
Better still...



# reduction trees

Better still...

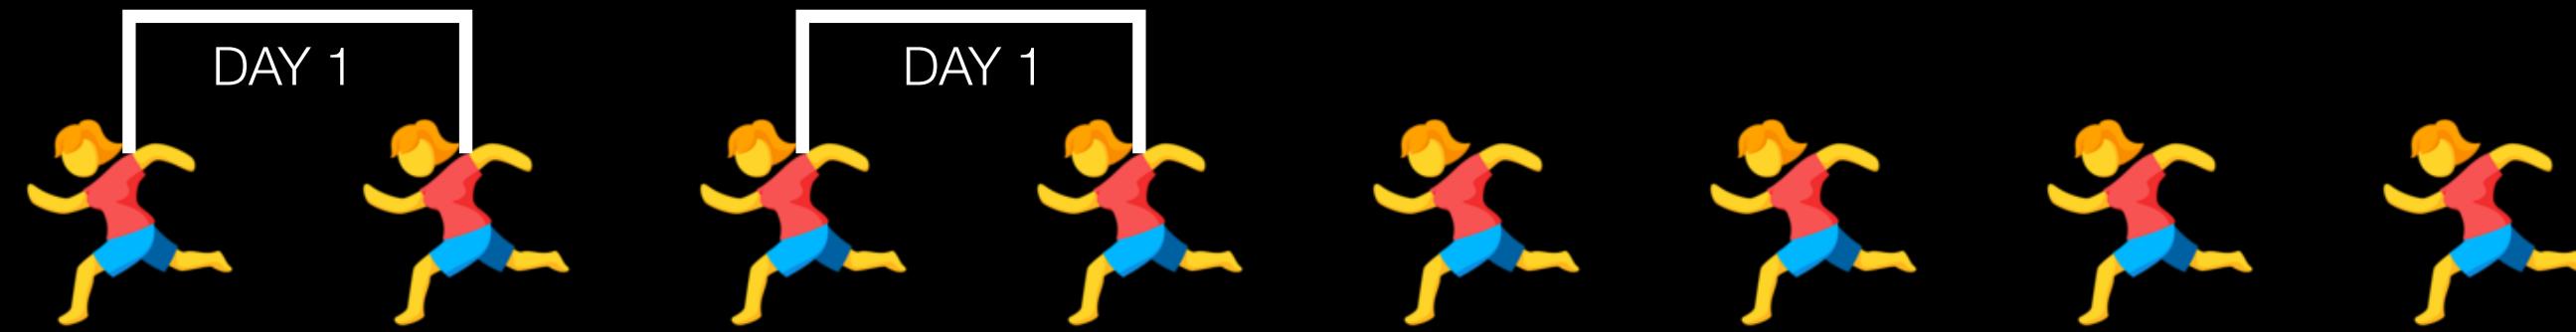
But hang on, do we have four courts at the same time?



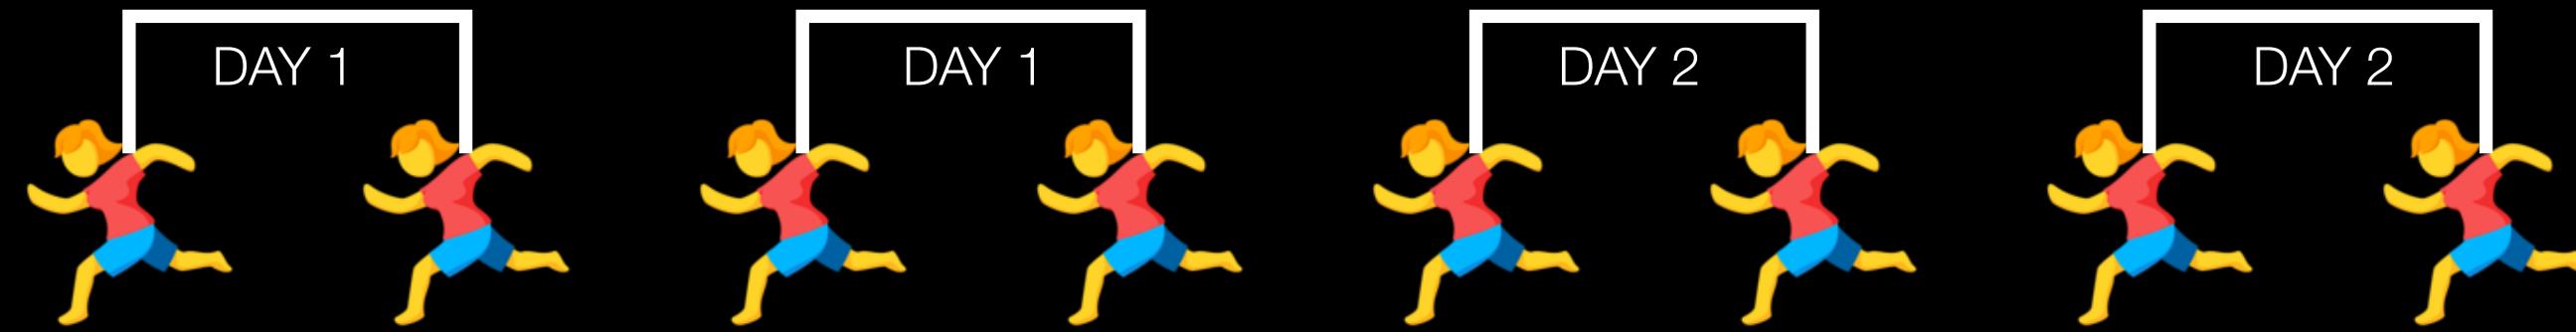
# reduction trees



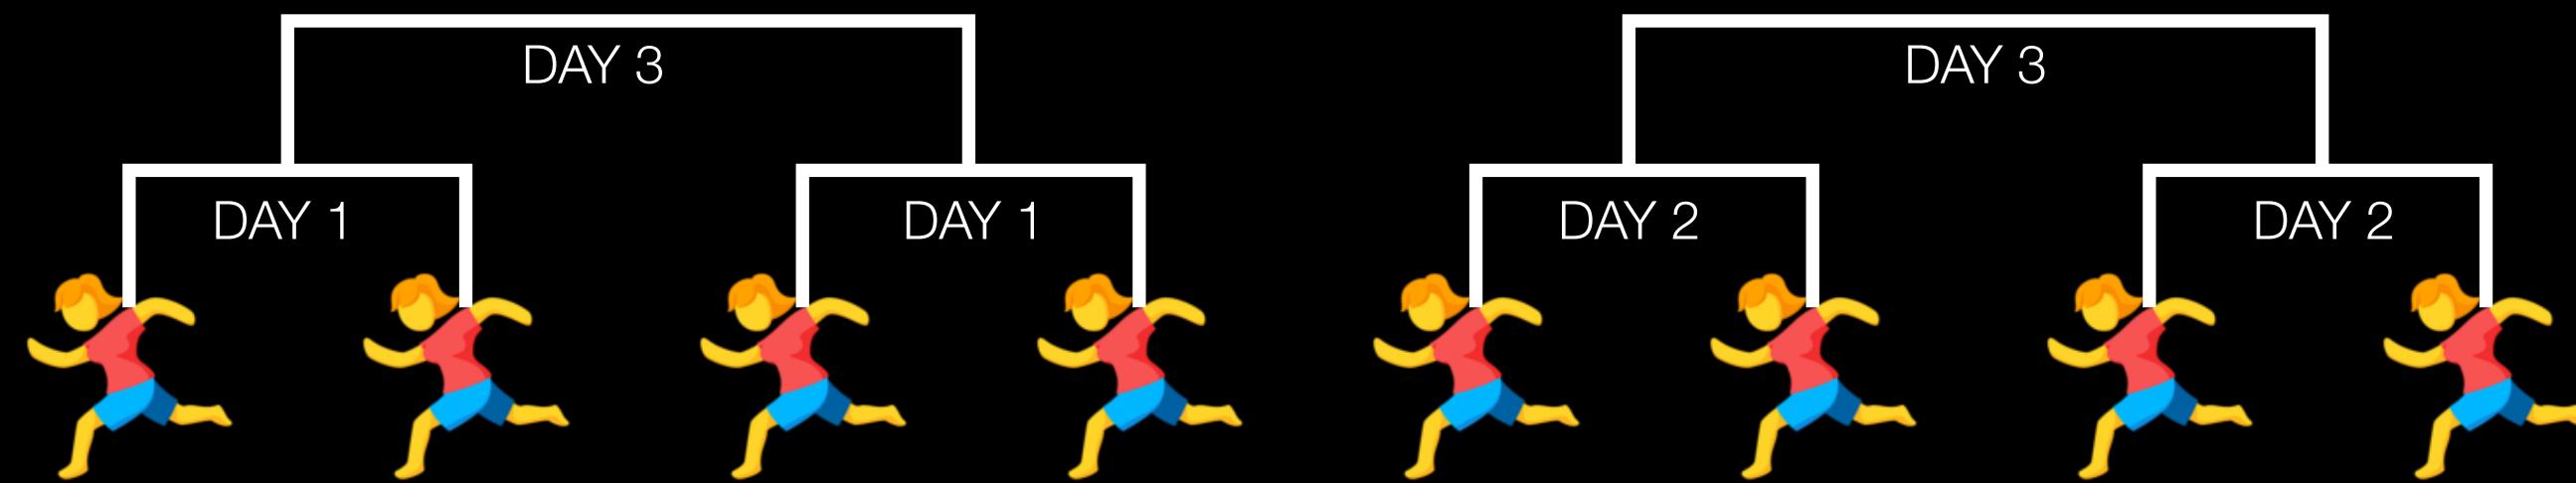
# reduction trees



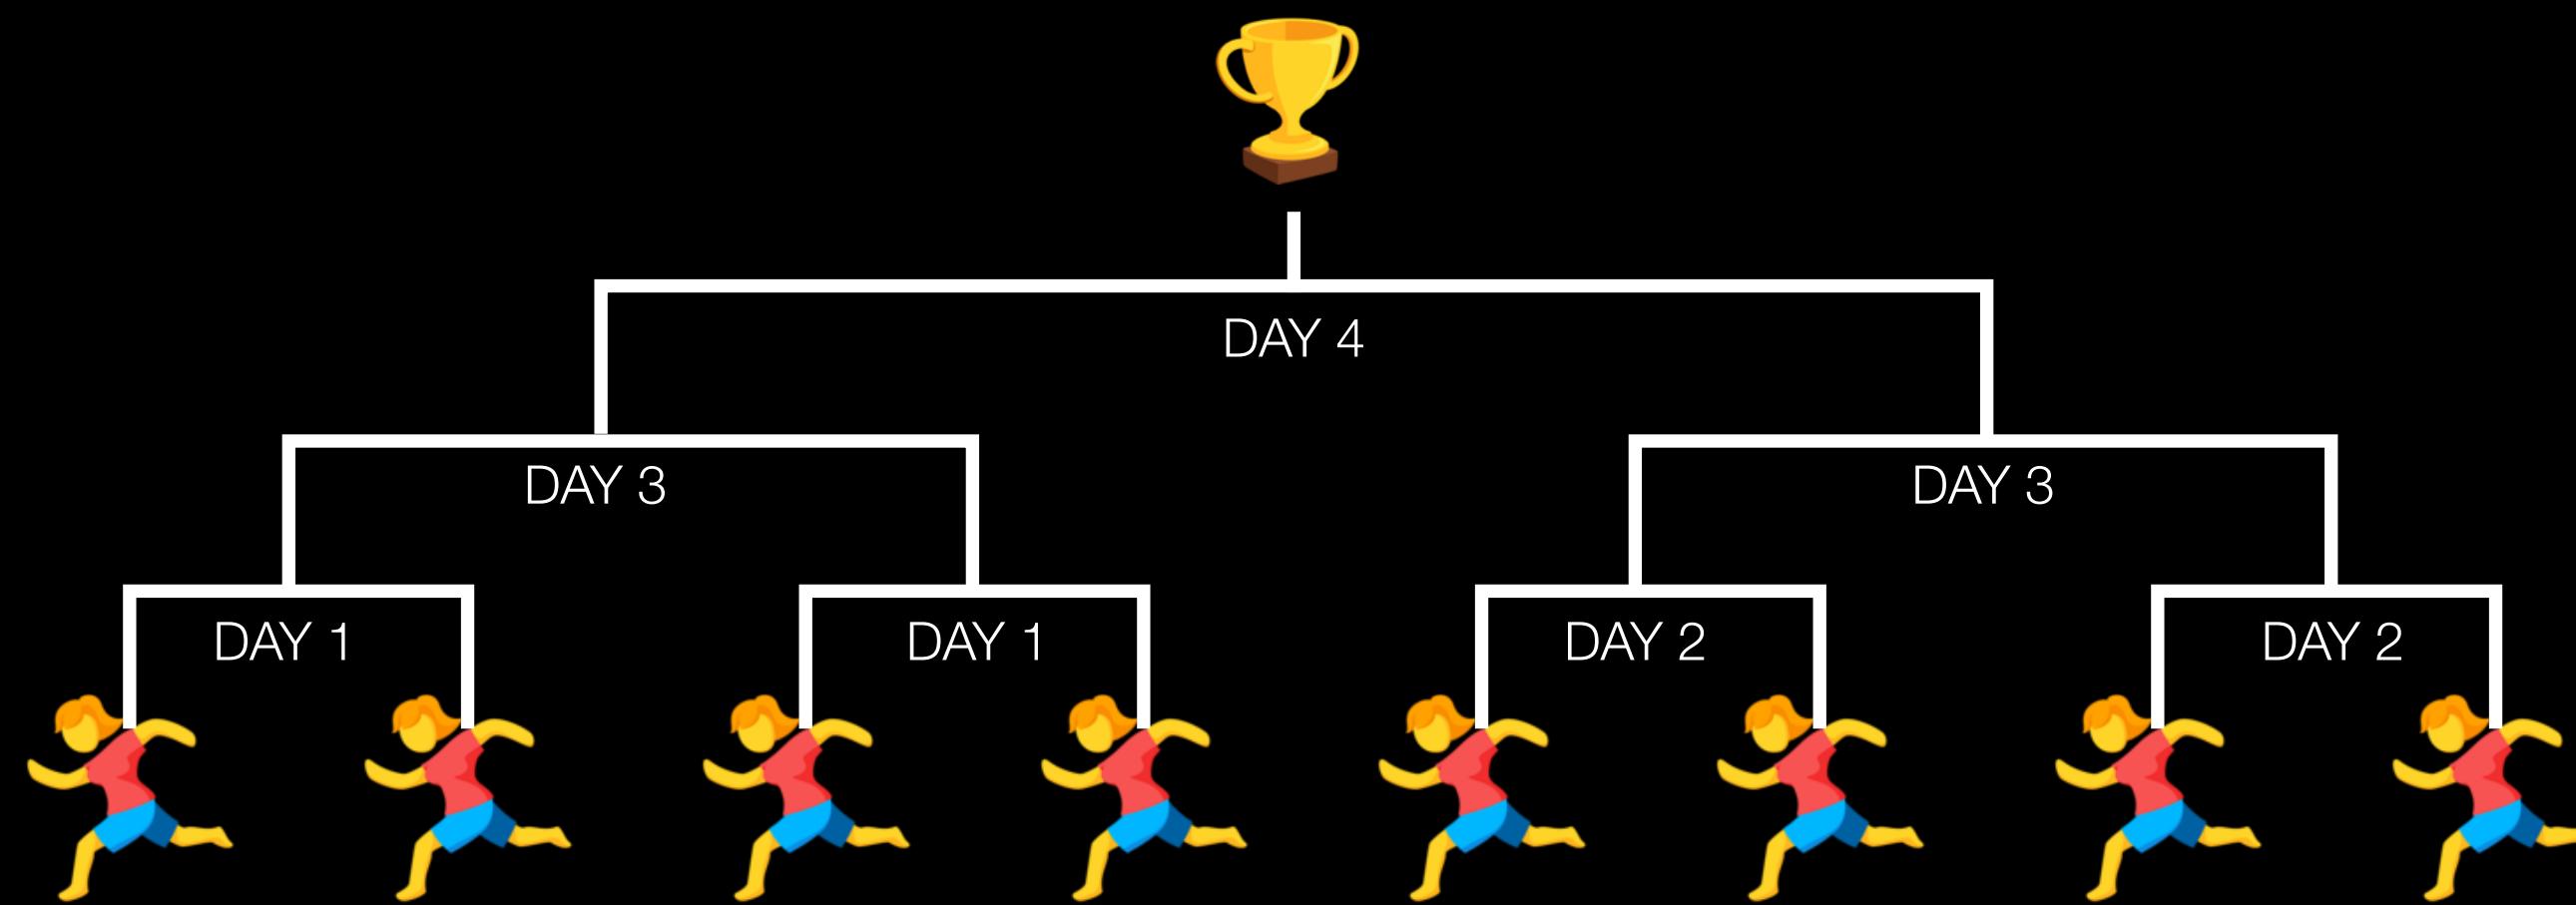
# reduction trees



# reduction trees

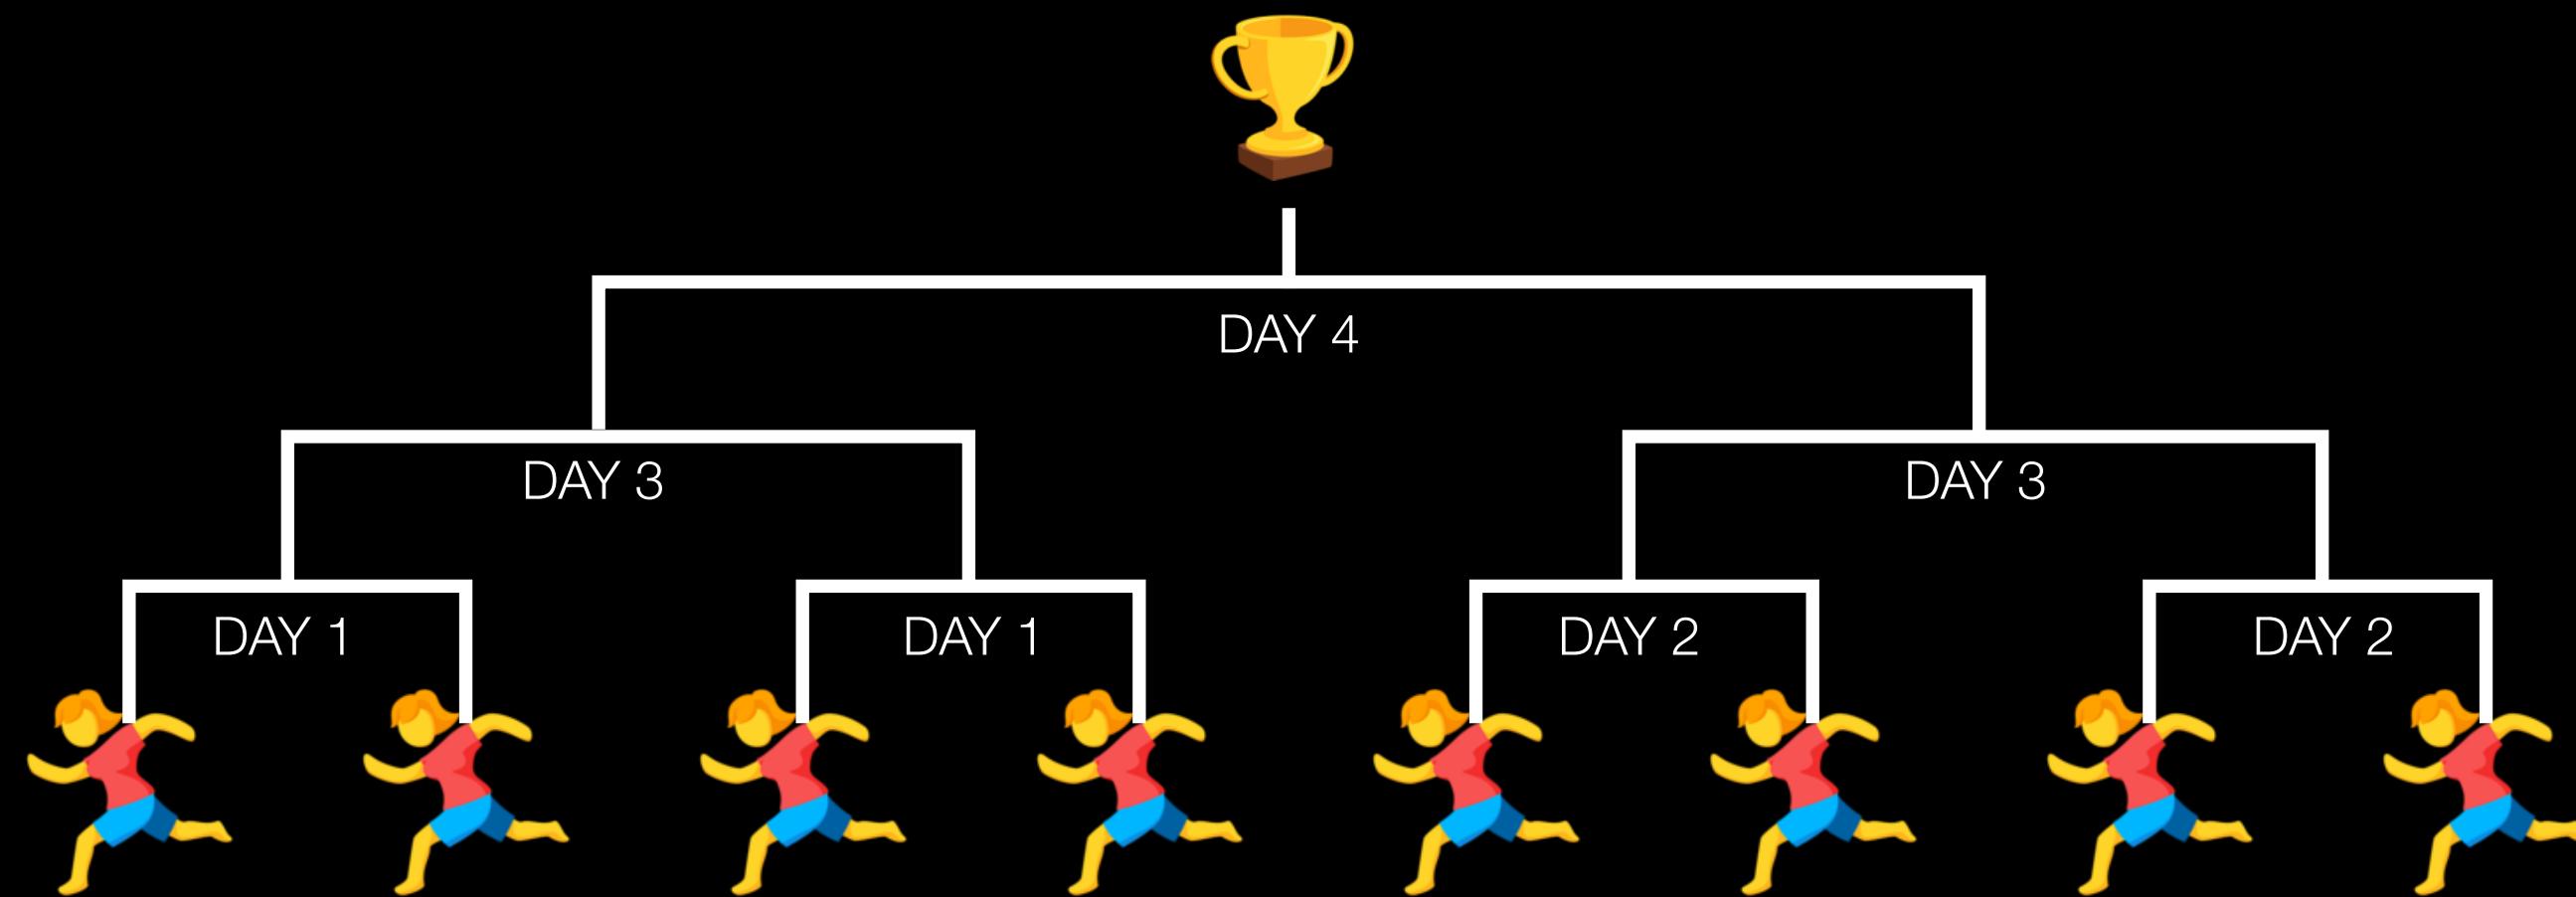


# reduction trees



# reduction trees

We can do this with two courts!



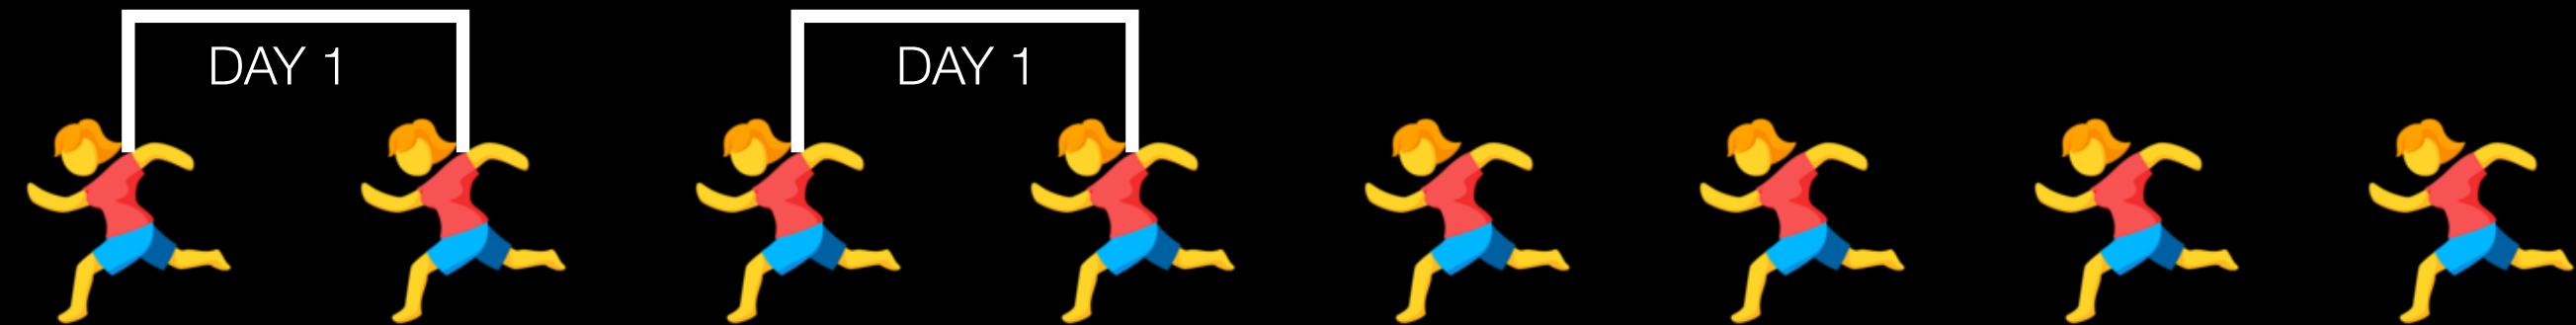
# reduction trees

Ugh it's getting expensive to transport players to and fro.  
Here's another option...



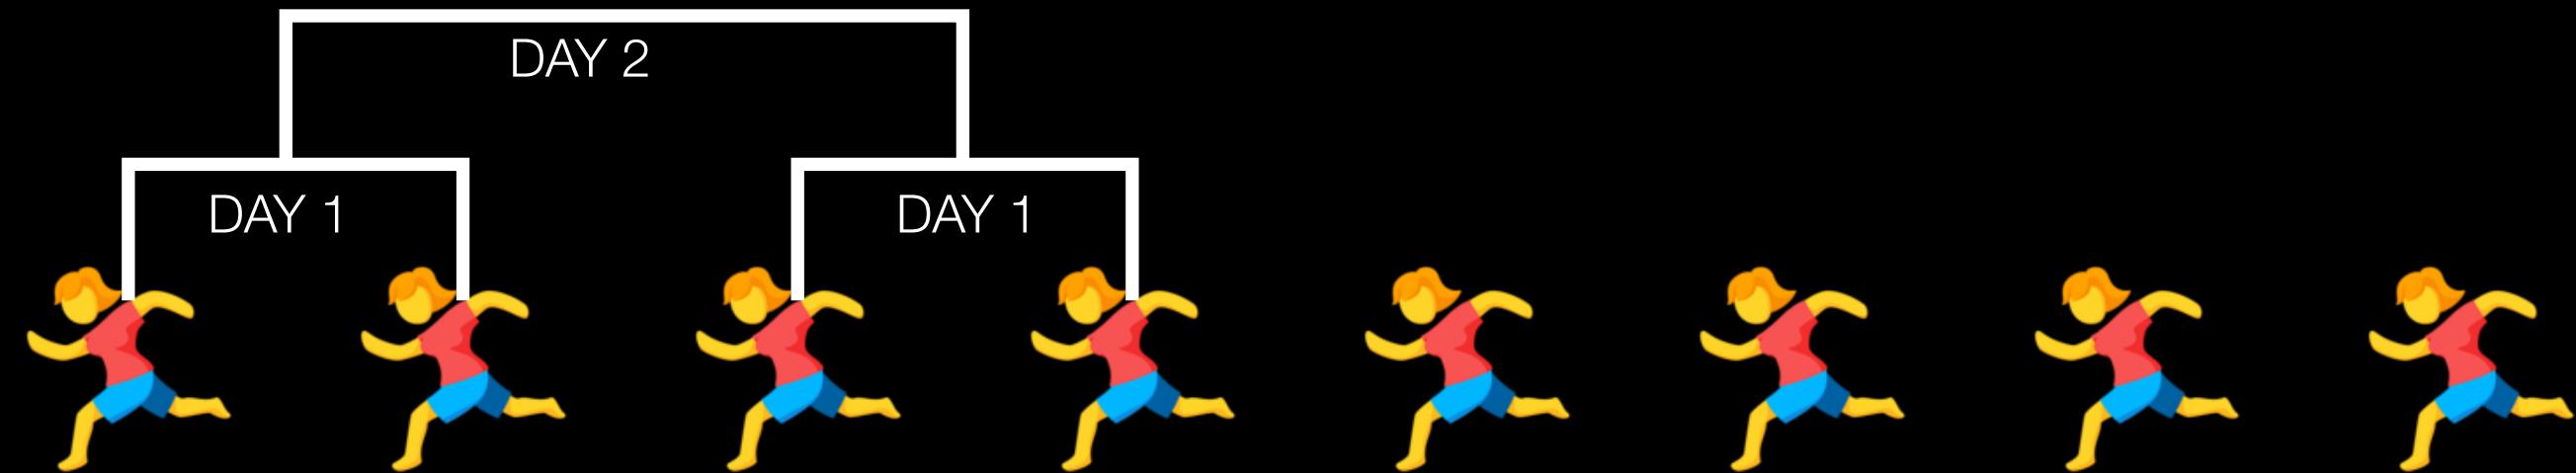
# reduction trees

Ugh it's getting expensive to transport players to and fro.  
Here's another option...



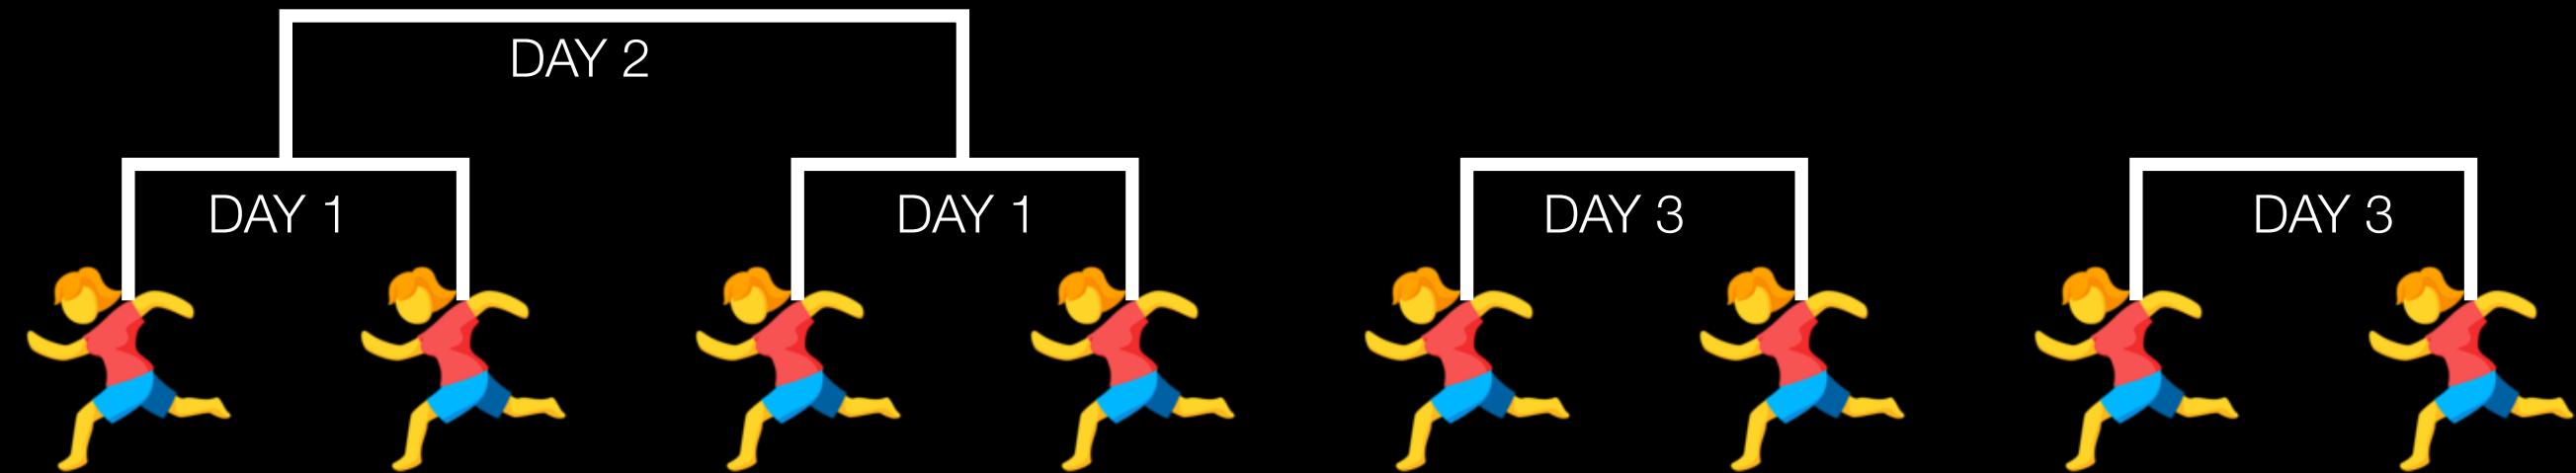
# reduction trees

Ugh it's getting expensive to transport players to and fro.  
Here's another option...



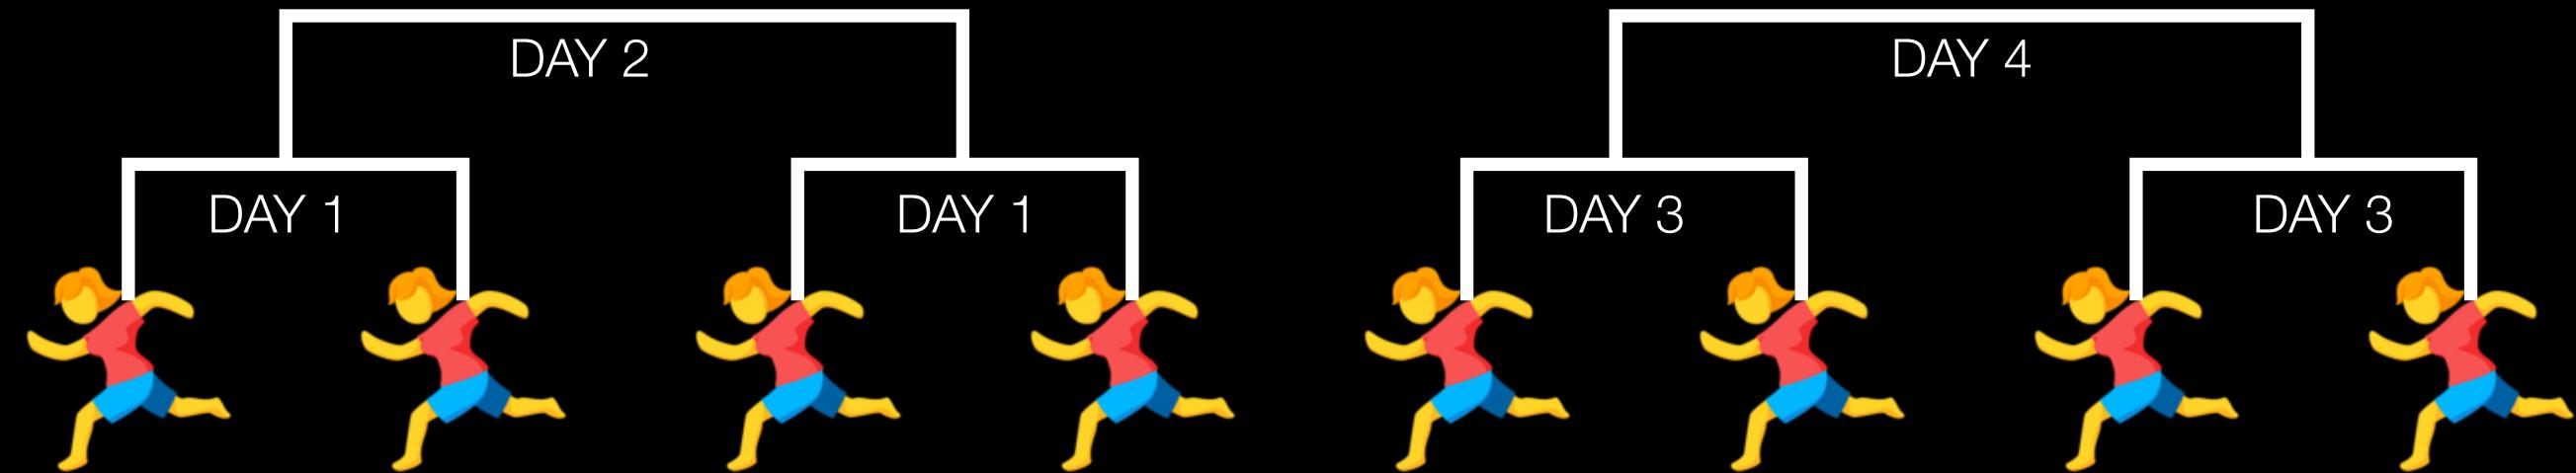
# reduction trees

Ugh it's getting expensive to transport players to and fro.  
Here's another option...



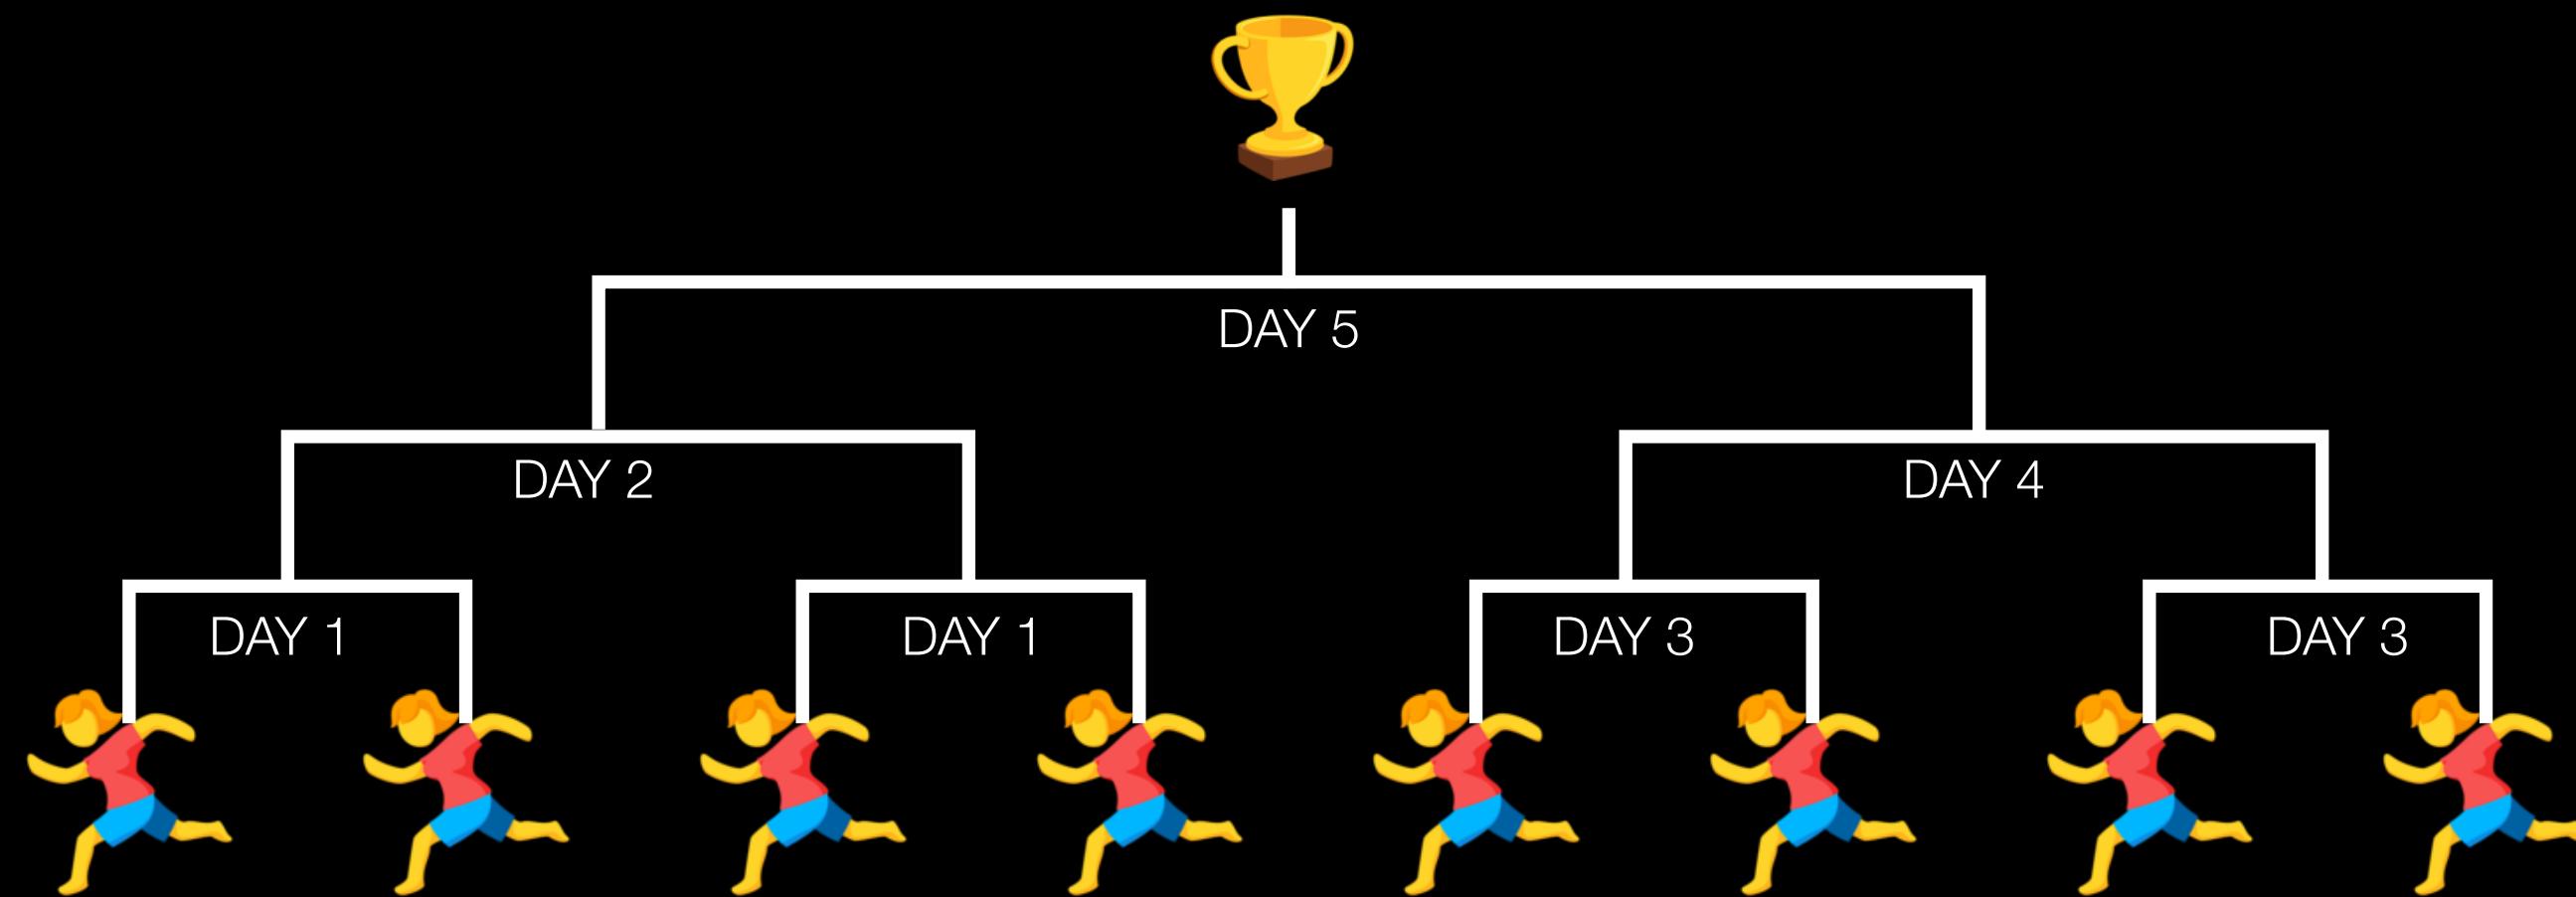
# reduction trees

Ugh it's getting expensive to transport players to and fro.  
Here's another option...



# reduction trees

Ugh it's getting expensive to transport players to and fro.  
Here's another option...



# reduction trees

This just in:  
the players are numbers;  
the matches are addition!

1    2    3    4    5    6    7    8

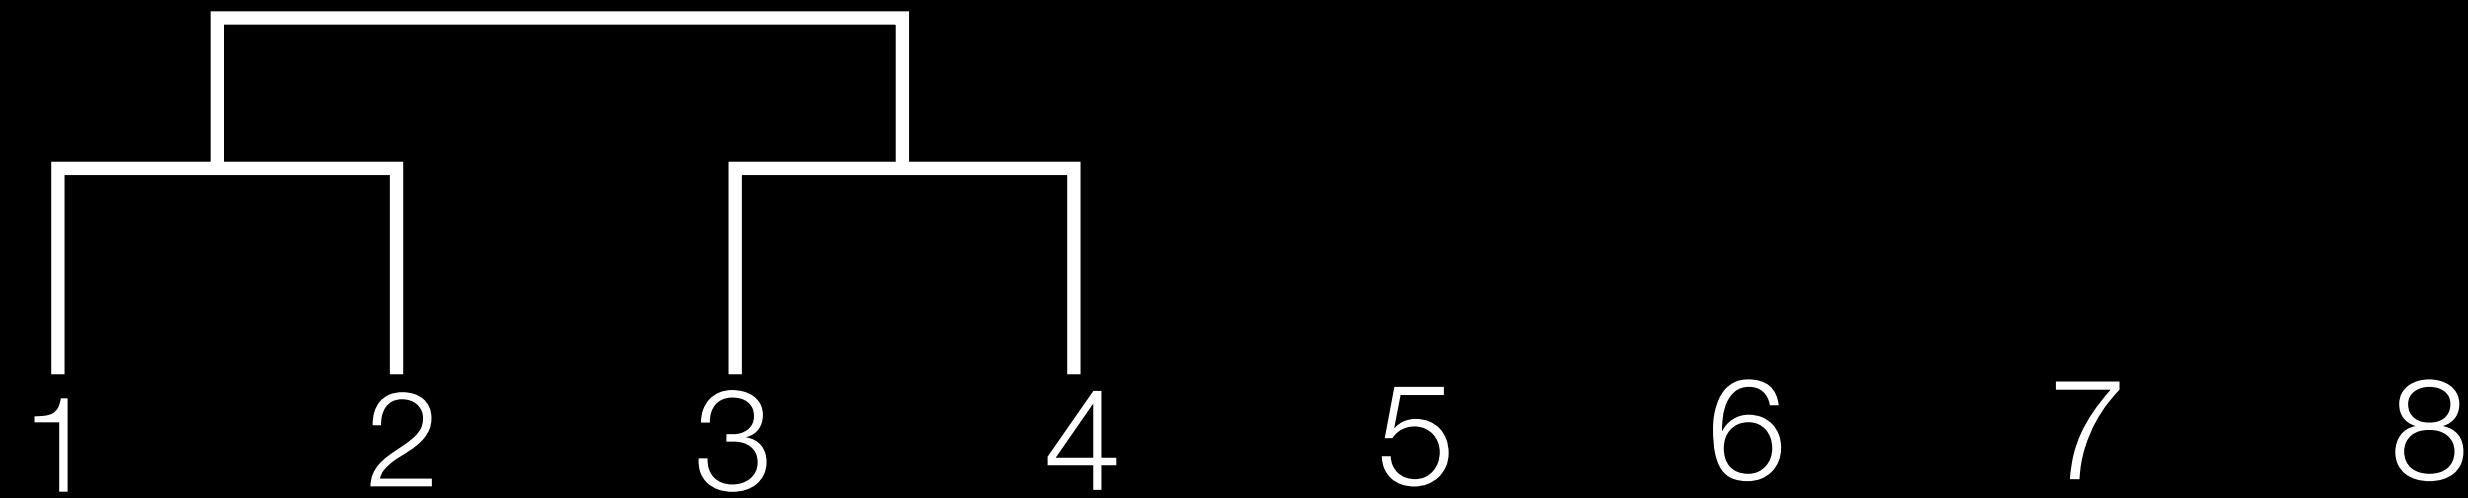
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



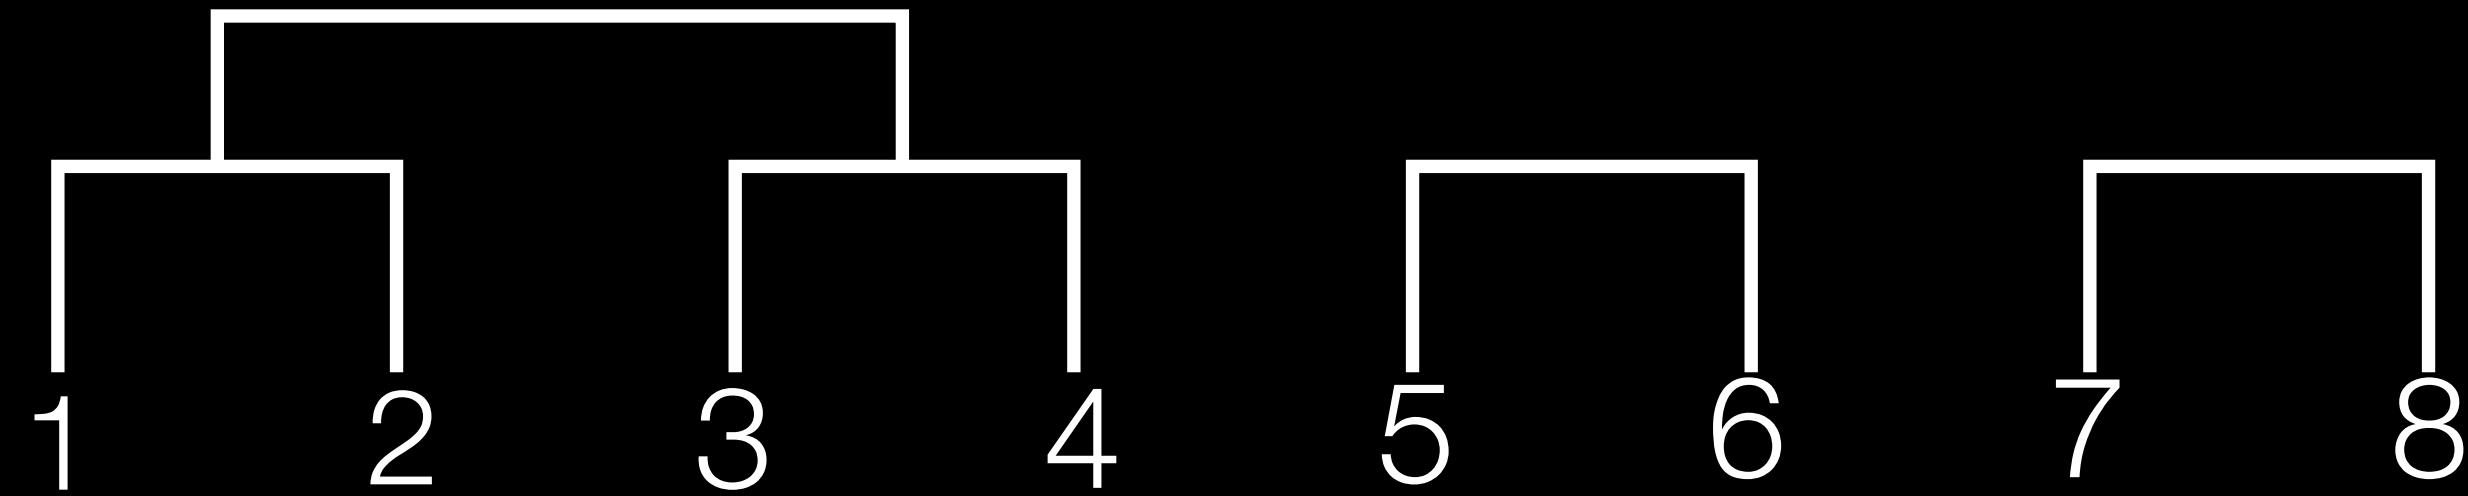
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



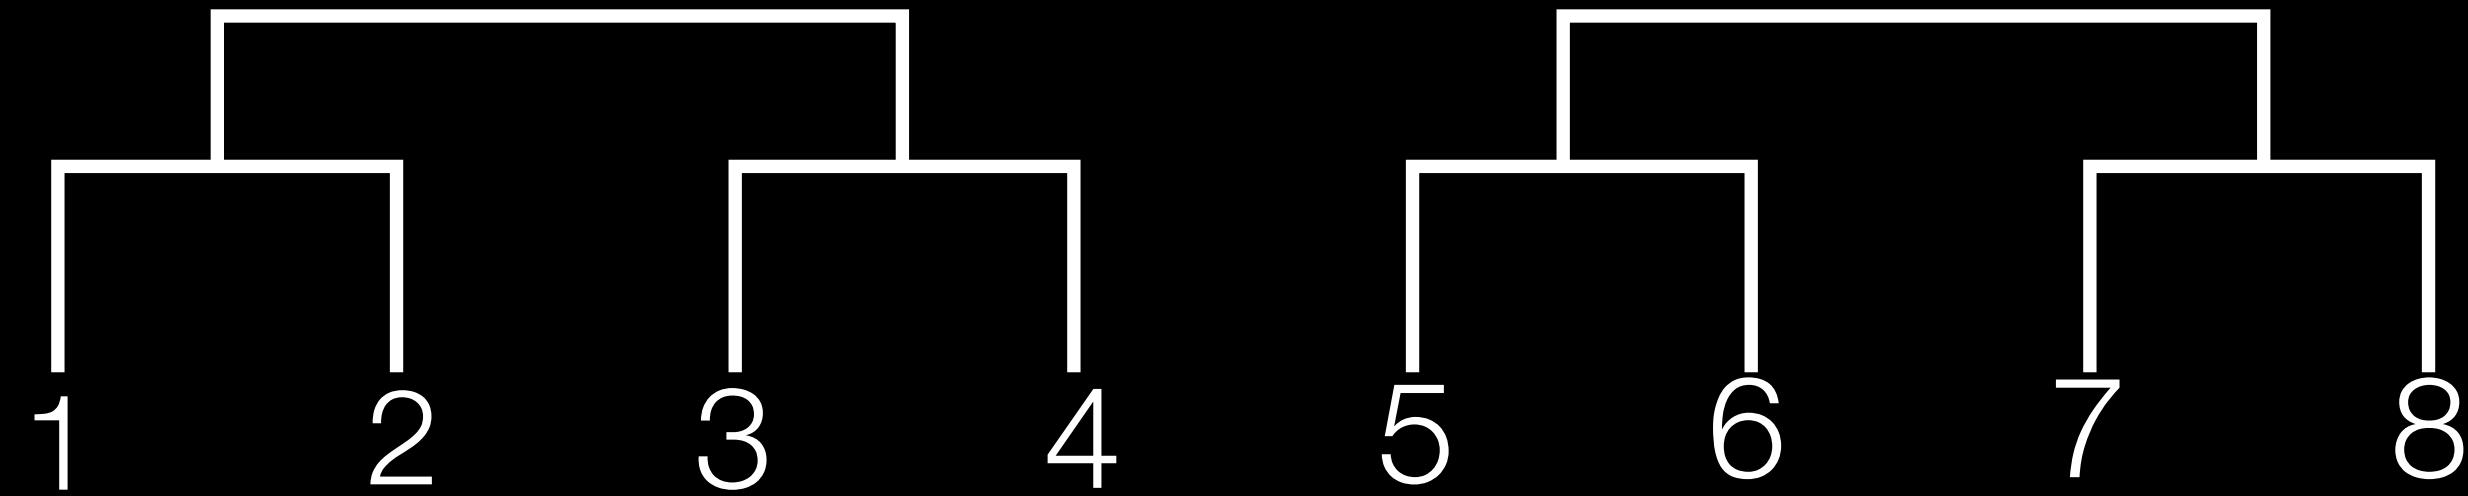
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



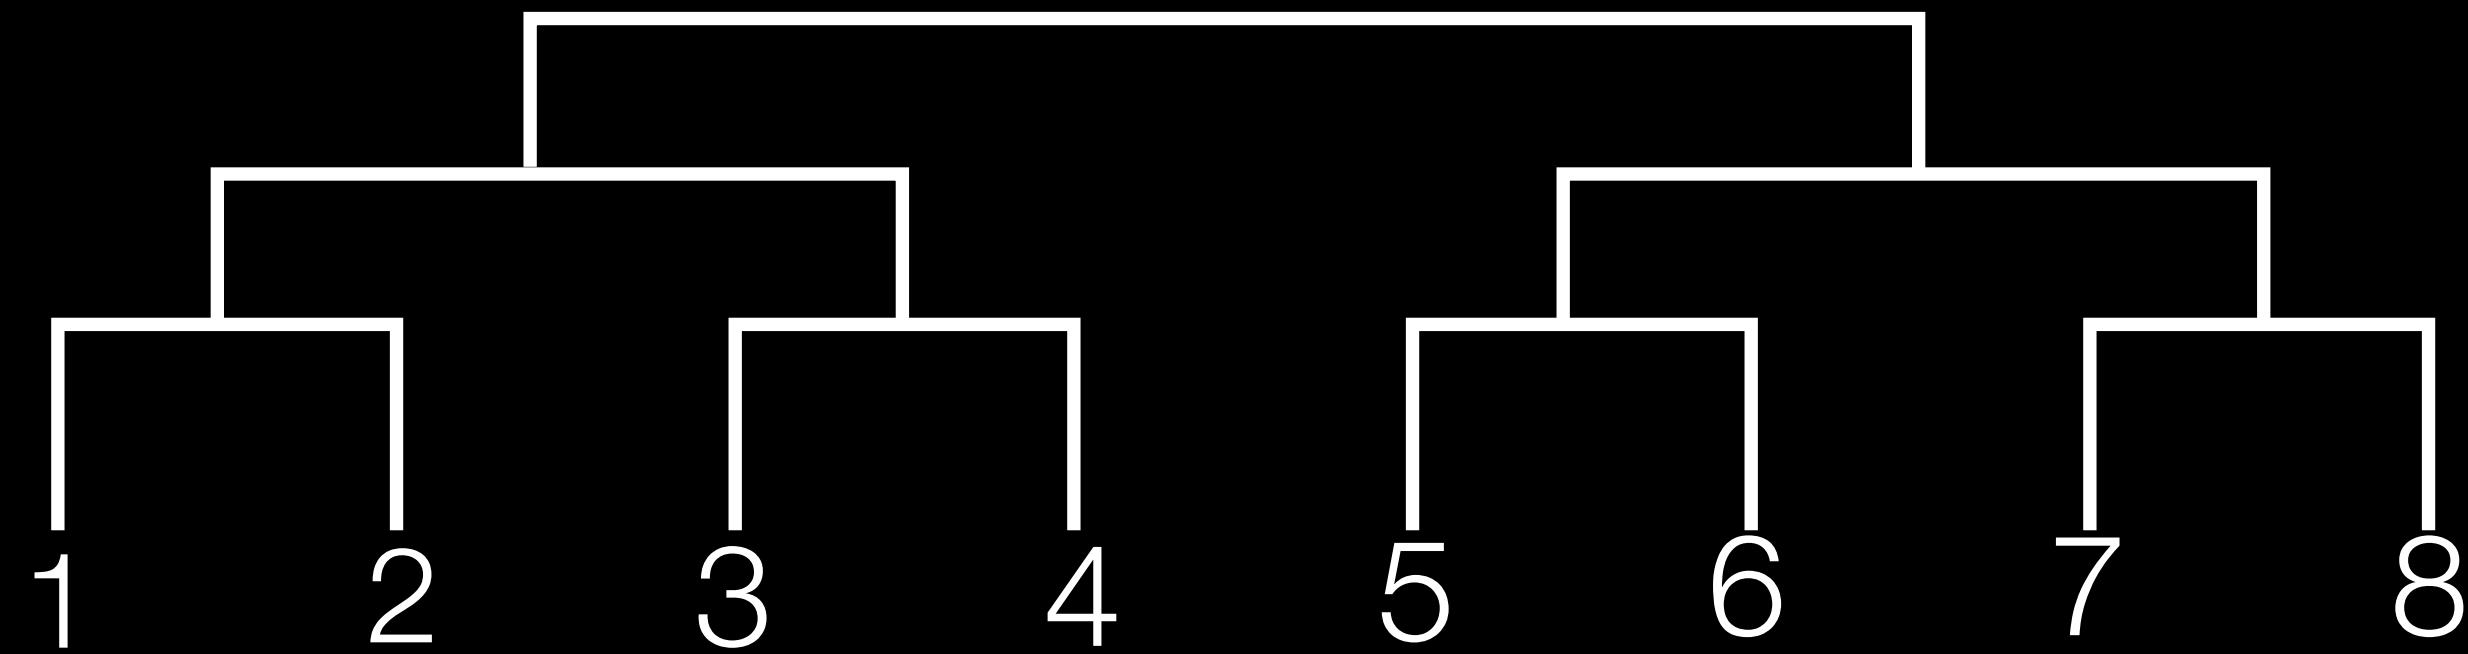
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



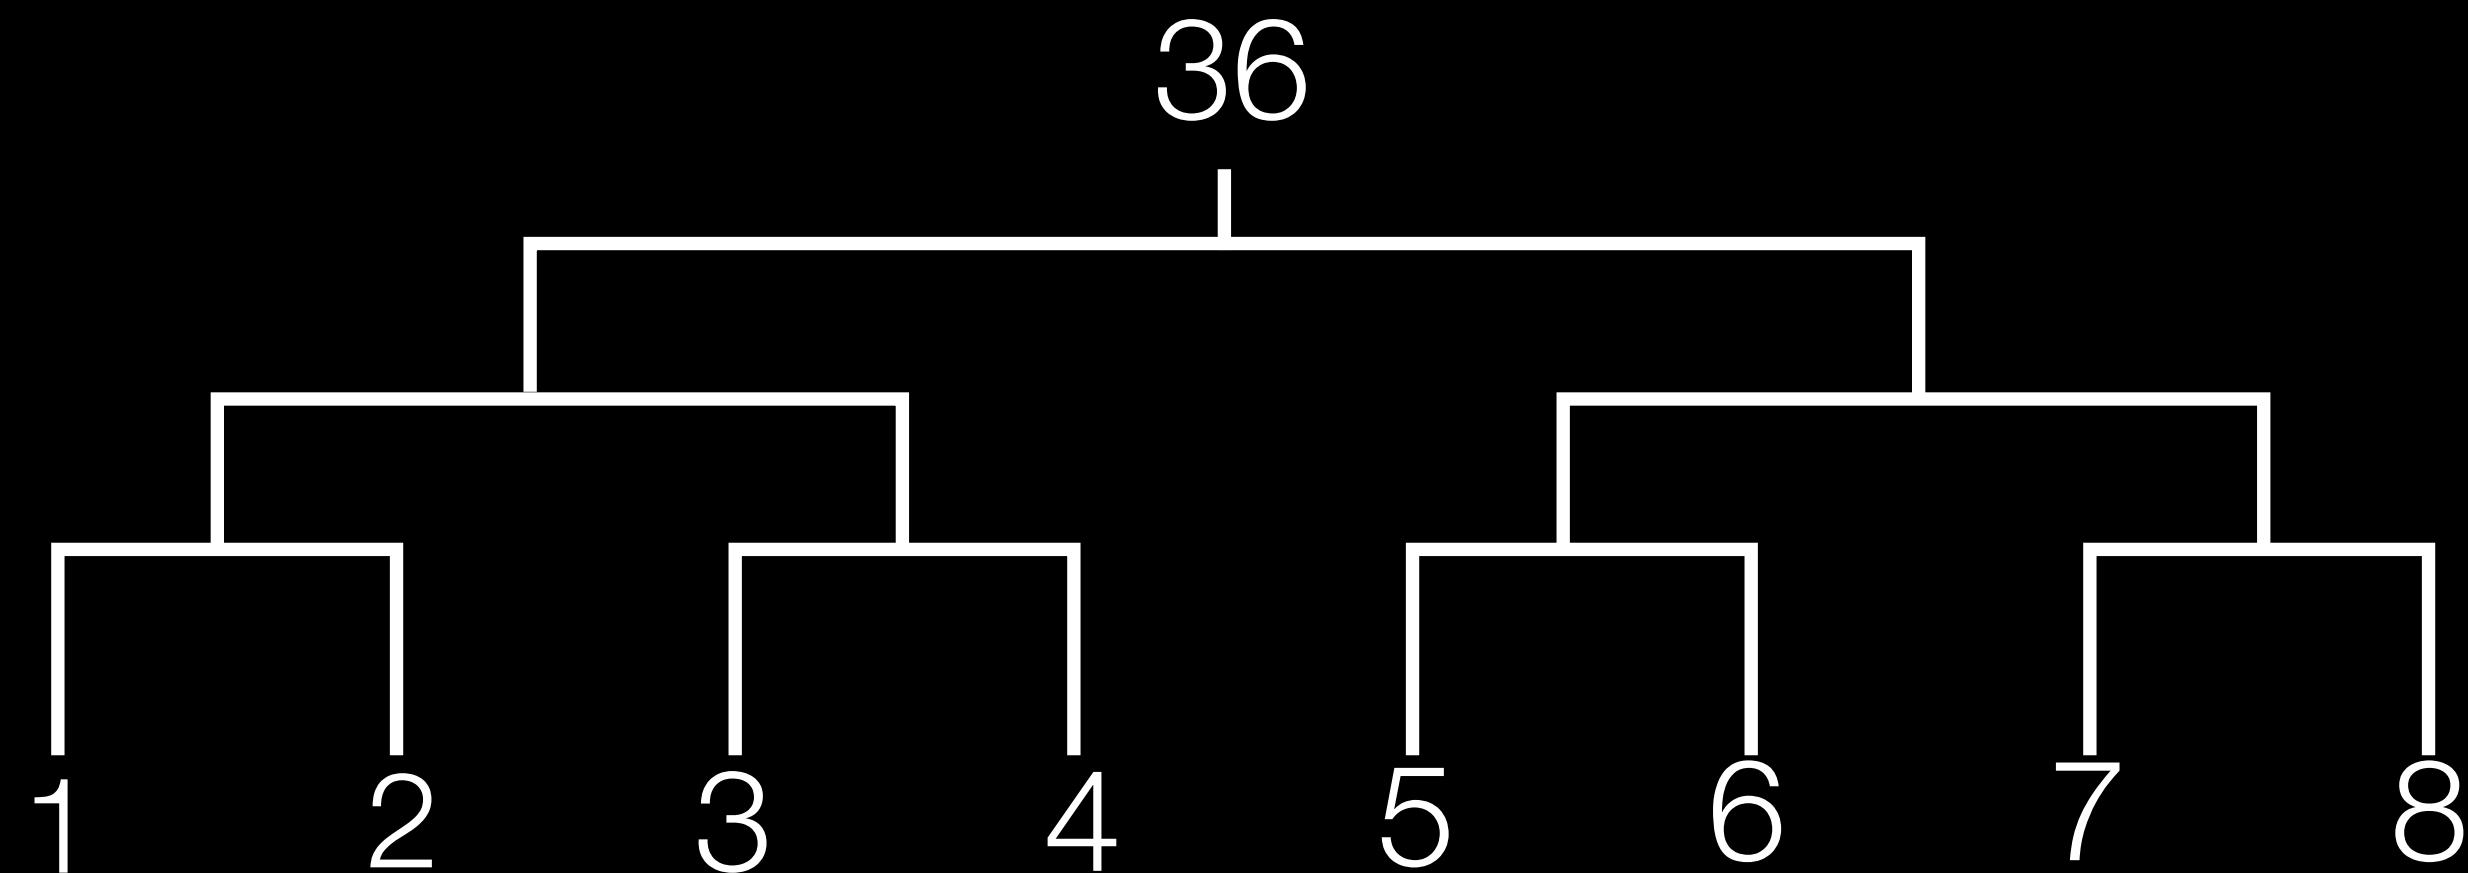
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



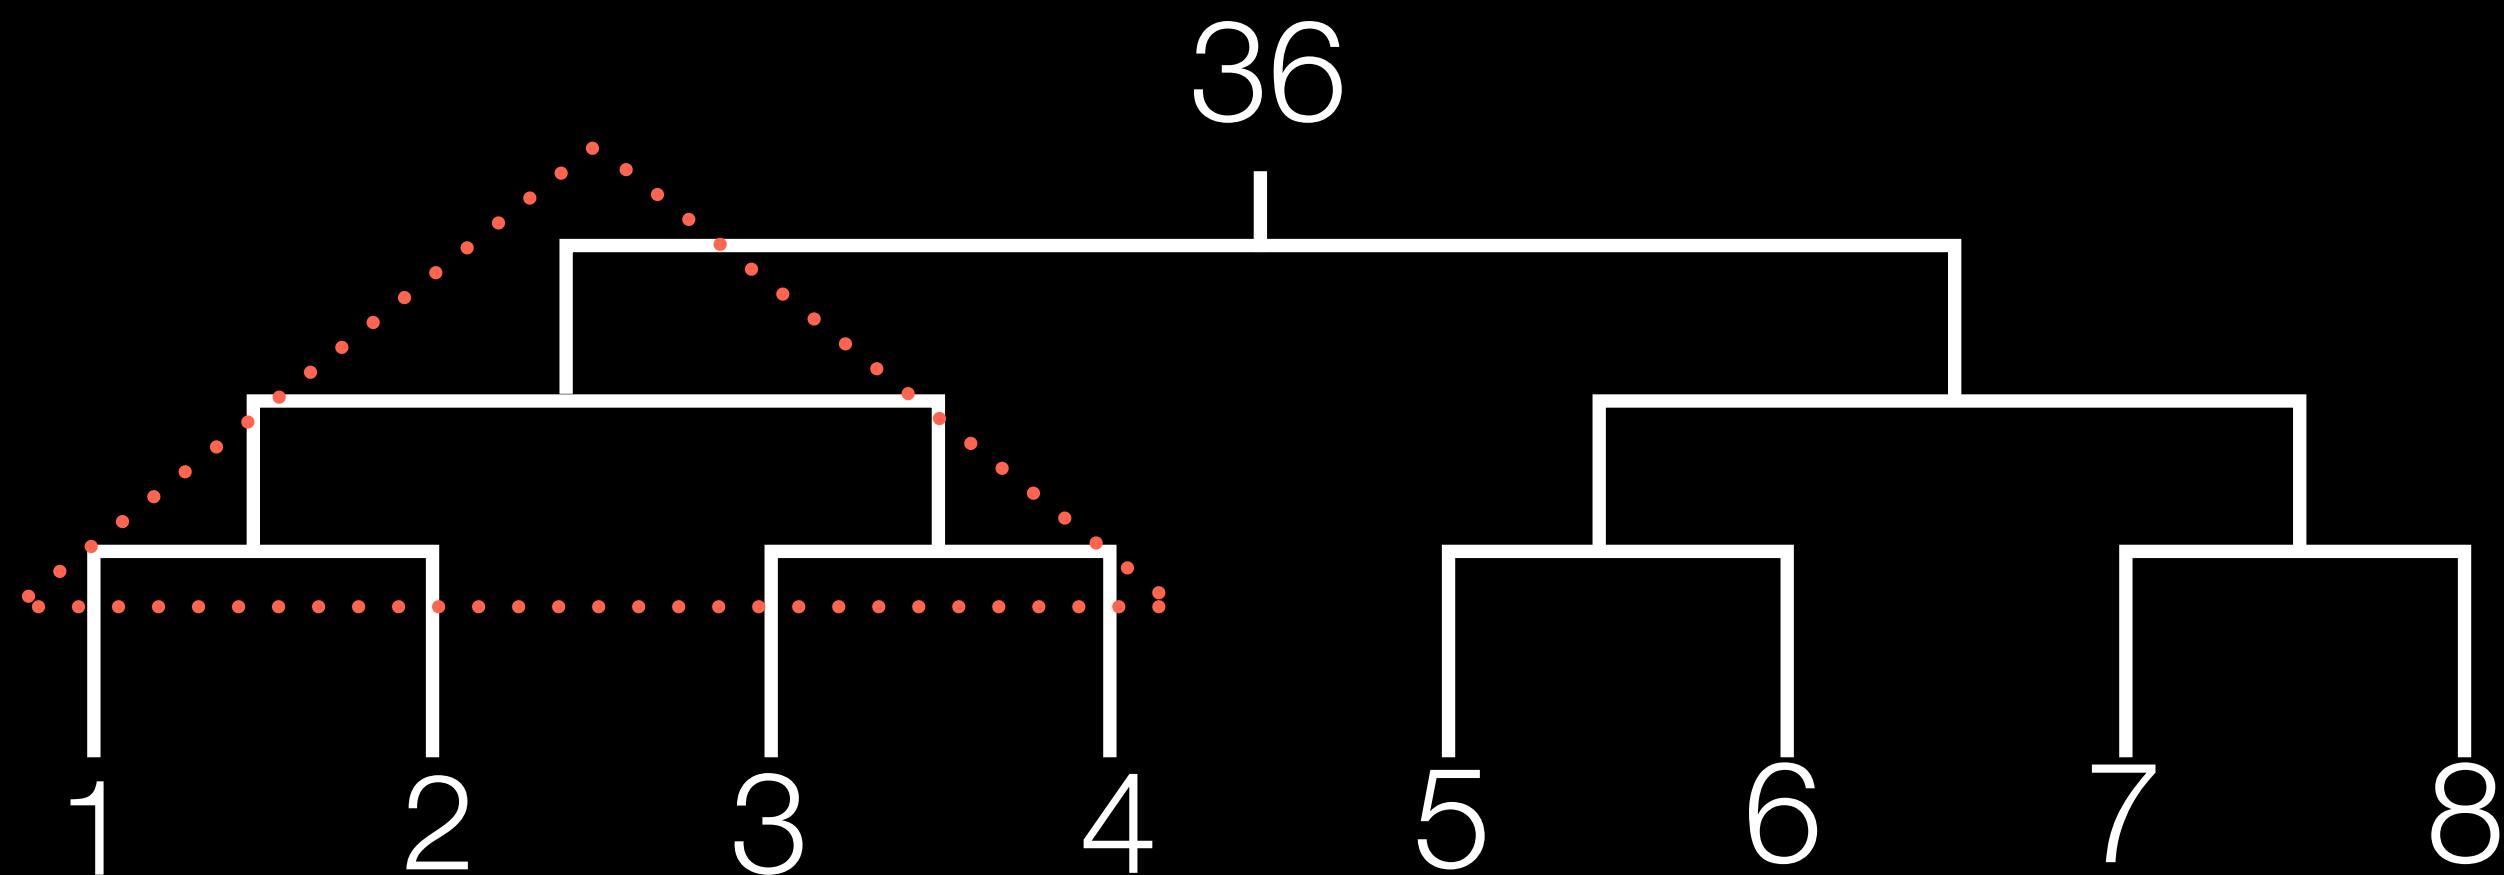
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



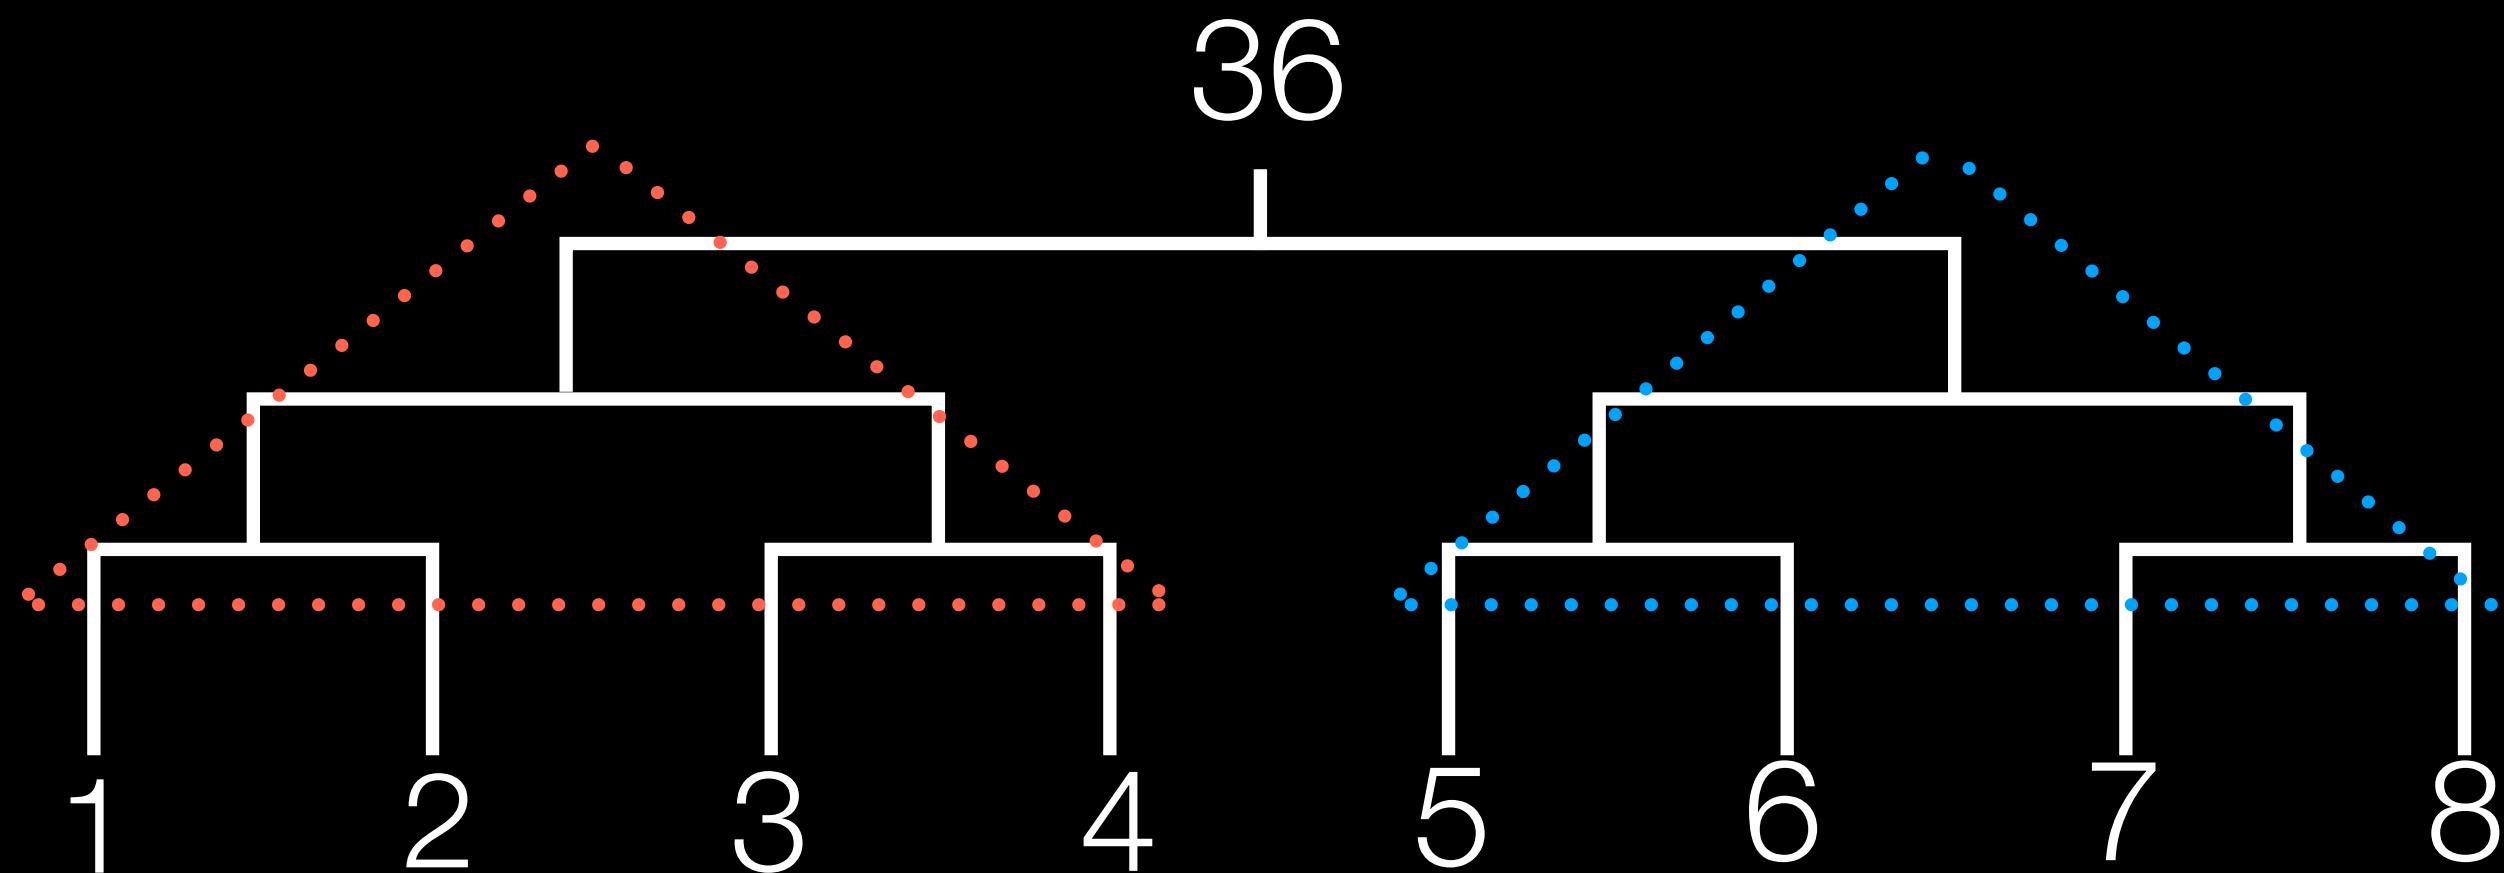
# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



# reduction trees

This just in:  
the players are numbers;  
the matches are addition!



# reduction trees

How would you handle 16 numbers?

1    2    3    4    5    6    7    8    9    10    11    12    13    14    15    16

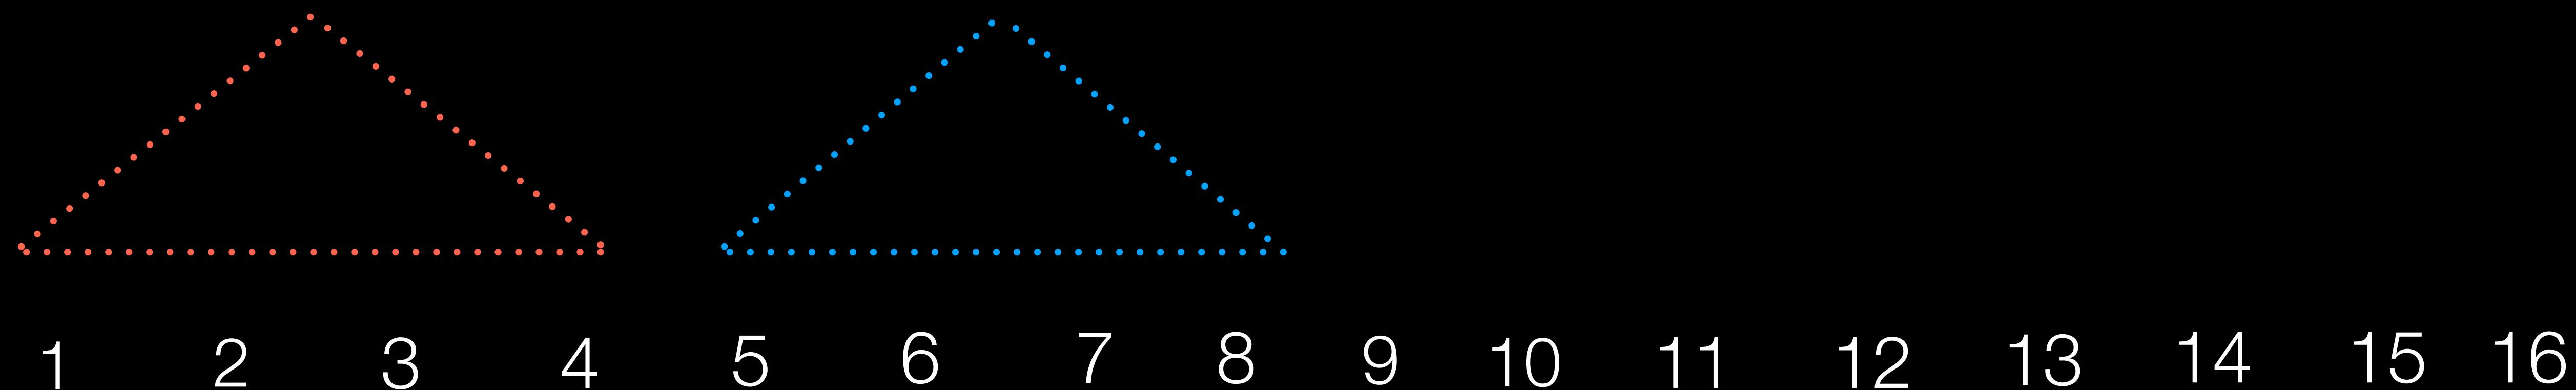
# reduction trees

How would you handle 16 numbers?



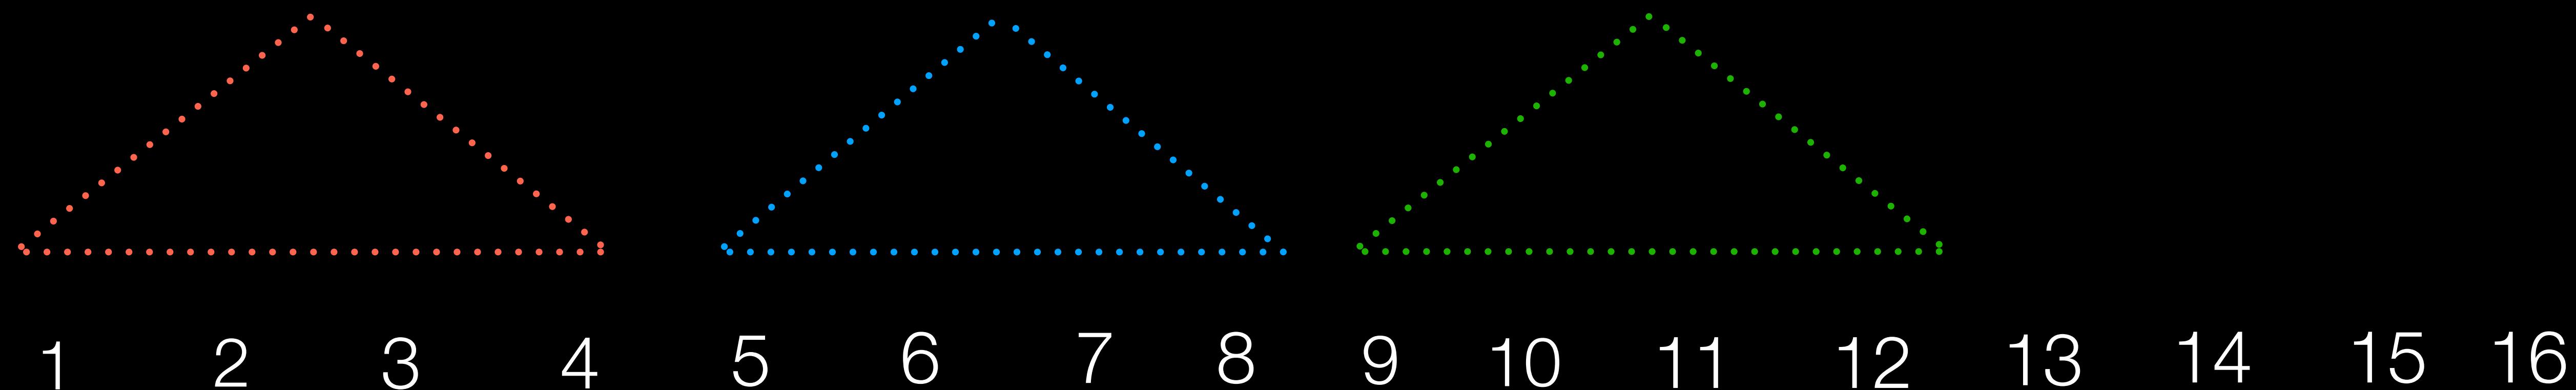
# reduction trees

How would you handle 16 numbers?



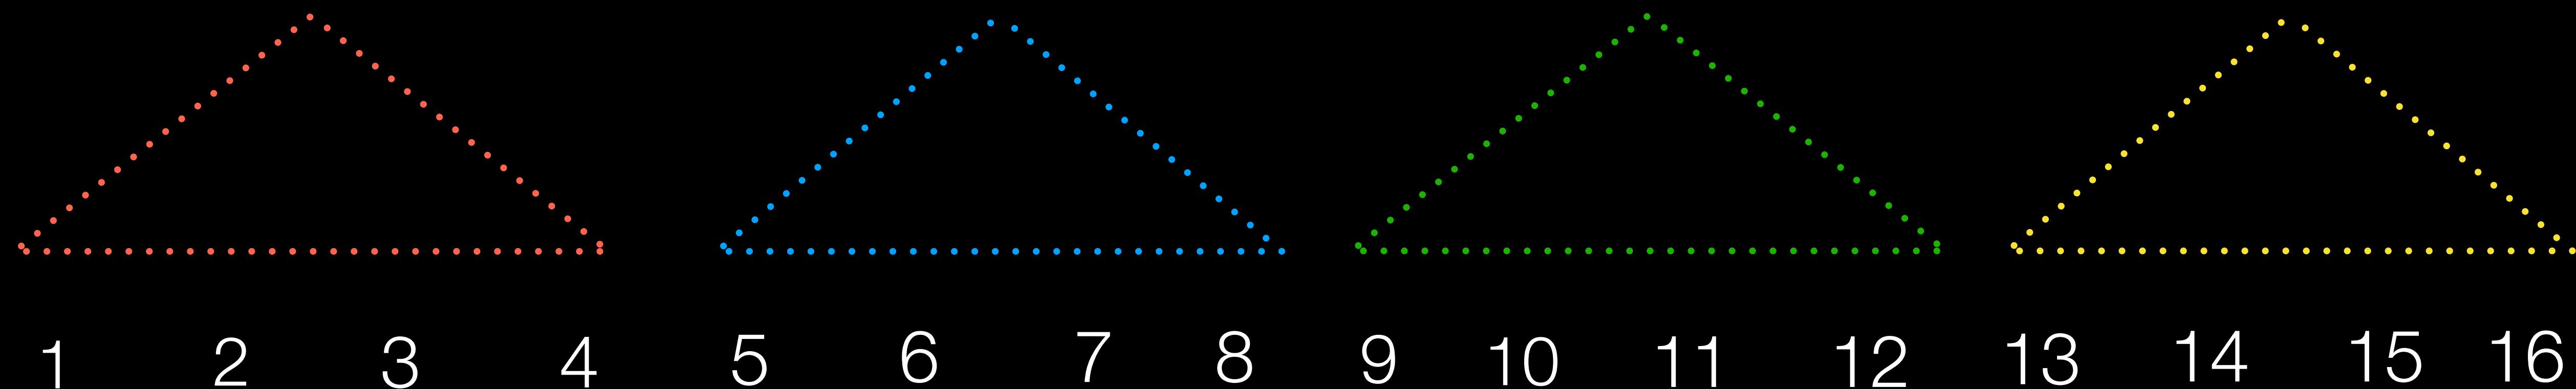
# reduction trees

How would you handle 16 numbers?



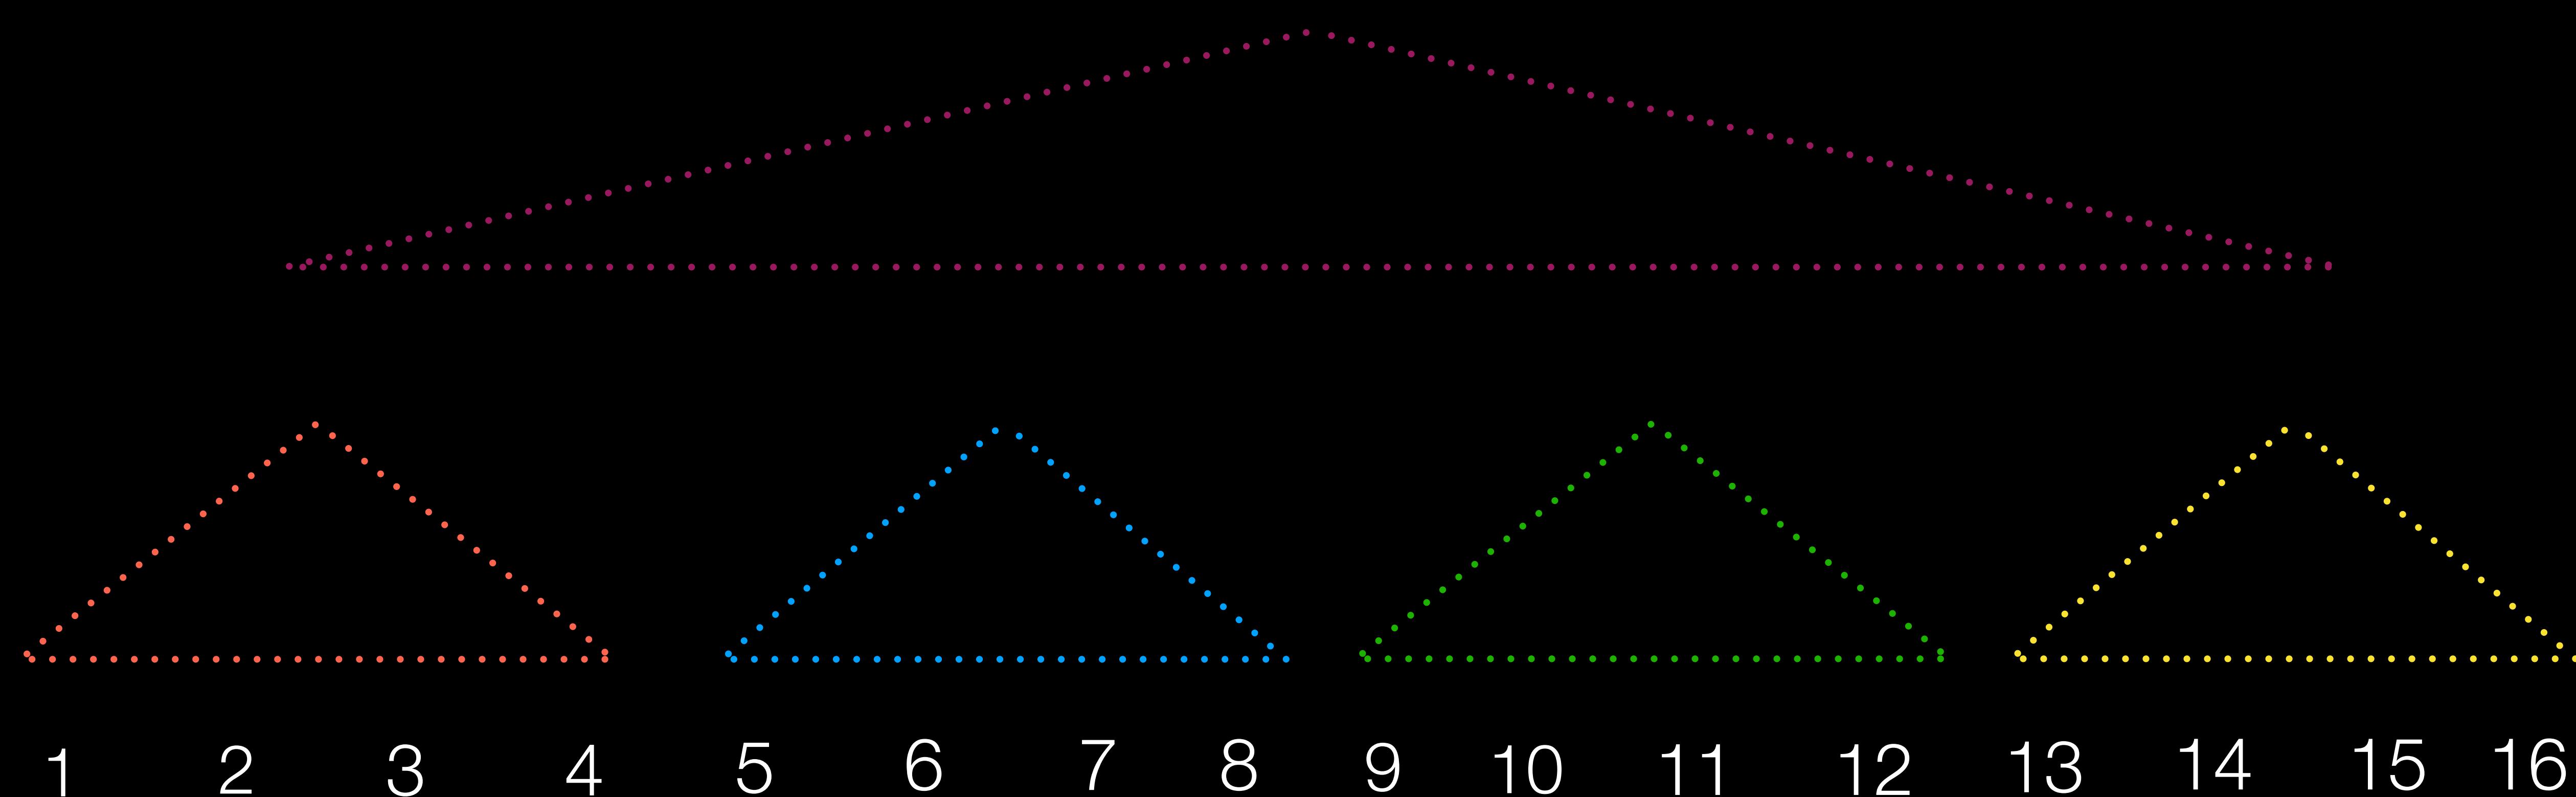
# reduction trees

How would you handle 16 numbers?



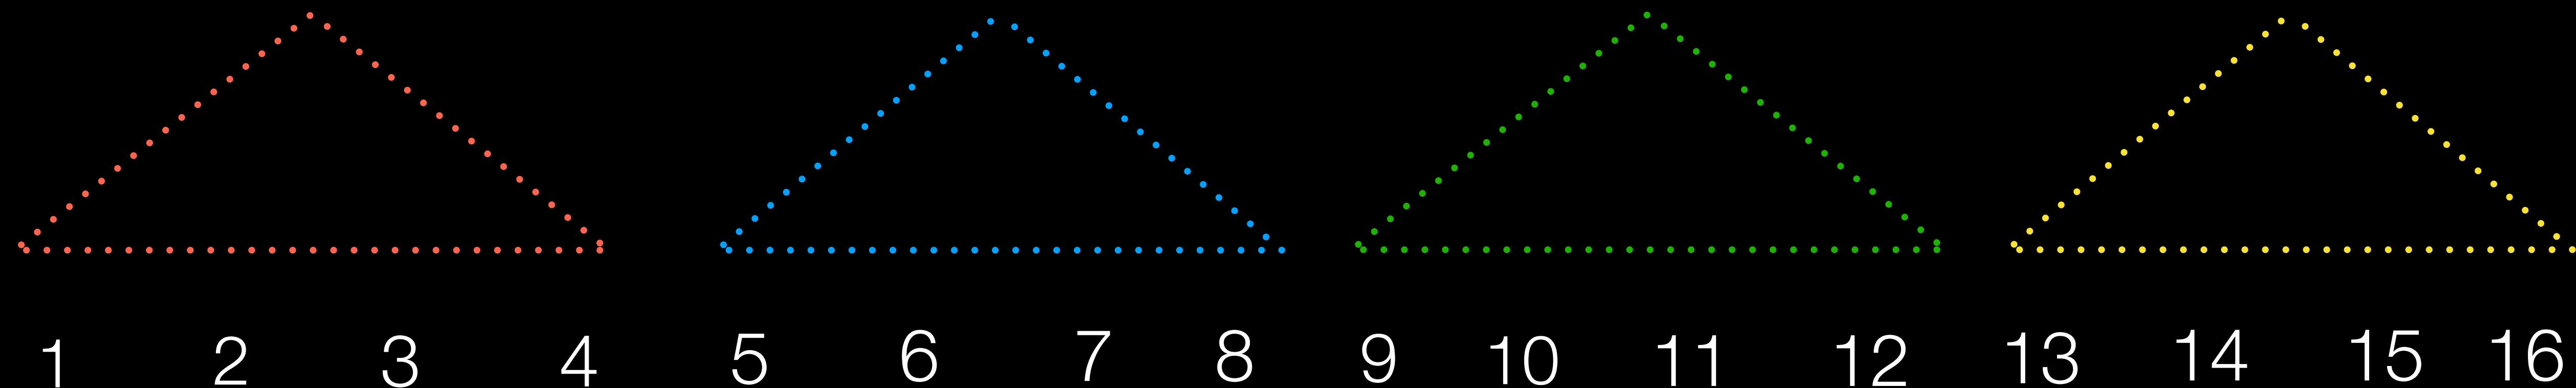
# reduction trees

How would you handle 16 numbers?



# reduction trees, feat. banking

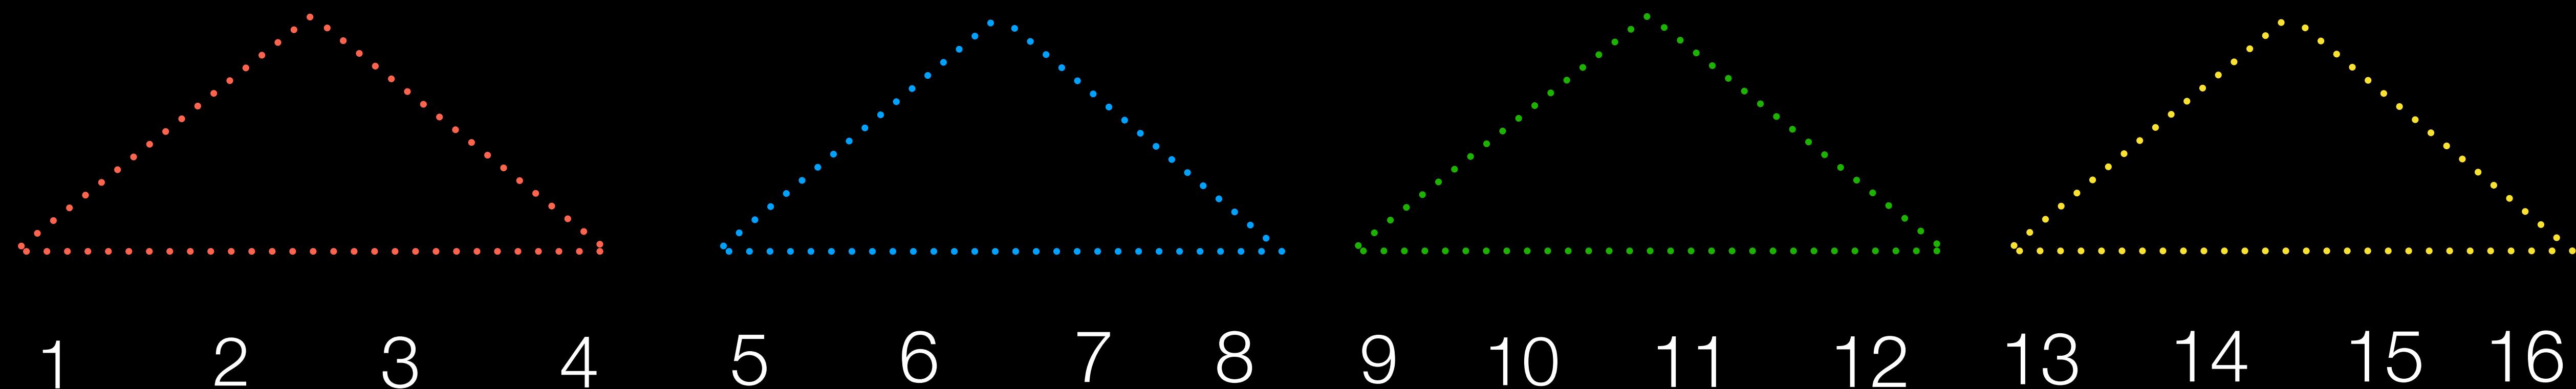
Where would you place the 16 numbers?



# reduction trees, feat. banking

Where would you place the 16 numbers?

mem := [ 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 ]



# reduction trees, feat. banking

Where would you place the 16 numbers?

mem := 

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----

m1 := 

1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16

m2 := 

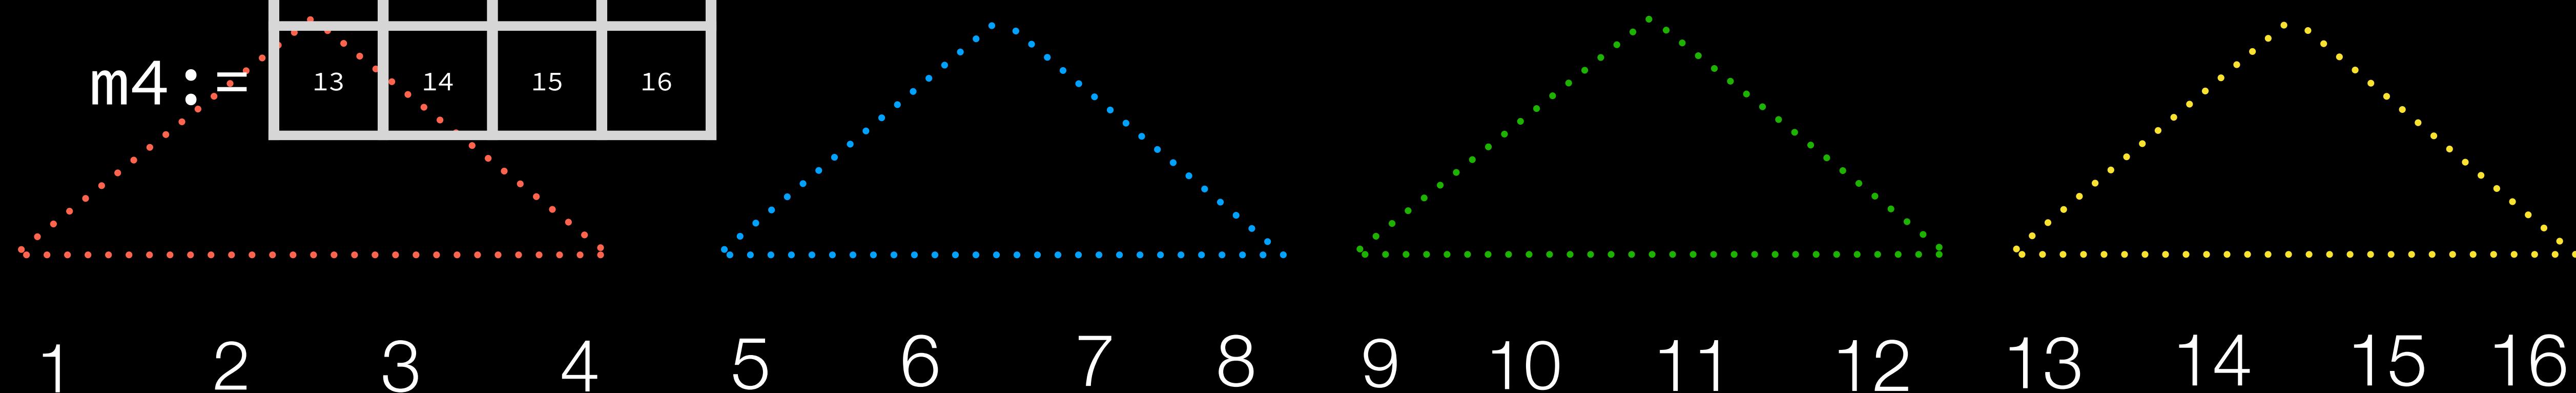
5	6	7	8
---	---	---	---

m3 := 

9	10	11	12
---	----	----	----

m4 := 

13	14	15	16
----	----	----	----

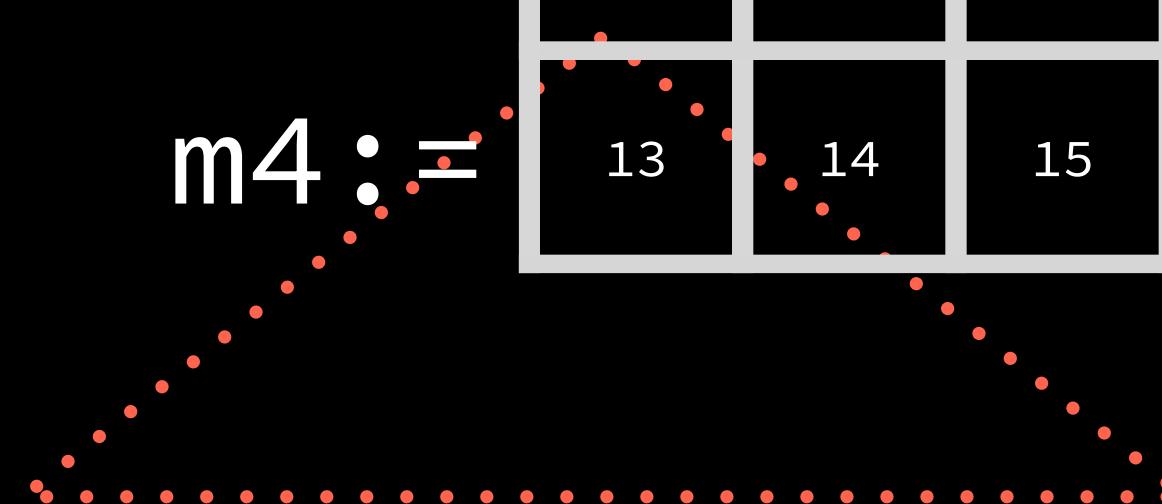


# reduction trees, feat. banking

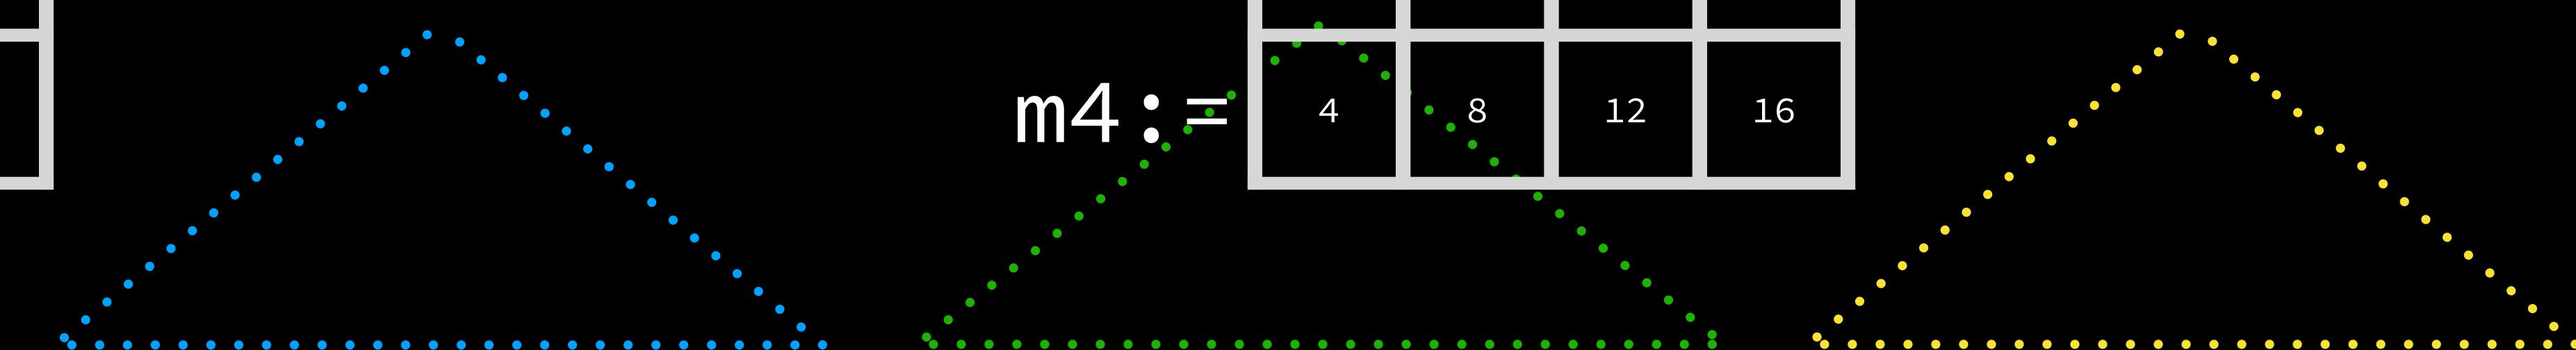
Where would you place the 16 numbers?

mem :=	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16
--------	---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----

m1 :=	1	2	3	4
m2 :=	5	6	7	8
m3 :=	9	10	11	12
m4 :=	13	14	15	16



m1 :=	1	5	9	13
m2 :=	2	6	10	14
m3 :=	3	7	11	15
m4 :=	4	8	12	16



1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16

# arbitrage

What if I want to parallelize the same memory, but with different banking factors?

# arbitrage

What if I want to parallelize the same memory, but with different banking factors?

```
avec := 

|   |   |   |   |   |   |   |   |   |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |
|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 |
|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|


```

```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

# arbitrage

What if I want to parallelize the same memory, but with different banking factors?

```
avec := 

|   |   |   |   |   |   |   |   |   |    |    |    |    |    |    |    |    |    |    |    |    |    |    |    |
|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|
| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 |
|---|---|---|---|---|---|---|---|---|----|----|----|----|----|----|----|----|----|----|----|----|----|----|----|


```

```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

We need to break `avec` into  $\text{lcm}(2, 3) = 6$  banks, and then arbitrate between those.

# arbitrage

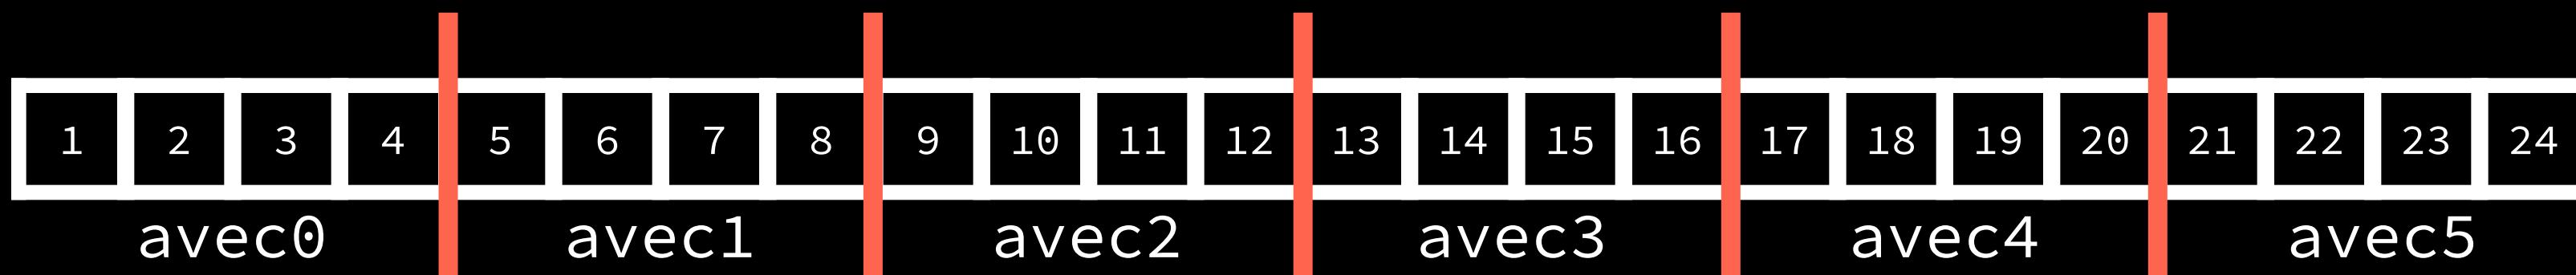
What if I want to parallelize the same memory, but with different banking factors?

```
avec := [1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]
```

```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

We need to break `avec` into  $\text{lcm}(2, 3) = 6$  banks, and then arbitrate between those.



# arbitrage

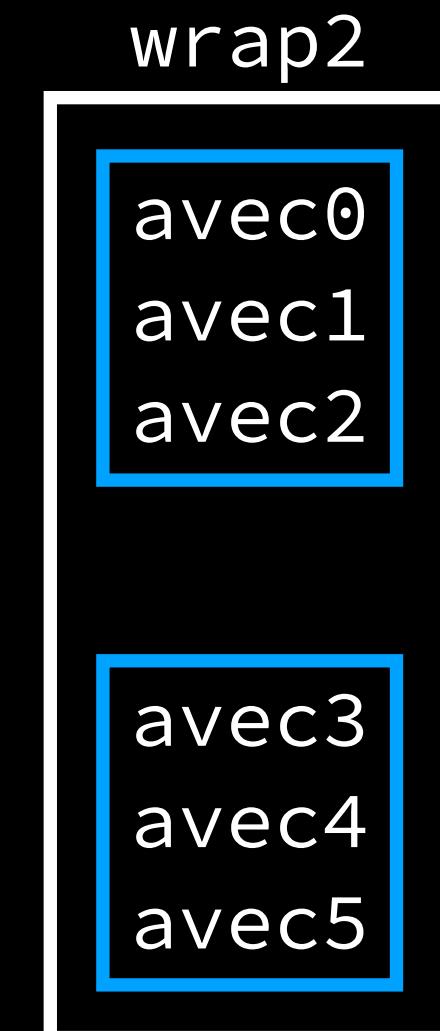
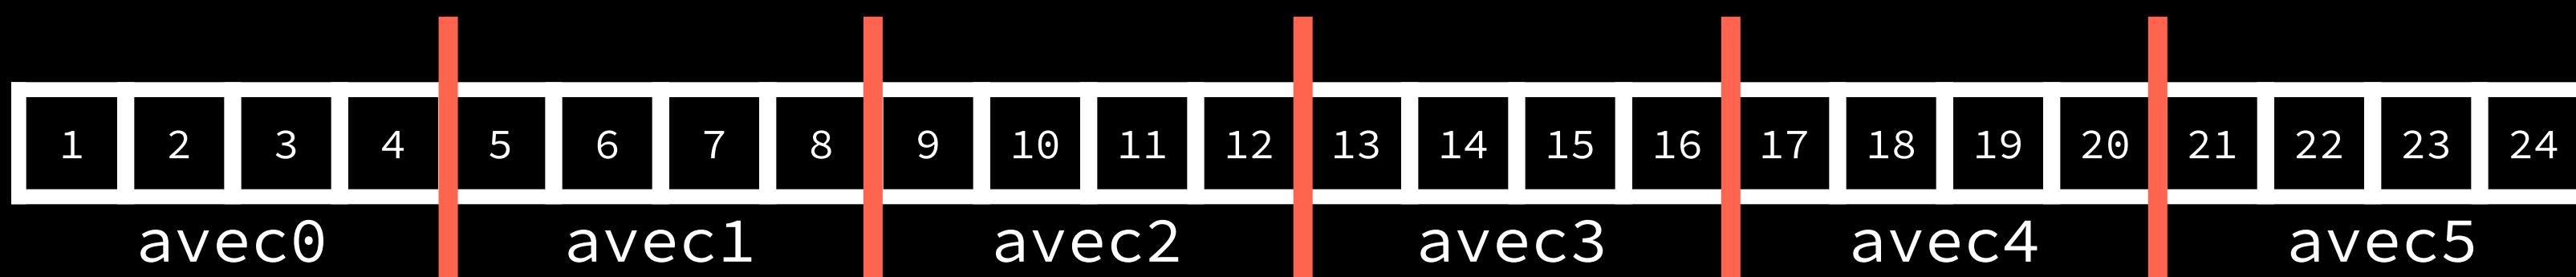
What if I want to parallelize the same memory, but with different banking factors?

avec := [1](#) [2](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) [9](#) [10](#) [11](#) [12](#) [13](#) [14](#) [15](#) [16](#) [17](#) [18](#) [19](#) [20](#) [21](#) [22](#) [23](#) [24](#)

```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

We need to break avec into  
 $\text{lcm}(2, 3) = 6$  banks,  
and then arbitrate between those.



# arbitrage

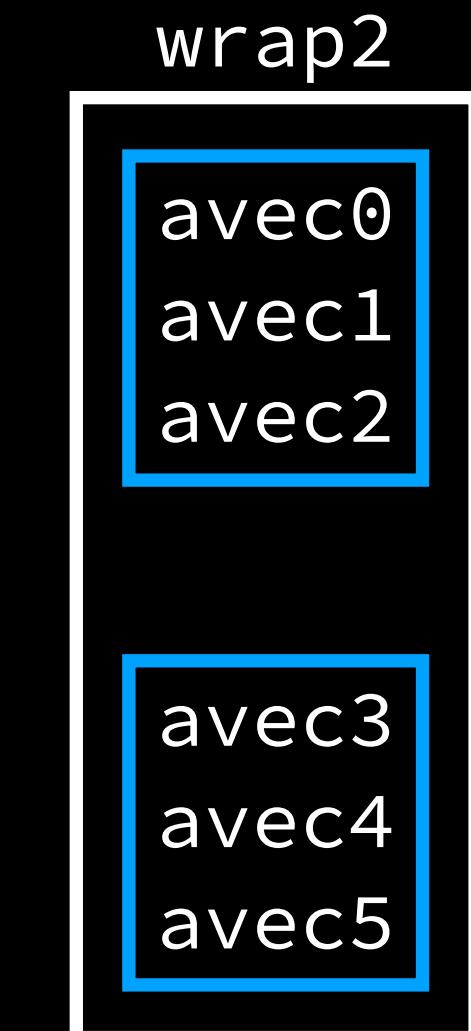
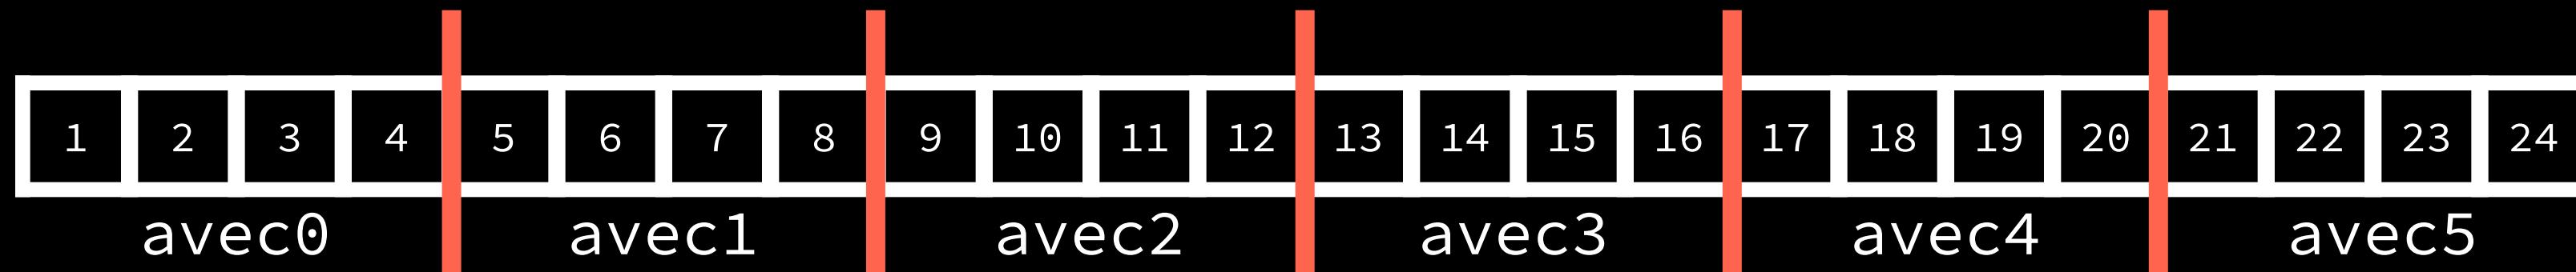
What if I want to parallelize the same memory, but with different banking factors?

```
avec := [1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]
```

```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

We need to break `avec` into  $\text{lcm}(2, 3) = 6$  banks, and then arbitrate between those.



$$0 \leq i < 2$$
$$0 \leq j < 12$$

# arbitrage

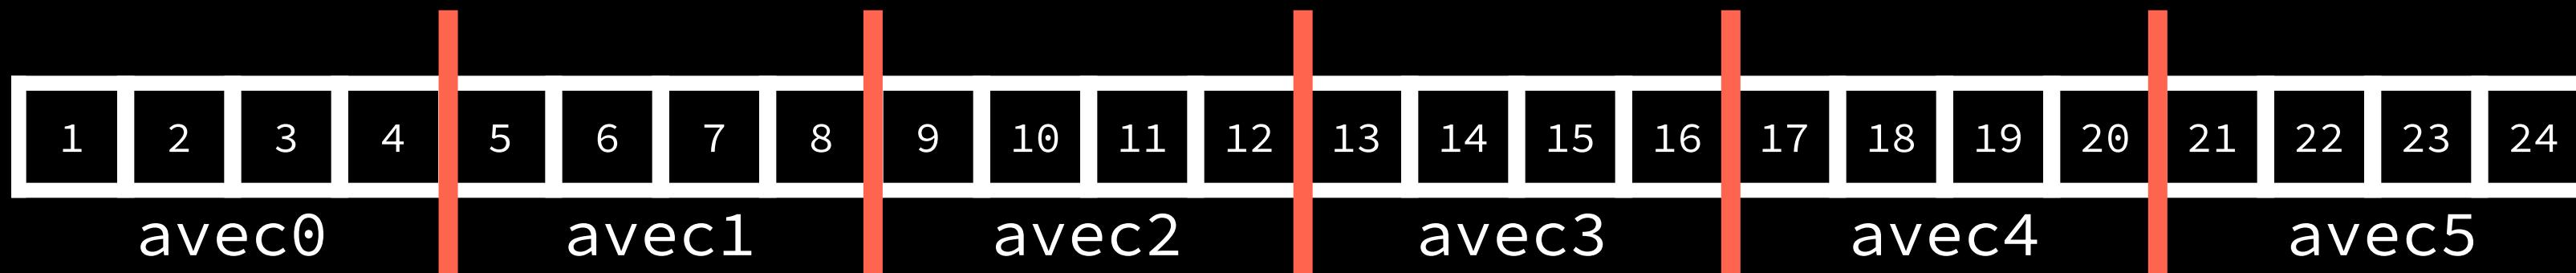
What if I want to parallelize the same memory, but with different banking factors?

```
avec := [1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]
```

```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

We need to break `avec` into  $\text{lcm}(2, 3) = 6$  banks, and then arbitrate between those.



wrap2

```
avec0  
avec1  
avec2
```

```
avec3  
avec4  
avec5
```

$0 \leq i < 2$

$0 \leq j < 12$

wrap3

```
avec0  
avec1
```

```
avec2  
avec3
```

```
avec4  
avec5
```

# arbitrage

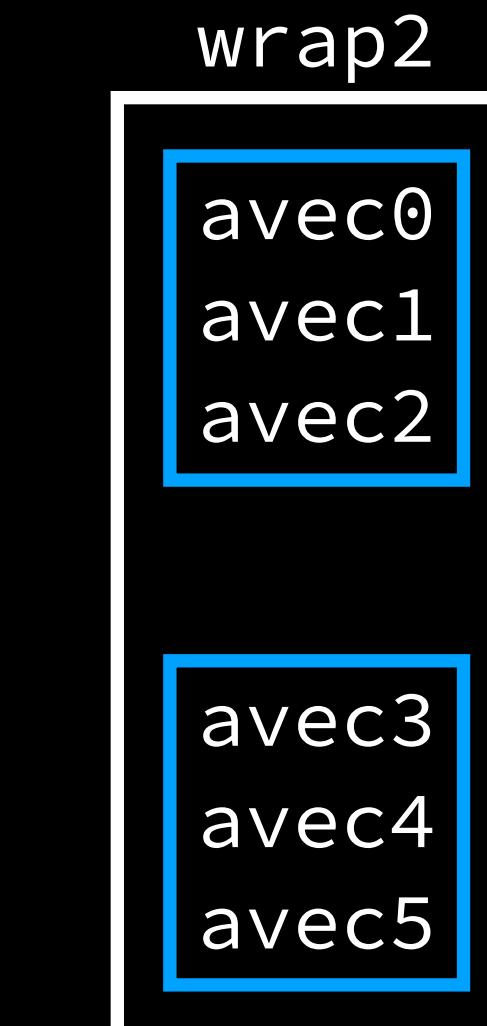
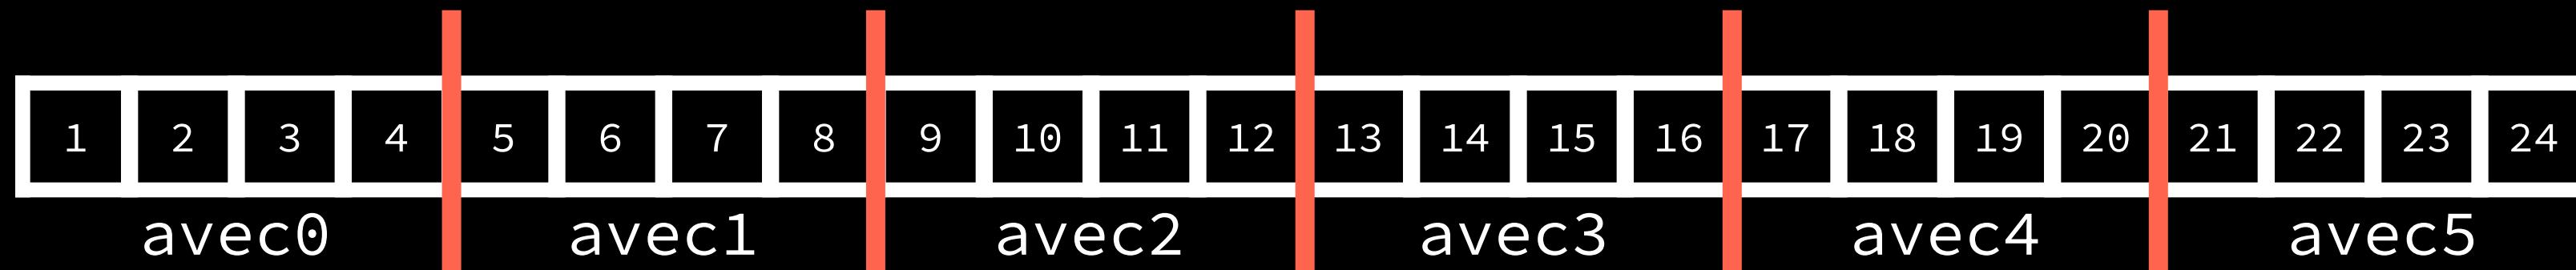
What if I want to parallelize the same memory, but with different banking factors?

```
avec := [1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24]
```

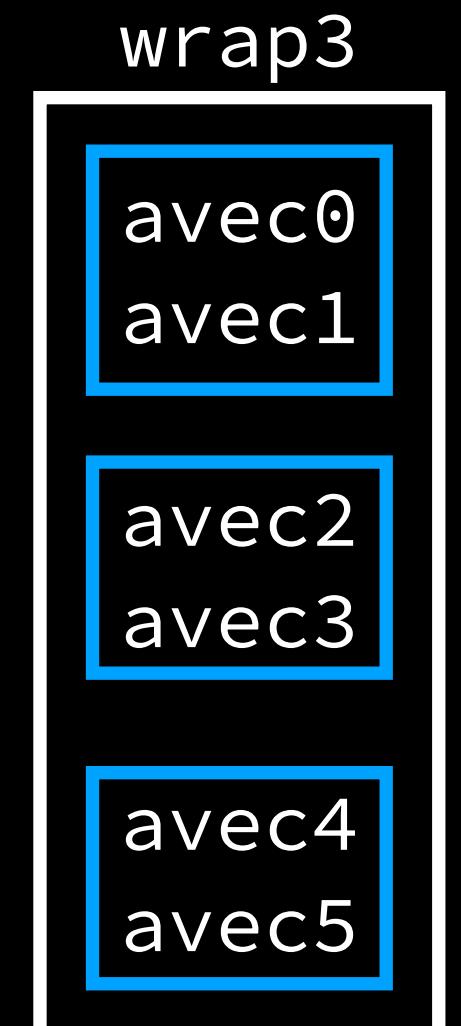
```
squares := map 3 (a <- avec) { a * a }
```

```
add_1 := map 2 (a <- avec) { a + 1 }
```

We need to break `avec` into  $\text{lcm}(2, 3) = 6$  banks, and then arbitrate between those.



$0 \leq i < 2$   
 $0 \leq j < 12$



$0 \leq i < 3$   
 $0 \leq j < 8$

make MrXL better!

# closing remarks

- you can work on calyx frontends/calyx/calyx backends
- here's people working on calyx!
- verification, optimization, lots of buzzwords :D